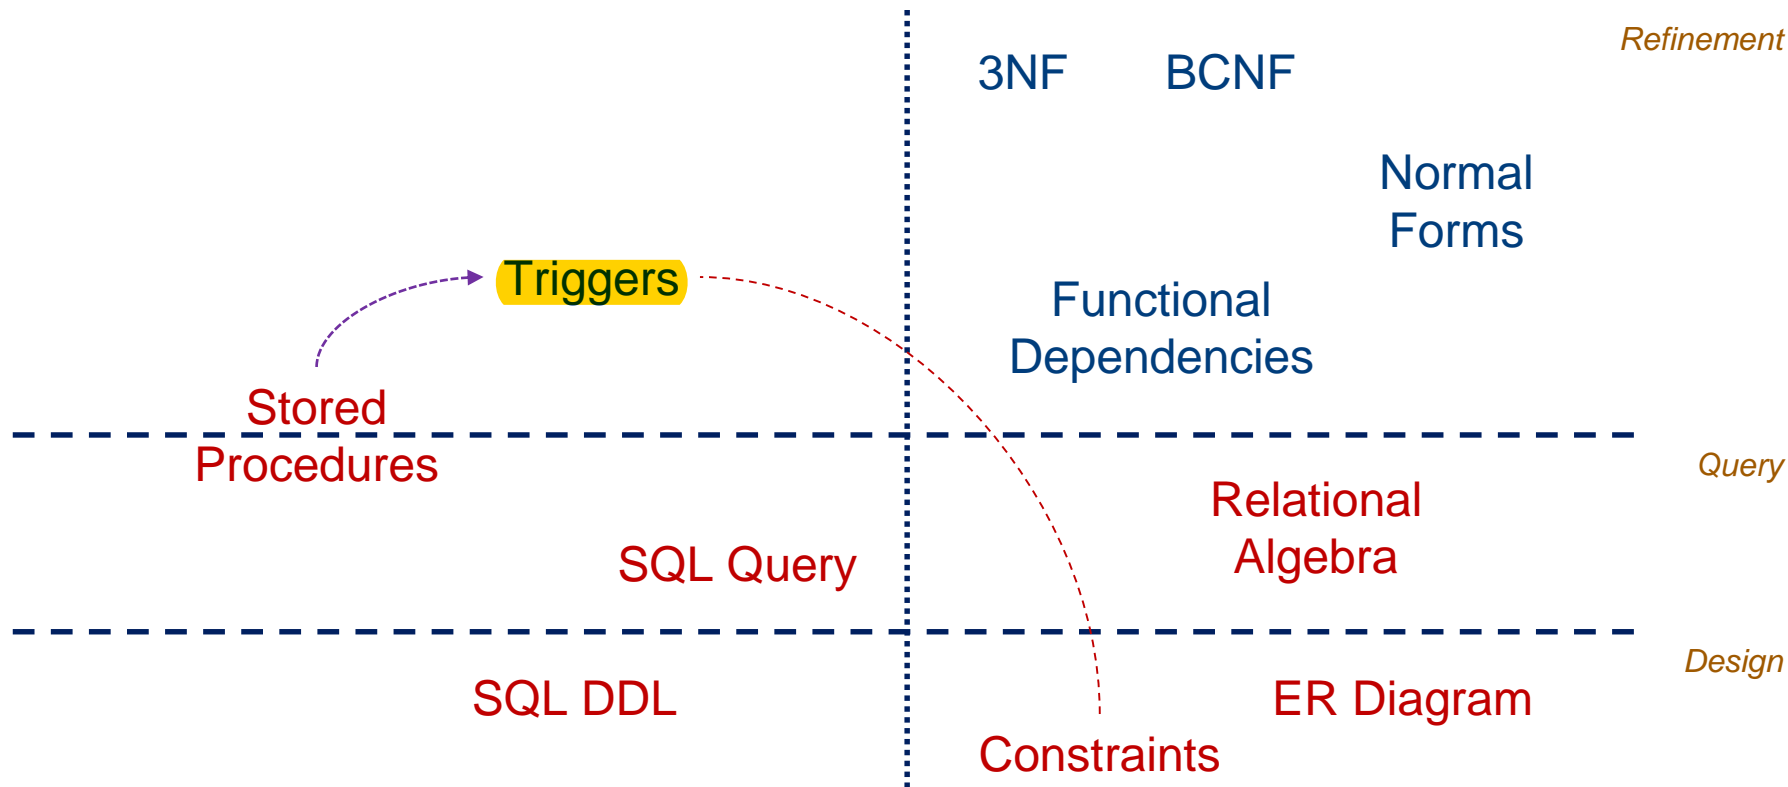


# **CS2102 Database Systems**

## Lecture 9 – Triggers

# Roadmap



# Previous Lectures

- **SQL Functions and Procedures**

- Named operations
- Expressive control structures (*Turing complete*)

- **Constraints**

- Key constraints
- Referential constraints
- Uniqueness
- Check

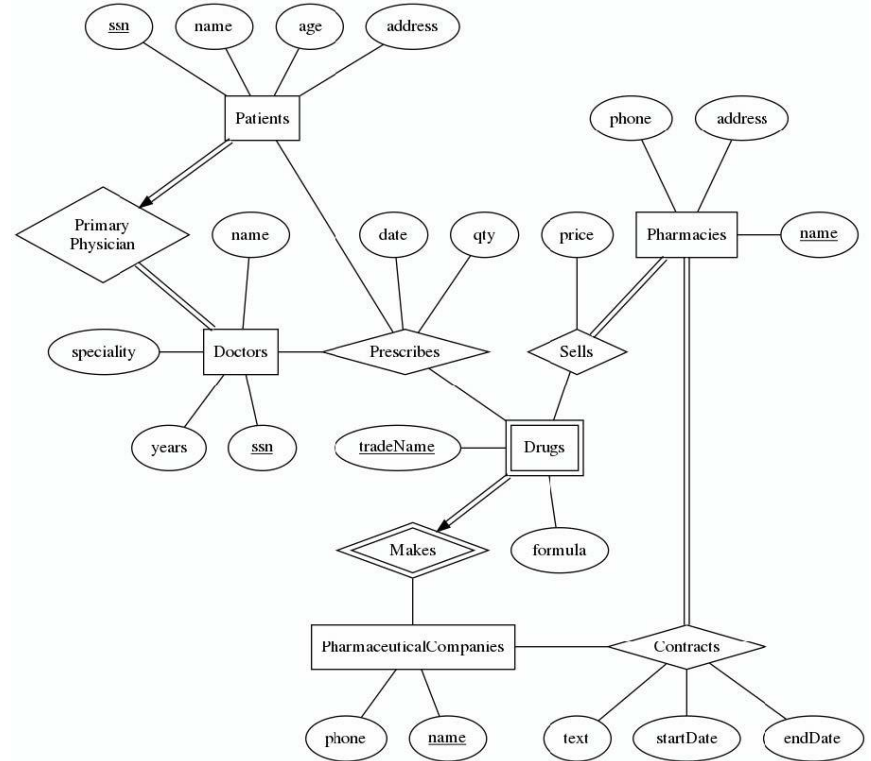
- What if there are constraints that cannot be handled by the constructs above?

# Difficult Constraints

- **Tutorial 03**

- There is exactly one contract between a pharmacy and a pharmaceutical company if and only if that pharmacy sells some drug that is made by that pharmaceutical company.

This is quite complex, instead, we will use a simpler motivating example but Trigger is powerful enough to handle this




# Difficult Constraints

- **ISA Hierarchy**

- A different difficult constraint that cannot be enforced via relational schema only is the ISA hierarchy with non-overlapping constraint (i.e., does not satisfy overlapping constraint).

The solution is:

- For each insertion to E1
- Check if the row is already present in E2
- Same for insert to E2, check on E1
- Checking can be done via trigger



since the uniqueness is not check across tables

# Triggers

# Triggers

- Other Motivating Example

Table "Scores"

<u>Name</u>	Mark
Alice	92
Bob	63
Cathy	58
David	47



Table "Scores\_Log"

<u>Name</u>	Date
Alice	2021-10-01
Bob	2021-10-09
Cathy	2021-10-12
David	2021-10-15

- Suppose we want to log when the marks are entered
  - This should be done automatically
    - The user of the database should not be bothered to write SQL statement to insert
  - This should be done *each time* an insertion occurs regardless of how it is done
  - We want to record the following data
    - The name of the student
    - The date of entry

# Triggers

- **Other Motivating Example**

What we want?

1. Automatic insertion
2. Regardless of how

Table "Scores"

<u>Name</u>	Mark
Alice	92
Bob	63
Cathy	58
David	47



Table "Scores\_Log"

<u>Name</u>	Date
Alice	2021-10-01
Bob	2021-10-09
Cathy	2021-10-12
David	2021-10-15

- Suppose we want to log when the marks are entered

- **Idea?**

- A procedure that enters both data

- ❖ What if the user forgot to use the procedure?

```
CREATE OR REPLACE PROCEDURE enter_data
    (Name TEXT, Mark INT)
AS $$

    INSERT INTO Scores VALUES (Name, Mark);
    INSERT INTO Scores_Log
        VALUES (Name, CURRENT_DATE);

$$ LANGUAGE sql;
```



# Triggers

- Other Motivating Example

Issues?

1. How to check insertion?
2. How to get the name?

Table "Scores"

<u>Name</u>	Mark
Alice	92
Bob	63
Cathy	58
David	47




Table "Scores\_Log"

<u>Name</u>	Date
Alice	2021-10-01
Bob	2021-10-09
Cathy	2021-10-12
David	2021-10-15

- Suppose we want to log when the marks are entered

- What we want?

- A procedure to insert into Scores\_Log
    - But called *automatically*

```
CREATE OR REPLACE PROCEDURE log_score()  
AS $$  
BEGIN  
    IF (there is an insertion into Scores) THEN  
        INSERT INTO Scores_Log  
            VALUES (Name, CURRENT_DATE);  
    END IF;  
END;  
$$ LANGUAGE plpgsql;
```

① How to check insertion?

② How to get the name?


# Triggers

- **Basic of Triggers**

- Trigger is an **event-condition-action** (ECA) rule
  - When **event** occurs
  - Test **condition**
  - If satisfied, execute **action**

- **Example**

- **Event** New tuple inserted into Scores
  - **Condition** Nothing (*always execute action*)
  - **Action** Insert into Scores\_Log
- } *Trigger*
- } *Trigger Function*



# Triggers

- Other Motivating Example

Trigger

```
CREATE TRIGGER score_log
AFTER INSERT ON Scores
FOR EACH ROW EXECUTE FUNCTION
log_score();
```

Table "Scores"

<u>Name</u>	Mark
Alice	92
Bob	63
Cathy	58
David	47




Table "Scores\_Log"

<u>Name</u>	Date
Alice	2021-10-01
Bob	2021-10-09
Cathy	2021-10-12
David	2021-10-15

- Suppose we want to log when the marks are entered

- **Trigger Function**

- Actions to run when event occurred and conditions satisfied

```
CREATE OR REPLACE FUNCTION log_score()
RETURNS TRIGGER AS $$
BEGIN

    INSERT INTO Scores_Log
        VALUES (NEW.Name, CURRENT_DATE);

    RETURN NULL;
END;
$$ LANGUAGE plpgsql;
```

Trigger Function

# Triggers

- The Trigger

Trigger

```
CREATE TRIGGER score_log  
AFTER INSERT ON Scores  
FOR EACH ROW EXECUTE FUNCTION  
log_score();
```

- This tells the database to

- Watch out for insertion on Score
- ~~Check for certain condition~~ *(we'll come back to this later)*
- Call the function log\_score() after each insertion of a tuple
- There are other options, but that will be discussed later on

Table "Scores"

<u>Name</u>	Mark
Alice	92
Bob	63
Cathy	58
David	47



Table "Scores\_Log"

<u>Name</u>	Date
Alice	2021-10-01
Bob	2021-10-09
Cathy	2021-10-12
David	2021-10-15

# Triggers

- The Trigger Function

- RETURNS **TRIGGER** indicates that this is a trigger function
  - Can only RETURNS **TRIGGER**
- **NEW** refers to the new row inserted into Scores
  - Only accessible by trigger functions
- Other accessible data?
  - **TG\_OP**
  - **TG\_TABLE\_NAME**
  - **OLD**
  - ...

Table "Scores"

<u>Name</u>	Mark
Alice	92
Bob	63
Cathy	58
David	47



Table "Scores\_Log"

<u>Name</u>	Date
Alice	2021-10-01
Bob	2021-10-09
Cathy	2021-10-12
David	2021-10-15

```
CREATE OR REPLACE FUNCTION log_score()  
RETURNS TRIGGER AS $$  
BEGIN  
  
    INSERT INTO Scores_Log  
        VALUES (NEW.Name, CURRENT_DATE);  
  
    RETURN NULL;  
END;  
$$ LANGUAGE plpgsql;
```

Trigger Function

# Triggers

- Syntax

Trigger

```
CREATE TRIGGER <trigger_name>  
<timing> <event> ON <table>  
FOR EACH <granularity>  
EXECUTE FUNCTION <function_name>();
```

only applicable to this table

when triggered, this  
trigger calls this function

## NOTE

<timing>, <event> and  
<granularity> are called  
trigger options.

This is what we will discuss  
next.

```
CREATE OR REPLACE FUNCTION <function_name>()  
RETURNS TRIGGER AS $$  
BEGIN
```

can only return TRIGGER

<code>

this code is then executed

```
END;  
$$ LANGUAGE plpgsql;
```

Trigger Function

# Triggers Options

- Trigger Options

- Events:

- `INSERT ON table`
    - `DELETE ON table`
    - `UPDATE [OF column] ON table`

*TG\_OP* {  
*'INSERT'*  
*'DELETE'*  
*'UPDATE'*

- Timing:

- `AFTER` or `BEFORE` *(the triggering event)*
    - `INSTEAD OF` *(the triggering event on views)*

- Granularity:

- `FOR EACH ROW` *(modified)*
    - `FOR EACH STATEMENT` *(that performs the modification)*

# Triggers Events

INSERT

DELETE

UPDATE



# Triggers Event

- **Another Motivating Example**

Table "Scores"

<u>Name</u>	Mark
Alice	92
Bob	63
Cathy	58
David	47




Table "Scores\_Log2"

<u>Name</u>	Op	Date
Alice	Insert	2021-10-01
Bob	Delete	2021-10-09
Cathy	Update	2021-10-12
David	Insert	2021-10-15

- Suppose we want to log when the marks are entered
  - This should be done automatically
    - The user of the database should not be bothered to write SQL statement to insert
  - This should be done *each time* an insertion occurs regardless of how it is done
  - We want to record the following data
    - The name of the student
    - The date of entry
    - The operation performed

# Triggers Event

- The Trigger Function

```
CREATE OR REPLACE FUNCTION log_score2() RETURNS TRIGGER AS $$
BEGIN
    IF (TG_OP = 'INSERT') THEN
        INSERT INTO Scores_Log2 VALUES (NEW.Name, 'Insert', CURRENT_DATE);
        RETURN NEW;
    ELSIF (TG_OP = 'DELETE') THEN
        INSERT INTO Scores_Log2 VALUES (OLD.Name, 'Delete', CURRENT_DATE);
        RETURN OLD;
    ELSIF (TG_OP = 'UPDATE') THEN
        INSERT INTO Scores_Log2 VALUES (NEW.Name, 'Update', CURRENT_DATE);
        RETURN NEW;
    END IF;
END;
$$ LANGUAGE plpgsql;
```

Trigger Function

## Question

Can we RETURN NULL like before?

*Return value has effect, we will discuss later.*

# Triggers Event

- The Trigger Function

```
CREATE OR REPLACE FUNCTION log_score2() RETURNS TRIGGER AS $$  
BEGIN  
    IF      (TG_OP = 'INSERT') THEN ...  
    ELSIF (TG_OP = 'DELETE') THEN ...  
    ELSIF (TG_OP = 'UPDATE') THEN ...  
    END IF;  
END;  
$$ LANGUAGE plpgsql;
```

In the case of ISA, most likely will use  
AFTER INSERT OR UPDATE

Trigger Function

- The Trigger

```
CREATE TRIGGER score_log2  
AFTER INSERT OR DELETE OR UPDATE ON Scores  
FOR EACH ROW EXECUTE FUNCTION log_score2();
```

Trigger

# Triggers Timing

AFTER  
BEFORE  
INSTEAD OF

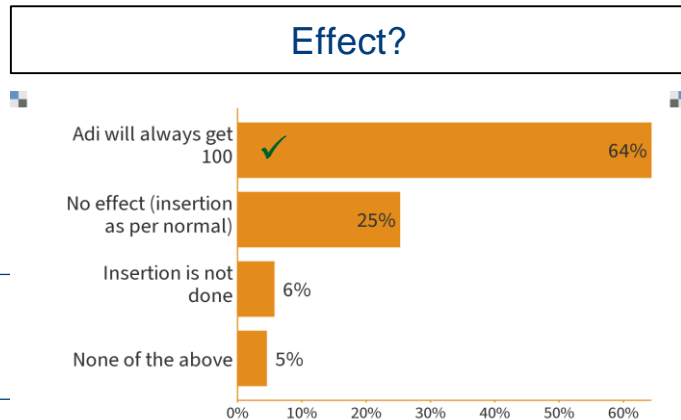
# Triggers Timing

Table "Scores"

<u>Name</u>	Mark
...	...

- The Trigger Function

```
CREATE OR REPLACE FUNCTION give_adi_full_mark() RETURNS TRIGGER AS $$  
BEGIN  
    IF (NEW.Name = 'Adi') THEN  
        NEW.Mark := 100;  
    END IF;    RETURN NEW;  
END;  
$$ LANGUAGE plpgsql;
```



Trigger Function

```
CREATE TRIGGER adi_should_get_full_mark  
BEFORE INSERT ON Scores  
FOR EACH ROW EXECUTE FUNCTION give_adi_full_mark();
```

**NOTE**  
NEW is a tuple (Adi, 100).  
Hence, the result.

Trigger

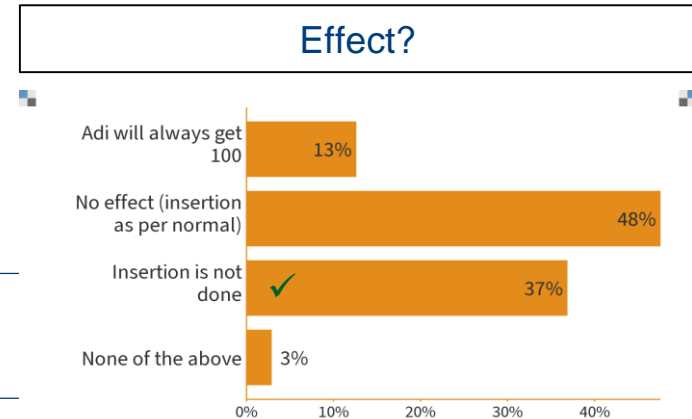
# Triggers Return Value

Table "Scores"

<u>Name</u>	Mark
...	...

- The Trigger Function

```
CREATE OR REPLACE FUNCTION give_adi_full_mark() RETURNS TRIGGER AS $$  
BEGIN  
    IF (NEW.Name = 'Adi') THEN  
        NEW.Mark := 100;  
    END IF;    RETURN NULL;  
END;  
$$ LANGUAGE plpgsql;
```



Trigger Function

- The Trigger

```
CREATE TRIGGER adi_should_get_full_mark  
BEFORE INSERT ON Scores  
FOR EACH ROW EXECUTE FUNCTION give_adi_full_mark();
```

**NOTE**  
NULL stops BEFORE trigger.  
*Hence, the result.*

Trigger

# Triggers Return Value

## ERRATA

Here, OLD is initially NULL.  
So it is the same as return NULL.

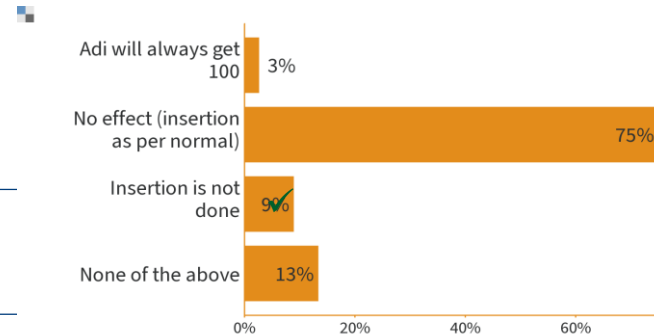
Table "Scores"

Name	Mark
...	...

- The Trigger Function

```
CREATE OR REPLACE FUNCTION give_adi_full_mark() RETURNS TRIGGER AS $$  
BEGIN  
    IF (NEW.Name = 'Adi') THEN  
        NEW.Mark := 100;  
    END IF;    RETURN OLD;  
END;  
$$ LANGUAGE plpgsql;
```

### Effect?



Trigger Function

- The Trigger

```
CREATE TRIGGER adi_should_get_full_mark  
BEFORE INSERT ON Scores  
FOR EACH ROW EXECUTE FUNCTION give_adi_full_mark();
```

## NOTE

OLD is NULL for BEFORE trigger.  
*Hence, the result.*

Trigger


# Triggers Return Value

Table "Scores"

<u>Name</u>	Mark
...	...

- The Trigger Function

```
CREATE OR REPLACE FUNCTION give_adi_full_mark() RETURNS TRIGGER AS $$  
BEGIN  
    IF (NEW.Name = 'Adi') THEN  
        OLD.Name := 'Adi';    OLD.Mark := 100;  
    END IF;    RETURN OLD;  
END;  
$$ LANGUAGE plpgsql;
```



Trigger Function

- The Trigger

```
CREATE TRIGGER adi_should_get_full_mark  
BEFORE INSERT ON Scores  
FOR EACH ROW EXECUTE FUNCTION give_adi_full_mark();
```

**NOTE**  
OLD is NULL for non-Adi.  
Old is (Adi, 100) for Adi.  
*Hence, the result.*

Trigger



# Triggers Return Value

- Transition Variables

- **NEW** The modified row *after* the triggering event
- **OLD** The modified row *before* the triggering event
- ❖ Not all make sense all the time

**NOTE**  
✗ is NULL

	NEW	OLD
INSERT	✓	✗
UPDATE	✓	✓
DELETE	✗	✓

- Effect of Return Value

	NULL tuple	non-NULL tuple t
BEFORE INSERT	No tuple inserted	Tuple <i>t</i> will be inserted
BEFORE UPDATE	No tuple updated	Tuple <i>t</i> will be the updated tuple
BEFORE DELETE	No deletion performed	Deletion proceeds as normal
AFTER INSERT	<i>Does not matter</i> <i>It's done already, cannot be changed now</i>	
AFTER UPDATE		
AFTER DELETE		

# Triggers Timing

- **INSTEAD OF Trigger**
  - This kind of trigger can only be defined on **VIEWS**
- **First Off...**
  - What is a **VIEW**?
    - A “*virtual*” table defined by a query to compute the view
    - Can be used in queries just like a regular table
  - Why use a **VIEW**?
    - Hide data and/or complexity from users
    - Logical data independence
    - Captures common query to be reused
  - ❖ Real database applications use **tons** of VIEW

# Triggers Timing

- **INSTEAD OF Trigger**
  - This kind of trigger can only be defined on **VIEWS**
- **Secondly ...**
  - Why do we want to modify a **VIEW**?
    - User sees it as a table and they can modify table
  - Does it even make sense?
    - In most cases, we cannot modify **VIEW** directly
    - But we can modify the underlying table

# Triggers Timing

- INSTEAD OF Trigger

- This kind of trigger can only be defined on **VIEWS**

- Modifying a VIEW

- Simple case

- ```
CREATE VIEW Students AS  
SELECT Name FROM Scores;
```

- ```
DELETE FROM Students WHERE Name = 'Alice';
```

  
**translates to**  

```
DELETE FROM Scores WHERE Name = 'Alice';
```

# Triggers Timing

- INSTEAD OF Trigger

- This kind of trigger can only be defined on **VIEWS**

- Modifying a VIEW

- Impossible case

- ```
CREATE VIEW A_Students AS  
SELECT Name, Mark FROM Scores WHERE Mark >= 70;
```
- ```
INSERT INTO A_Students VALUES ('Bryan', 69);
```

  
insertion does not affect the VIEW, only the underlying table

# Triggers Timing

- INSTEAD OF Trigger

- This kind of trigger can only be defined on **VIEWS**

- Modifying a VIEW

- Too many possible case
  - `CREATE VIEW Top_Marks(Mark) AS  
SELECT MAX(Mark) FROM Scores;`
  - `UPDATE Top_Marks SET Mark = 80;`  
what if more than one people have top marks?

# Triggers Timing

- **INSTEAD OF Trigger**
  - This kind of trigger can only be defined on **VIEWS**
- **OK, so what to do here?**
  - Let's say we want to modify all people with the maximum mark
  - We modify the underlying table instead
    - Use INSTEAD OF trigger

```
CREATE TRIGGER modify_top_mark  
INSTEAD OF UPDATE ON Top_Marks  
FOR EACH ROW EXECUTE FUNCTION update_top_mark();
```

Trigger

# Triggers Timing

Table "Scores"

<u>Name</u>	Mark
...	...

- The Trigger Function

```
CREATE OR REPLACE FUNCTION update_top_mark() RETURNS TRIGGER AS $$  
BEGIN  
    UPDATE Scores SET Mark = NEW.Mark WHERE Mark = OLD.Mark;  
    RETURN NEW;  
END;  
$$ LANGUAGE plpgsql;
```

Trigger Function

### Return Value

- **NULL**
  - Ignore all operations on current row
- **Non-NULL**
  - Signals the database to proceed as normal

- The Trigger

```
CREATE TRIGGER modify_top_mark  
INSTEAD OF UPDATE ON Top_Marks  
FOR EACH ROW EXECUTE FUNCTION update_top_mark();
```

Trigger



# Triggers Granularity

FOR EACH ROW

FOR EACH STATEMENT

# Triggers Granularity

- **Row-Level Trigger**

*(FOR EACH ROW)*

- Executes the trigger function for every tuple encountered

- **Statement-Level Trigger**

*(FOR EACH STATEMENT)*

- Executes the trigger function only once
- Why would we want to do this?
  - If we already prevent an operation, then just once is enough
  - There are other contextual data available, but we will not discuss this



will just do it once - until the ;

```
CREATE TRIGGER modify_top_mark  
INSTEAD OF UPDATE ON Top_Marks  
FOR EACH ROW EXECUTE FUNCTION update_top_mark();
```

# Triggers Granularity

Table "Scores\_Log"

<u>Name</u>	Date
...	...

- The Trigger Function

```
CREATE OR REPLACE FUNCTION show_warning() RETURNS TRIGGER AS $$  
BEGIN  
    RAISE NOTICE 'You are not supposed to delete from log...';  
    RETURN NULL;  
END;  
$$ LANGUAGE plpgsql;
```

Trigger Function

- The Trigger

```
CREATE TRIGGER warn_delete  
BEFORE DELETE ON Scores_Log  
FOR EACH STATEMENT EXECUTE FUNCTION show_warning();
```

Trigger

## Effect?

- **RAISE NOTICE**
  - Database will give you prompt whenever a deletion is attempted
- **RETURN NULL**
  - Unfortunately, still perform deletion

# Triggers Granularity

Table "Scores\_Log"

<u>Name</u>	Date
...	...

- The Trigger Function

```
CREATE OR REPLACE FUNCTION show_warning() RETURNS TRIGGER AS $$  
BEGIN  
    RAISE EXCEPTION 'You are not supposed to delete from log...';  
    RETURN NULL;  
END;  
$$ LANGUAGE plpgsql;
```

Trigger Function

- Statement-Level Trigger *(FOR EACH STATEMENT)*

- Ignores the values returned by the trigger functions
- RETURN NULL would not make the database omit the subsequent operation
- What to do?
  - RAISE EXCEPTION

# Triggers Granularity

- **Granularity and Timing**

- `INSTEAD OF` is only allowed on row-level granularity
- `AFTER` or `BEFORE` are allowed on both row-level as well as statement-level granularity

Granularity

Timing	Row-Level	Statement-Level
AFTER	Tables	Tables and View
BEFORE	Tables	Tables and View
INSTEAD OF	Views	-

# Triggers Condition

Event

*Trigger*

Condition

*... we are here ...*

Action

*Trigger Function*

# Triggers Condition

so the check is not in the function - but in the trigger



Table "Scores"

<u>Name</u>	Mark
Adi	100

- The Trigger Function

```
CREATE OR REPLACE FUNCTION give_adi_full_mark() RETURNS TRIGGER AS $$
BEGIN
  IF (NEW.Name = 'Adi') THEN
    NEW.Mark := 100;
  END IF;
  RETURN NEW;
END;
$$ LANGUAGE plpgsql;
```

Trigger Function

### Observation

Trigger function only cares about the case when name is *Adi*

We can move this condition to the trigger definition

- The Trigger

```
CREATE TRIGGER adi_should_get_full_mark
BEFORE INSERT ON Scores
FOR EACH ROW WHEN (NEW.Name = 'Adi')
EXECUTE FUNCTION give_adi_full_mark();
```

Trigger

# Triggers Condition

- What Can We Use?

- In general, the condition in `WHEN()` could be more complicated
- Subject to the following requirements:
  - **NO** `SELECT` in `WHEN()`
  - **NO** `OLD` in `WHEN()` for `INSERT`
  - **NO** `NEW` in `WHEN()` for `DELETE`
  - **NO** `WHEN()` for `INSTEAD OF`

but can be used to chain - by using `AND`

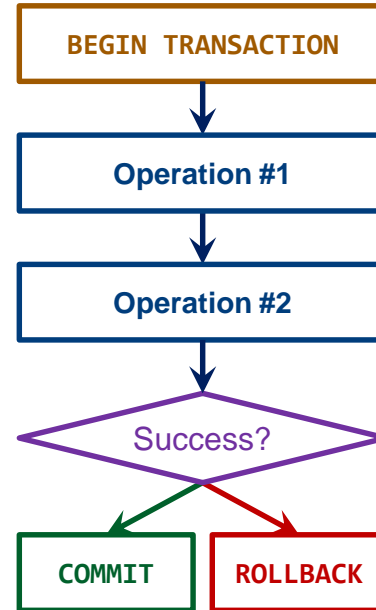
```
CREATE TRIGGER adi_should_get_full_mark
BEFORE INSERT ON Scores
FOR EACH ROW WHEN (NEW.Name = 'Adi')
EXECUTE FUNCTION give_adi_full_mark();
```



# Triggers Condition

- **Deferred Trigger**

- Triggers happen at the end of either statement or at the end of transaction
- **Operation** consisting of multiple statements may leave the database in an intermediate *inconsistent* state
- In such cases, we want the trigger to check consistency constraint only at the end of a transaction
- ❖ This is called a **deferred trigger**



# Triggers Condition


- **Deferred Trigger**

- Example:
  - Customer may have multiple account
  - Total balance for all accounts must be at least 150

Table "Account"

<u>AID</u>	Name	Balance
1	Alice	100
2	Alice	100

- Task: Transfer money from one account to another



# Triggers Condition

- **Deferred Trigger**

- Naïve approach:

1. Deduct amount from account 1
2. Add amount to account 2


Table "Account"

<u>AID</u>	Name	Balance
1	Alice	100
2	Alice	100

- Solution:

1. Put the two update into one transaction
2. Defer trigger to check only at the end of transaction

*May violate constraint!*



# Triggers Condition

Table "Account"

<u>AID</u>	Name	Balance
1	Alice	100
2	Alice	100

- Deferred Trigger

```
CREATE CONSTRAINT TRIGGER balance_check  
AFTER INSERT OR UPDATE OR DELETE ON Account  
DEFERRABLE INITIALLY DEFERRED  
FOR EACH ROW EXECUTE FUNCTION check_balance();
```

- **CONSTRAINT** and **DEFERRABLE** *together* indicate that the trigger can be deferred
- **INITIALLY DEFERRED** indicates that by default, the trigger is deferred
  - In other words, only check at the end of transaction
  - Other option is **INITIALLY IMMEDIATE**
    - *i.e.*, the trigger is not deferred by default
- Deferred triggers only work with **AFTER** and **FOR EACH ROW**



# Triggers Condition

Table "Account"

<u>AID</u>	Name	Balance
1	Alice	100
2	Alice	100

- Deferred Trigger

```
CREATE CONSTRAINT TRIGGER balance_check  
AFTER INSERT OR UPDATE OR DELETE ON Account  
DEFERRABLE INITIALLY DEFERRED  
FOR EACH ROW EXECUTE FUNCTION check_balance();
```

- Now we can do the following

```
BEGIN TRANSACTION;  
  
UPDATE Account SET Balance = Balance - Amount WHERE AID = Account1;  
UPDATE Account SET Balance = Balance + Amount WHERE AID = Account2;  
  
COMMIT;
```

# Triggers Condition

Table "Account"

<u>AID</u>	Name	Balance
1	Alice	100
2	Alice	100

- **Deferred Trigger**

```
CREATE CONSTRAINT TRIGGER balance_check  
AFTER INSERT OR UPDATE OR DELETE ON Account  
DEFERRABLE INITIALLY IMMEDIATE  
FOR EACH ROW EXECUTE FUNCTION check_balance();
```

- **What if the trigger is INITIALLY IMMEDIATE?**

- Change it on the fly

```
BEGIN TRANSACTION;  
SET CONSTRAINTS balance_check DEFERRED;  
UPDATE Account SET Balance = Balance - Amount WHERE AID = Account1;  
UPDATE Account SET Balance = Balance + Amount WHERE AID = Account2;  
  
COMMIT;
```

# Final Note on Triggers

- **Multiple Triggers**

- There can be multiple triggers defined for the same event on the same table
  - There need to be an *order of activation*
    - BEFORE statement-level triggers
    - BEFORE row-level triggers
    - AFTER row-level triggers
    - AFTER statement-level triggers
  - Within each category, triggers are activated in *alphabetic* order
  - If BEFORE row-level trigger returns NULL, then subsequent triggers on the same row are omitted

- **Universality of Triggers?**

- Our discussions are based on PostgreSQL syntax and implementation only

# Questions?