

Memory Management

# Disjoint Memory Schemes

Lecture 8

# Overview

## ■ Paging Scheme

- ❑ Basic Idea
- ❑ Logical Address Translation
- ❑ Hardware Support
- ❑ Protection & Page Sharing

## ■ Segmentation Scheme

- ❑ Motivation
- ❑ Basic Idea
- ❑ Logical Address Translation
- ❑ Hardware Support


## ■ Segmentation with Paging Scheme

# Disjoined Memory Space with Paging

- In lecture 7, memory management was discussed with two assumptions:
  1. Process occupies a contiguous physical memory region
  2. Physical memory is large enough to hold the whole process
- Let us remove assumption(1) in this lecture:
  - Process memory space can now be in **disjoint physical memory locations**
  - Via a main mechanism known as **paging**

# Paging Scheme: Basic Idea

- **Physical memory** is split into regions of fixed size (decided by hardware)
  - known as **physical frames**
- **Logical memory** of a process is split into regions of the **same size**
  - known as **logical pages**
- At execution time, pages of a process are loaded into **any available** memory frame
  - Logical memory space remain contiguous
  - Occupied physical memory region can be disjointed!



# Page Table: Lookup Mechanism

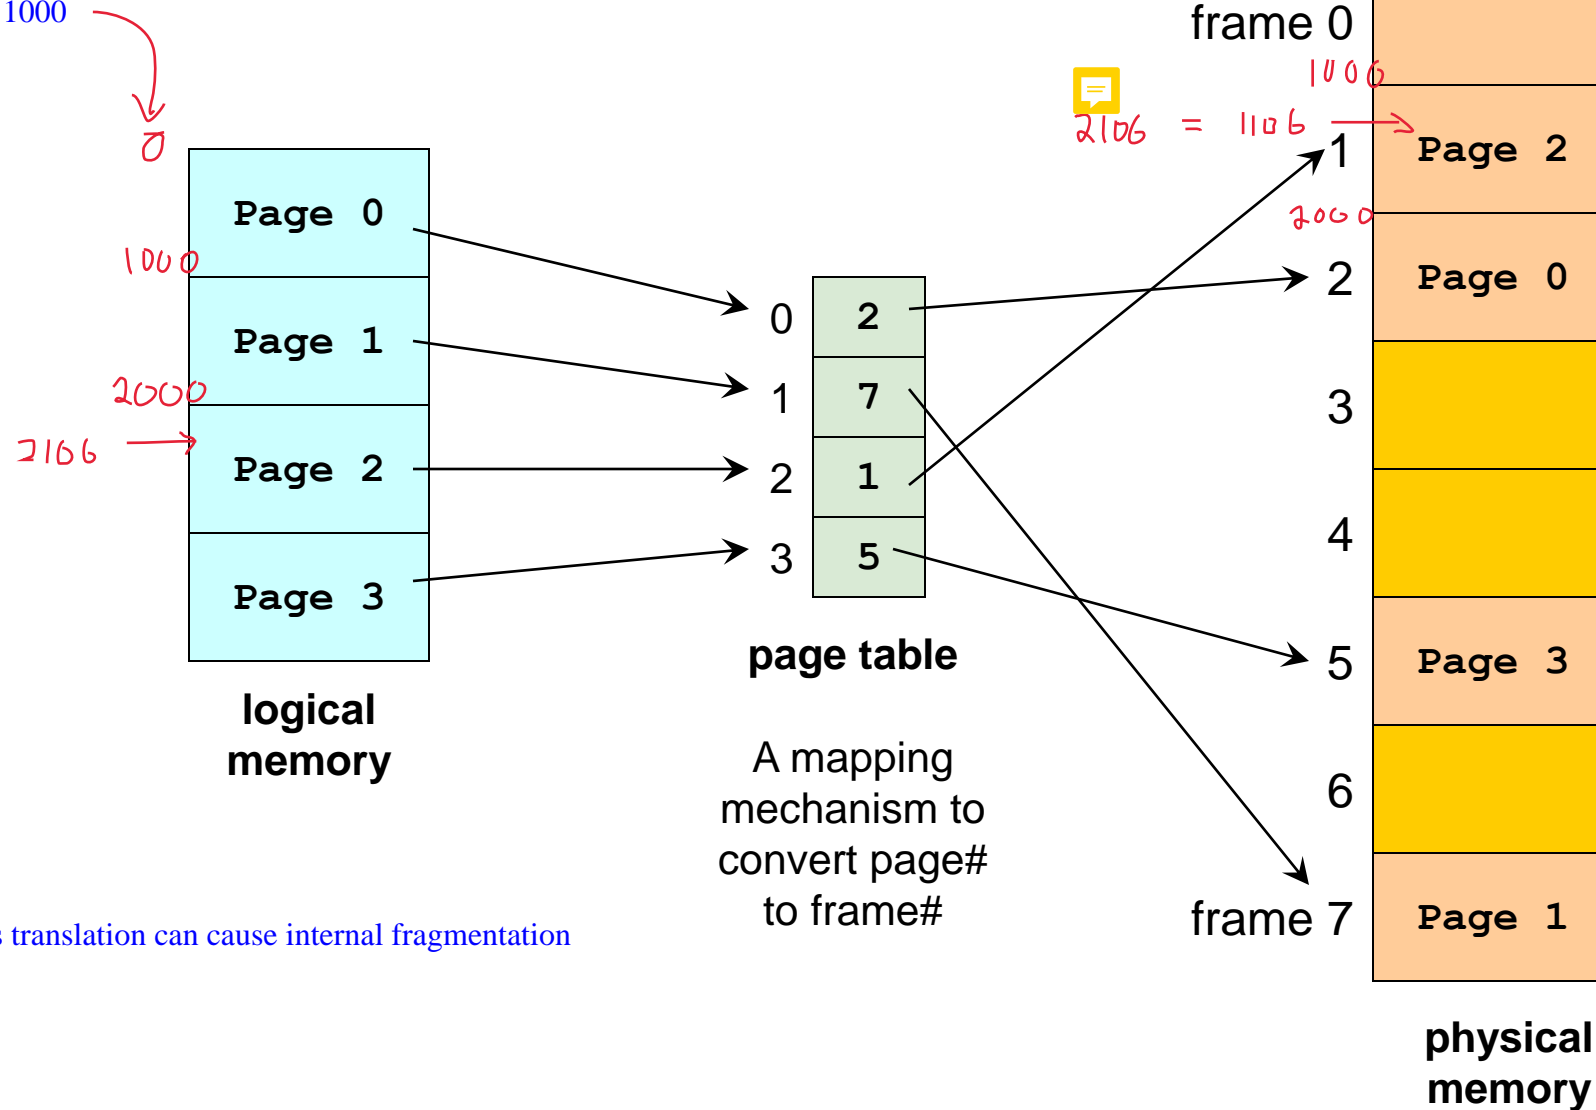
- In contiguous memory allocation, it is very simple to keep track of the usage of a process:
  - Minimally, **starting address** and **size of process**
- Under paging scheme:
  - Logical page  $\leftrightarrow$  physical frame mapping is no longer straightforward
  - Need a lookup table to provide the translation
  - This structure is known as a **Page Table**

# Paging: Illustration

Logical Address Translation

Page Size = 1000

Load 2106



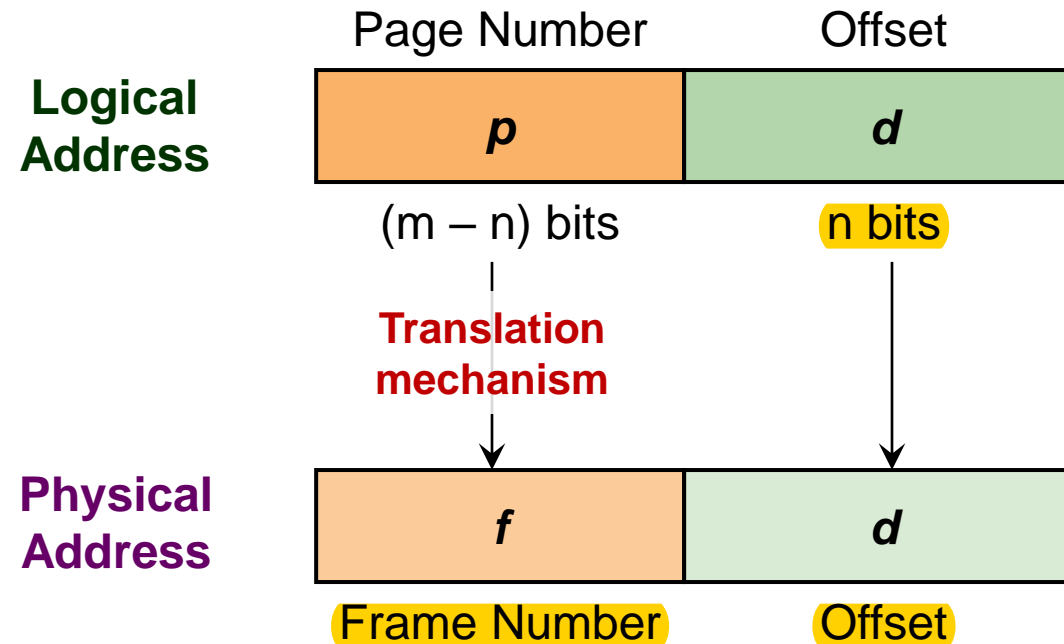
using logical address translation can cause internal fragmentation

# Logical Address Translation

- Program code uses logical memory address
  - However, to actually access the value, physical memory address is needed
- To locate a value in physical memory in the paging scheme, we need to know:
  - **F**, the physical frame number
  - **Offset**, displacement from the beginning of the physical frame
- Physical Address  
**= F x size of physical frame + Offset**

# Logical Address Translation: Essential Tricks

- Two important design decisions simplify address translation calculation
  1. Keep frame size (page size) as a power-of-2  *$2^n$  allows for  $F \ll n$*
  2. Physical frame size == Logical page size

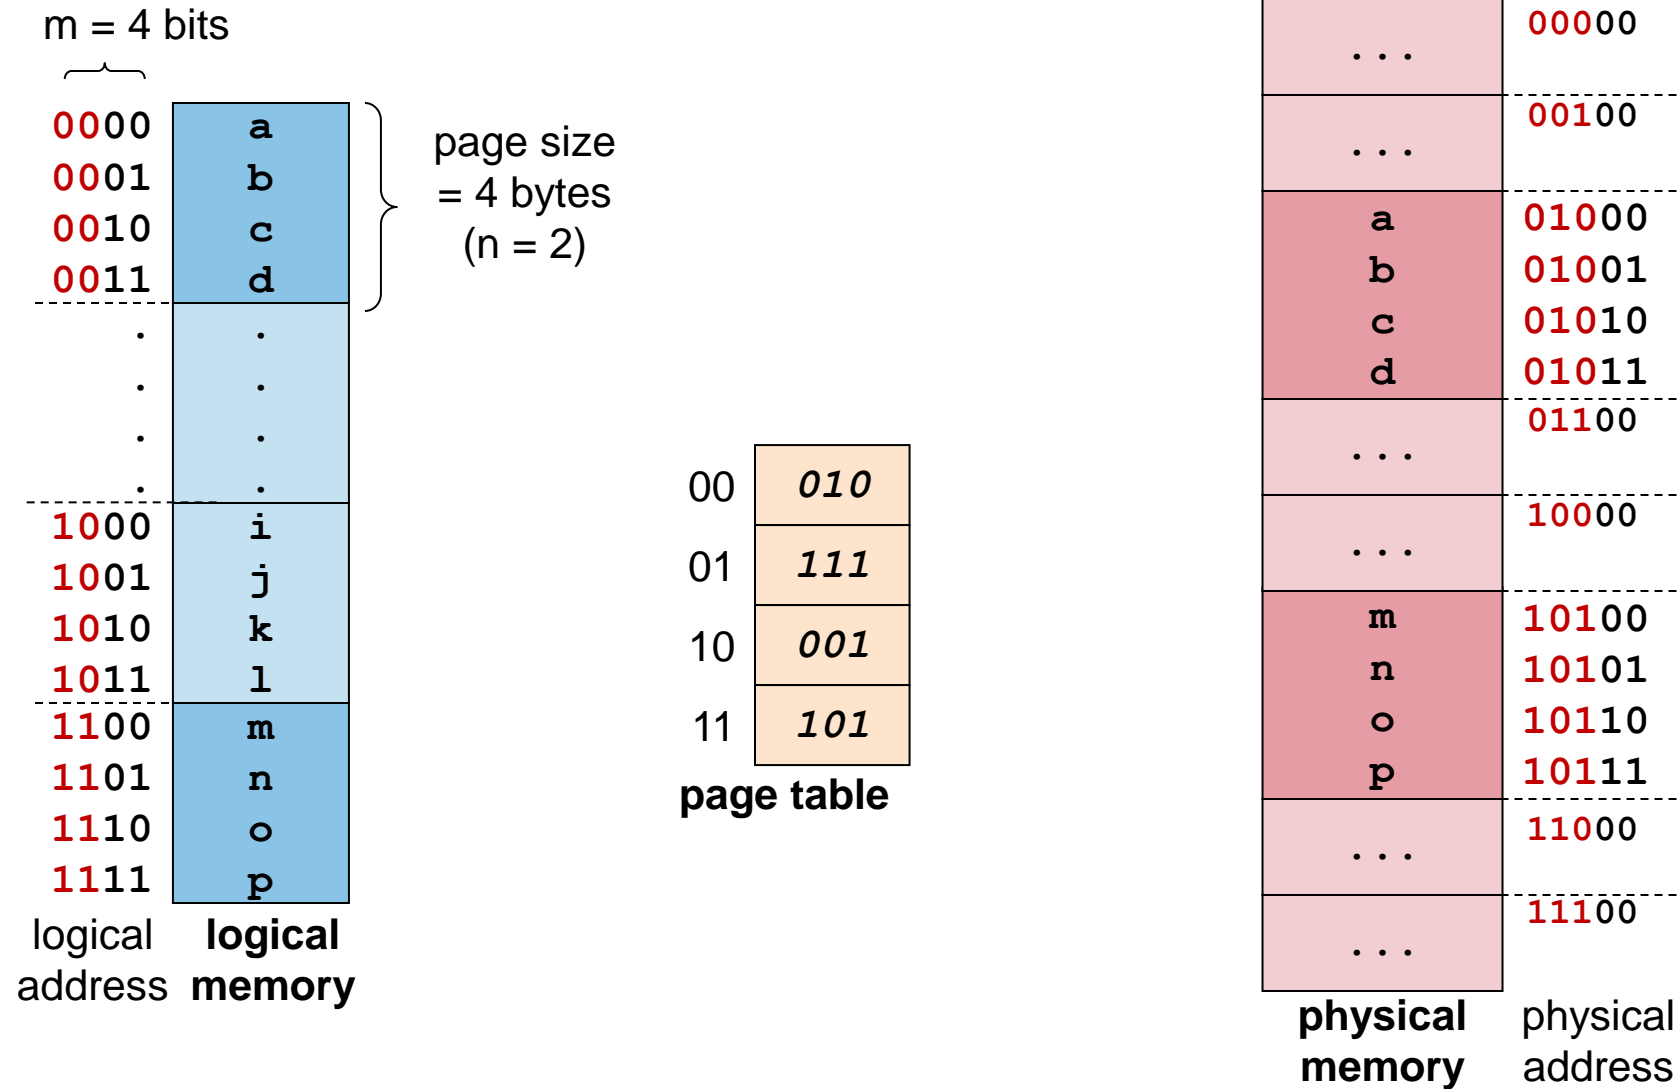




# Logical Address Translation: Formula

- Given:
  - Page/Frame size of  $2^n$
  - $m$  bits of logical address
- Logical Address  $LA$ :
  - $p$  = Most significant  $m-n$  bits of  $LA$
  - $d$  = Remaining  $n$  bits of  $LA$
- Use  $p$  to find frame number  $f$ 
  - from mapping mechanism like page table
- Physical Address  $PA$ :
  - $PA = f * 2^n + d$

# Example: 4 Logical Pages, 8 Physical Frames




# Paging: Observations

- Paging removes **external fragmentation**
  - ❑ No left-over physical memory region
  - ❑ All free frame can be used with no wastage
- Paging can still have **internal fragmentation**
  - ❑ Logical memory space **may not be multiple of page size**
- Clear separation of logical and physical address space
  - ❑ Allow great flexibility
  - ❑ Simple address translation

# Implementing Paging Scheme

- Common pure-software solution:

- OS stores page table alongside process information (e.g., PCB)
- **Improved understanding: Memory context of a process == Page table**  
 each process requires its own page table

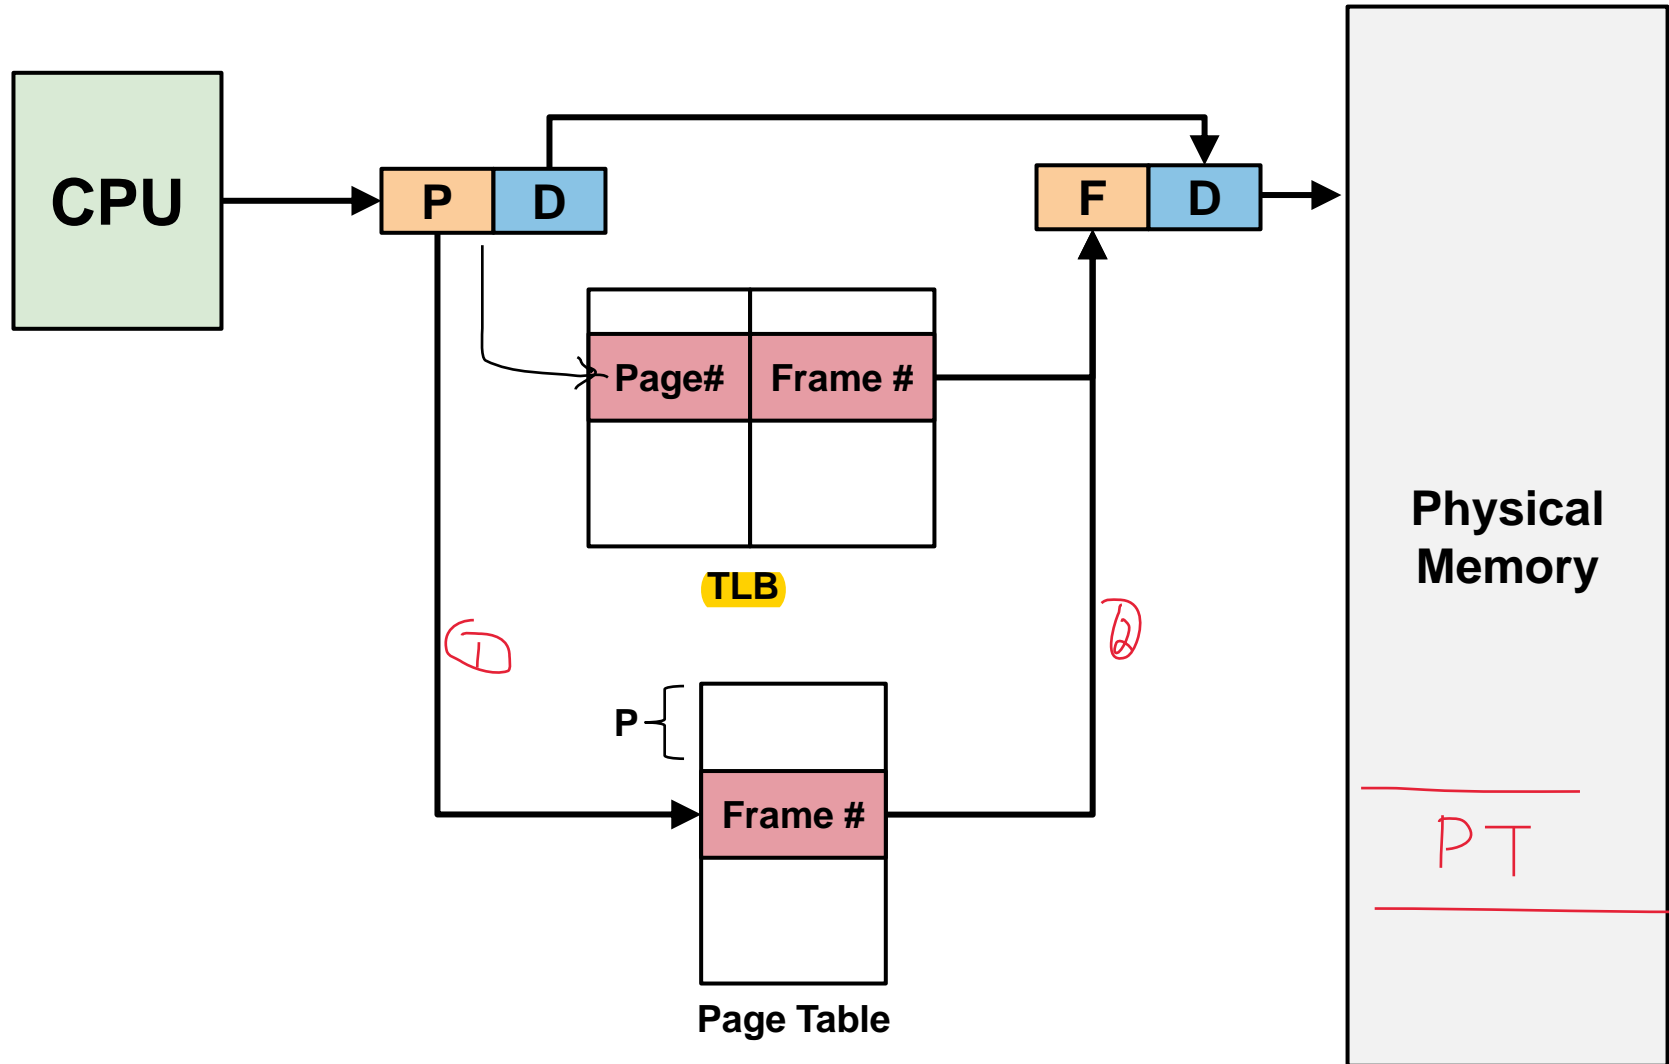
- Issues:

- **Require two memory accesses for every memory reference**
  - 1<sup>st</sup> access is to read the indexed page table entry to get frame number
  - 2<sup>nd</sup> access is to access the actual memory item

# Paging Scheme: **Hardware Support**

- Modern processors provide specialized on-chip component to support paging
  - Known as **Translation Look-Aside Buffer (TLB)**
  - TLB acts as a **cache** for *a few* page table entries
- Logical address translation with TLB:
  - Use page number to search TLB associatively
  - Entry found (**TLB-Hit**):
    - Frame number is retrieved to generate physical address
  - Entry not found (**TLB-Miss**):
    - Access the full page table
    - Retrieved frame number is used to generate physical address and update TLB

# Translation Look-Aside Buffer: Illustration



# TLB: Impact on Memory Access Time

- Suppose:

- TLB access takes 1ns
- Main memory access takes 50ns
- What is the average memory access time if TLB contains 40% of the whole page table?

- Memory access time

= TLB hit + TLB miss

= 40% x (1ns + 50ns) + 60%(1ns + 50ns + 50ns)

= 81ns

- Note:

- Overhead of filling in TLB entry and impact of cache ignored.

# TLB and Process Switching


- **Improved Understanding:** TLB is part of the **hardware context** of a process
- When a context switch occurs:
  - TLB entries are flushed
    - So that new process will not get incorrect translation
- Hence, when a process resumes running:
  - Will encounter many **TLB misses** to fill the TLB
    - It is possible to place some entries initially, e.g., the code pages, to reduce TLB misses



# Paging Scheme: Protection

- The basic paging scheme can be easily extended to provide memory protection between processes using:

- Access-Right Bits
- Valid Bit



- Access Right Bits:

- Each Page Table Entry includes several bits to indicate:
  - Whether the page is writable, readable, executable
  - E.g., page containing code should be executable, page containing data should be readable and writable, etc.
- Memory access is checked against the access right bits

# Paging Scheme: Protection (cont)

## ■ Observation:

- ❑ The logical memory range is usually the same for **all** processes
- ❑ However, not all processes utilize the whole range
  - ➔ Some pages are out-of-range for a particular process

## ■ Valid bit:

different processes will have different values

- ❑ Attached to each page table entry to indicate:
  - Whether the page is valid for the process to access
- ❑ OS will set the valid bits when a process is running
- ❑ **Memory access is checked against this bit:**
  - Out-of-range access will be caught by OS

# Paging Scheme: Page sharing

- Page table can allow several processes to share the same physical memory frame
  - Use the same physical frame number in the page table entries
- Possible usage:
  - Shared code page:
    - Some code are being used by many processes, e.g., C standard library, system calls, etc.
  - Implement Copy-On-Write:
    - As discussed earlier in "Process Abstraction" lecture, parent and child process can share a page until one tries to change a value in it

# Paging Scheme: **Page sharing**

Page 0
Page 1
Page 2
Page 3

**Process P's**  
logical  
memory

Page 0
Page 1
Page 2
Page 3

**Process Q's**  
logical  
memory

0	2
1	7
2	1
3	5

**Process P's**  
page table

0	
1	
2	
3	

**Process Q's**  
page table

frame 0	Page 1
1	Page 2
2	Page 0
3	Page 2
4	Page 0
5	Page 3
6	Page 3
frame 7	Page 1

**physical  
memory**

# Paging Scheme: **Copy-On-Write**

Page 0
Page 1
Page 2
Page 3

**Process P's**  
logical  
memory

Page 0
Page 1
Page 2
Page 3

**Process Q's**  
logical  
memory

0	2
1	7
2	1
3	5

**Process P's**  
page table

0	
1	
2	
3	

**Process Q's**  
page table

frame 0	
1	<b>Page 2</b>
2	Page 0
3	
4	
5	Page 3
6	
frame 7	Page 1

**physical  
memory**

Why memory error is often called ***segmentation fault***?

# SEGMENTATION SCHEME

# Segmentation Scheme: **Motivation**

- Memory space of a process is treated as a single entity so far
  - However, there are **several memory regions** with different usage in a process
- For example, in a typical C program:
  - User Code Region
  - Global Variables Region
    - Static Data (persistent as long as program executes)
  - Heap Region:
    - Dynamic data (persistent as long as user doesn't free)
  - Stack Region
  - Library Code Region
  - etc.

# Segmentation Scheme: Motivation (cont)

- Some regions may grow/shrink at execution time:
  - E.g., the stack region, the heap region, the library code region
- It is hard to place different regions in a contiguous memory space and still allow them grow/shrink freely
  - Also hard to check whether a memory access in a region is within its bound

thus the motivation for discontiguous memory will allow for the stack to grow dynamically during run time  
- and still maintain the protection for these regions




# Segmentation Scheme: Basic Idea

- Separate regions into multiple **memory segments**
  - Logical memory space of a process is now a collection of **segments**
- Each memory segment:
  - Has a name
    - For ease of reference
  - Has a limit
    - Indicate the range of the segment
- All memory reference is now specified as:
  - **Segment name + Offset**

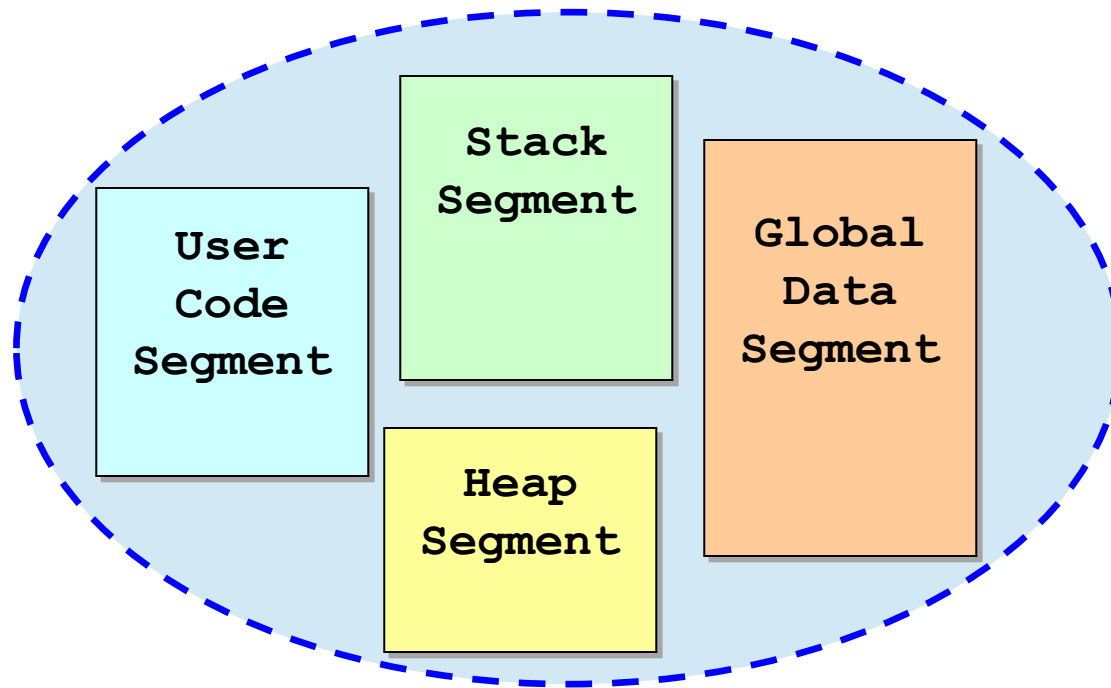
similar to the base + limit concept

Memory Access  
Example:

"*Heap*" + 245

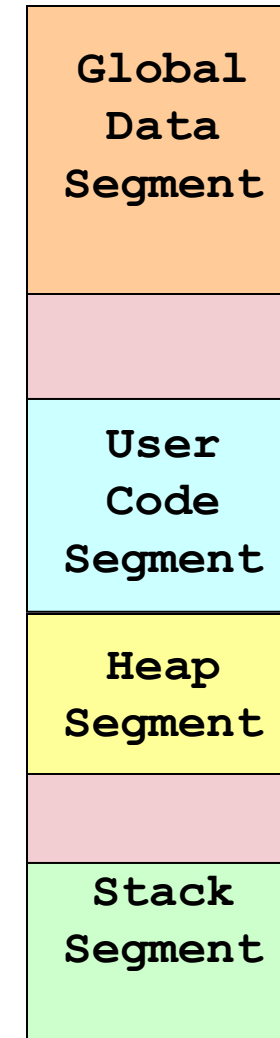


# Segmentation: Illustration



Logical Address Space

(Segments can be of different sizes)



Physical Memory

(Segments are disjoint in RAM)

# Segmentation: Logical Address Translation

- Each segment is mapped to a contiguous physical memory region
  - with a **base address** and a **limit**
- The segmentation name is usually represented as a single number
  - Known as **segment id**
- Logical address **<SegID, Offset>**:
  - **SegID** is used to look up **<Base, Limit>** of the segment in a **segment table**
  - Physical Address  $PA = Base + Offset$
  - **Offset** < **Limit** for valid access

# Logical Address Translation: Illustration

Assume:

User Code Segment = 0

Global Data Segment = 1

Heap Segment = 2

Stack Segment = 3

segmentation fault  
<1, 1600>

Since offset is larger than the limit

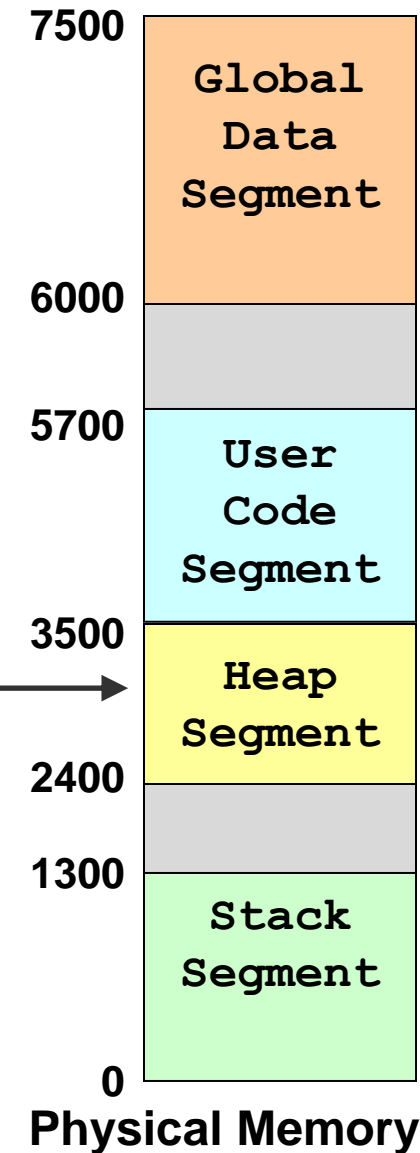
Memory Access  
< Segment Id, Offset >

	Base	Limit
0	3500	2200
1	6000	1500
2	2400	1100
3	0	1300

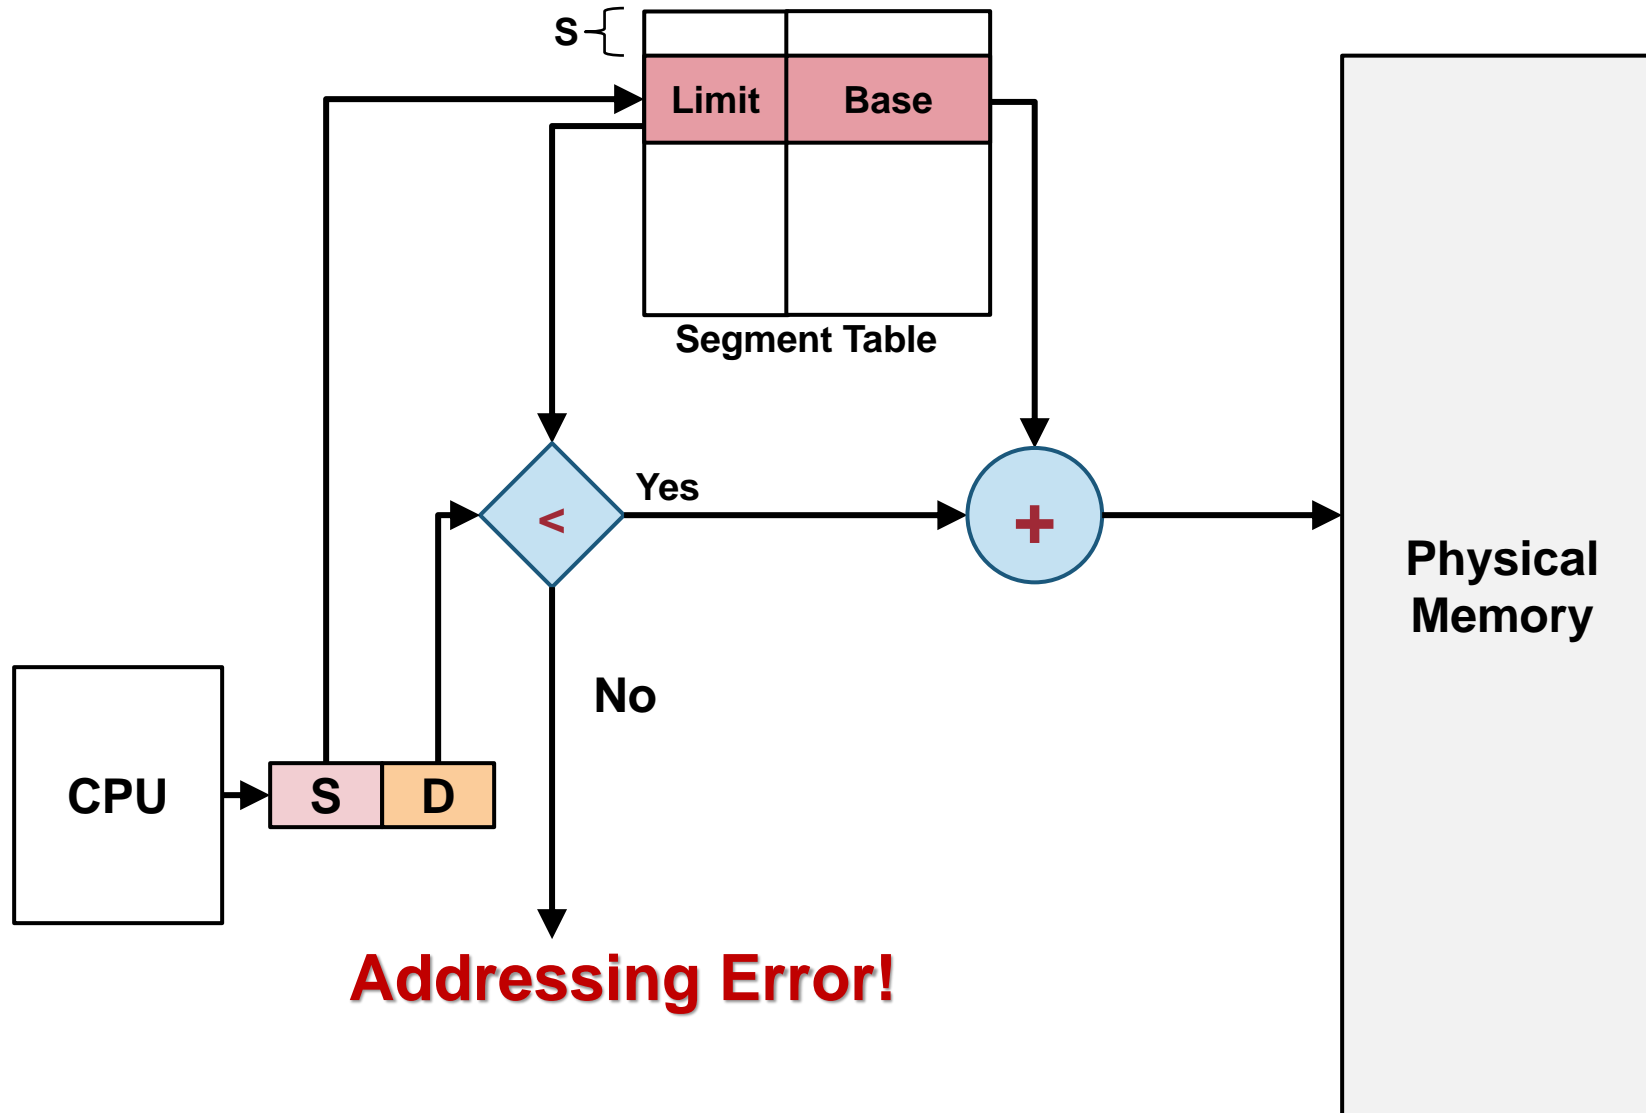
segment table

< 2, 500 >

2900



# Segmentation: Hardware Support



# Segmentation: Summary

## ■ Pros:

- ❑ Each segment is an independent contiguous memory space
  - ➔ Can grow / shrink independently
  - ➔ Can be protected / shared independently

## ■ Cons:

- ❑ Segmentation requires **variable-size contiguous memory regions**
  - ➔ can cause **external fragmentation**

## ■ Important observation:

- ❑ Segmentation is **not** the same as paging
- ❑ Each of them trying to solve a different problem

P is good; S is also good; Let's do S+P then!

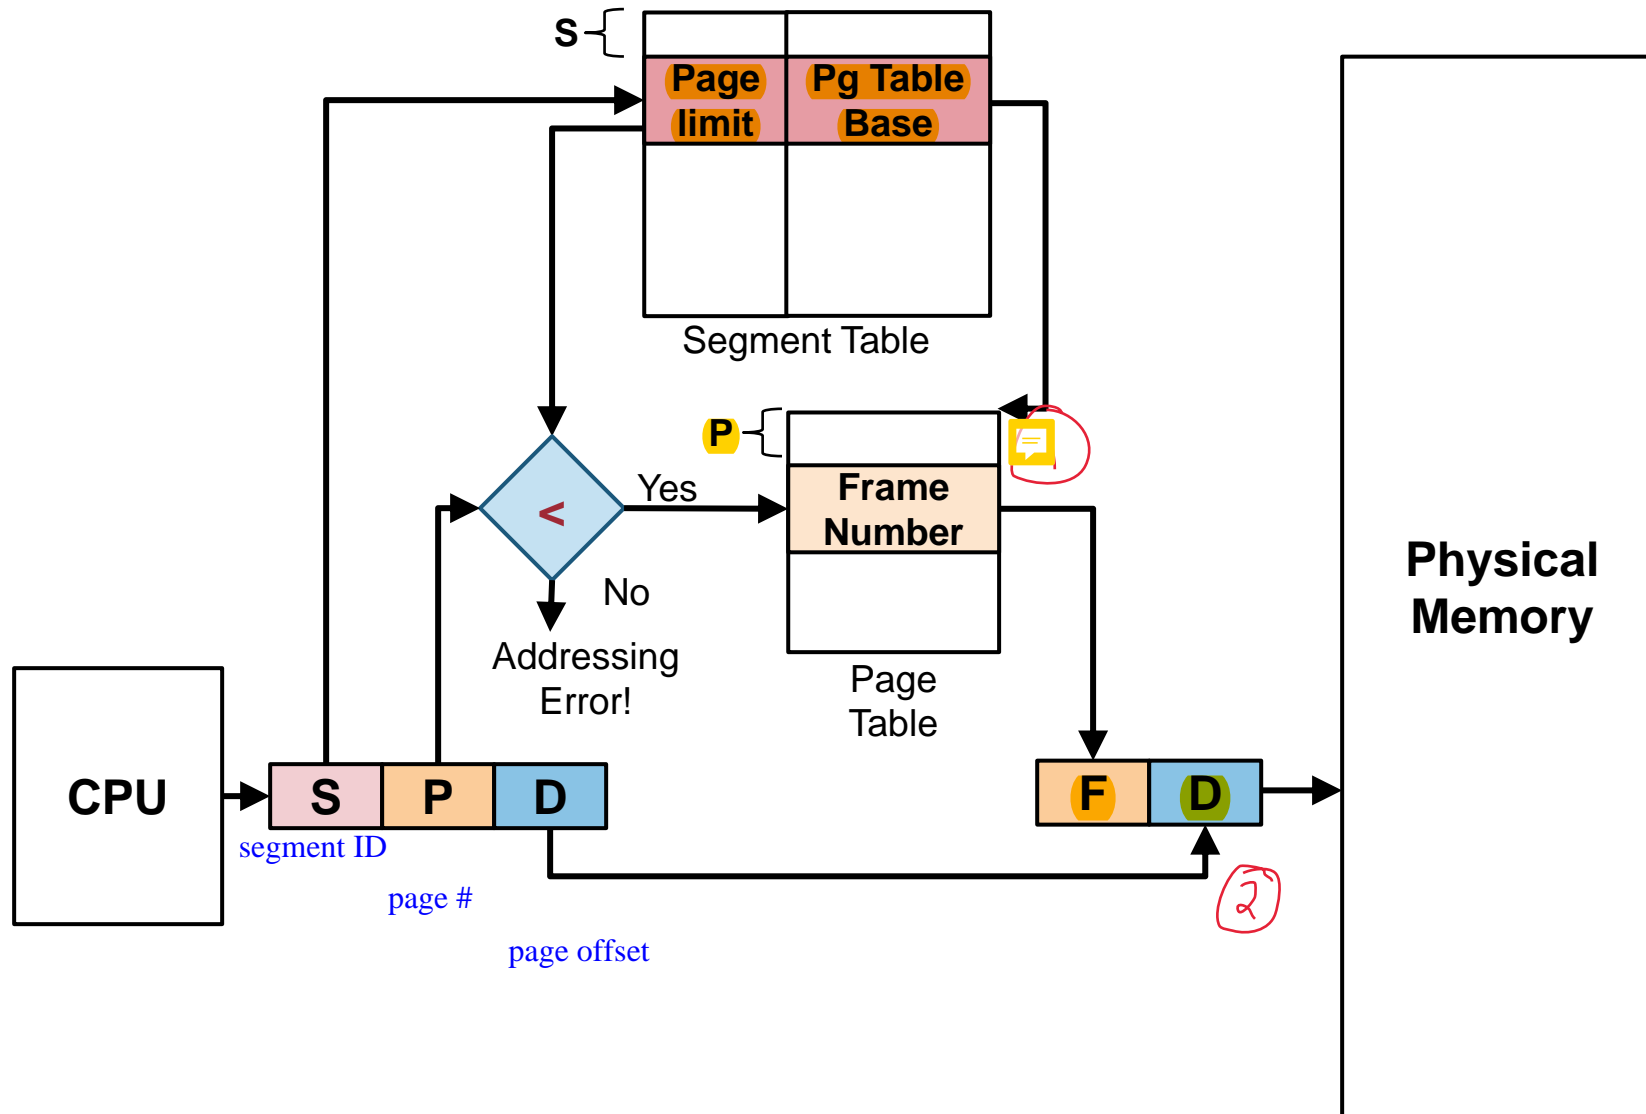
# SEGMENTATION WITH PAGING

# Segmentation with Paging: Basic Idea

- The intuitive next step is to combine segmentation and paging
- Basic Idea:
  - ❑ Each segment is now composed of several pages instead of a contiguous memory region
  - ➔ Essentially, each segment has a page table
  - ❑ Segment can grow by allocating new pages then add to its page table, and similarly for shrinking





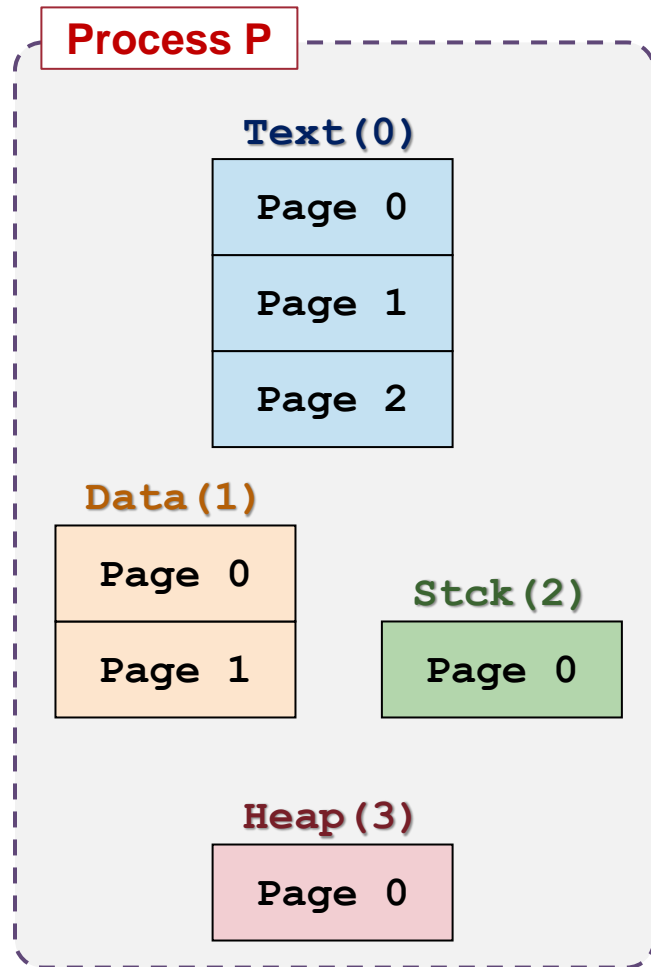
# Segmentation with Paging: Illustration



# Segmentation with Paging

in this example each page is 4 byte

<1, 6>



# Summary

- Discussed two popular memory management schemes
  - Both allow the logical address space to be in disjoint physical regions
- Paging split the logical address into fixed size pages
  - Store in fixed size physical memory frames
- Segmentation split the logical address into variable size segments according to their usage
  - Stored in variable-sized physical partitions
- Combining segmentation and paging