

The background of the image consists of numerous large, 3D-rendered blue numbers of various sizes and orientations, creating a dense, abstract pattern.

CS3223 Tutorial 6

Gary Lim

Admin

- ❖ How was the midterms? :p

Admin

- ❖ How was the midterms? :p
- ❖ **Project**
 - ❖ Submit report and code by **18 Mar 2022**, Week 9 Friday
 - ❖ Project demo in Week 10 (book a slot on the Google docs)
 - ❖ Support partition-based join, aggregates (SUM, COUNT, AVG, MIN, MAX), GROUP BY, DISTINCT
 - ❖ Experiments and discussions

Chapter Review

- ❖ Query Optimization
 - ❖ Plan Space
 - ❖ Search algorithm: Systems R

Query Optimization

- ❖ Find the **best** query evaluation plan (QEP) for a given query

Efficient, less I/O

these query evaluation plans will return the same answers

Query Optimization

- ❖ Find the **best** query evaluation plan (QEP) for a given query
- ❖ Consists of the following:
 - ❖ **Plan space** (set of all semantically equivalent plans)
 - ❖ **Enumeration / Search algorithm** (efficient search, able to find a ‘decent’ plan)
 - ❖ **Cost model** (basis for comparison between plans)

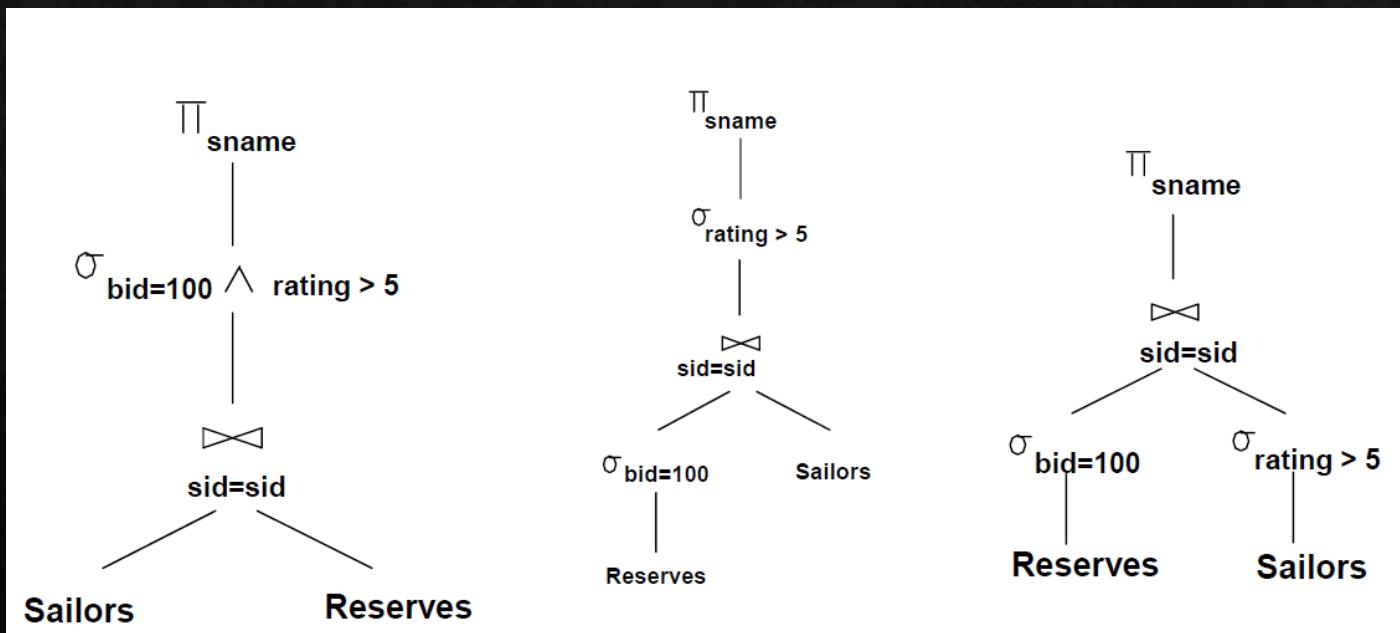
Query Optimization

- ❖ Common Heuristics?

Plan Space

- ◇ Common Heuristics
 - ◇ Push down **selections** and **projections**
 - ◇ Smallest relation first
 - ◇ Avoid cross products
 - ◇ Focus on a certain type of plans (e.g. left deep trees)

Selection pushed down

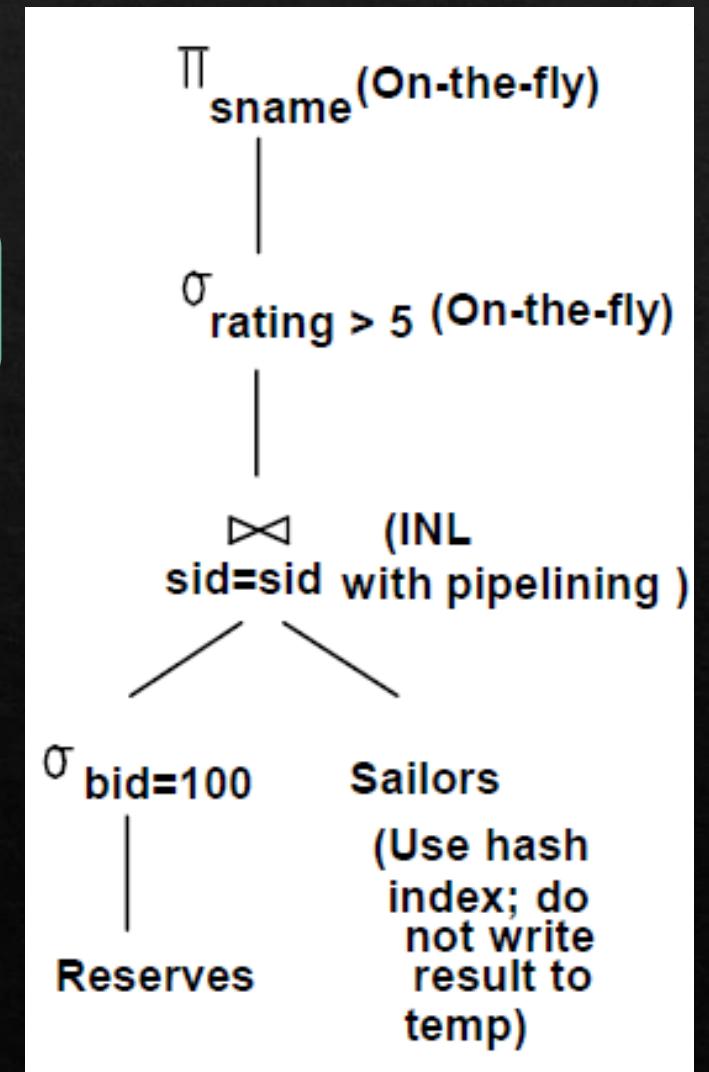


Plan Space

- ❖ Common Heuristics
 - ❖ Push down **selections** and **projections**
 - ❖ Smallest relation first
 - ❖ Avoid cross products
 - ❖ Focus on a certain type of plans (e.g. left deep trees)

Not always ideal to push down selection

Why?



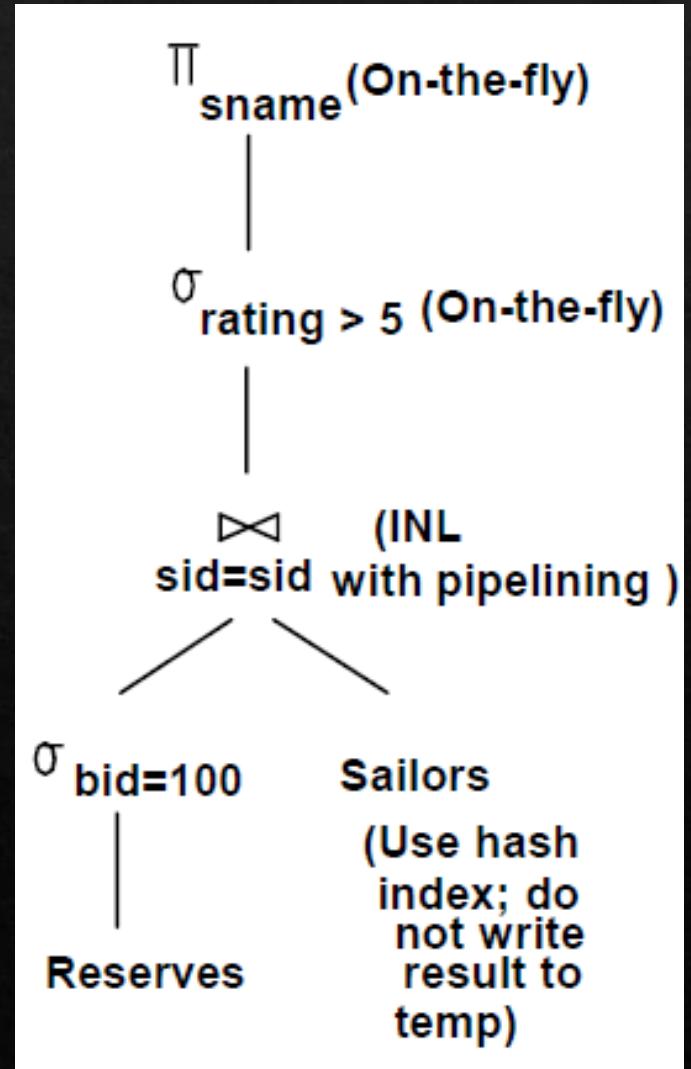
Plan Space

- ❖ Common Heuristics
 - ❖ Push down **selections** and **projections**
 - ❖ Smallest relation first
 - ❖ Avoid cross products
 - ❖ Focus on a certain type of plans (e.g. left deep trees)

Not always ideal to push down selection

Why?

If we push down $\sigma_{rating > 5}$ we need to perform a full index scan on the hash index for sailors, and write out the intermediate result, and cannot use index nested loop join anymore



Plan Space

Left-deep tree

- ❖ Right child has to be a base table

Right-deep Tree

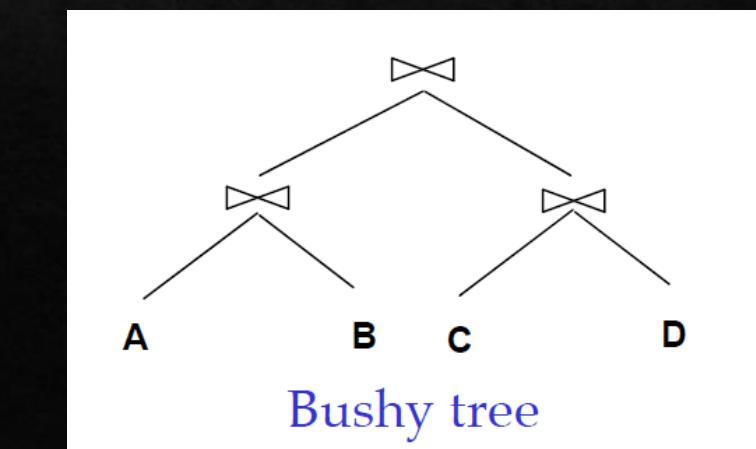
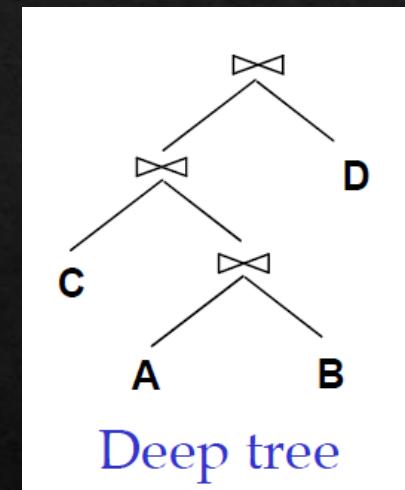
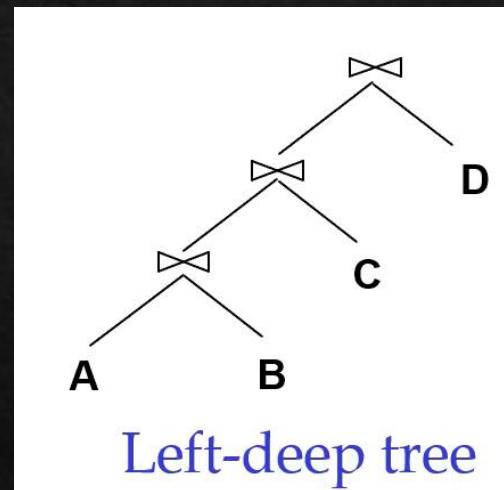
- ❖ Left child has to be a base table

Deep Tree

- ❖ One of the children is a base table

Bushy Tree

- ❖ No restriction



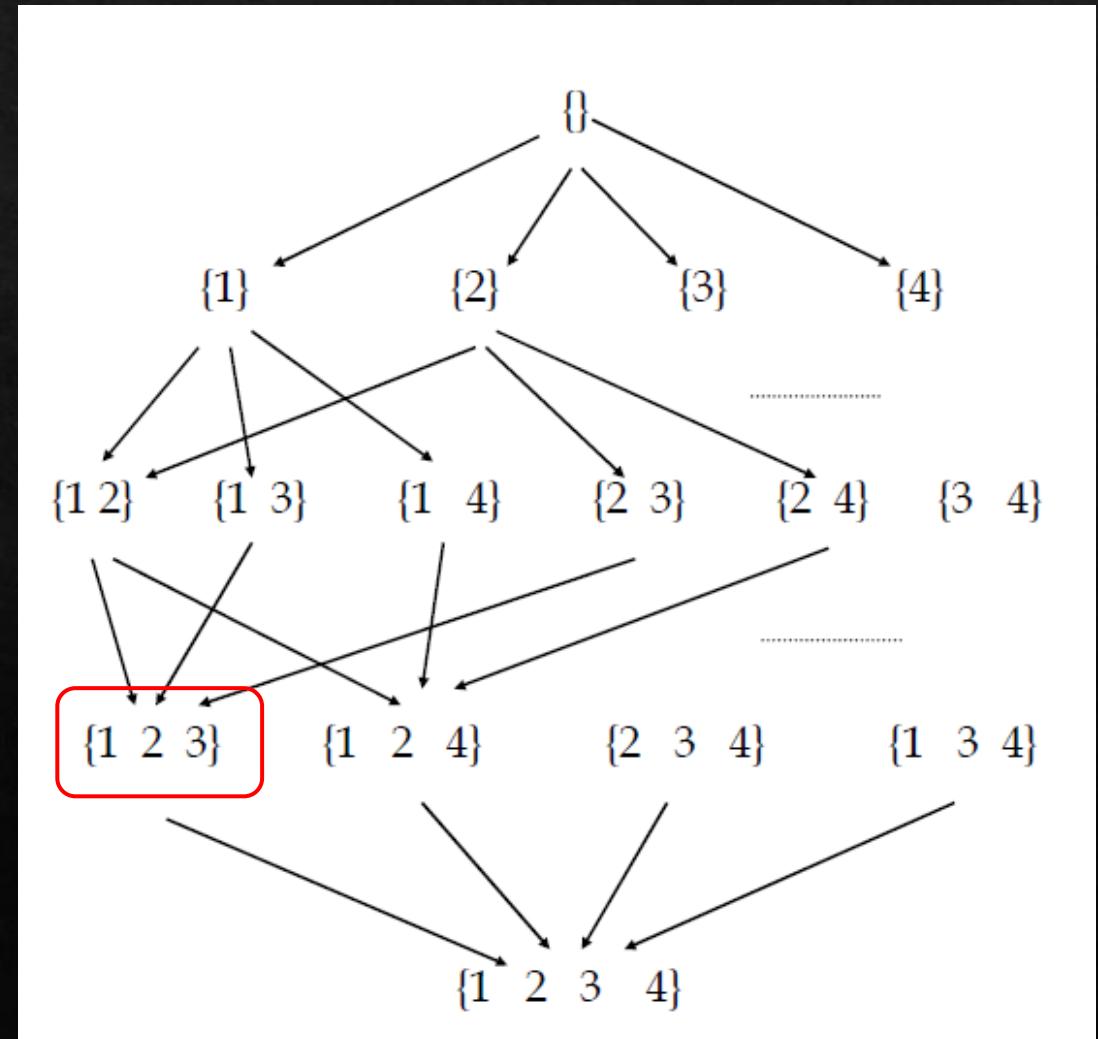
Search Algorithms

Search Algorithm	Search Space	Plan Quality	Optimization Overhead
Exhaustive	Complete	✓ ✓	✗ ✗
Greedy	Small	✓/✗	✓ Polynomial
Randomized	Depends	✓/✗	Depends on how long it is ran
Systems R	Almost complete	✓	✗ O(2^k)

Systems R

- ❖ Only left-deep trees
- ❖ **Principal of Optimality**
 - ❖ If two plans differ only in a subplan, then the plan with the better subplan is also the better plan (allows you to solve with DP due to overlapping subproblems)
- ❖ Number of subproblems: 2^k where k is number of relations

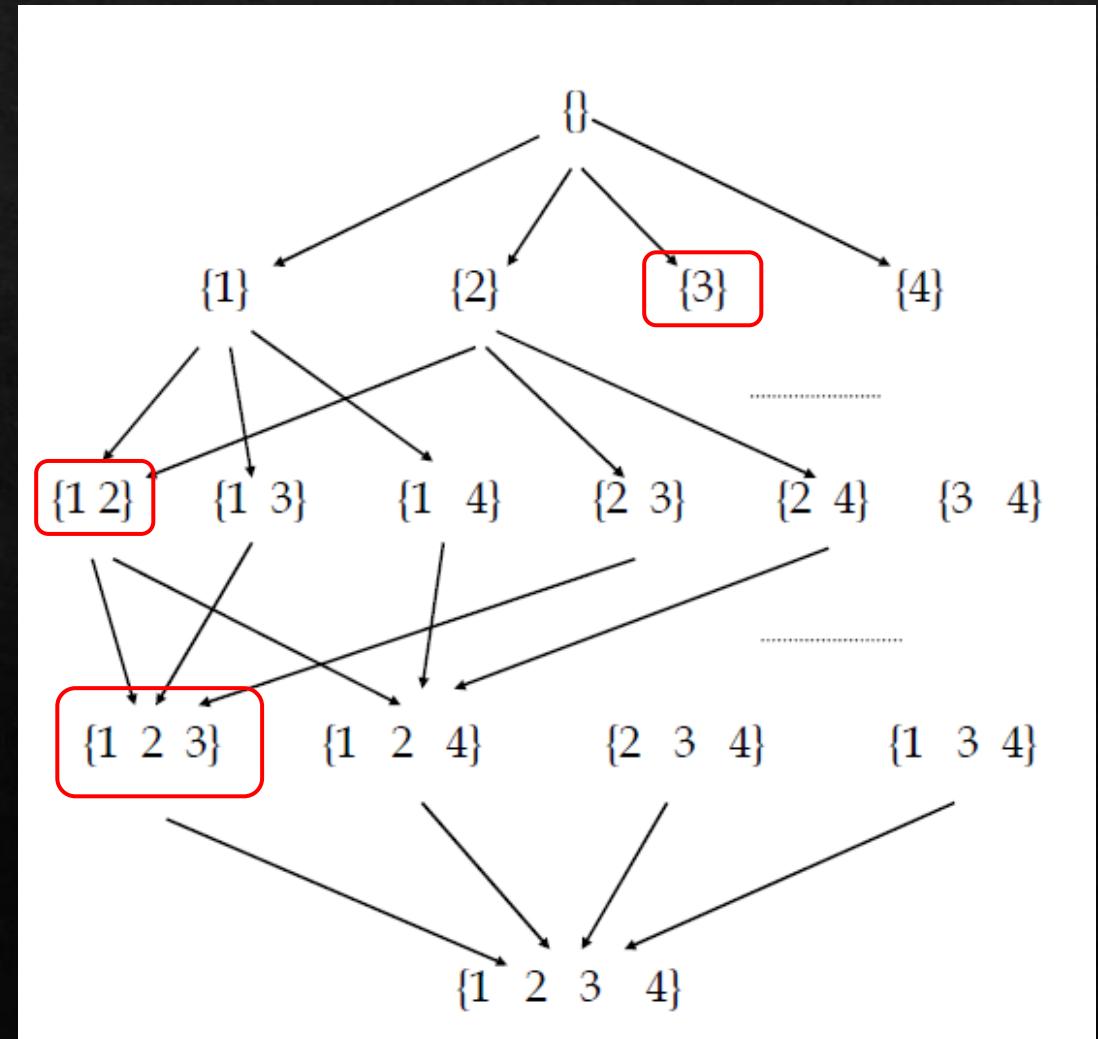
$$\sum_{k=0}^n \binom{n}{k} = 2^n$$



Systems R

- ❖ Only left-deep trees
- ❖ **Principal of Optimality**
 - ❖ If two plans differ only in a subplan, then the plan with the better subplan is also the better plan (allows you to solve with DP due to overlapping subproblems)
- ❖ Number of subproblems: 2^k where k is number of relations

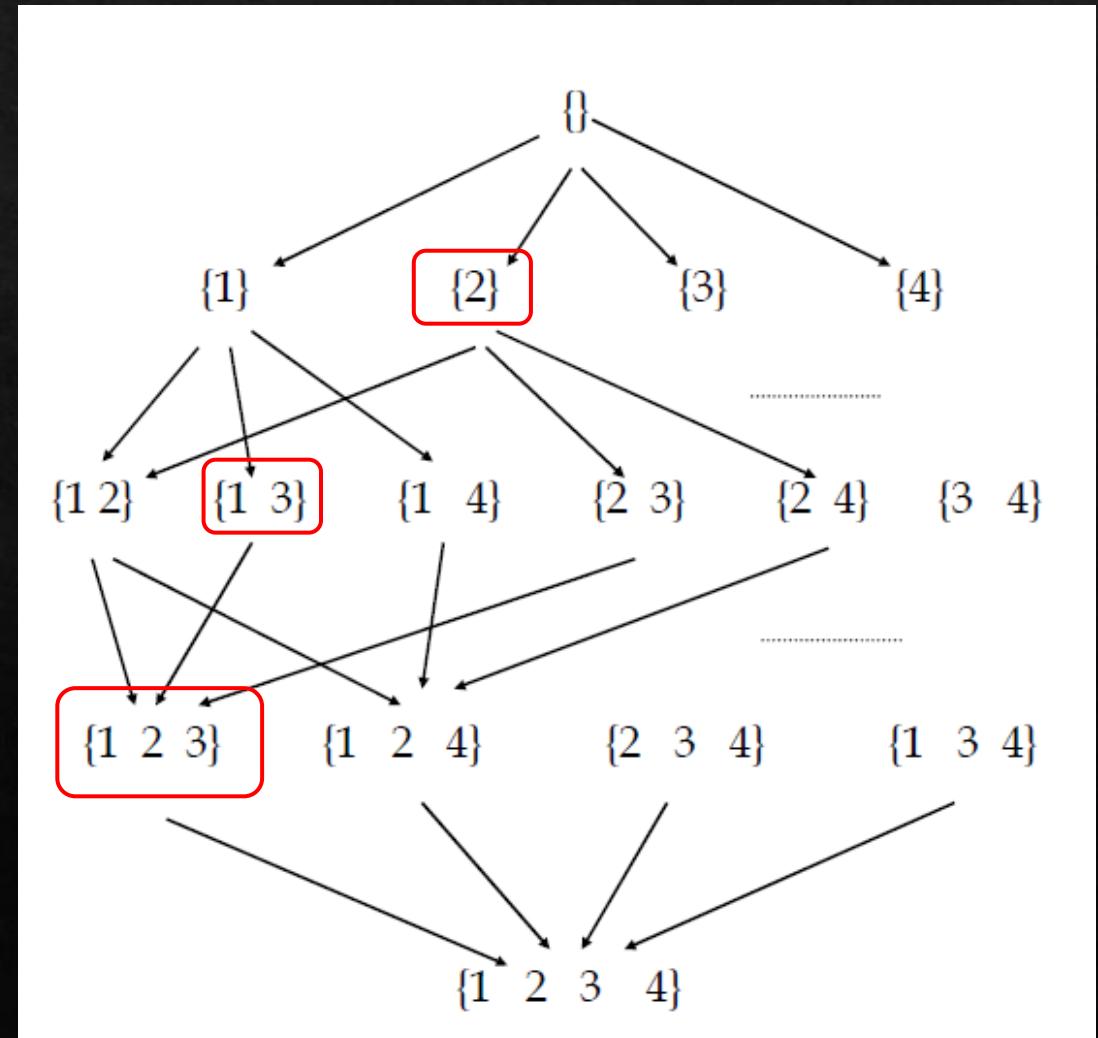
$$\sum_{k=0}^n \binom{n}{k} = 2^n$$



Systems R

- ❖ Only left-deep trees
- ❖ **Principal of Optimality**
 - ❖ If two plans differ only in a subplan, then the plan with the better subplan is also the better plan (allows you to solve with DP due to overlapping subproblems)
- ❖ Number of subproblems: 2^k where k is number of relations

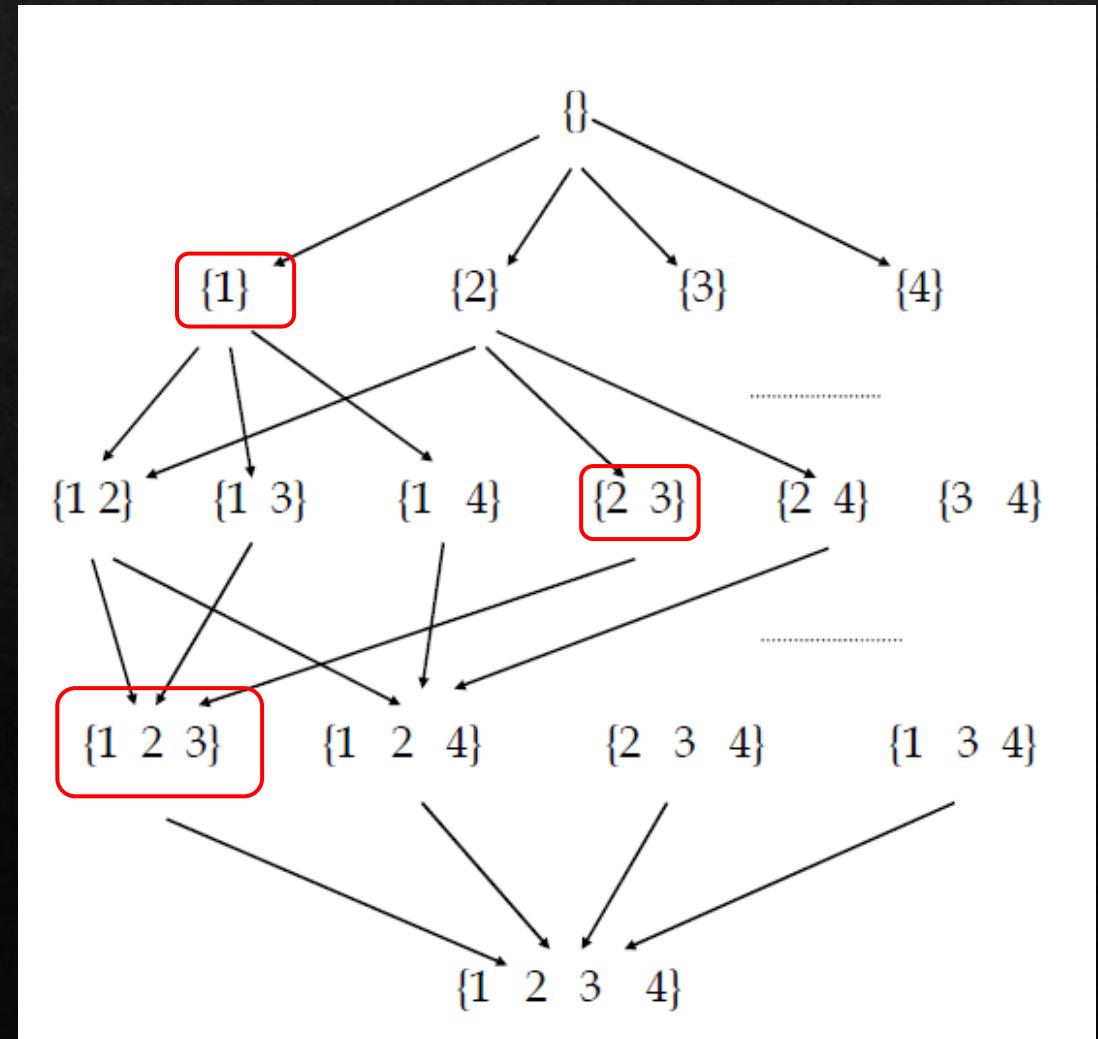
$$\sum_{k=0}^n \binom{n}{k} = 2^n$$



Systems R

- ❖ Only left-deep trees
- ❖ **Principal of Optimality**
 - ❖ If two plans differ only in a subplan, then the plan with the better subplan is also the better plan (allows you to solve with DP due to overlapping subproblems)
- ❖ Number of subproblems: 2^k where k is number of relations

$$\sum_{k=0}^n \binom{n}{k} = 2^n$$



Tutorial Questions

Q1

- ❖ R1(A, B, C) – 10,000 tuples
- ❖ R2(C, D, E) – 15000 tuples
- ❖ R3(E, F) – 7500 tuples
- ❖ 1 page \leftarrow 10 tuples of R1

Q1a

- ❖ R1(A, B, C) – 10,000 tuples
- ❖ R2(C, D, E) – 15000 tuples
- ❖ R3(E, F) – 7500 tuples
- ❖ 1 page \leftarrow 10 tuples of R1

Estimate the result size of $R1 \bowtie R2 \bowtie R3$

Q1a

- ❖ R1(**A**, B, **C**) – 10,000 tuples
- ❖ R2(**C**, D, **E**) – 15,000 tuples
- ❖ R3(**E**, F) – 7,500 tuples
- Since A, C and E are **primary keys**, there are
 - 10,000 distinct values of R1.A,
 - 15,000 distinct values of R2.C
 - 7,500 distinct values of R3.E
- Assuming **uniform distribution**,
 - all 10,000 values of **R1.C** are evenly distributed
 - all 15,000 values of **R2.E** are evenly distributed

Estimate the result size of $R1 \bowtie R2 \bowtie R3$

Q1a

- ❖ R1(A, B, **C**) – 10,000 tuples
- ❖ R2(**C**, D, **E**) – 15,000 tuples
- ❖ R3(**E**, F) – 7,500 tuples
- Since A, C and E are **primary keys**, there are
 - 10,000 distinct values of R1.A,
 - 15,000 distinct values of R2.C
 - 7,500 distinct values of R3.E
- Assuming **uniform distribution**,
 - **all 10,000 values of R1.C are evenly distributed**
 - all 15,000 values of R2.E are evenly distributed

Estimate the result size of $R1 \bowtie R2 \bowtie R3$

- ❖ R1 \bowtie R2
 - ❖ Every R1 tuple has a matching R2 tuple (**10,000 tuples**)

Q1a

- ❖ R1(**A**, B, **C**) – 10,000 tuples
 - ❖ R2(**C**, D, **E**) – 15,000 tuples
 - ❖ R3(**E**, F) – 7,500 tuples
- Since A, C and E are **primary keys**, there are
 - 10,000 distinct values of R1.A,
 - 15,000 distinct values of R2.C
 - 7,500 distinct values of R3.E
 - Assuming **uniform distribution**,
 - all 10,000 values of **R1.C** are evenly distributed
 - **all 15,000 values of R2.E are evenly distributed**

Estimate the result size of $R1 \bowtie R2 \bowtie R3$

- ❖ $R1 \bowtie R2$
 - ❖ Every R1 tuple has a matching R2 tuple (**10,000 tuples**)
- ❖ $R1 \bowtie R2 \bowtie R3$
 - ❖ Every R3 tuple has 2 matching R2 tuples; (still **10,000 tuples**)

Q1a

- ◊ R1(A, B, **C**) – 10,000 tuples ₁
 - ◊ R2(**C**, D, **E**) – 15,000 tuples ₁
₂
 - ◊ R3(**E**, F) – 7,500 tuples ₁
- Since A, C and E are **primary keys**, there are
 - 10,000 distinct values of R1.A,
 - 15,000 distinct values of R2.C
 - 7,500 distinct values of R3.E
 - Assuming **uniform distribution**,
 - all 10,000 values of **R1.C** are evenly distributed
 - **all 15,000 values of R2.E are evenly distributed**

Estimate the result size of $R1 \bowtie R2 \bowtie R3$

- ◊ $R1 \bowtie R2$
 - ◊ Every R1 tuple has a matching R2 tuple (**10,000 tuples**)
- ◊ $R1 \bowtie R2 \bowtie R3$
 - ◊ Every R3 tuple has 2 matching R2 tuples; (still **10,000 tuples**)

Q1b

- ❖ R1(A, B, **C**) – 10,000 tuples
- ❖ R2(**C**, D, **E**) – 15,000 tuples
- ❖ R3(**E**, F) – 7,500 tuples

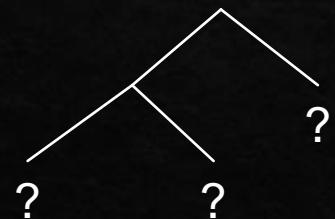
List all left-deep tree plans, excluding those with cross products.

Q1b

- ◊ R1(A, B, **C**) – 10,000 tuples
- ◊ R2(**C**, D, **E**) – 15,000 tuples
- ◊ R3(**E**, F) – 7,500 tuples

List all **left-deep tree** plans, excluding those with cross products.

Recall: For left-deep trees, right child must be base table



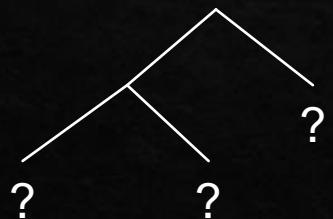
3! ways

Q1b

- ◊ R1(A, B, **C**) – 10,000 tuples
- ◊ R2(**C**, D, **E**) – 15,000 tuples
- ◊ R3(**E**, F) – 7,500 tuples

List all left-deep tree plans, **excluding those with cross products.**

Exclude $R1 \bowtie R3$



$$3! - 2 = 4$$

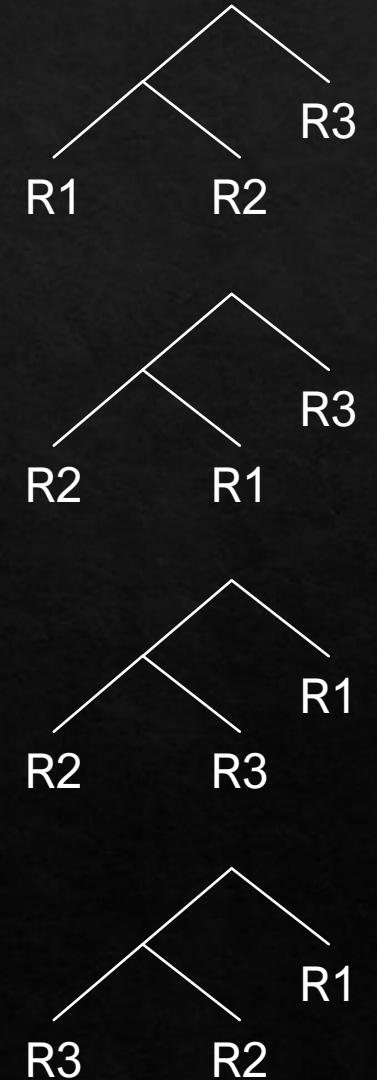
Q1b

- ◊ R1(A, B, **C**) – 10,000 tuples
- ◊ R2(**C**, D, **E**) – 15,000 tuples
- ◊ R3(**E**, F) – 7,500 tuples

List all left-deep tree plans, excluding those with cross products.

1. $(R1 \bowtie R2) \bowtie R3$
2. $(R2 \bowtie R1) \bowtie R3$
3. $(R2 \bowtie R3) \bowtie R1$
4. $(R3 \bowtie R2) \bowtie R1$

Depending on join algorithm,
 $R1 \bowtie R2$ can have a different
cost from $R2 \bowtie R1$ (e.g. block
nested join, hash-join)



Q1c

- ❖ R1(A, B, C) – 10,000 tuples
 - Using block-nested loop joins
 - 100 buffer pages
- ❖ R2(C, D, E) – 15,000 tuples
- ❖ R3(E, F) – 7,500 tuples
- ❖ R1 \bowtie R2(A, B, C, D, E) – 10,000 tuples
- ❖ R2 \bowtie R3(C, D, E, F) – 15,000 tuples

main focus is that must calculate how many tuples there can be in a page

- since we know the size of an attribute
- then will need to know how many attributes will exist in the relation

Q1c

- ◊ R1(A, B, C) – 10,000 tuples
- ◊ R2(C, D, E) – 15,000 tuples
- ◊ R3(E, F) – 7,500 tuples
- ◊ R1 \bowtie R2(A, B, C, D, E) – 10,000 tuples
- ◊ R2 \bowtie R3(C, D, E, F) – 15,000 tuples

- Using block-nested loop joins
- 100 buffer pages
- 10 tuples of R1 per page (30 attributes per page)
- 10 tuples of R2 per page
- 15 tuples of R3 per page
- 6 tuples of R1 \bowtie R2 per page
- 7 tuples of R2 \bowtie R3 per page

Attributes have
the same size

Cost of (R1 \bowtie R2) \bowtie R3

$$T: R1 \bowtie R2: \frac{10000}{10} + \left\lceil \frac{\frac{10000}{10}}{98} \right\rceil \times \frac{15000}{10} + \left\lceil \frac{10000}{6} \right\rceil = 19167$$

read write R1 \bowtie R2

} 29834 I/O

$$T \bowtie R3 : \left\lceil \frac{10000}{6} \right\rceil + \left\lceil \frac{\left\lceil \frac{10000}{6} \right\rceil}{98} \right\rceil \times \frac{7500}{15} = 10667$$

read R1 \bowtie R2 and R3

Q1c

- ◊ R1(A, B, C) – 10,000 tuples
- ◊ R2(C, D, E) – 15,000 tuples
- ◊ R3(E, F) – 7,500 tuples
- ◊ R1 \bowtie R2(A, B, C, D, E) – 10,000 tuples
- ◊ R2 \bowtie R3(C, D, E, F) – 15,000 tuples

- Using block-nested loop joins
- 100 buffer pages
- 10 tuples of R1 per page (30 attributes per page)
- 10 tuples of R2 per page
- 15 tuples of R3 per page
- 6 tuples of R1 \bowtie R2 per page
- 7 tuples of R2 \bowtie R3 per page

Attributes have
the same size

Cost of (R2 \bowtie R1) \bowtie R3

$$T: R2 \bowtie R1: \frac{15000}{10} + \left\lceil \frac{\frac{15000}{10}}{98} \right\rceil \times \frac{10000}{10} + \left\lceil \frac{10000}{6} \right\rceil = 19167$$

read write R1 \bowtie R2

} 29834 I/O

$$T \bowtie R3 : \left\lceil \frac{10000}{6} \right\rceil + \left\lceil \frac{\left\lceil \frac{10000}{6} \right\rceil}{98} \right\rceil \times \frac{7500}{15} = 10667$$

read R1 \bowtie R2 and R3

Q1c

- ◆ R1(A, B, C) – 10,000 tuples
- ◆ R2(C, D, E) – 15,000 tuples
- ◆ R3(E, F) – 7,500 tuples
- ◆ R1 \bowtie R2(A, B, C, D, E) – 10,000 tuples
- ◆ R2 \bowtie R3(C, D, E, F) – 15,000 tuples

- Using block-nested loop joins
- 100 buffer pages
- 10 tuples of R1 per page (30 attributes per page)
- 10 tuples of R2 per page
- 15 tuples of R3 per page
- 6 tuples of R1 \bowtie R2 per page
- 7 tuples of R2 \bowtie R3 per page

Attributes have
the same size

Cost of (R2 \bowtie R3) \bowtie R1

$$T: R2 \bowtie R3: \frac{15000}{10} + \left\lceil \frac{\frac{15000}{10}}{98} \right\rceil \times \frac{7500}{15} + \left\lceil \frac{15000}{7} \right\rceil = 11643$$

read write R2 \bowtie R3

35786 I/O

$$T \bowtie R1 : \left\lceil \frac{15000}{7} \right\rceil + \left\lceil \frac{\left\lceil \frac{15000}{7} \right\rceil}{98} \right\rceil \times \frac{10000}{10} = 24143$$

read R2 \bowtie R3 and R1

Q1c

- ◊ R1(A, B, C) – 10,000 tuples
- ◊ R2(C, D, E) – 15,000 tuples
- ◊ R3(E, F) – 7,500 tuples
- ◊ R1 \bowtie R2(A, B, C, D, E) – 10,000 tuples
- ◊ R2 \bowtie R3(C, D, E, F) – 15,000 tuples

- Using block-nested loop joins
- 100 buffer pages
- 10 tuples of R1 per page (30 attributes per page)
- 10 tuples of R2 per page
- 15 tuples of R3 per page
- 6 tuples of R1 \bowtie R2 per page
- 7 tuples of R2 \bowtie R3 per page

Attributes have
the same size

Cost of (R3 \bowtie R2) \bowtie R1

$$T: R3 \bowtie R2 : \frac{7500}{15} + \left\lceil \frac{\frac{7500}{15}}{98} \right\rceil \times \frac{15000}{10} + \left\lceil \frac{15000}{7} \right\rceil = 11643$$

read write R2 \bowtie R3

$$T \bowtie R1 : \left\lceil \frac{15000}{7} \right\rceil + \left\lceil \frac{\left\lceil \frac{15000}{7} \right\rceil}{98} \right\rceil \times \frac{10000}{10} = 24143$$

read R2 \bowtie R3 and R1

35786 I/O

Q1

- ◊ $R_3 \bowtie (R_1 \bowtie R_2)$ actually has a lower cost but this is not within the search space of left deep trees
- ◊ Left-deep trees search space does not guarantee optimal plan

Q2

- ❖ Discuss the motivation for the “avoid cross products” heuristics in multi-join query optimization.
- ❖ Does it always lead to better plans? Justify your answer.

Q2

- ❖ Discuss the motivation for the “avoid cross products” heuristics in multi-join query optimization.
 - ❖ Does it always lead to better plans? Justify your answer.
-
- ❖ Avoid large intermediate results (**more efficient query plan**)
 - ❖ Excluding cross-product plans also significantly **reduces the plan space**
 - + join is one of the most expensive operations in the dbms

Q2

- ❖ Discuss the motivation for the “avoid cross products” heuristics in multi-join query optimization.
- ❖ **Does it always lead to better plans? Justify your answer.**

Q2

- ❖ Not necessarily. Consider this counter example:
 - ❖ R1(A, B): tuple size = 10 bytes, 10 tuples
 - ❖ R2(D, E): tuple size = 10 bytes, 10 tuples
 - ❖ R3(B, E, F, G, H, ...): tuple size = 1000 bytes, 10000 tuples
- ❖ Query: $R1 \bowtie R2 \bowtie R3$
 - ❖ Non-cross product plans:
 - ❖ $(R1 \bowtie R3) \bowtie R2$
 - ❖ $(R2 \bowtie R3) \bowtie R1$
 - ❖ Cross product plans:
 - ❖ $(R1 \bowtie R2) \bowtie R3$

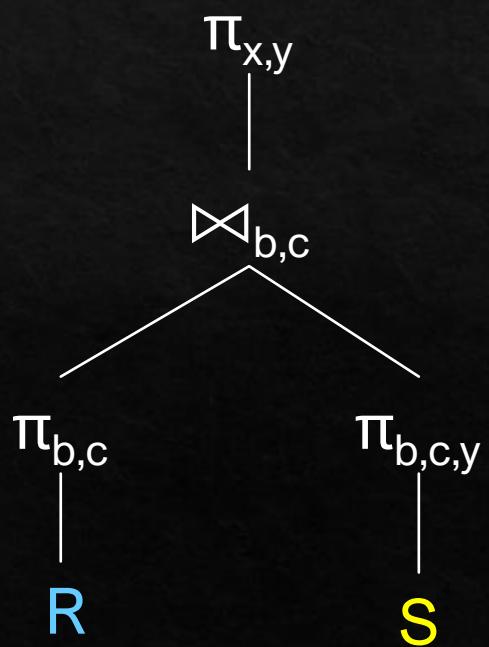
Large intermediate
result!

Q3

- ❖ $\pi_L(R(a, b, c) \bowtie S(b, c, d, e))$ where L is
 - ❖ $(x, y), x = b + c, y = c + d$
 - ❖ $(a, b, z), z = a + d$
- ❖ Push down projection as far as possible

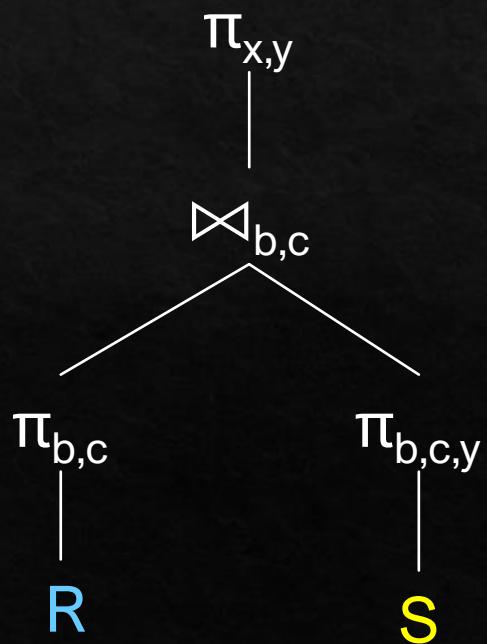
Q3a

- ◊ $\pi_L(R(a, b, c) \bowtie S(b, c, d, e))$ where L is
 - ◊ $(x, y), x = b + c, y = c + d$



Q3a

- ◇ $\pi_L(R(a, b, c) \bowtie S(b, c, d, e))$ where L is
 - ◇ (x, y) , $x = b + c$, $y = c + d$



Why is it wrong to
push x down?

Q3a

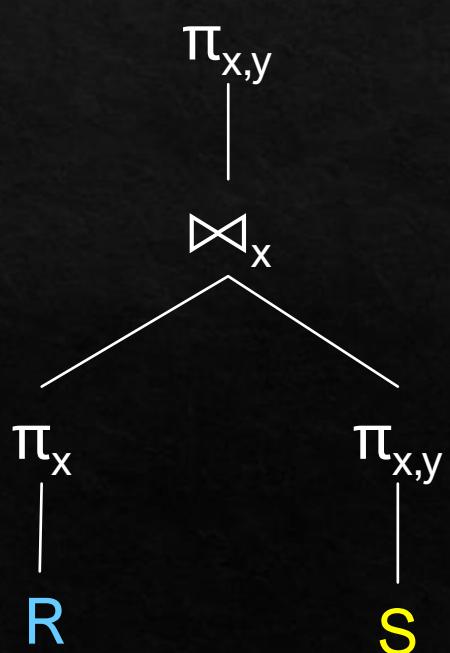
- ◇ $\pi_L(R(a, b, c) \bowtie S(b, c, d, e))$ where L is
- ◇ $(x, y), x = b + c, y = c + d$

R

b	c
1	3

S

b	c	d
1	3	7
2	2	7



Why is it wrong to push x down?

$\pi_x(R)$

x
4

$\pi_{x,y}(S)$

x	y
4	10
4	9

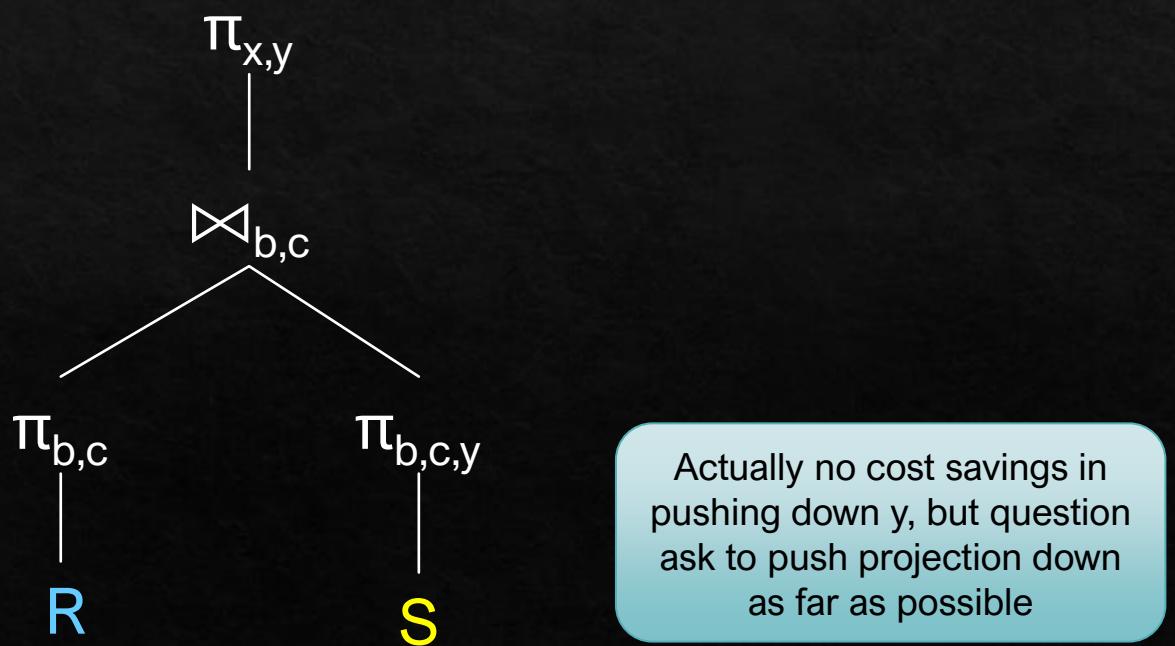
$\pi_{x,y}$

x	y
4	10
4	9

This is wrong! We need to keep b and c for the join

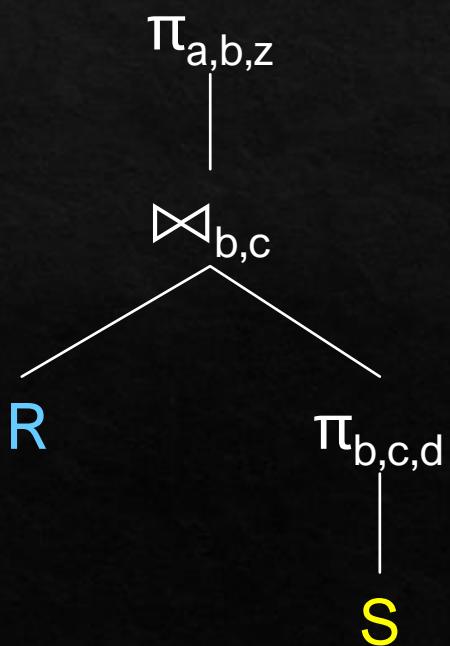
Q3a

- ◊ $\pi_L(R(a, b, c) \bowtie S(b, c, d, e))$ where L is
 - ◊ $(x, y), x = b + c, y = c + d$



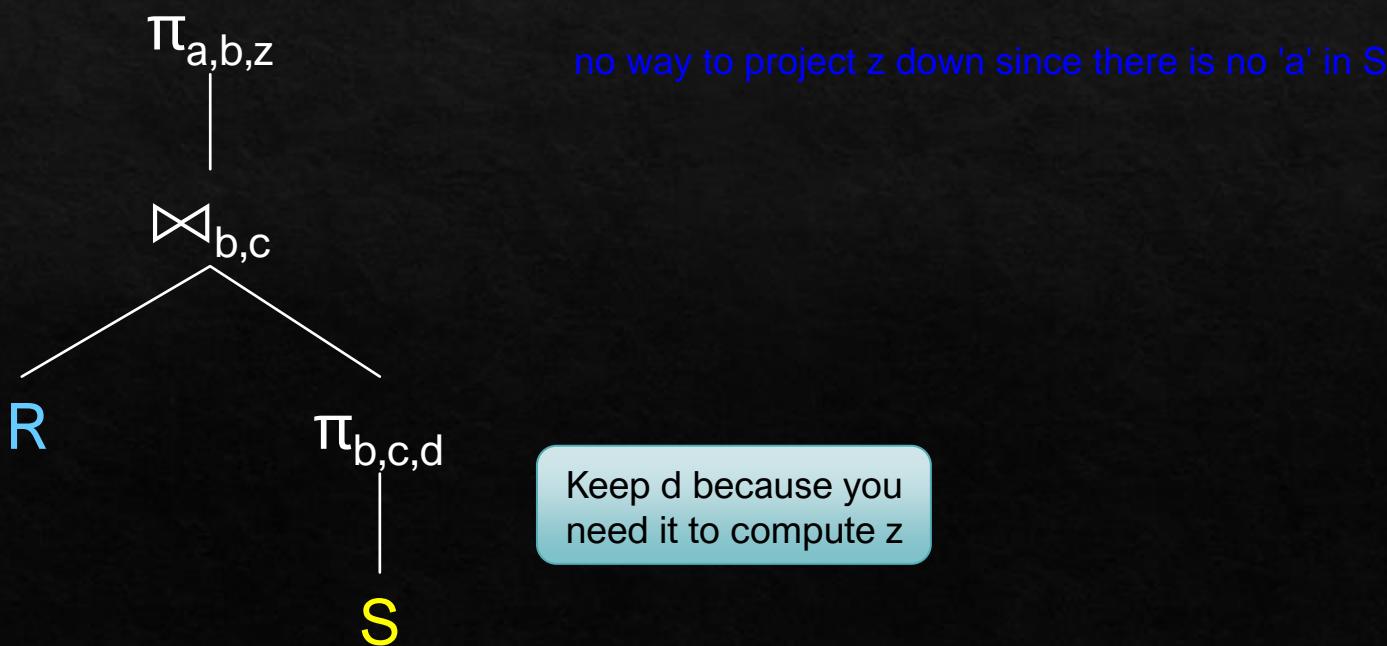
Q3b

- ◊ $\pi_L(R(a, b, c) \bowtie S(b, c, d, e))$ where L is
 - ◊ $(a, b, z), z = a + d$



Q3b

- ◊ $\pi_L(R(a, b, c) \bowtie S(b, c, d, e))$ where L is
 - ◊ $(a, b, z), z = a + d$

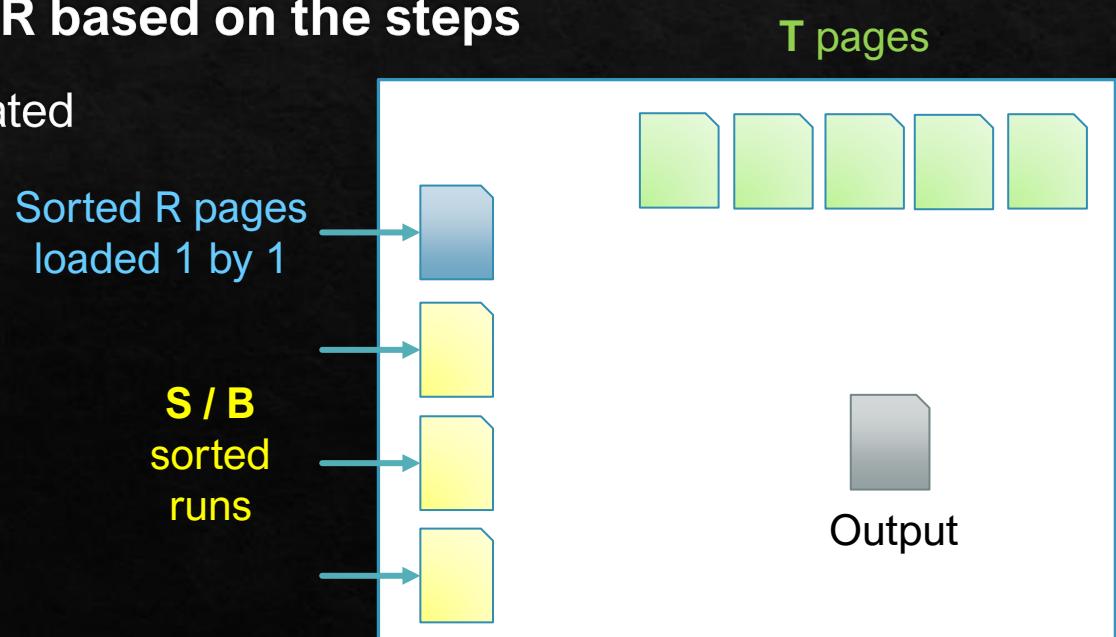


Q4

- ❖ Wish to join $R(a, b)$, $S(b, c)$ and $T(c, d)$ with R , S and T pages respectively
- ❖ B buffers (exclude output buffer)
- ❖ R is already sorted on b
- ❖ **Provide an inequality in terms of B , S , T and R based on the steps**

Q4

- ❖ Wish to join $R(a, b)$, $S(b, c)$ and $T(c, d)$ with R , S and T pages respectively
- ❖ B buffers (exclude output buffer)
- ❖ R is already sorted on b
- ❖ **Provide an inequality in terms of B , S , T and R based on the steps**
 - ❖ Generate sorted runs for S : S/B sorted runs created
 - ❖ Load T into memory entirely: T buffers needed
 - ❖ Merge R and S and T
- ❖ $1 + T + \left\lceil \frac{S}{B} \right\rceil \leq B$



Q4

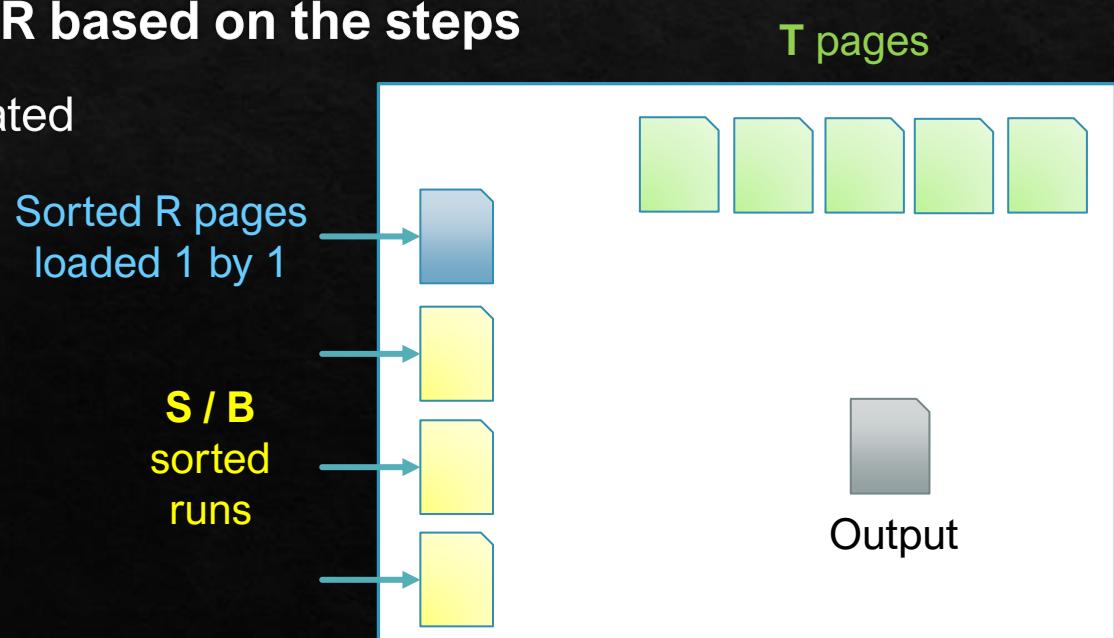
- ❖ Wish to join $R(a, b)$, $S(b, c)$ and $T(c, d)$ with R , S and T pages respectively
- ❖ B buffers (exclude output buffer)
- ❖ R is already sorted on b
- ❖ **Provide an inequality in terms of B , S , T and R based on the steps**

- ❖ Generate sorted runs for S : S/B sorted runs created
- ❖ Load T into memory entirely: T buffers needed
- ❖ Merge R and S and T

$$1 + T + \left\lceil \frac{S}{B} \right\rceil \leq B$$

$$B^2 \geq S + BT$$

Ignore the 1, multiply B
on both sides



Extra

- ❖ Let δ denote duplicate elimination. Which of the following statements is TRUE?
 - ❖ $\delta(R \times S) = \delta(R) \times \delta(S)$
 - ❖ $\delta(R \bowtie S) = \delta(R) \bowtie \delta(S)$
 - ❖ $\delta(\sigma_s(R)) = \sigma_s(\delta(R))$

Extra

- ❖ All true.
- ❖ To prove equality, try to show that LHS is a subset of RHS, and RHS is a subset of LHS

End