

The background of the image consists of numerous large, 3D-rendered blue numbers of various sizes and orientations, creating a dense, textured pattern.

CS3223 Tutorial 8

Gary Lim

Chapter Review

- ❖ ACID Transactions
- ❖ Serializability
- ❖ Recoverable, Cascadeless, Strict

ACID Transactions

- ❖ **Atomicity:**
- ❖ **Consistency:**
- ❖ **Isolation:**
- ❖ **Durability:**

ACID Transactions

- ❖ **Atomicity:** All actions in a Xact happen, or none of them happen
 - ❖ Guaranteed by using COMMIT after all actions has finished, or ABORT if some actions are not allowed to be executed
- ❖ **Consistency:**
- ❖ **Isolation:**
- ❖ **Durability:**

ACID Transactions

- ❖ **Atomicity:** All actions in a Xact happen, or none of them happen
 - ❖ Guaranteed by using COMMIT after all actions has finished, or ABORT if some actions are not allowed to be executed
- ❖ **Consistency:** If each Xact is consistent, and the DB was in a consistent state initially, then it should end up being consistent
- ❖ **Isolation:**
- ❖ **Durability:**

ACID Transactions

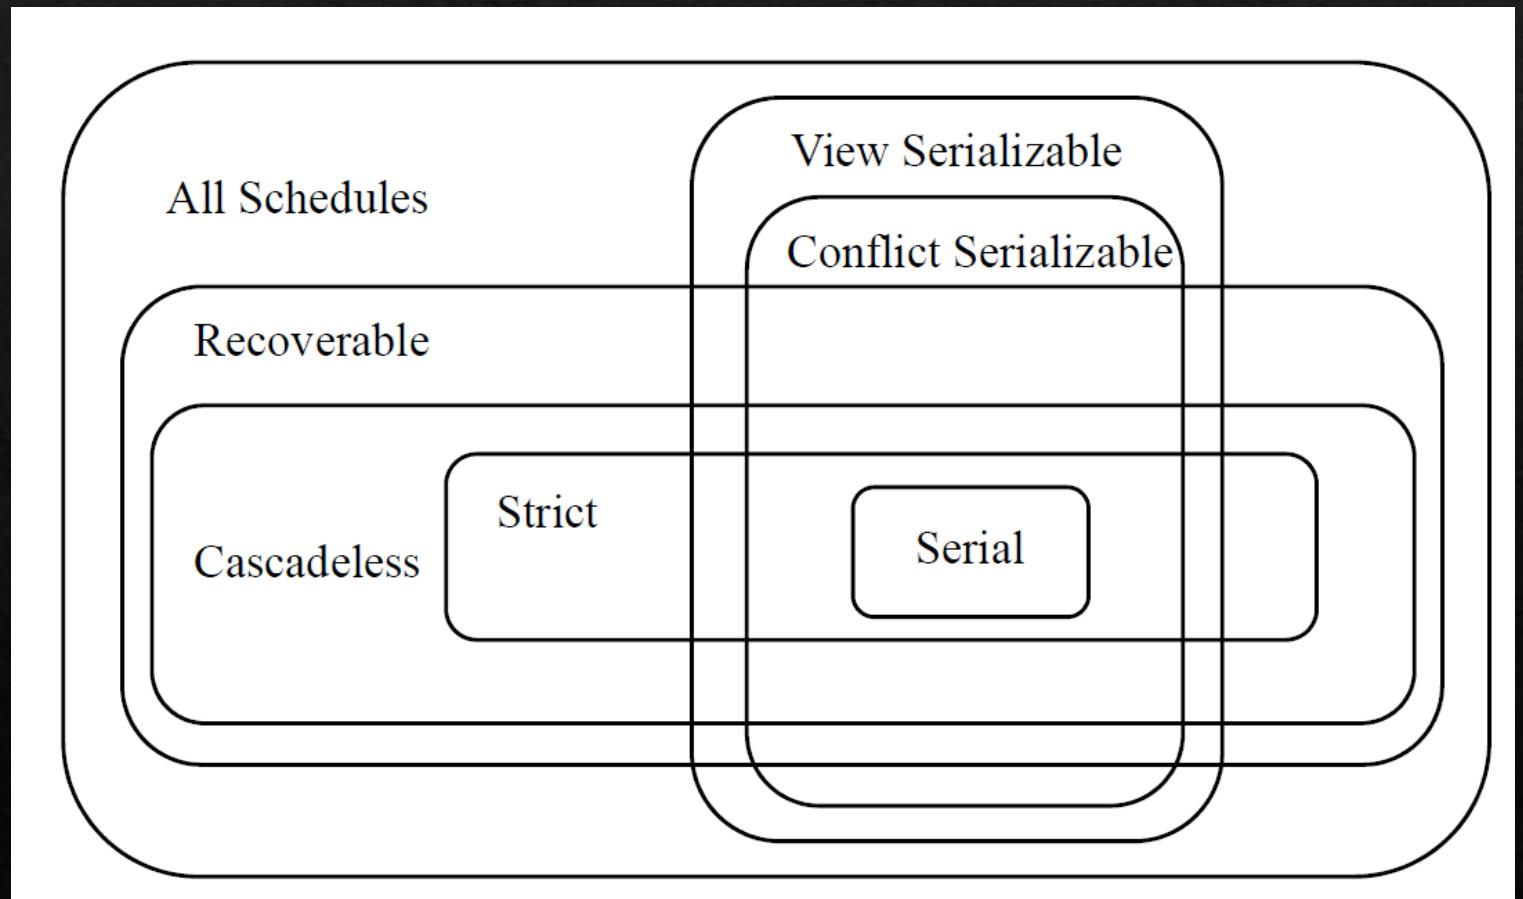
- ❖ **Atomicity:** All actions in a Xact happen, or none of them happen
 - ❖ Guaranteed by using COMMIT after all actions has finished, or ABORT if some actions are not allowed to be executed
- ❖ **Consistency:** If each Xact is consistent, and the DB was in a consistent state initially, then it should end up being consistent
- ❖ **Isolation:** Execution of one Xact is isolated from that of another Xacts
 - ❖ Net effect must be identical to executing all Xacts in some serial order
- ❖ **Durability:**

ACID Transactions

- ❖ **Atomicity:** All actions in a Xact happen, or none of them happen
 - ❖ Guaranteed by using COMMIT after all actions has finished, or ABORT if some actions are not allowed to be executed
- ❖ **Consistency:** If each Xact is consistent, and the DB was in a consistent state initially, then it should end up being consistent
- ❖ **Isolation:** Execution of one Xact is isolated from that of another Xacts
 - ❖ Net effect must be identical to executing all Xacts in some serial order
- ❖ **Durability:** If a Xact commits, its effects persist
 - ❖ Recovery schemes

Schedules

- ❖ Serial
- ❖ Conflict Serializable
- ❖ View Serializable
- ❖ Strict
- ❖ Cascadeless
- ❖ Recoverable



Serializability

- ◊ **Serial**
- ◊ Conflict Serializable
- ◊ View Serializable
- ◊ Strict
- ◊ Cascadeless
- ◊ Recoverable

Serial

- ◊ $S_1 = R_1(X), R_1(Y), W_1(Z), W_3(X), W_3(Z), R_2(Y), R_2(X), W_2(Z)$
- ◊ **No interleaving** of actions from different transactions
 - ◊ All transactions execute one after another (as though in an atomic manner)

T1	T2	T3
$R_1(X)$ $R_1(Y)$ $W_1(Z)$		$W_3(X)$ $W_3(Z)$
	$R_2(Y)$ $R_2(X)$ $W_2(Z)$	

Serializability

- ◊ **Serial**
- ◊ Conflict Serializable
- ◊ View Serializable
- ◊ Strict
- ◊ Cascadeless
- ◊ Recoverable

Serializable

- ◊ $S_1 = R_1(X), R_1(Y), W_3(X), W_3(Z), R_2(Y), W_1(Z), R_2(X), W_2(Z)$
- ◊ **effect is identical** to that of some complete **serial** schedule
 - ◊ $R_1(X) \rightarrow W_3(X) \rightarrow R_2(X)$ // preserved
 - ◊ $W_2(Z)$ performs last write on Z
 - ◊ Net effect is the same as serial schedule!

Equivalent serial schedule:

- ◊ $S_2 = R_1(X), R_1(Y), W_1(Z), W_3(X), W_3(Z), R_2(Y), R_2(X), W_2(Z)$

T1	T2	T3
$R_1(X)$ $R_1(Y)$		$W_3(X)$ $W_3(Z)$
$W_1(Z)$	$R_2(Y)$ $R_2(X)$ $W_2(Z)$	

- ◊ Serial
- ◊ Conflict Serializable
- ◊ View Serializable
- ◊ Strict
- ◊ Cascadeless
- ◊ Recoverable

Serializability

Serializable BUT NOT Conflict-serializable

- ◊ $S_1 = R_1(X), R_1(Y), W_3(X), W_3(Z), R_2(Y), W_1(Z), R_2(X), W_2(Z)$
- ◊ **effect is identical** to that of some complete **serial** schedule
 - ◊ $R_1(X) \rightarrow W_3(X) \rightarrow R_2(X)$ // preserved
 - ◊ $W_2(Z)$ performs last write on Z
 - ◊ Net effect is the same as serial schedule!
- ◊ Not conflict serializable
 - ◊ $R_1(X) \rightarrow W_3(X)$
 - ◊ $W_3(Z) \rightarrow W_1(Z)$

T1	T2	T3
$R_1(X)$ $R_1(Y)$		
$W_1(Z)$	$R_2(Y)$ $R_2(X)$ $W_2(Z)$	$W_3(X)$ $W_3(Z)$

Serializability

- ◊ Serial
- ◊ **Conflict Serializable**
- ◊ View Serializable
- ◊ Strict
- ◊ Cascadeless
- ◊ Recoverable

Conflict Serializable

- ◊ $S_1 = R_1(X), R_1(Y), W_3(X), W_1(Z), R_2(Y), W_3(Z), R_2(X), W_2(Z)$
- ◊ it is conflict equivalent to some serial schedule
 - ◊ $R_1(X) \rightarrow W_3(X) \rightarrow R_2(X)$ // preserved
 - ◊ $W_1(Z) \rightarrow W_3(Z) \rightarrow W_2(Z)$ // preserved

Conflict equivalent serial schedule:

- ◊ $S_2 = R_1(X), R_1(Y), W_1(Z), W_3(X), W_3(Z), R_2(Y), R_2(X), W_2(Z)$

T1	T2	T3
$R_1(X)$ $R_1(Y)$		$W_3(X)$
$W_1(Z)$	$R_2(Y)$ $R_2(X)$	$W_3(Z)$ $W_2(Z)$

Serializability

- ◊ Serial
- ◊ Conflict Serializable
- ◊ **View Serializable**
- ◊ Strict
- ◊ Cascadeless
- ◊ Recoverable

View serializable

- ◊ $S_1 = R_1(X), R_1(Y), W_3(X), W_3(Z), R_2(Y), W_1(Z), R_2(X), W_2(Z)$

◊ **View equivalent** to some serial schedule

- ◊ T_i reads the same initial value of A in both schedules
- ◊ T_i reads the value of A written by T_j in both schedules
- ◊ The same exact T_i performs the final write on A

View equivalent serial schedule:

- ◊ $S_2 = R_1(X), R_1(Y), W_1(Z), W_3(X), W_3(Z), R_2(Y), R_2(X), W_2(Z)$

T1	T2	T3
$R_1(X)$ $R_1(Y)$		$W_3(X)$ $W_3(Z)$
$W_1(Z)$	$R_2(Y)$ $R_2(X)$	$W_2(Z)$

Serializability

- ◊ Serial
- ◊ Conflict Serializable
- ◊ **View Serializable**
- ◊ Strict
- ◊ Cascadeless
- ◊ Recoverable

View serializable

- ◊ $S_1 = R_1(X), R_1(Y), W_3(X), W_3(Z), R_2(Y), W_1(Z), R_2(X), W_2(Z)$

View equivalent to some serial schedule

- ◊ T_i reads the same initial value of A in both schedules
- ◊ T_i reads the value of A written by T_j in both schedules
- ◊ The same exact T_i performs the final write on A

View equivalent serial schedule:

- ◊ $S_2 = R_1(X), R_1(Y), W_1(Z), W_3(X), W_3(Z), R_2(Y), R_2(X), W_2(Z)$

T1	T2	T3
$R_1(X)$ $R_1(Y)$		
$W_1(Z)$	$R_2(Y)$ $R_2(X)$ $W_2(Z)$	$W_3(X)$ $W_3(Z)$

Final write on Z

$1 \rightarrow 2$
 $3 \rightarrow 2$

Serializability

- ◊ Serial
- ◊ Conflict Serializable
- ◊ **View Serializable**
- ◊ Strict
- ◊ Cascadeless
- ◊ Recoverable

View serializable

- ◊ $S_1 = R_1(X), R_1(Y), W_3(X), W_3(Z), R_2(Y), W_1(Z), R_2(X), W_2(Z)$
- ◊ **View equivalent** to some serial schedule
 - ◊ T_i reads the same initial value of A in both schedules
 - ◊ T_i reads the value of A written by T_j in both schedules
 - ◊ The same xact T_i performs the final write on A

View equivalent serial schedule:

- ◊ $S_2 = R_1(X), R_1(Y), W_1(Z), W_3(X), W_3(Z), R_2(Y), R_2(X), W_2(Z)$

T1 reads initial value
(before T3 writes)

1→3

	T1	T2	T3
T1	$R_1(X)$ $R_1(Y)$		
T2		$R_2(Y)$ $R_2(X)$	$W_3(X)$ $W_3(Z)$
T3			$W_2(Z)$

T2 reads
from T3
3→2

Other Schedule Types

strict is like a cascadeless but for both read and write

- ◊ Serial
- ◊ Conflict Serializable
- ◊ View Serializable
- ◊ **Strict**
- ◊ Cascadeless
- ◊ Recoverable

Strict

- ◊ $R_1(X), R_1(Y), R_2(Y), W_3(X), W_3(Z), c3, R_2(X), W_1(Z), c1, W_2(Z), c2$
- ◊ A schedule S is a **strict** schedule if for every $W_i(O)$ in S, O is **not read** or **written** by another Xact until T_i either **aborts** or **commits**
 - ◊ If $W_1(Z) \rightarrow W_2(Z)$, then $W_1(Z), c1 \dots W_2(Z)$
 - ◊ T2 only writes to Z after T1 commits $W_1(Z)$
 - ◊ If $W_3(X) \rightarrow R_2(X)$, then $W_3(X), c3 \dots R_2(X)$
 - ◊ T2 only reads X after T3 commits $W_3(X)$

T1	T2	T3
$R_1(X)$ $R_1(Y)$		
	$R_2(Y)$	
		$W_3(X)$ $W_3(Z)$ commit
	$R_2(X)$	
	$W_1(Z)$ commit	
		$W_2(Z)$ commit

Other Schedule Types

- ◊ Serial
- ◊ Conflict Serializable
- ◊ View Serializable
- ◊ Strict
- ◊ **Cascadeless**
- ◊ Recoverable

Cascadeless (but not strict)

- ◊ $R_1(X), R_1(Y), R_2(Y), W_3(X), W_3(Z), c3, R_2(X), W_1(Z), W_2(Z), c2, c1$
- ◊ A schedule S is a cascadeless schedule if whenever T_i reads from T_j in S, Commit_j must precede this read action
 - ◊ If $W_4(Z) \rightarrow W_2(Z)$, then $W_4(Z), c1 \dots W_2(Z)$
 - ◊ T2 only writes to Z after T1 commits $W_4(Z)$
 - ◊ If $W_3(X) \rightarrow R_2(X)$, then $W_3(X), c3 \dots R_2(X)$
 - ◊ T2 only reads X after T3 commits $W_3(X)$

Drop the requirement
on write after write

T1	T2	T3
$R_1(X)$ $R_1(Y)$		
	$R_2(Y)$	
		$W_3(X)$ $W_3(Z)$ commit
		$R_2(X)$
	$W_1(Z)$	$W_2(Z)$ commit
		commit

Other Schedule Types

- ◊ Serial
- ◊ Conflict Serializable
- ◊ View Serializable
- ◊ Strict
- ◊ Cascadeless
- ◊ **Recoverable**

Recoverable (but not cascadeless)

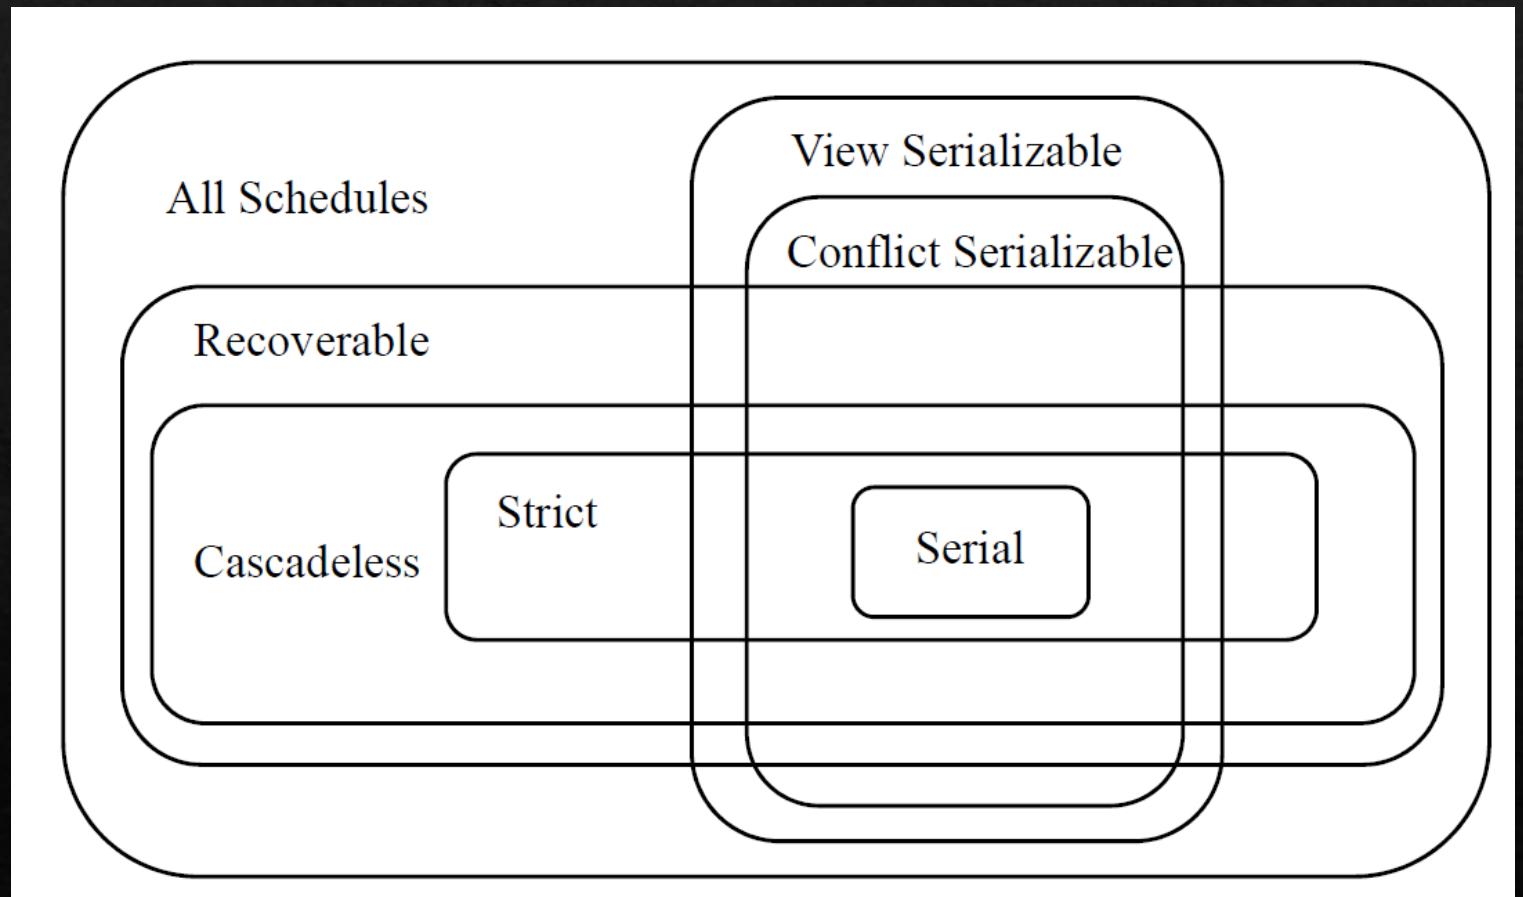
- ◊ $R_1(X), R_1(Y), R_2(Y), W_3(X), W_3(Z), R_2(X), c3, W_1(Z), W_2(Z), c2, c1$
- ◊ A schedule S is said to be a recoverable schedule if for every Xact T that commits in S, T must commit after T' if T reads from T'
 - ◊ If $W_3(X) \rightarrow R_2(X)$, then $c3 \dots c2$
 - ◊ T2 only commits after T3 commits

 recoverable only checks commits

T1	T2	T3
$R_1(X)$		
$R_1(Y)$	$R_2(Y)$	
		$W_3(X)$
		$W_3(Z)$
	$R_2(X)$	
		commit
	$W_1(Z)$	
		$W_2(Z)$
		commit
		commit

Schedules

- ❖ Serial
- ❖ Conflict Serializable
- ❖ View Serializable
- ❖ Strict
- ❖ Cascadeless
- ❖ Recoverable



Q1

- ❖ a. $\overbrace{W_1(X), R_2(X)}^{\text{1}} \overbrace{W_1(X), c_2, c_1}^{\text{2}}$
 - ❖ S1. Not view serializable cos $1 \rightarrow 2$, and $2 \rightarrow 1$. Not recoverable because to be recoverable, T2 must commit after T1 (since T2 read from T1), but this is not true
- ❖ b. $R_1(X), R_1(Y), \overbrace{W_1(X), R_2(Y)}^{\text{2}} \overbrace{W_3(Y), W_1(X), R_2(Y)}^{\text{3}}, c_3, c_2, c_1$
 - ❖ S4. Not view serializable. Based on object Y, $2 \rightarrow 3$ (T2 reads initial value of Y), $3 \rightarrow 2$ (T2 reads Y from T3). Recoverable because T2 read from T3, and T2 committed after T3 commits. Not cascadeless: T2 did not read Y after T3 commits $W_3(Y)$. i.e. $W_3(Y), c_3, \dots R_2(Y)$
- ❖ c. $\overbrace{W_1(X), R_2(Y)}^{\text{1}}, \overbrace{R_1(Y), R_2(X)}^{\text{2}}, c_1, c_2$
 - ❖ S6. Conflict serializable because we can have $W_1(X), R_1(Y), R_2(Y), R_2(X)$. Recoverable because T2 read from T1, and T2 committed after T1 commits. Not cascadeless, same reason as part b
- ❖ d. $R_1(X), R_2(X), W_1(X), c_1, W_2(X), c_2$
 - ❖ S10. Not view serializable because $1 \rightarrow 2$ (T1 performs initial read, before T2 writes X, and because T2 must perform final write on X), and $2 \rightarrow 1$ (because T2 performs initial read, before T1 writes X). Strict because whenever T2 writes to X, it happens only after T1 has committed.

Q2

Give an example of a schedule with two or more transaction with the following three properties:

- ❖ T1 commits before T2 starts.
- ❖ The schedule is conflict serializable.
- ❖ In any equivalent serial schedule, T2 must come before T1 (there may be other transactions between T2 and T1).

Q2

Enforce this

- ❖ $T_2 \rightarrow T_3$ (on some object X)
- ❖ $T_3 \rightarrow T_1$ (on some object Y)
- ❖ $S_1: W_3(X), R_1(X), c1, R_2(Y), W_3(Y)$
- ❖ $S_2: R_2(Y), W_3(X), W_3(Y), R_1(X)$ // conflict equivalent serial schedule

Other alternatives

- ❖ $W_3(A), W_1(A), c1, W_2(B), c2, W_3(B), c3$

Q3a

- ❖ Consider the following transactions

T0	read(A) read(B) if A = 0 then B \leftarrow B+1 write(B)	T1:	read(B) read(A) if B = 0 then A \leftarrow A+1 write(A)
-----------	--	------------	--

- ❖ Show that every serial execution involving these two transactions preserves the consistency of the database.

Q3a

- ❖ Consider the following transactions

T0	read(A) read(B) if A = 0 then B \leftarrow B+1 write(B)	T1:	read(B) read(A) if B = 0 then A \leftarrow A+1 write(A)
-----------	--	------------	--

- ❖ Show that every serial execution involving these two transactions preserves the consistency of the database.
- ❖ Show that T0 \rightarrow T1 and T1 \rightarrow T0 (symmetrical case) preserves the consistency requirement

Q3a

◇ If $T_0 \rightarrow T_1$

T0	read(A) read(B) if A = 0 then B \leftarrow B+1 write(B)	A = 0 B = 0 B = 1 A = 0, B = 1
T1	read(B) read(A) if B = 0 then A \leftarrow A+1 write(A)	B = 1 A = 0 - A = 0, B = 1

◇ If $T_1 \rightarrow T_0$

T1	read(B) read(A) if B = 0 then A \leftarrow A+1 write(A)	B = 0 A = 0 A = 1 A = 1, B = 0
T0	read(A) read(B) if A = 0 then B \leftarrow B+1 write(B)	A = 1 B = 0 - A = 1, B = 0

consistency requirement

- **A = 0 or B = 0**
- A = B = 0 be the initial values

Q3b

Show a concurrent execution of T0 and T1 which produces a non-serializable schedule.

❖ S_1 :

Q3b

Show a concurrent execution of T0 and T1 which produces a non-serializable schedule.

- ❖ $S_1: R_0(A), R_0(B), R_1(B), R_1(A), W_0(B), W_1(A)$
-
- The diagram consists of two curved arrows. The top arrow originates from the number '1' and points to the number '0'. The bottom arrow originates from the number '0' and points to the number '1'. These arrows are positioned above the schedule S_1 , indicating a cyclic dependency between the two transaction instances.

Q3c

Is there a concurrent execution of T0 and T1 which produces a serializable schedule?

Q3c

- ❖ $R_0(A)$, $R_0(B)$, $R_1(B)$, $R_1(A)$, $W_0(B)$, $W_1(A)$
- ❖ To make it serializable, we must somehow have $0 \rightarrow 1$ (or $1 \rightarrow 0$)
- ❖ Let's try $0 \rightarrow 1$, by moving $W_0(B)$
- ❖ $R_0(A)$, $R_0(B)$, $W_0(B)$, $R_1(B)$, $R_1(A)$, $W_1(A)$
- ❖ But this is not a concurrent execution that we are looking for (no interleaving executions)
- ❖ It's the same outcome for $1 \rightarrow 0$
- ❖ No, we cannot achieve a concurrent execution that produces a serializable schedule

Q3c

- ❖ There is no parallel execution resulting in a serializable schedule. From part(a), we know that a serializable schedule results in A=0 OR B=0. Suppose we start with T0 read(A). Then when the schedule ends, no matter when we run the steps of T1, B=1. Now suppose we start executing T1 prior to completion of T0. Then T1 read(B) will give B a value of 0. So, when T1 completes, A=1. Thus, B=1 AND A=1. We get similar logic for starting with T1 read(B).

T0	T1
read(A) read(B) if A = 0 then B \leftarrow B+1 write(B) // B=1	read(B) // B=0 read(A) if B = 0 then A \leftarrow A+1 write(A) // A=1

As long as the yellow line happens before the blue line, the data will not be consistent

Q4

Consider the schedule S:

- ❖ $R_2(B), W_2(A), R_1(A), R_3(A), W_1(B), W_2(B), W_3(B)$

Is S view-serializable?



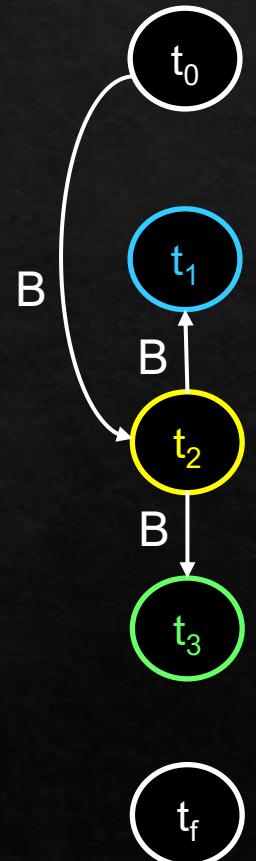
Q4

Consider the schedule S:

- ◊ R₂(B), W₂(A), R₁(A), R₃(A), W₁(B), W₂(B), W₃(B)

Is S view-serializable?

- ◊ T2 reads initial value of B (before anyone writes): 2→1, 2→3



Q4

Consider the schedule S:

- ◇ $R_2(B)$, $W_2(A)$, $R_1(A)$, $R_3(A)$, $W_1(B)$, $W_2(B)$, $W_3(B)$

Is S view-serializable?

- ◇ T2 reads initial value of B (before anyone writes): $2 \rightarrow 1$, $2 \rightarrow 3$
- ◇ T1 reads value of A from T2: $2 \rightarrow 1$
- ◇ T3 reads value of A from T2: $2 \rightarrow 3$



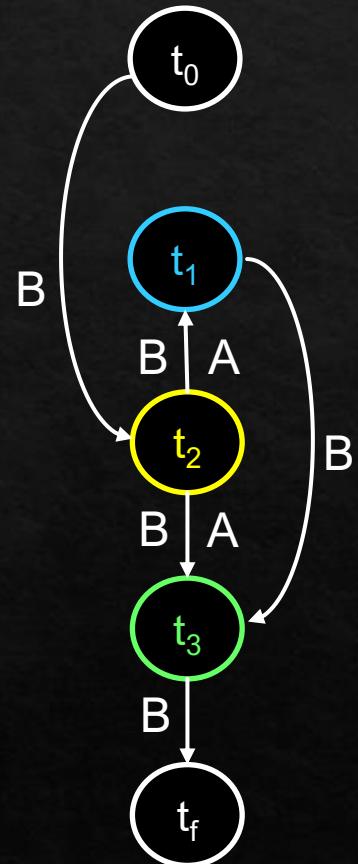
Q4

Consider the schedule S:

- ◇ $R_2(B), W_2(A), R_1(A), R_3(A), \underline{W_1(B)}, \underline{W_2(B)}, \underline{W_3(B)}$

Is S view-serializable?

- ◇ T2 reads initial value of B (before anyone writes): $2 \rightarrow 1, 2 \rightarrow 3$
- ◇ T1 reads value of A from T2: $2 \rightarrow 1$
- ◇ T3 reads value of A from T2: $2 \rightarrow 3$
- ◇ T3 performs final write on B (and no writes after that): $1 \rightarrow 3, 2 \rightarrow 3$



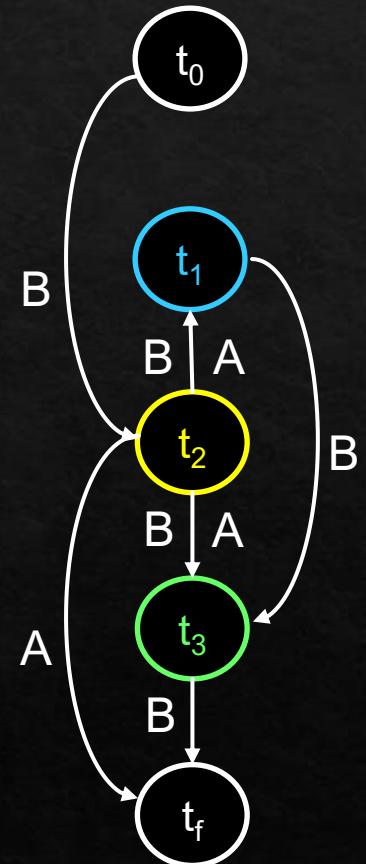
Q4

Consider the schedule S:

- ◇ $R_2(B)$, $W_2(A)$, $R_1(A)$, $R_3(A)$, $W_1(B)$, $W_2(B)$, $W_3(B)$

Is S view-serializable?

- ◇ T2 reads initial value of B (before anyone writes): $2 \rightarrow 1$, $2 \rightarrow 3$
- ◇ T1 reads value of A from T2: $2 \rightarrow 1$
- ◇ T3 reads value of A from T2: $2 \rightarrow 3$
- ◇ T3 performs final write on B (and no writes after that): $1 \rightarrow 3$, $2 \rightarrow 3$
- ◇ T2 performs final write on A



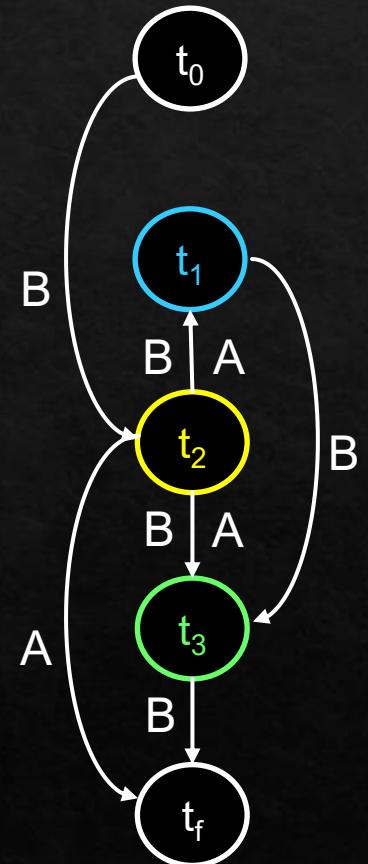
Q4

Consider the schedule S:

- ◇ $R_2(B), W_2(A), R_1(A), R_3(A), W_1(B), W_2(B), W_3(B)$

Is S view-serializable?

- ◇ T2 reads initial value of B (before anyone writes): $2 \rightarrow 1, 2 \rightarrow 3$
- ◇ T1 reads value of A from T2: $2 \rightarrow 1$
- ◇ T3 reads value of A from T2: $2 \rightarrow 3$
- ◇ T3 performs final write on B (and no writes after that): $1 \rightarrow 3, 2 \rightarrow 3$
- ◇ View equivalent serial schedule: $2 \rightarrow 1 \rightarrow 3$
- ◇ $R_2(B), W_2(A), W_2(B), R_1(A), W_1(B), R_3(A), W_3(B)$



Q4

Is S conflict serializable?

Blind writes

- ◊ $R_2(B)$, $W_2(A)$, $R_1(A)$, $R_3(A)$, $W_1(B)$, $W_2(B)$, $W_3(B)$
- ◊ S is not conflict serializable.
 - ◊ $W_1(B)$, $W_2(B)$ implies $1 \rightarrow 2$ but
 - ◊ $W_2(A)$, $R_1(A)$ suggests $2 \rightarrow 1$)
- ◊ But it is still view serializable.
 - ◊ View serializable schedule accepts blind writes (writes that are not read afterwards)

End