


Key Exchange Protocols

Key Exchange Protocols

- Why do we need KE protocols?
 - A way to establish secrets
 - Even if we have pre-established secrets, we want to use refresh “keys” per session
- Remark on (Perfect) Forward Secrecy
 - Protect Encrypted information, even if long-term key (for client and server) is compromised
 - Idea: Generate session keys, and throw away messages used to generate sessions keys...


Key Exchange: NS protocol

- Needham Schroder [1977] – Kerberos



- K_s and K_c are pre-established secrets between (C,S)
- Think about these 2 goals:
 - Is (N_c, N_s) a "good shared key"? [Key Secrecy]
 - Should A believe it's talking to B? [Entity Authentication]
 - Should B be sure if it's talking to A?

Lowe's Attack on NS protocol [1995]



- *Attacker knows N .*
- C thinks its talking to I, but *S thinks its talking to C*

Key Exchange: Diffie Hellman [1976]



Pick fresh a

$$g^a \pmod{p}$$



$$g^b \pmod{p}$$



$$K = g^{ba} \pmod{p}$$



Pick fresh b

$$K = g^{ab} \pmod{p}$$

- Is it secure (authenticated) key-exchange?
- Assuming: **CDH** is hard, **DLOG** is hard

Diffie Hellman Problem
(Given g, g^a, g^b . Find: g^{ab})

Discrete Log Problem
(Given g, g^a . Find: a)

Analysis of DHKE protocol

- Think about these 2 goals:
 - Is K a “good key” to talk with B ? [Key Secrecy]
 - Should A be sure it’s talking to B ? [Entity Authentication]

a new value of a here can be chosen whenever want to refresh the secret key

Pick fresh a $\xrightarrow{g^a \pmod{p}}$

$\xleftarrow{g^b \pmod{p}}$ Pick fresh b

$$K = g^{ba} \pmod{p}$$

$$K = g^{ab} \pmod{p}$$

Analysis of DHKE protocol

- DH does not achieve entity authentication



Pick fresh a

g^a

g^{r1}

g^b

Pick fresh b

g^{r2}

$$K1 = g^{r2 * a} \pmod{p}$$

$$K2 = g^{r1 * b} \pmod{p}$$

Adding Signatures to DHE: Station-to-Station (STS) Protocol

Pick fresh a

$$M1 = g^a$$



Pick fresh b

$$K = g^{ab} \pmod{p}$$

$$M2 = g^b, \{Sig_B(M2, M1)\}_K$$



$$K = g^{ba} \pmod{p}$$


$$\{Sig_A(M1, M2)\}_K$$



- Is this protocol secure?
 - Is K a “good key” to talk with B ? [Key Secrecy]
 - Should B be sure it’s talking to A ? [Entity Authentication]
 - Should A be sure it’s talking to B ? [Entity Authentication]

Security Properties: A Hierarchy of Authentication Specifications [Lowe'97]

We are ready for Secure Channels...



A Secure Channel establishes **between 2 programs**,
a data channel that has:

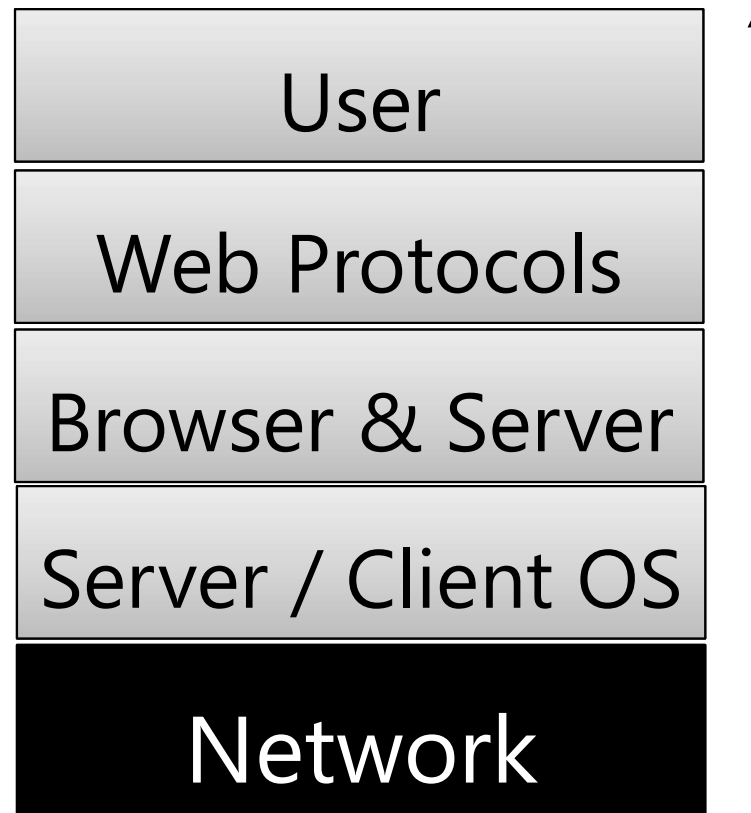
- **C**onfidentiality
- **I**ntegrity
- **A**uthentication

against a computationally-bounded "network attacker"

[Dolev-Yao-1983]


Recap:

Threat Model At Different Layers



Recap:

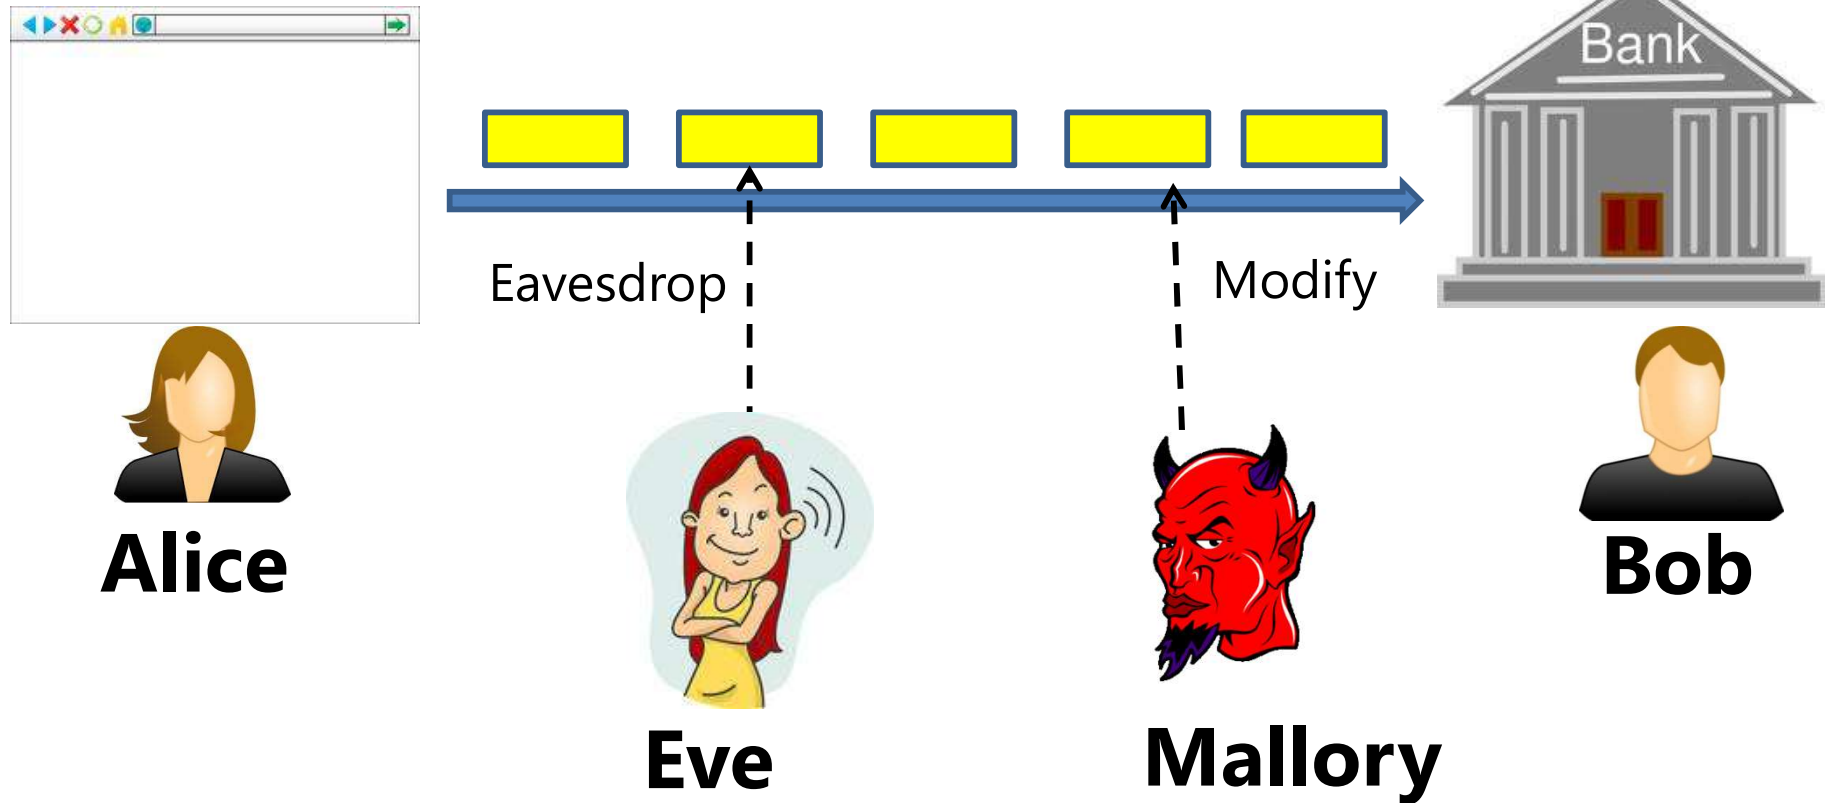
Story So Far



- Attacks: BGP, TCP/IP, DNS
- Threat Model:
 - Attackers (Eve & Mallory)
 - Assumptions
 - Desired Security Property:
 - The “CIA” of secure channels
- Solution: Secure Channels
- Building Blocks: Crypto primitives

Definition: Network Attacker


<http://bobbank.com>



Intercept ALL Traffic between Alice and Bob!

- Eve is assumed to only eavesdrop on traffic
- Mallory can listen and tamper with traffic

Definition: A Secure Channel



A Secure Channel is a data communication protocol established **between 2 programs** which preserves data:

- **C**onfidentiality
- **I**ntegrity
- **A**uthentication

against a computationally-bounded "network attacker"

[Dolev-Yao-1983]

* Note that availability is not a goal. So, denial-of-service attacks are permitted by the threat model.

* Integrity is also referred to as "authenticity" sometimes. Not to be confused with "authentication"

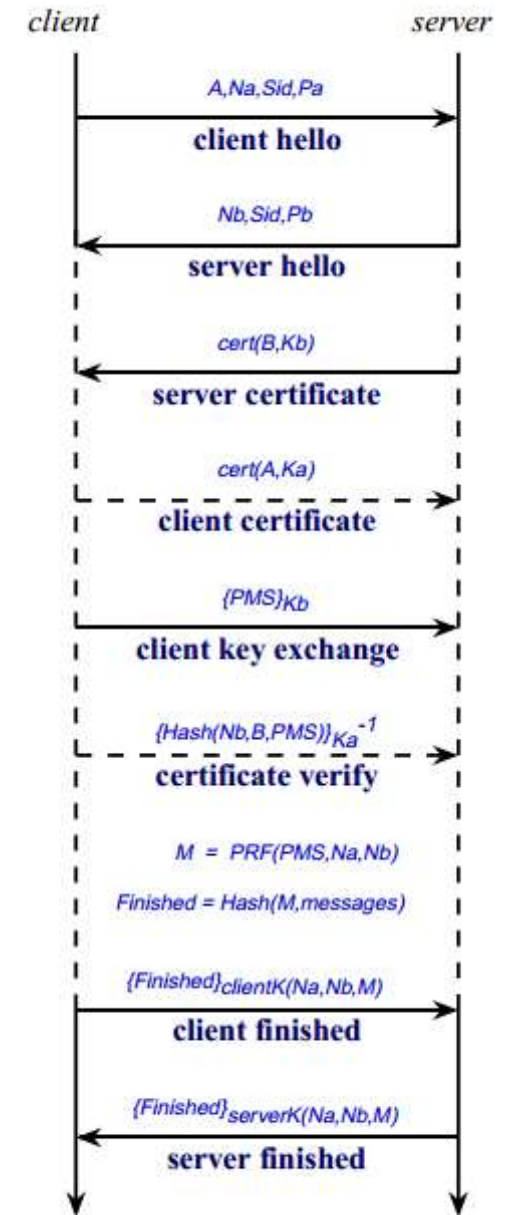
Secure Channels:

The Case of HTTPS (SSL/TLS)

Secure Channel on the Internet:

SSL + HTTP = HTTPS

- Used in HTTPS, SMTP, fax, ...
- Originated in Netscape SSL 2.0 [1993]
- Revised name to TLS
- Many versions:
 - TLS 1.2 [2008]
 - TLS 1.3 [2018]



A Very High-Level of TLS



Negotiation Phase



Key Exchange (RSA, DHE, ...)



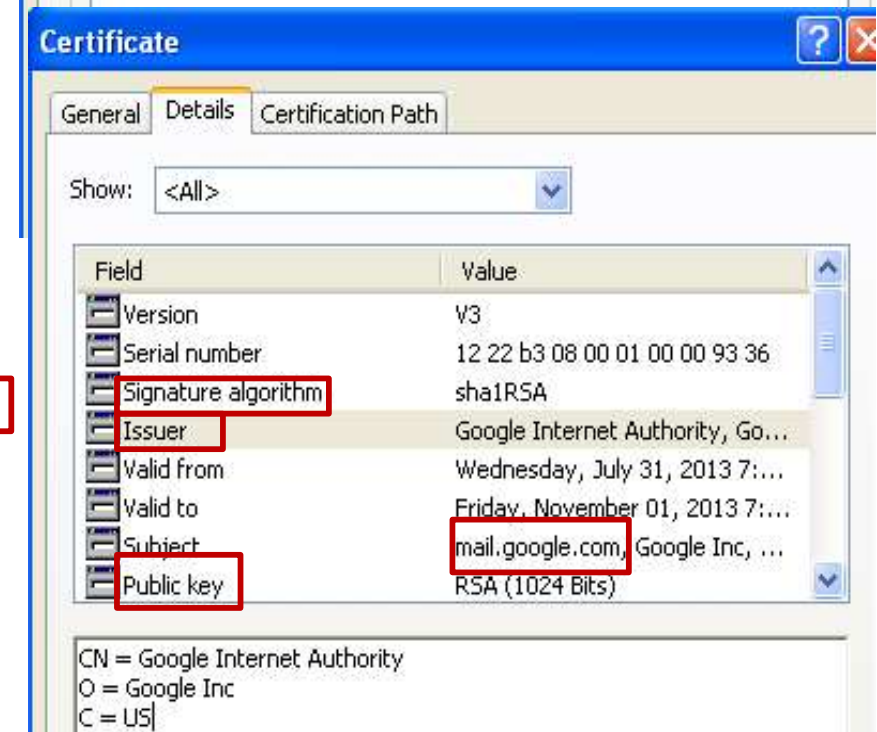
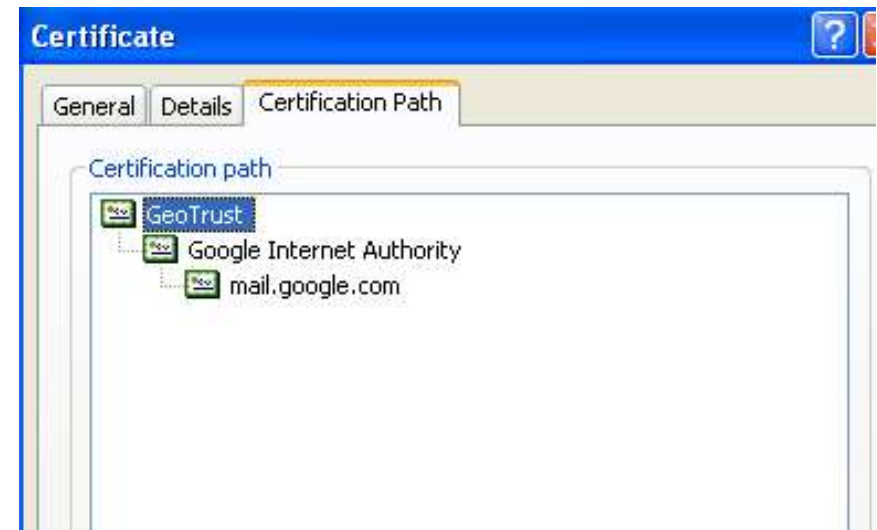
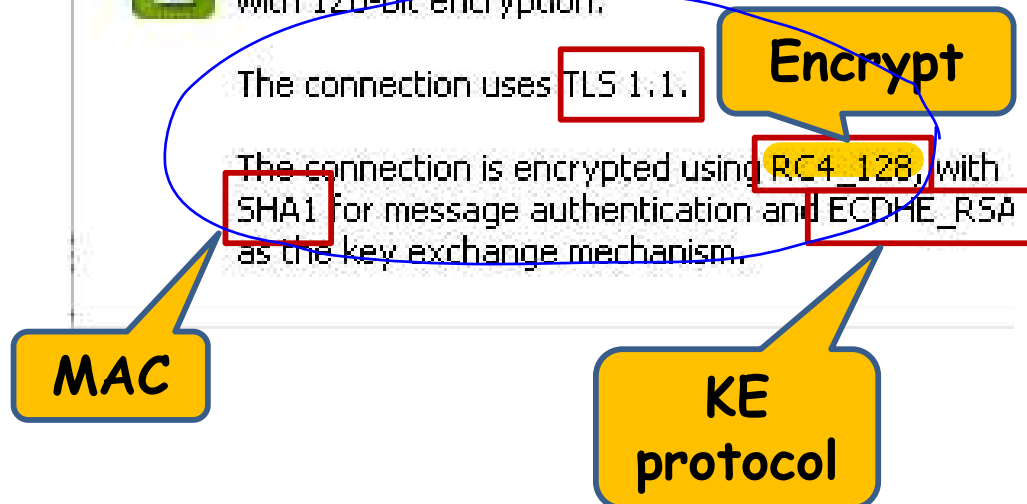
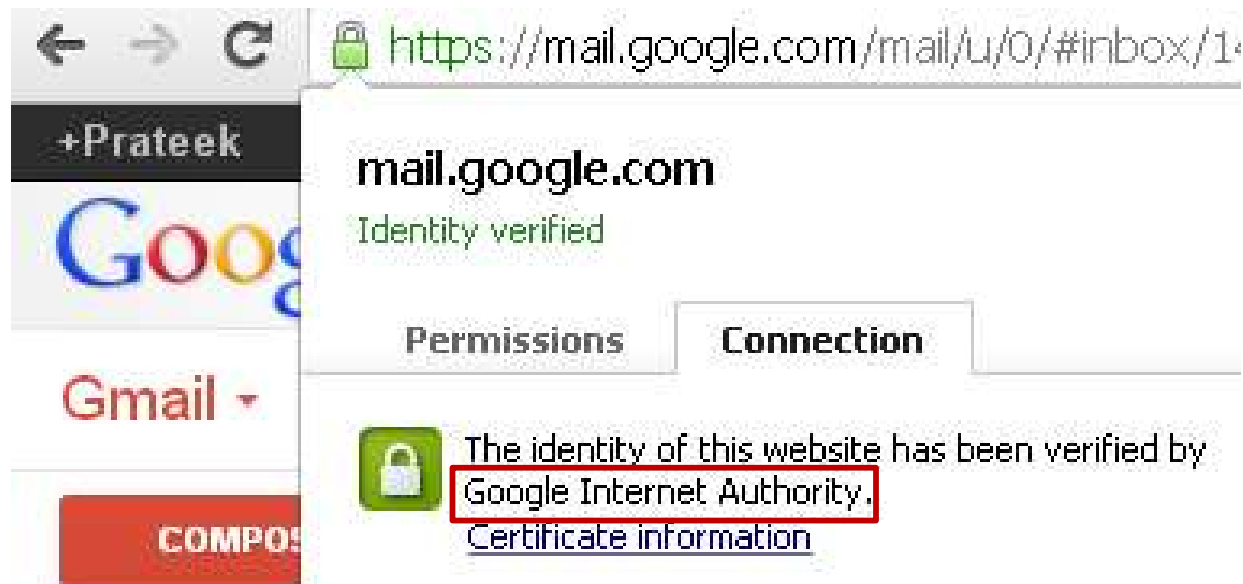
Symmetric Key Encrypted Session



Re-negotiation



HTTPS

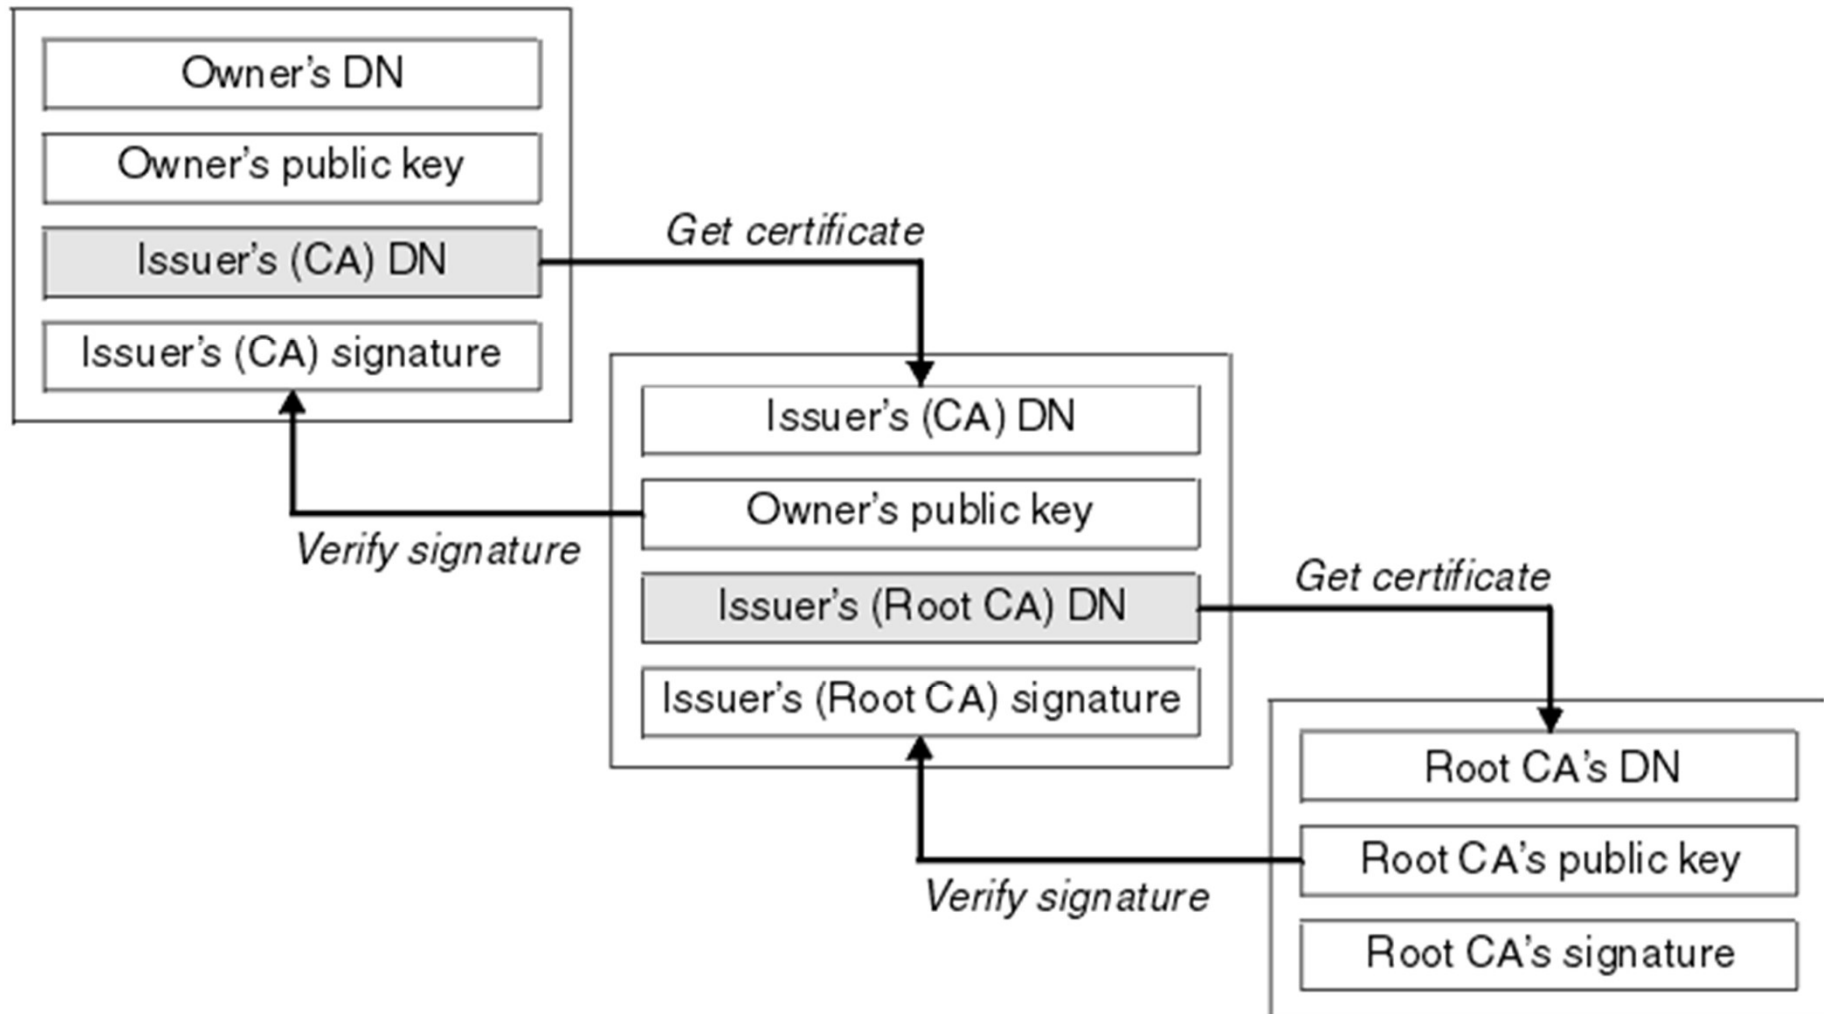


Chain Of Trusted Certificates

- Root CAs (e.g. GeoTrust)
 - Can designate Intermediate CA
 - E.g. Google Internet Authority
 - Restricted to signing certs for its subdomains
- Where does the browser start trusting?
 - Root CA's certificates are baked in your browser
 - ~50
- Who are the root CAs for the web?
 - Symantec (GeoTrust) – 38%
 - Comodo – 20%
 - GoDaddy – 13%, GlobalSign – 10%

Chain Of Trusted Certificates

- How Does the Chain Get Verified?

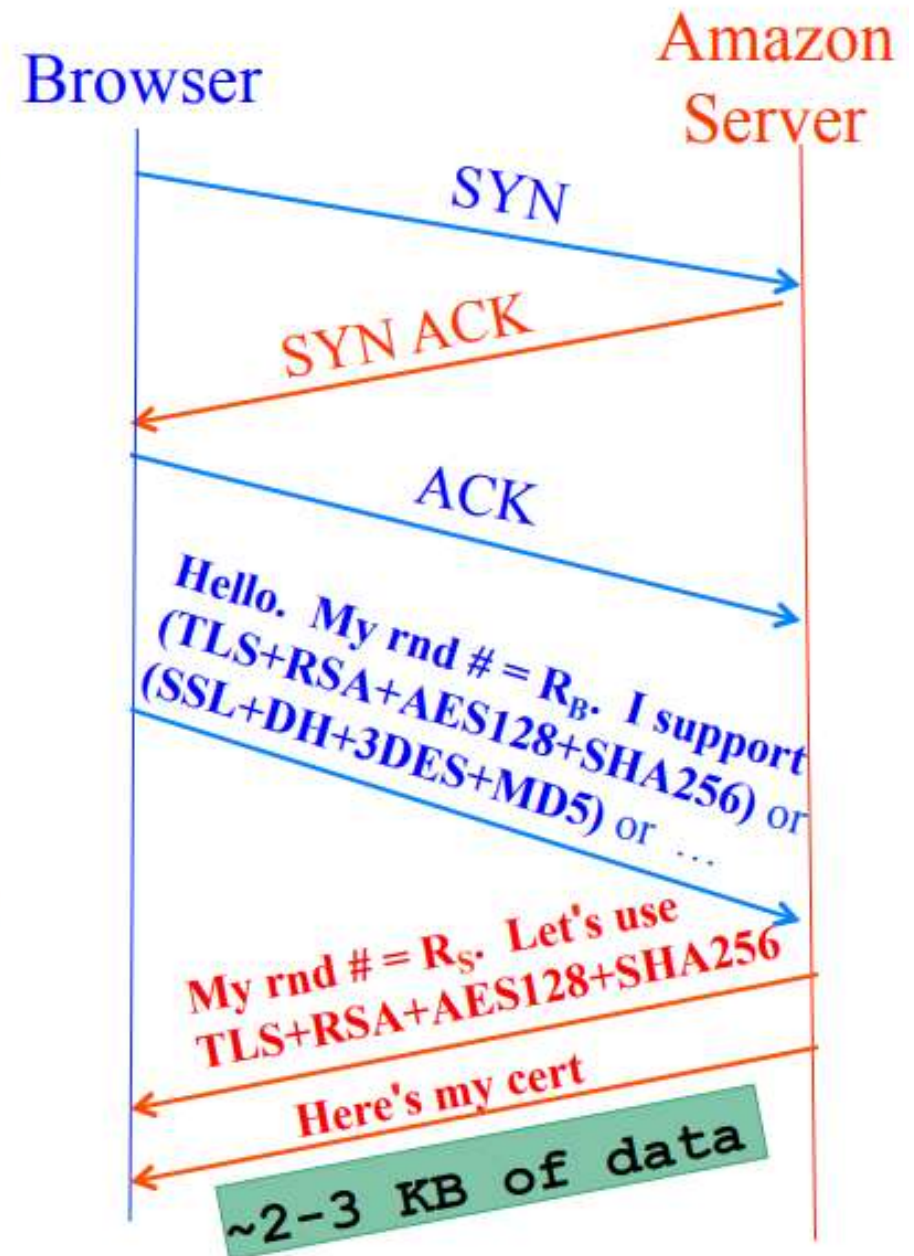


HTTPS (SSL/TLS):

Details of the Protocol


HTTPS Connection (TLS / SSL)

- Browser (client) connects via TCP to Amazon's **HTTPS** server
- Client picks 256-bit random number R_B , sends over list of crypto protocols it supports
- Server picks 256-bit random number R_S , selects *cipher suite* to use for this session
- Server sends over its certificate
- (all of this is in the clear)
- **Client now validates cert**



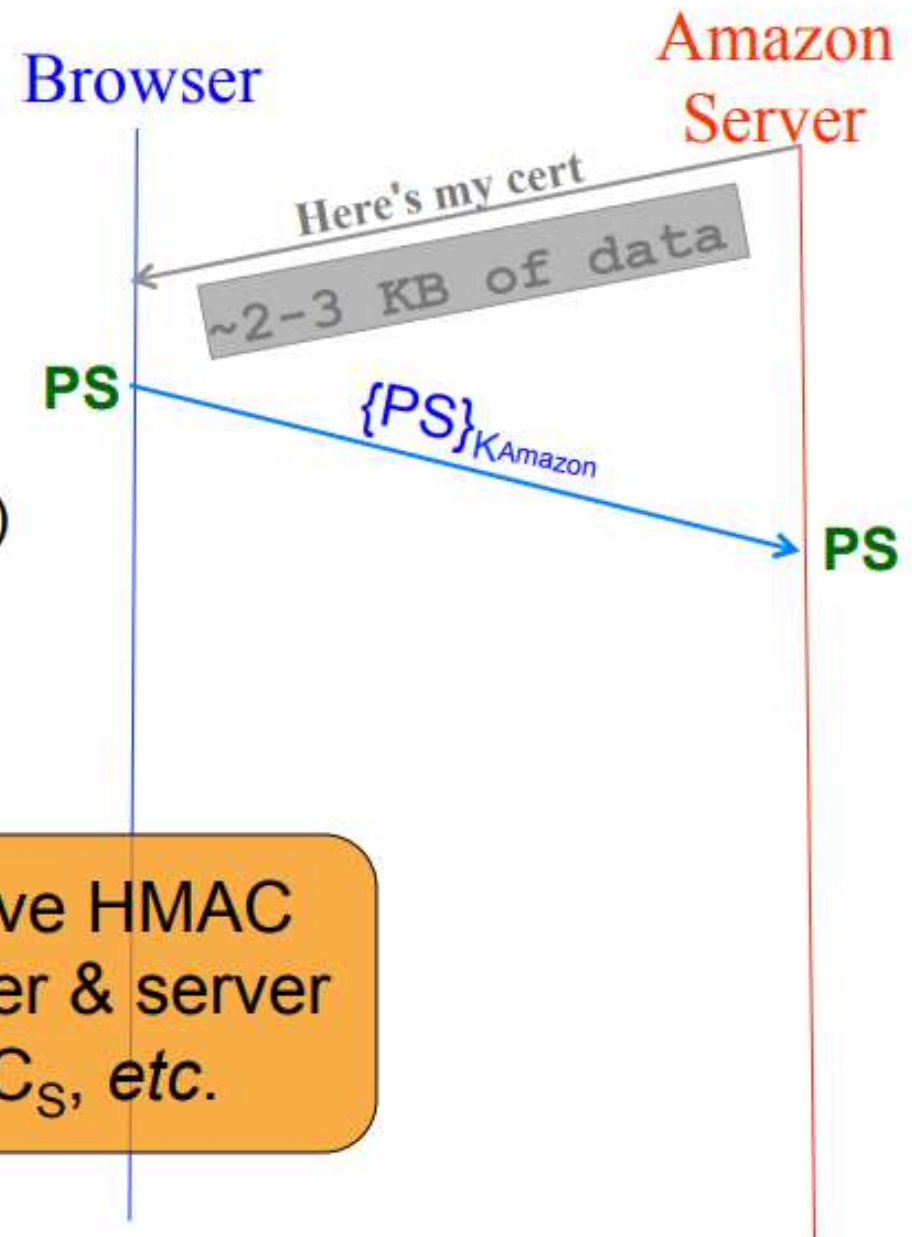
HTTPS Connection (TLS / SSL)

- For RSA, browser constructs long (368 bits) “Premaster Secret” **PS**
- Browser sends **PS** encrypted using Amazon's public RSA key K_{Amazon}
- Using **PS**, R_B , and R_S , browser & server derive symm. *cipher keys* (C_B , C_S) & MAC *integrity keys* (I_B , I_S)
 - One pair to use in each direction




HTTPS Connection (TLS / SSL)

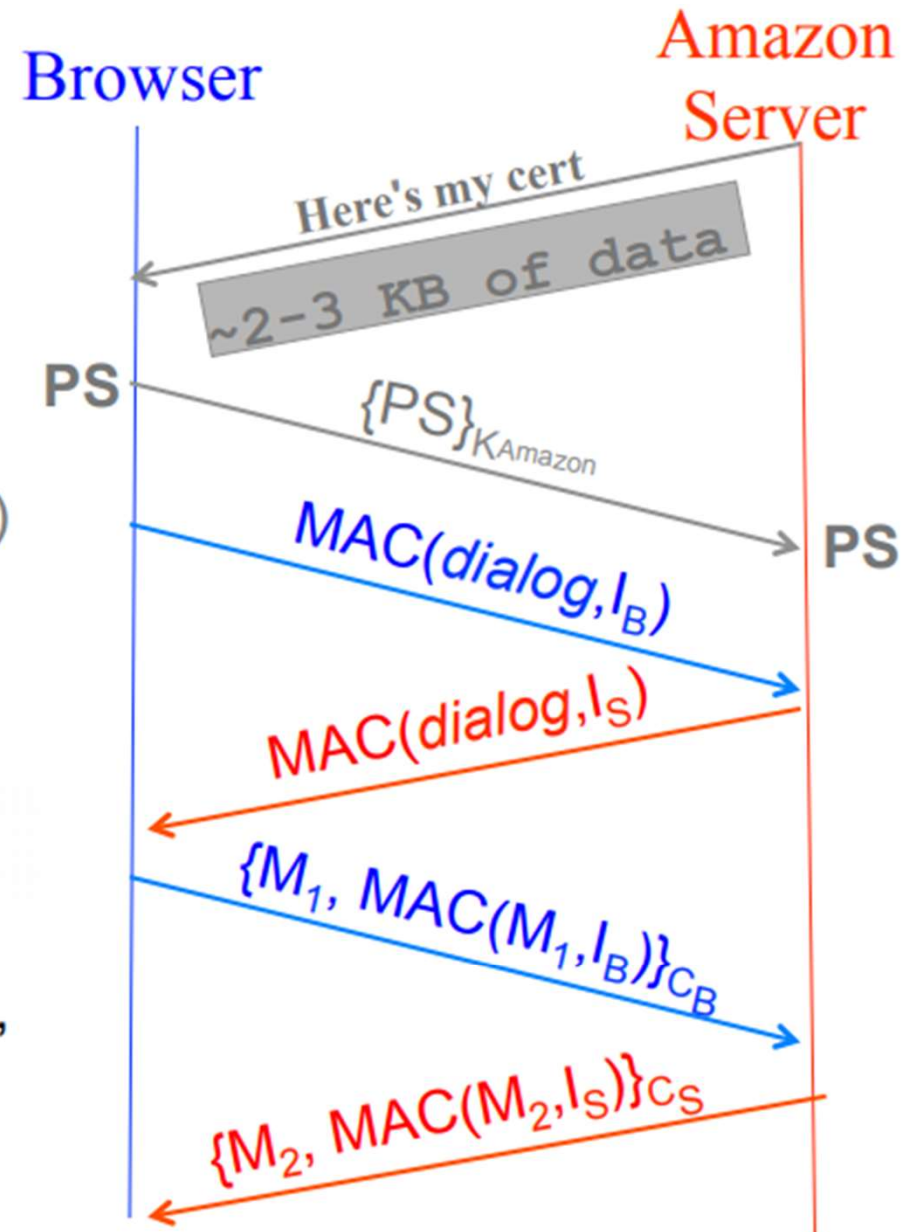
- For RSA, browser constructs long (368 bits) “Premaster Secret” **PS**
- Browser sends **PS** encrypted using Amazon's public RSA key K_{Amazon}
- Using **PS**, R_B , and R_S , browser & server derive symm. cipher keys (C_B , C_S) & MAC integrity keys (I_B , I_S)
 - One pair to use in each direction



PS is used as the key for iterative HMAC invocations on $R_B \parallel R_S$. Browser & server use the output to generate C_B , C_S , etc.

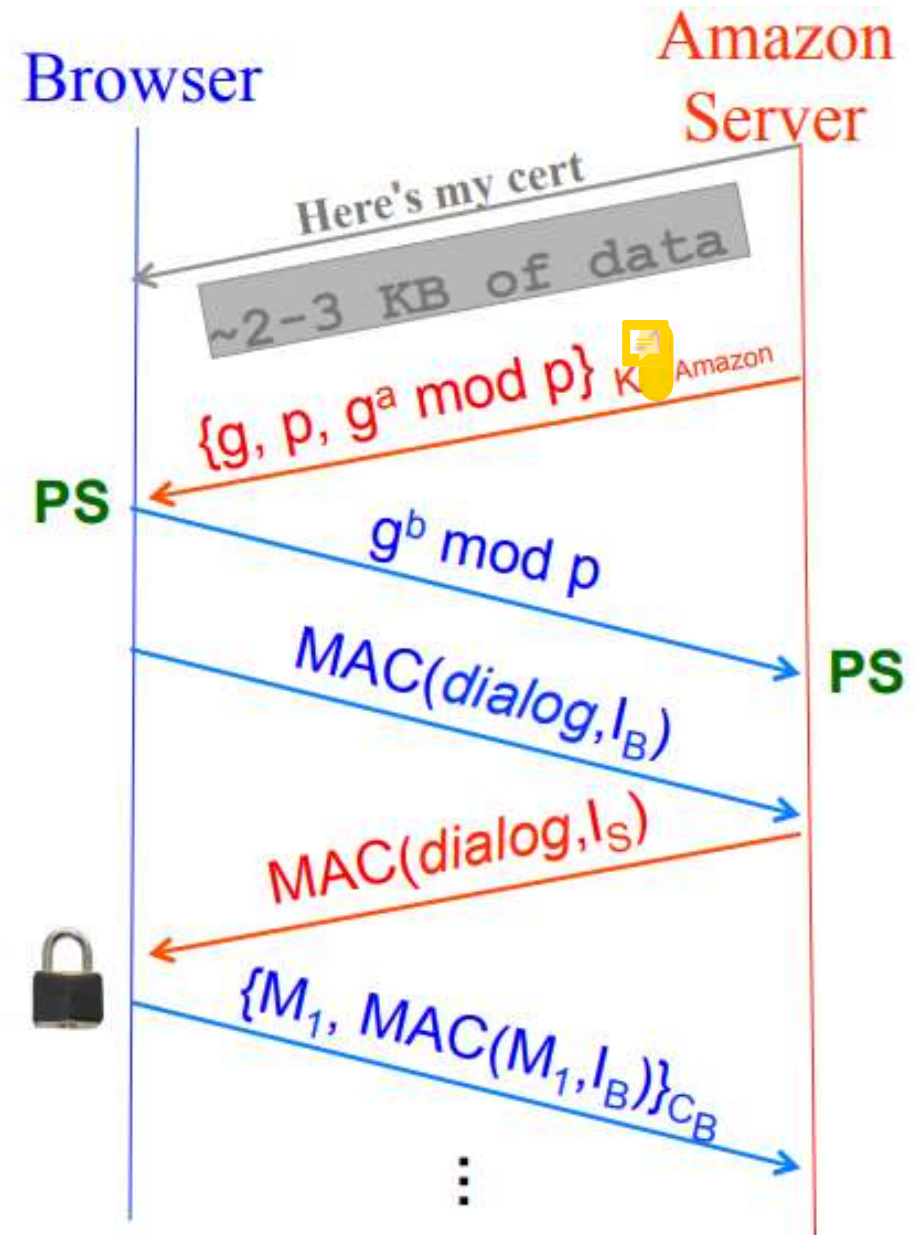
HTTPS Connection (TLS / SSL)

- For RSA, browser constructs long (368 bits) “Premaster Secret” **PS**
- Browser sends PS encrypted using Amazon's public RSA key K_{Amazon}
- Using PS, R_B , and R_S , browser & server derive symm. *cipher keys* (C_B , C_S) & MAC *integrity keys* (I_B , I_S)
 - One pair to use in each direction
- Browser & server exchange MACs computed over entire dialog so far
- If good MAC, Browser displays 
- All subsequent communication encrypted w/ symmetric cipher (e.g., **AES128**) cipher keys, MACs
 - Messages also numbered to thwart replay attacks



HTTPS Connection (TLS / SSL): Alternative via Diffie-Hellman KE

- For Diffie-Hellman, server generates random **a**, sends public params and $g^a \bmod p$
 - **Signed** with server's public key
- Browser verifies signature
- Browser generates random **b**, computes **PS** = $g^{ab} \bmod p$, sends to server
- Server also computes **PS** = $g^{ab} \bmod p$
- Remainder is as before: from **PS**, **R_B**, and **R_S**, browser & server derive symm. cipher keys (**C_B**, **C_S**) and MAC integrity keys (**I_B**, **I_S**), etc...



HTTPS (SSL/TLS): Security Analysis

Security Analysis of HTTPS

- We will assume that:
 - Cryptographic primitives are secure
 - Interacting end points are uncompromised
 - A "Perfect" Protocol achieves its stated / defined properties both in design and implementation
 - Attacker can do anything within its defined power

Important Principles:

- (1) State threat model, else there's no security argument!
- (2) Assumptions can fail, but that's not a flawed argument
- (3) Choose reasonable assumptions in your threat model

Assumptions in the threat model

- User is using a secure channel
- Crypto primitives are secure
- TLS protocol design is secure
- TLS protocol implementation is secure
- Certificate issuers are uncompromised
- Users check browser UI correctly
- Alice & Bob's secrets are secure
- Entities are authenticated correctly

Question (I)

- You visit <https://gmail.com> on a WiFi network with no cert errors. Can you safely assume that your email will reach Gmail safe from **ALL** network attackers?
 - Assume that HTTPS is a “perfect” secure channel
- **IMP:** A **valid security argument** should accept the threat model. You aren’t allowed to:
 - treat assumptions of the threat model as false
 - make additional assumptions just to arrive at a conclusion

Question (I) is Important!

- You visit <https://gmail.com> on a WiFi network with no cert errors. Can you safely assume that your email will reach Gmail safe from ALL network attackers?
- Stick to the threat model – assume the assumptions hold and then check:
 - (A) Defeats DNS Cache Poisoning?
 - (B) Defeats BGP Route Hijacking?
 - (C) Defeats TCP / IP attacks?

Yes, it does!

Question (II)

- You visit <http://evil.com>



- Should you switch to <https://evil.com>?

Question (III)

- You are installing an OS update from <http://microsoft.com>
- Is there a problem here? Fix?

SSL / TLS Protocol (Simplified)

CS3235 Lecture Notes

The basic TLS/SSL handshake protocol is presented below. Note that we use $\{X\}_{K^{-1}}$ to denote an encrypted message X with private key K^{-1} and $\{X\}_K$ is the encrypted message with the public key K . In the Diffie-Hellman (DH) below we denote g as the group generator over a multiplicative group of integers modulo a prime p ; values a and b are secrets. Assume that the browser, the CA and the CA keys are not compromised.

1. Client $\xrightarrow{\text{ClientHello}}$ Server: Client sends a **ClientHello** message containing random number R_b and supported cipher suites $Cipher_c$. The ciphersuite will select the key exchange protocol (e.g. RSA, Diffie-Hellman (DH)), the encryption algorithm for the data and the MAC algorithm.
2. Server $\xrightarrow{\text{ServerHello}}$ Client: Server sends a **ServerHello** message random number R_s , supported cipher suites $Cipher_s$.
3. Server $\xrightarrow{\text{Certificate}}$ Client: Server sends its **Certificate** containing the server's public key K_s signed with K_{auth}^{-1} where $auth$ is some signing authority (typically a CA) and K_{auth}^{-1} is the authority's private key. If the certificate is invalid, client aborts.
4. Server $\xrightarrow{\text{ServerKeyExchange}}$ Client:
If DH is selected, the server sends **ServerKeyExchange** message $\{g, p, g^a \bmod p\}_{K_s^{-1}}$
5. Server $\xrightarrow{\text{ServerHelloDone}}$ Client: Server signals end of handshake negotiation
6. Client $\xrightarrow{\text{ClientKeyExchange}}$ Server:
If DH is selected, client sends $g^b \bmod p$.
If RSA is selected, client sends $\{PS\}_{K_s}$, where PS is called PreMasterSecret and is a random value generated by the client.
After the **ClientKeyExchange** the client and server will derive a shared key: $K_{sess} \leftarrow$ derived from DH messages 4, 6 via a standard Diffie Hellman Key Exchange procedure, or $K_{sess} = PS$ if RSA is selected. Then the client and server derive the cipher keys CK_c, CK_s and integrity keys I_c, I_s deterministically using function $f()$:
 $(CK_c, CK_s, I_c, I_s) \leftarrow f(R_b, R_s, K_{sess})$
7. Client $\xrightarrow{\text{ChangeCipherSpec}}$ Server: Client sends $\text{MAC}(\text{msg}, I_c)$
8. Server $\xrightarrow{\text{ChangeCipherSpec}}$ Client: Server sends $\text{MAC}(\text{msg}, I_s)$

After the last step all data takes the form:

- Client \rightarrow Server: $\{M1, \text{MAC}(M1, I_c)\}_{CK_c}$
- Server \rightarrow Client: $\{M2, \text{MAC}(M2, I_s)\}_{CK_s}$