


# Process sandboxing & Inline Reference Monitors

Prateek Saxena

# Second Line of Defense

- First Line of Defense
  - Directly prevent the attack from happening
- Second Line of Defense
  - Assume that attack happens, minimize the impact



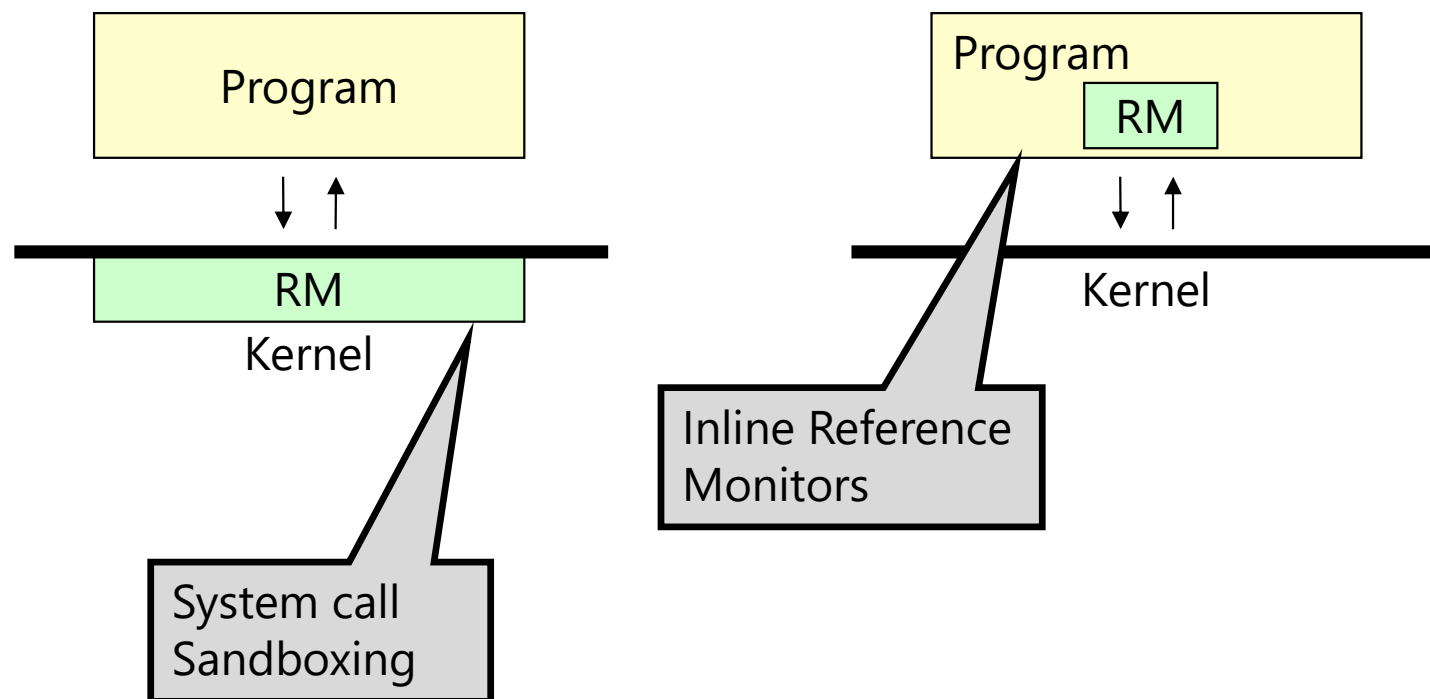
Vs.



# Reference Monitors

Reference Monitor: A piece of code that *checks all* **references** to an **object**

Syscall Sandbox: A reference monitor for protecting OS resource objects from an app



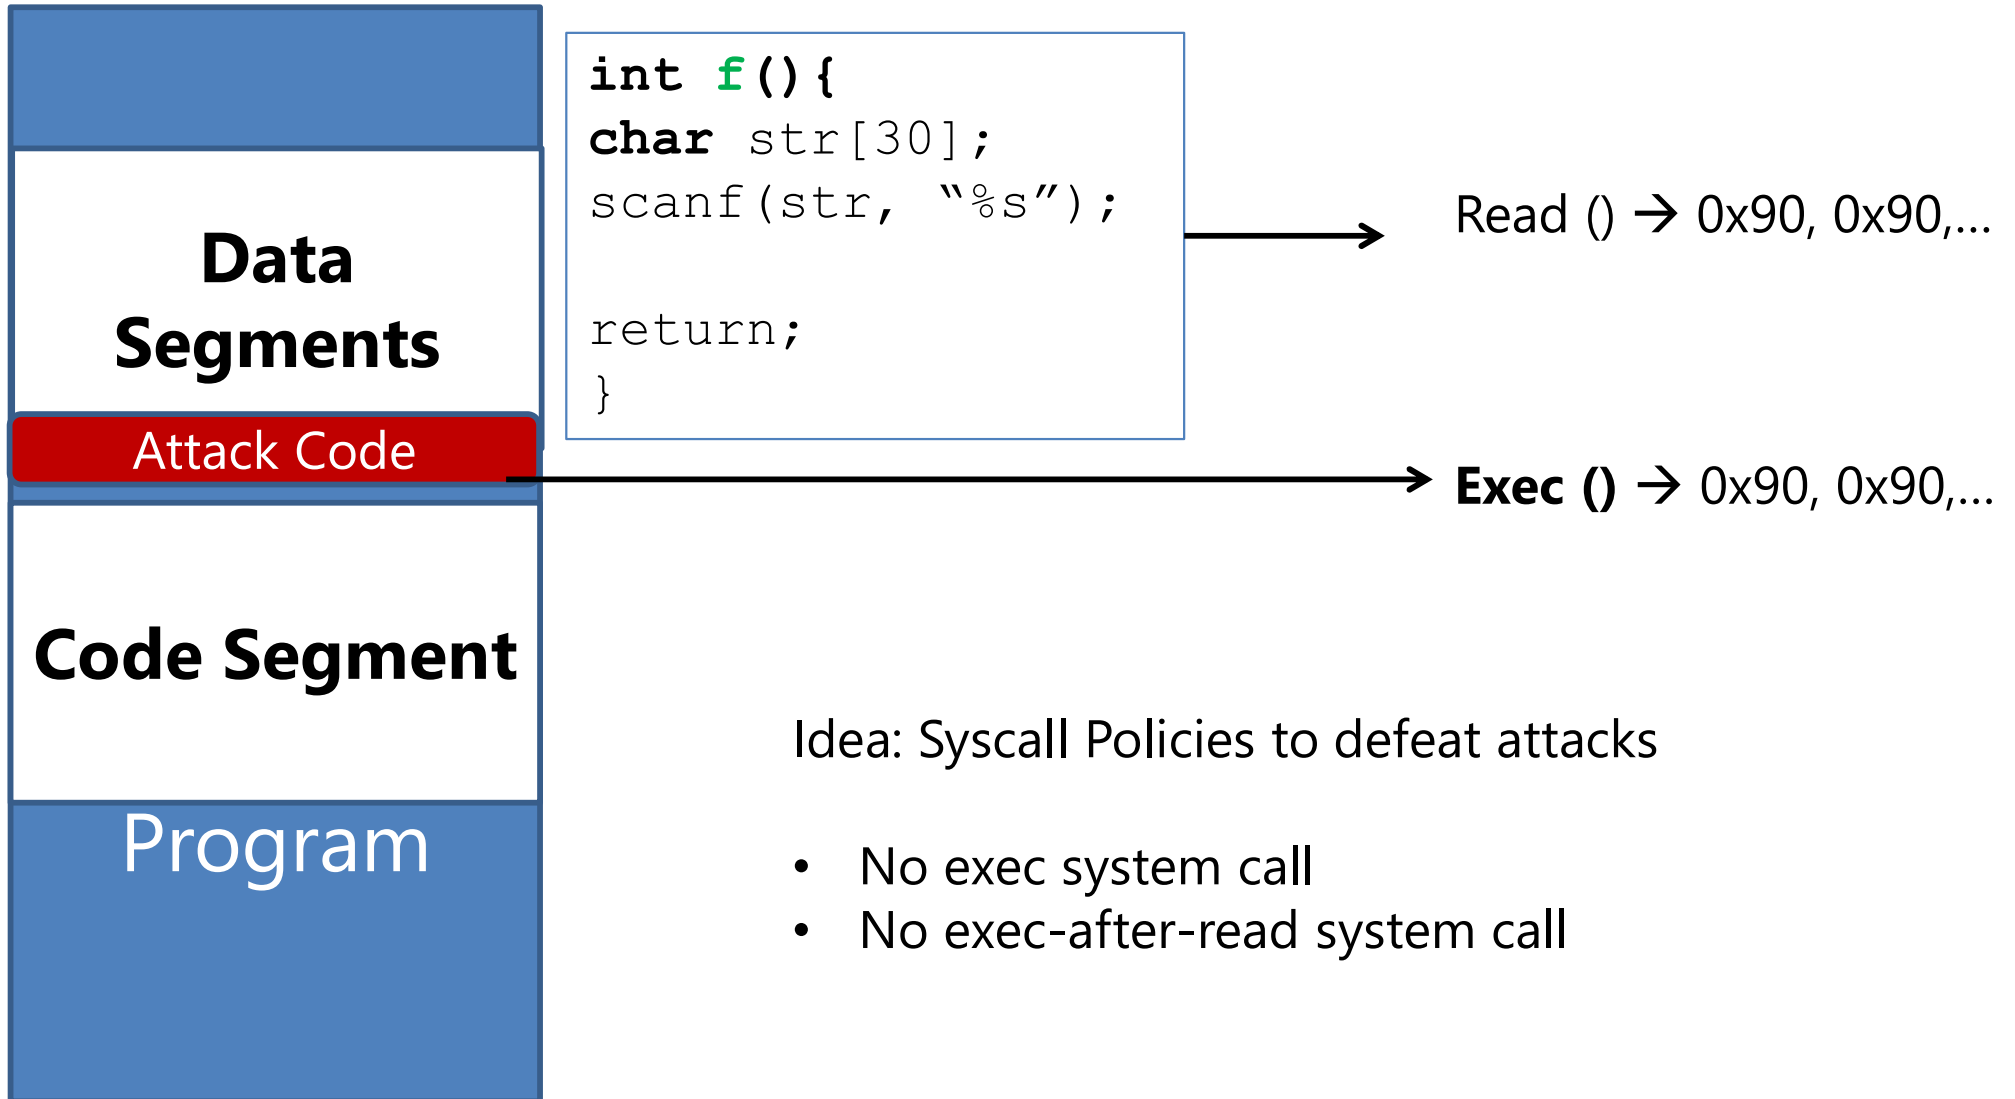
# 3 Security Principles

- Separation of Concerns:
  - Separate the **policy** from its **enforcement**
- Minimize Trusted Code Base (TCB)
  - Reduce what one needs to trust
  - Separate **verifier** from the **enforcement**
- Least Privilege
  - Give each component only the privileges necessary

# Policy vs. Enforcement Mechanism

- Access Control Policies (last lecture)
- Enforcement:
  - Process sandboxing (Today)
  - Inline Reference Monitors (Today)
  - Virtualization
  - Hardware-based isolation / Trusted Execution Env.

# Process Sandboxing

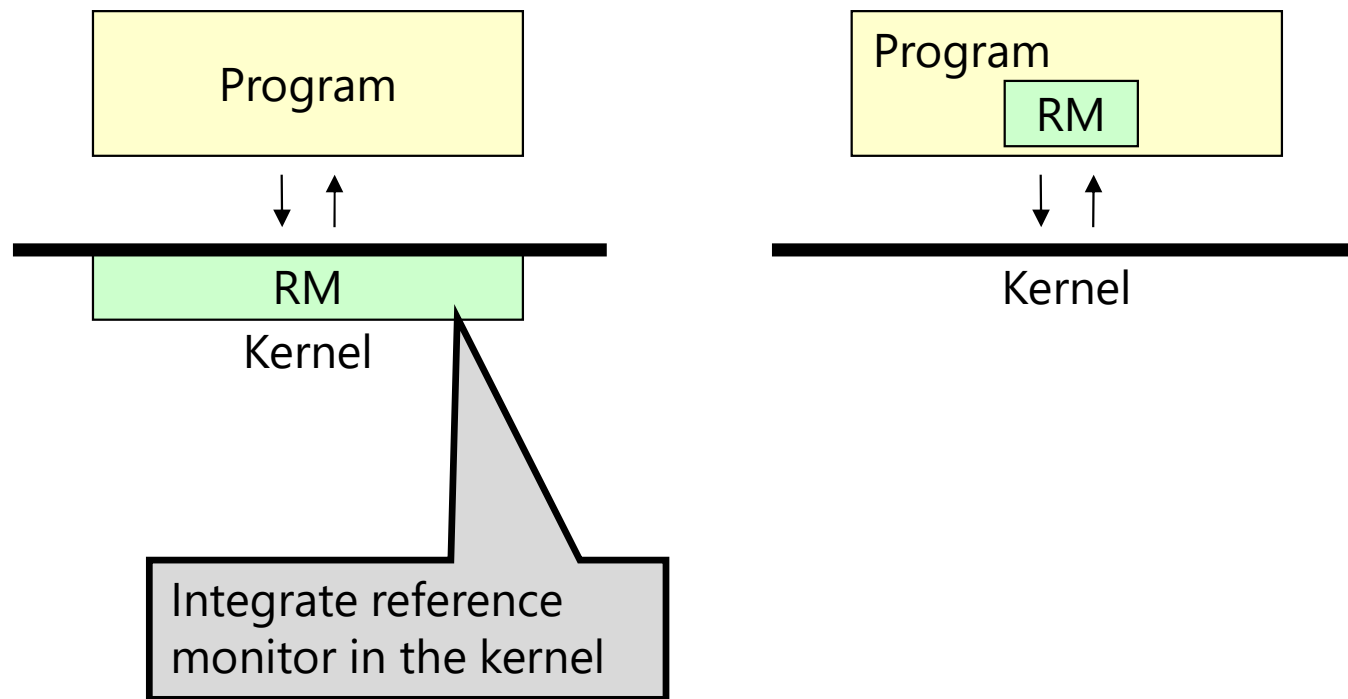


# **Enforcement Mechanisms:** Process Isolation / Sandboxing

# System Call Sandboxing

Reference Monitor: A piece of code that *checks all references* to an **object**

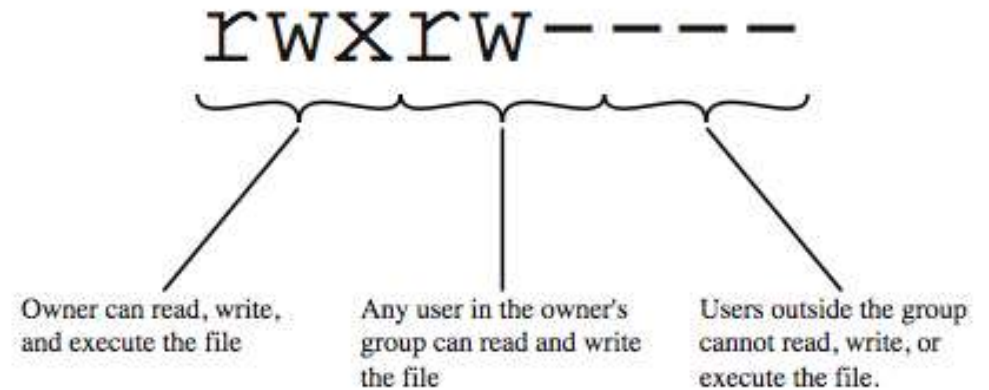
Syscall Sandbox: A reference monitor for protecting OS resource objects from an app





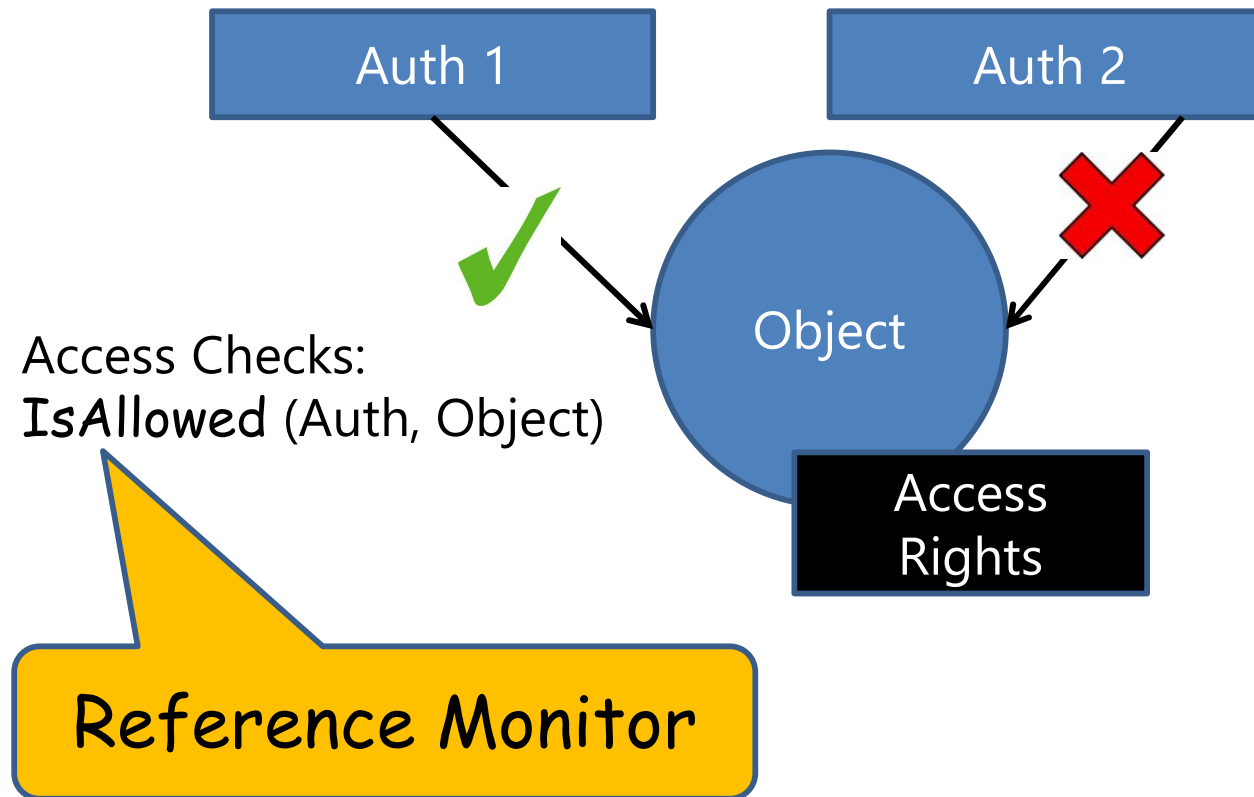
# Kernelized Syscall Sandbox (I): Access Control Lists

## Access Control List



	BIBLIOG	TEMP	F	HELP.TXT	C COMP	LINKER	SYS CLOCK	PRINTER
USER A	ORW	ORW	ORW	R	X	X	R	W
USER B	R	-	-	R	X	X	R	W
USER S	RW	-	R	R	X	X	R	W
USER T	-	-	-	R	X	X	R	W
SYS_MGR	-	-	-	RW	OX	OX	ORW	O
USER_SVCS	-	-	-	O	X	X	R	W

# Kernelized Syscall Sandbox (I): Access Control Lists




# Challenge: Ambient Authority

```
$ cp foo.txt bar.txt
```

The “cp” program has authority to write to **any** file on the system.

This is not in line with “Principle of Least Privilege”

# Kernelized Syscall Sandbox (II): Capabilities



The "cp" program has no authority, by default.  
It can only use "capabilities" it is given (e.g. UNIX file handles)

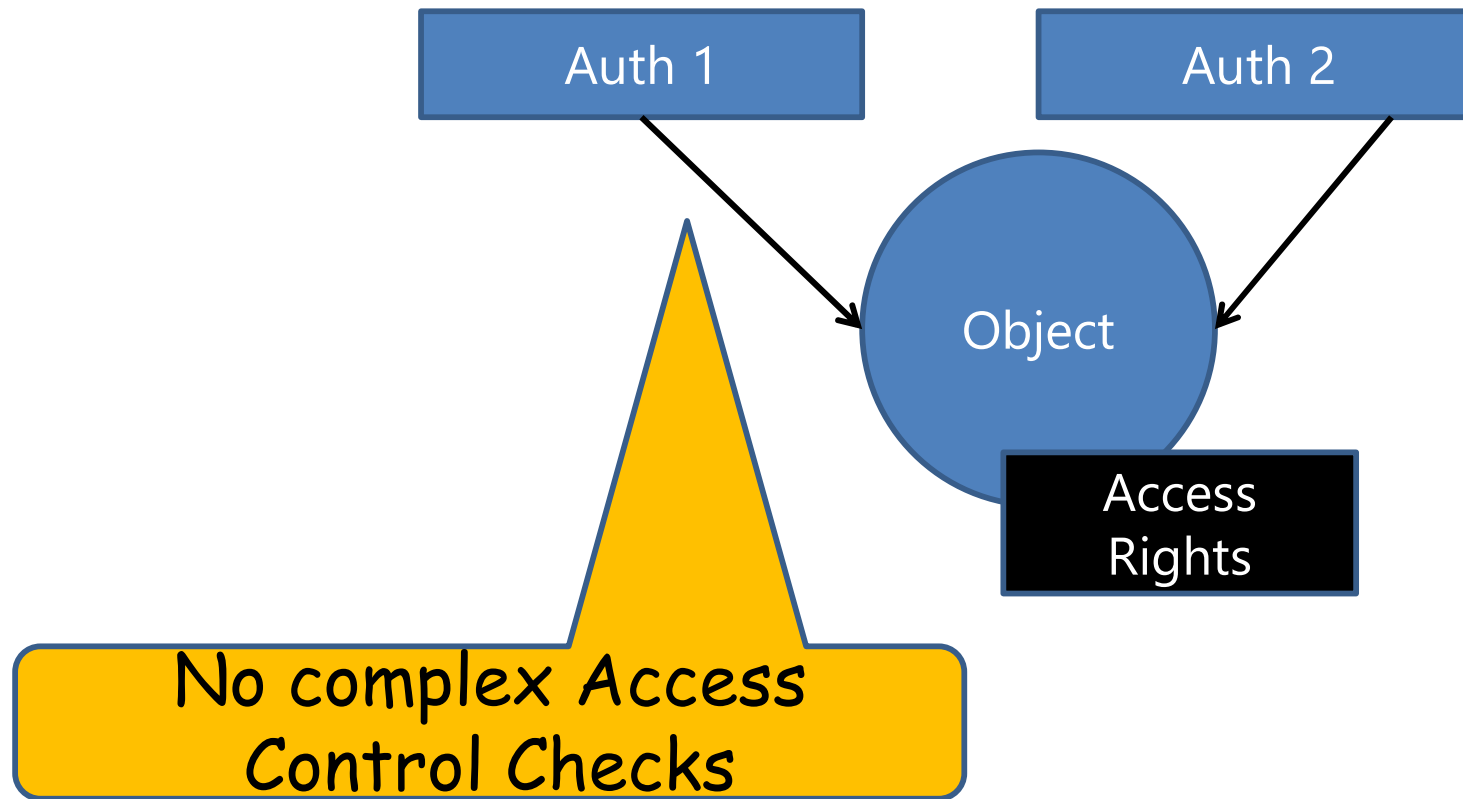
Definition of a **Capability**:

- An *identifier* which, when presented, provides certain access rights

Properties of a **Capability**:

- Unforgeable: Can't manufacture without explicitly getting it.

# Kernelized Syscall Sandbox (II): Capabilities



# Kernelized Syscall Sandbox (II): Capabilities

Access Control List

	BIBLIOG	TEMP	F	HELP.TXT	C_COMP	LINKER	SYS_CLOCK	PRINTER	
USER A	ORW	ORW	ORW	R	X	X	R	W	
USER B	R	-	-	R	X	X	R	W	
USER S	RW	-	R	R	X	X	R	W	
USER T	-	-	-	R	X	X	R	W	
SYS_MGR	-	-	-	RW	OX	OX	ORW	O	
USER_SVCS	-	-	-	O	X	X	R	W	

Capabilities

# Access Control Lists vs. Capabilities

## ACL

- Pros:
  - When the checks are simple and centralized, easier to implement ACL
  - Works well when rights change
- Cons:
  - Ambient Authority
  - Incomplete mediation:
    - Missing access control checks

## Capabilities

- Pros:
  - Eliminates access check logic
  - No pre-specification of who is allowed to access, i.e., can follow the natural flow of access rights
  - No ambient authority
    - Recall Least Privilege
- Cons:
  - Unsuitable when access rights change frequently
  - Capabilities can leak!

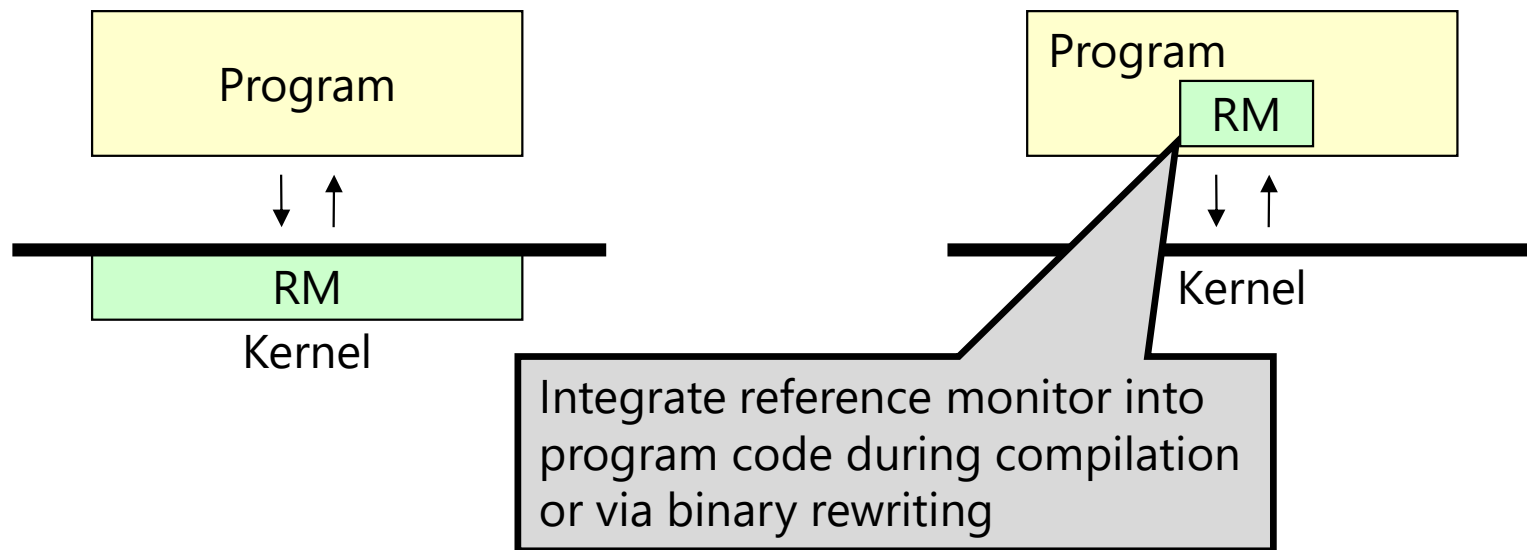
# **Inline Reference Monitors**



# Inline Reference Monitors

Reference Monitor: A piece of code that *checks all* **references** to an **object**

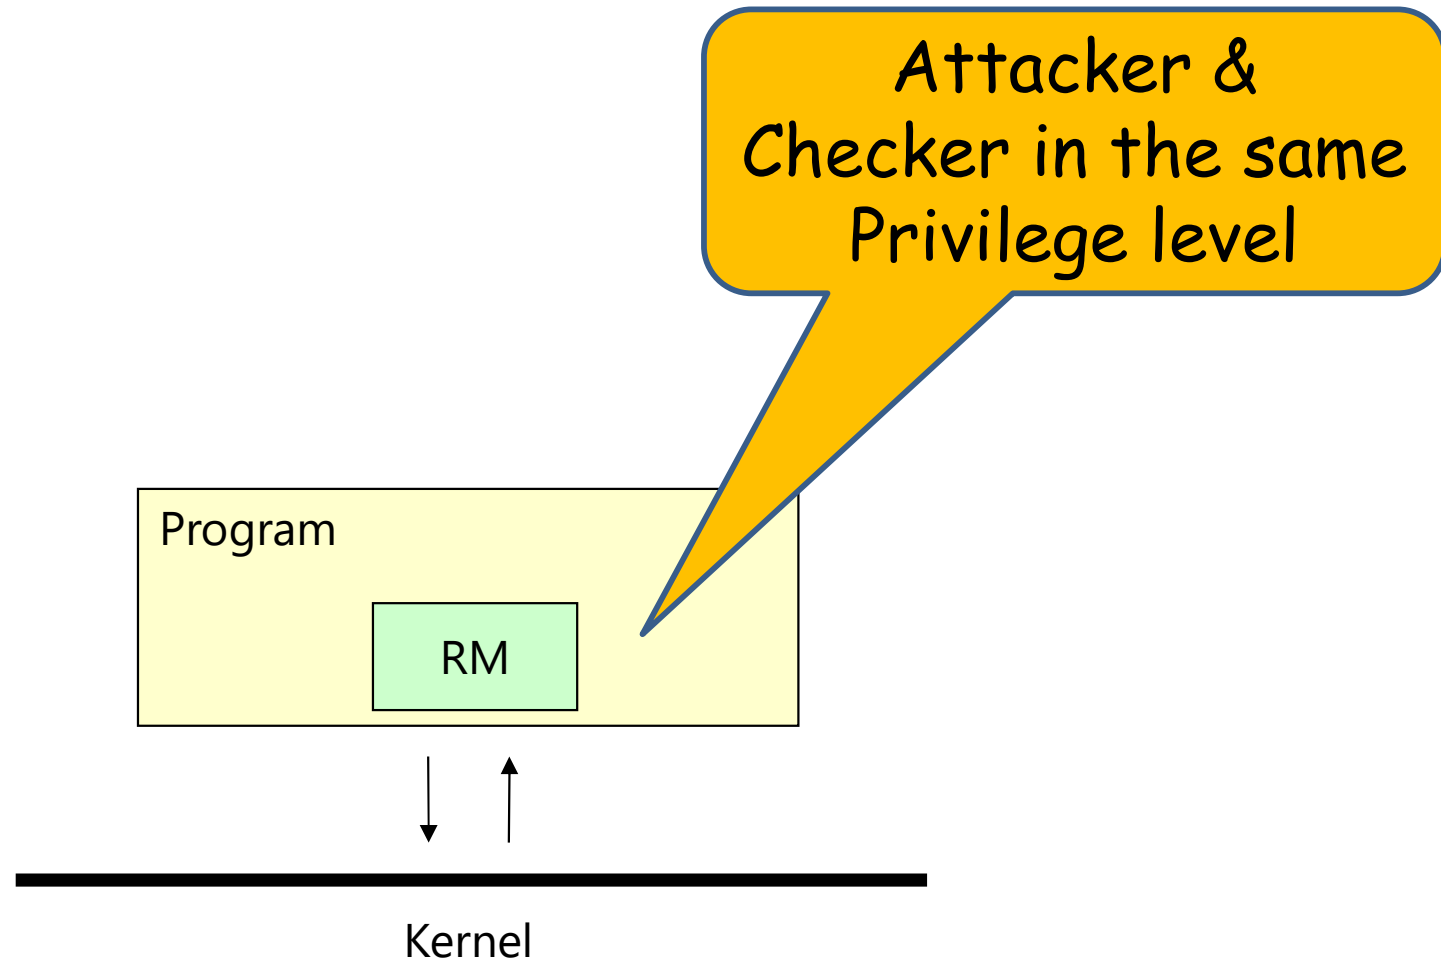
Syscall Sandbox: A reference monitor for protecting OS resource objects from an app



# Inline Reference Monitors Can Check...

- Complete Memory Safety
  - “Access memory objects in an intended way”
- Fault Isolation
  - “Each module only accesses pre-determined data / code”
- No foreign code
  - “Execute only predetermined code”
- Control Flow Integrity
  - “Control transfers are to legitimate points only”
- System Call Sandboxing
  - “Access only a subset of system calls”
- (Code) Pointers / Data Integrity
  - “Ensure (code) pointers / data have valid values”
- Data Flow Integrity...

# Challenges in Inline / Wrapper-based Enforcement

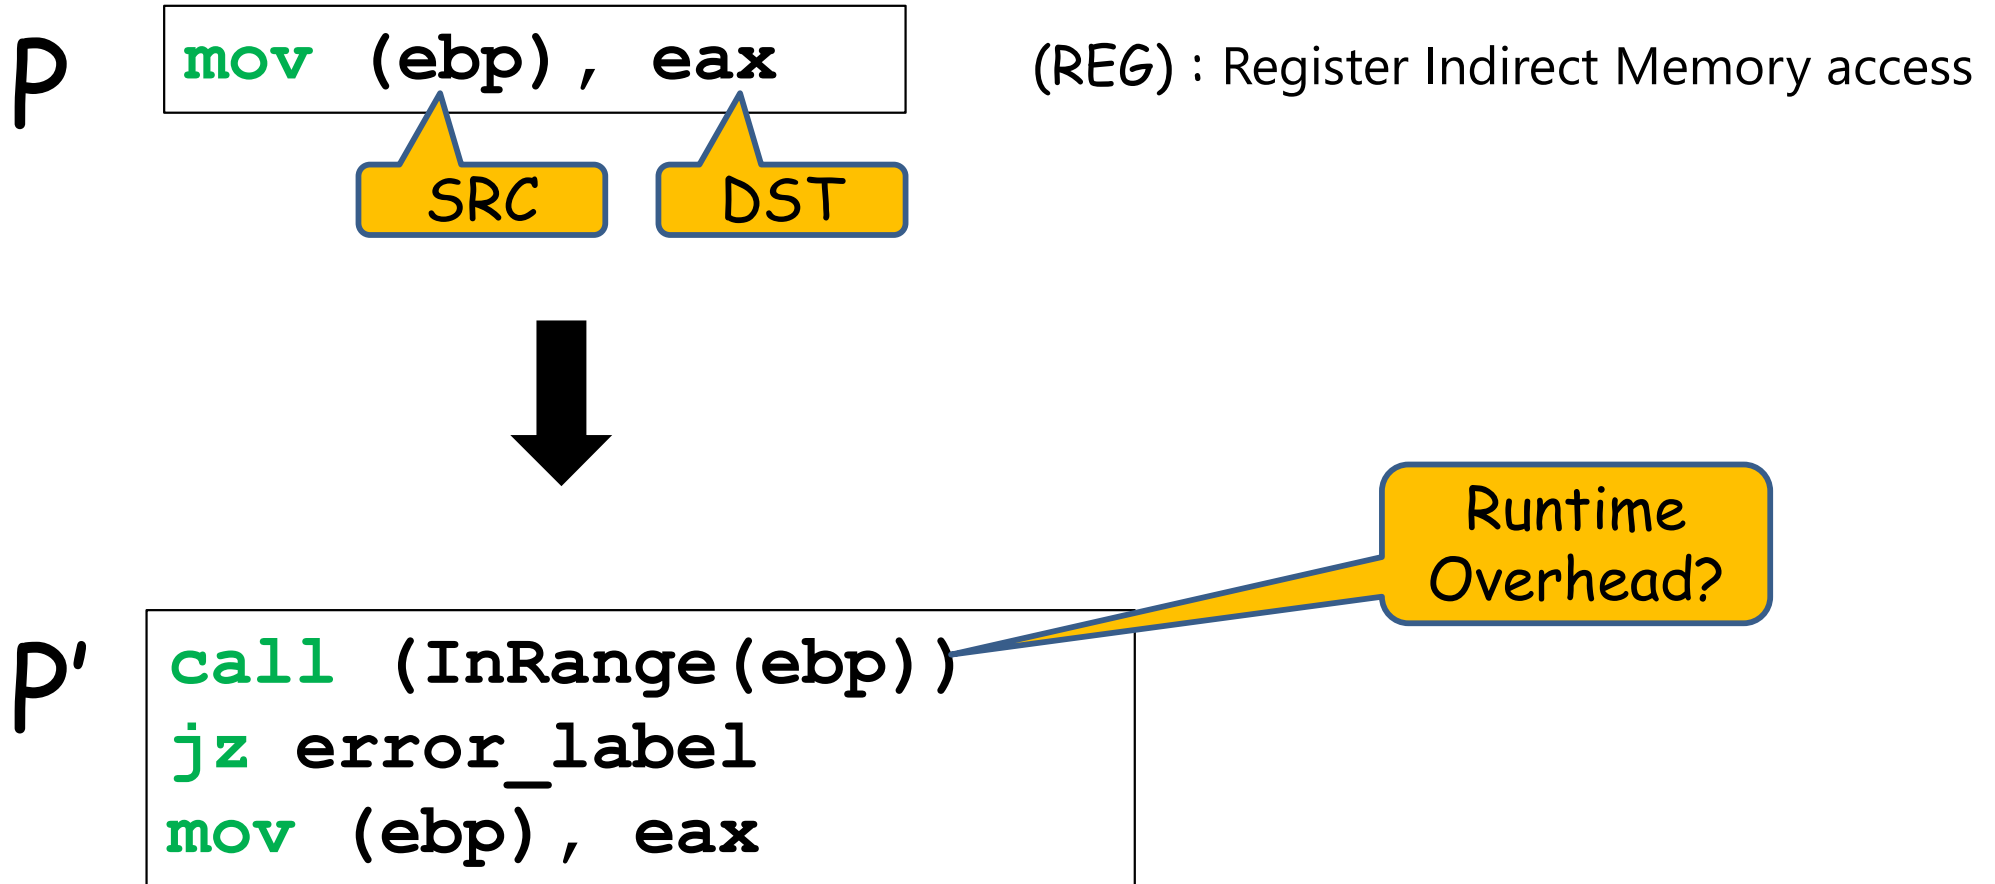


# **Inline Reference Monitors:** Software Fault Isolation

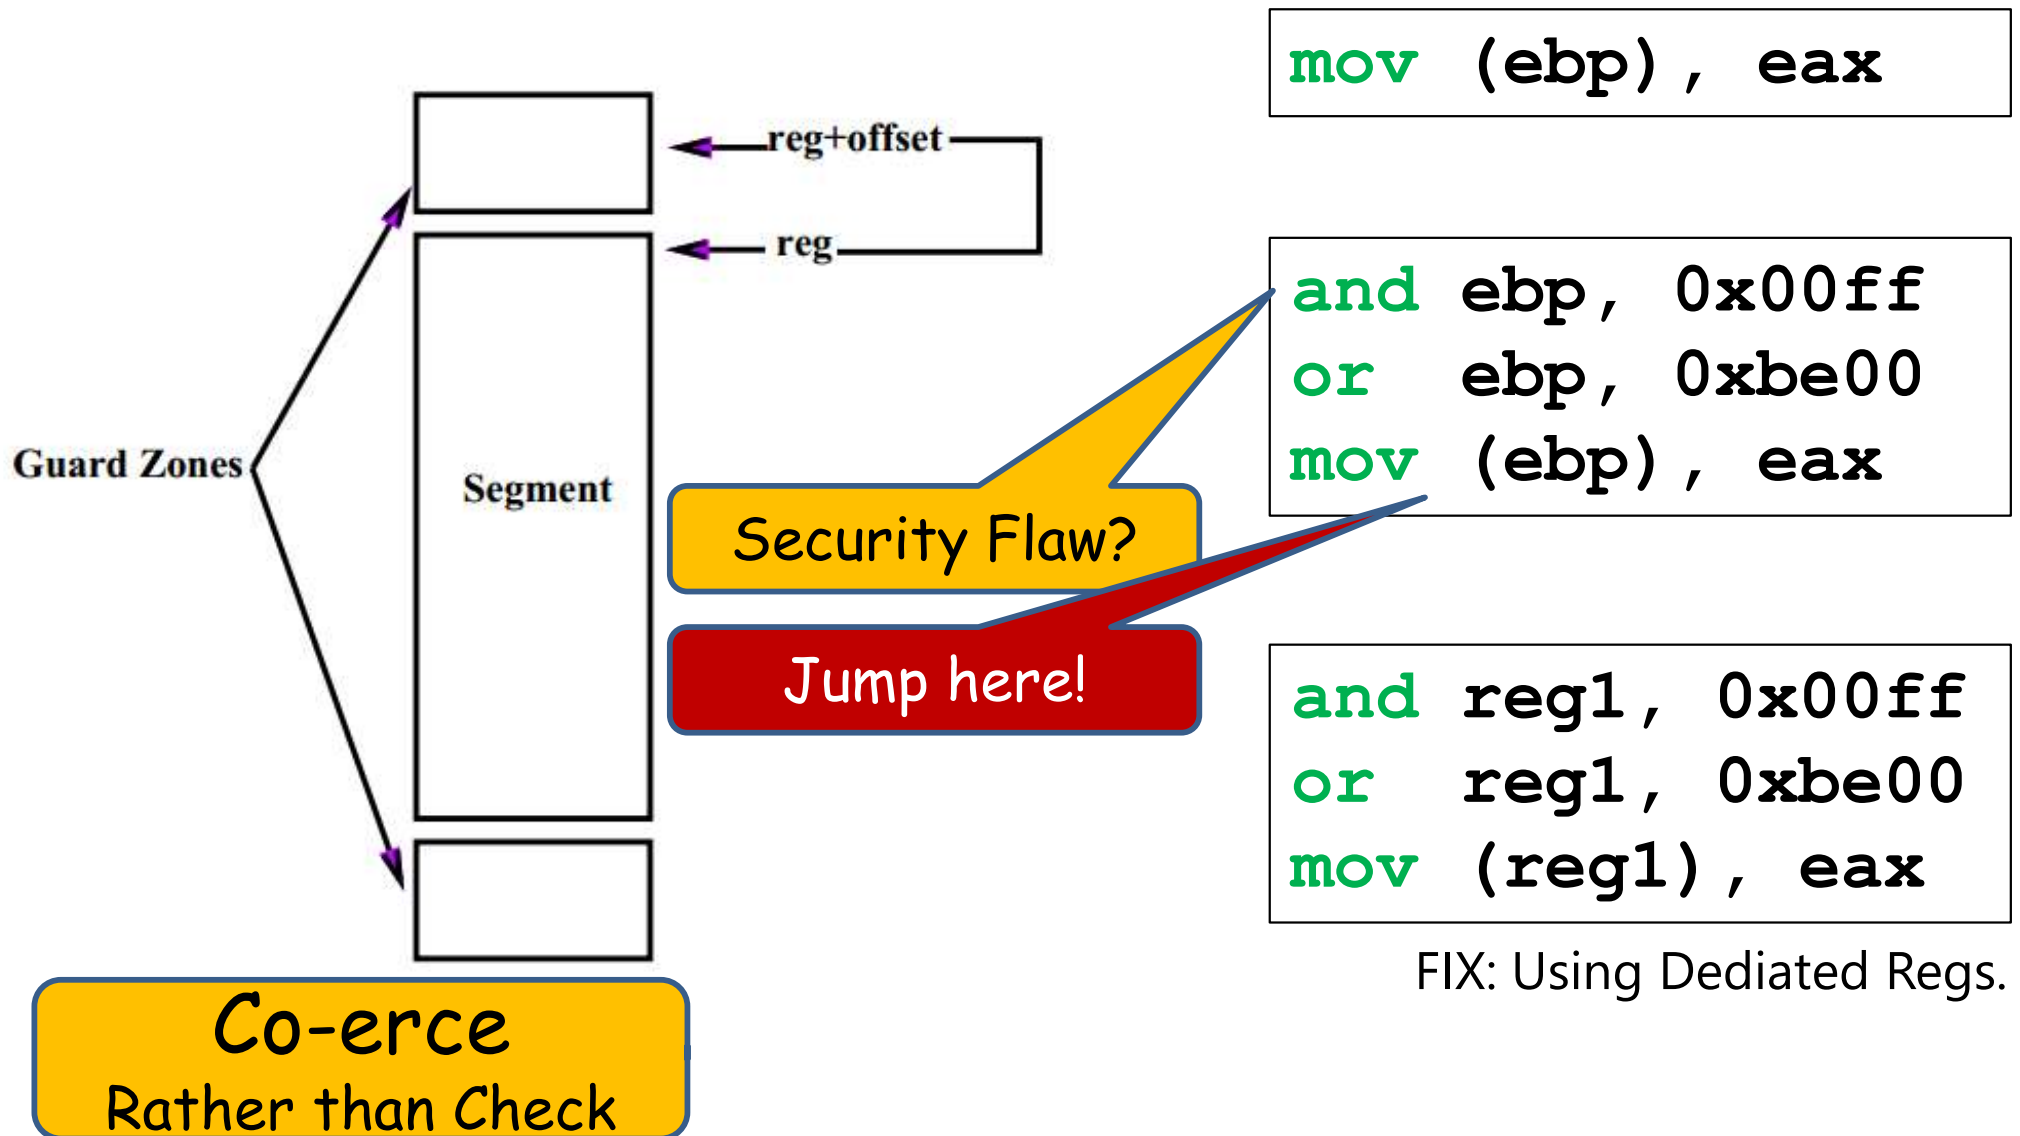
# Software Fault Isolation (SFI)

- Goal: **Fault Isolation**
  - Confine read/write to certain region M
  - This goal is also called “address sandboxing”
- Attacker controls all memory values in M
- Mechanism: Inline instrumentation of D
- Limit all memory accesses to region M
- Take an example: Let M be [0xbe00, 0xbeff]

# Naïve SFI Implementation



# Fast SFI Implementation



# Verifying Correctness of Fast-SFI

1. Check if these IRM instructions exist before memory access
2. All memory accesses use the dedicated register
3. The dedicated registers are used only in IRM instructions

```
and reg1, 0x00ff  
or  reg1, 0xbe00  
mov (reg1), eax
```



# 3 Security Principles


- Separation of Concerns:
  - Separate the **policy** from its **enforcement**

- Minimize Trusted Code Base (TCB)
  - Reduce what one needs to trust
  - Separate **verifier** from the **enforcement**

- Least Privilege
  - Give each component only the privileges necessary

# SFI Has a Small TCB...

- Goal of Software Fault Isolation:
  - Address Sandboxing
    - "Access memory segments statically verified"



- Trusted Computing Base (TCB):
  - "The trusted codebase for ensuring security properties"
- Smaller the TCB, the better the design

# **Aiding Syscall Sandboxing: Privilege Separation**


# Takeaways: 3 Security Principles

- Separation of Concerns:
  - Separate the **policy** from its **enforcement**
- Minimize Trusted Code Base (TCB)
  - Reduce what one needs to trust
  - Separate **verifier** from the **enforcement**
- Least Privilege
  - Give each component only the privileges necessary


# Problem: Bundling of Functionality



# Problem: Bundling of Functionality



# Solution: Privilege Separation



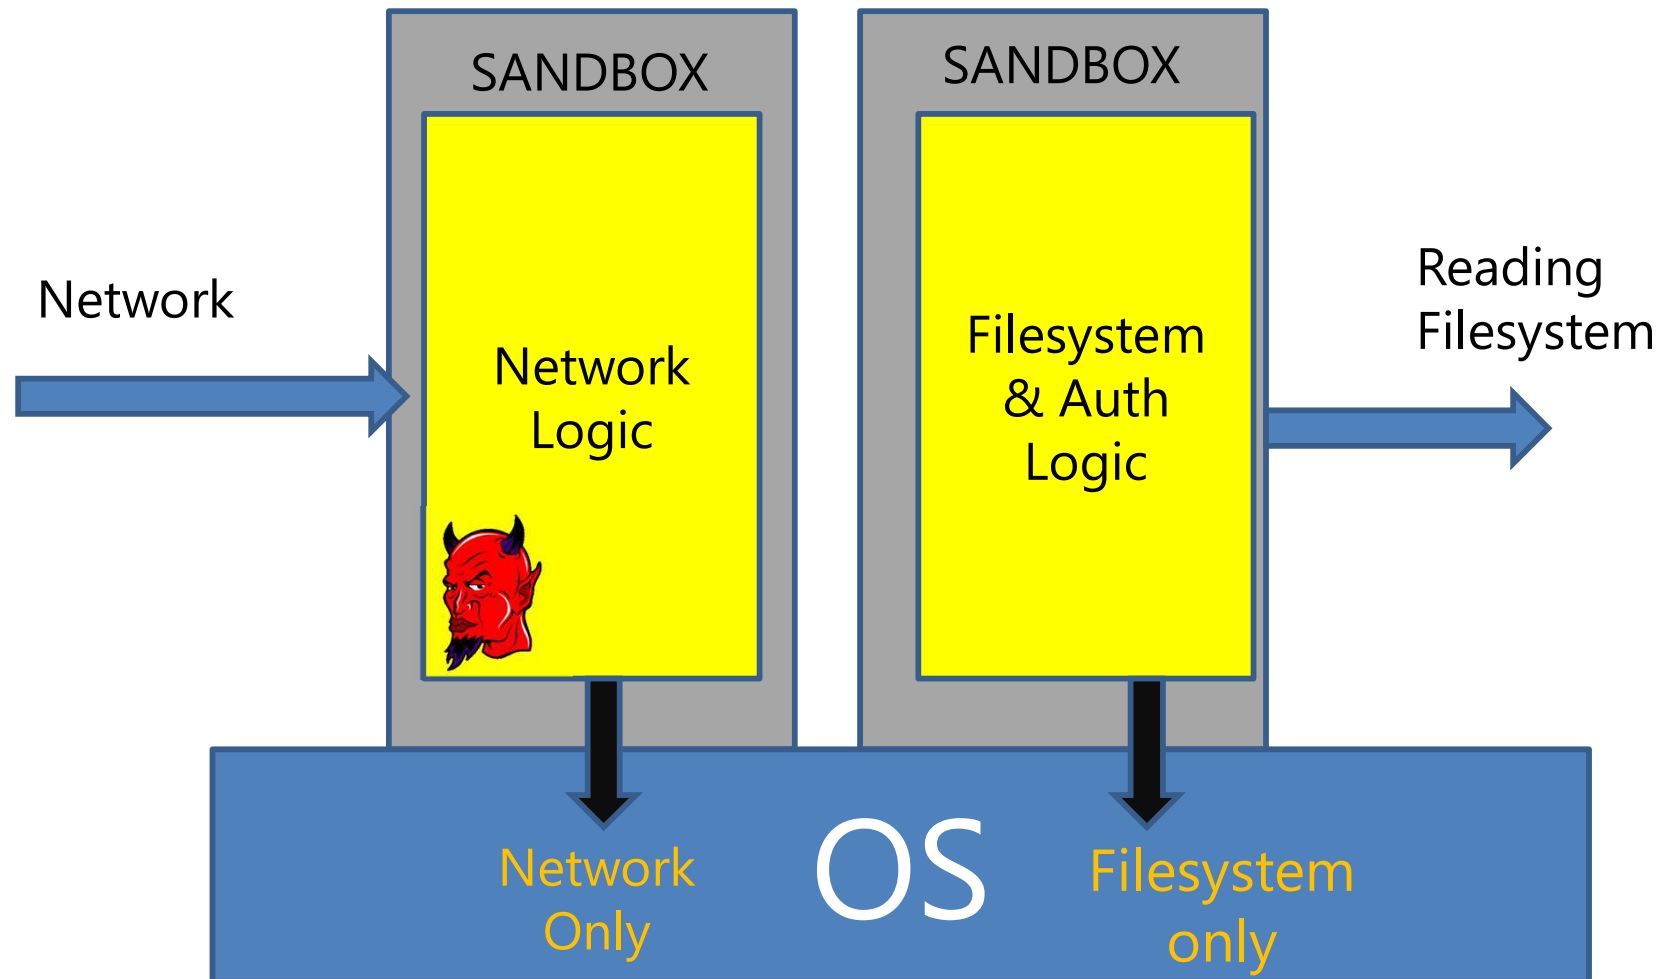


Courtesy: John Mitchell



# Principle of Least Privilege

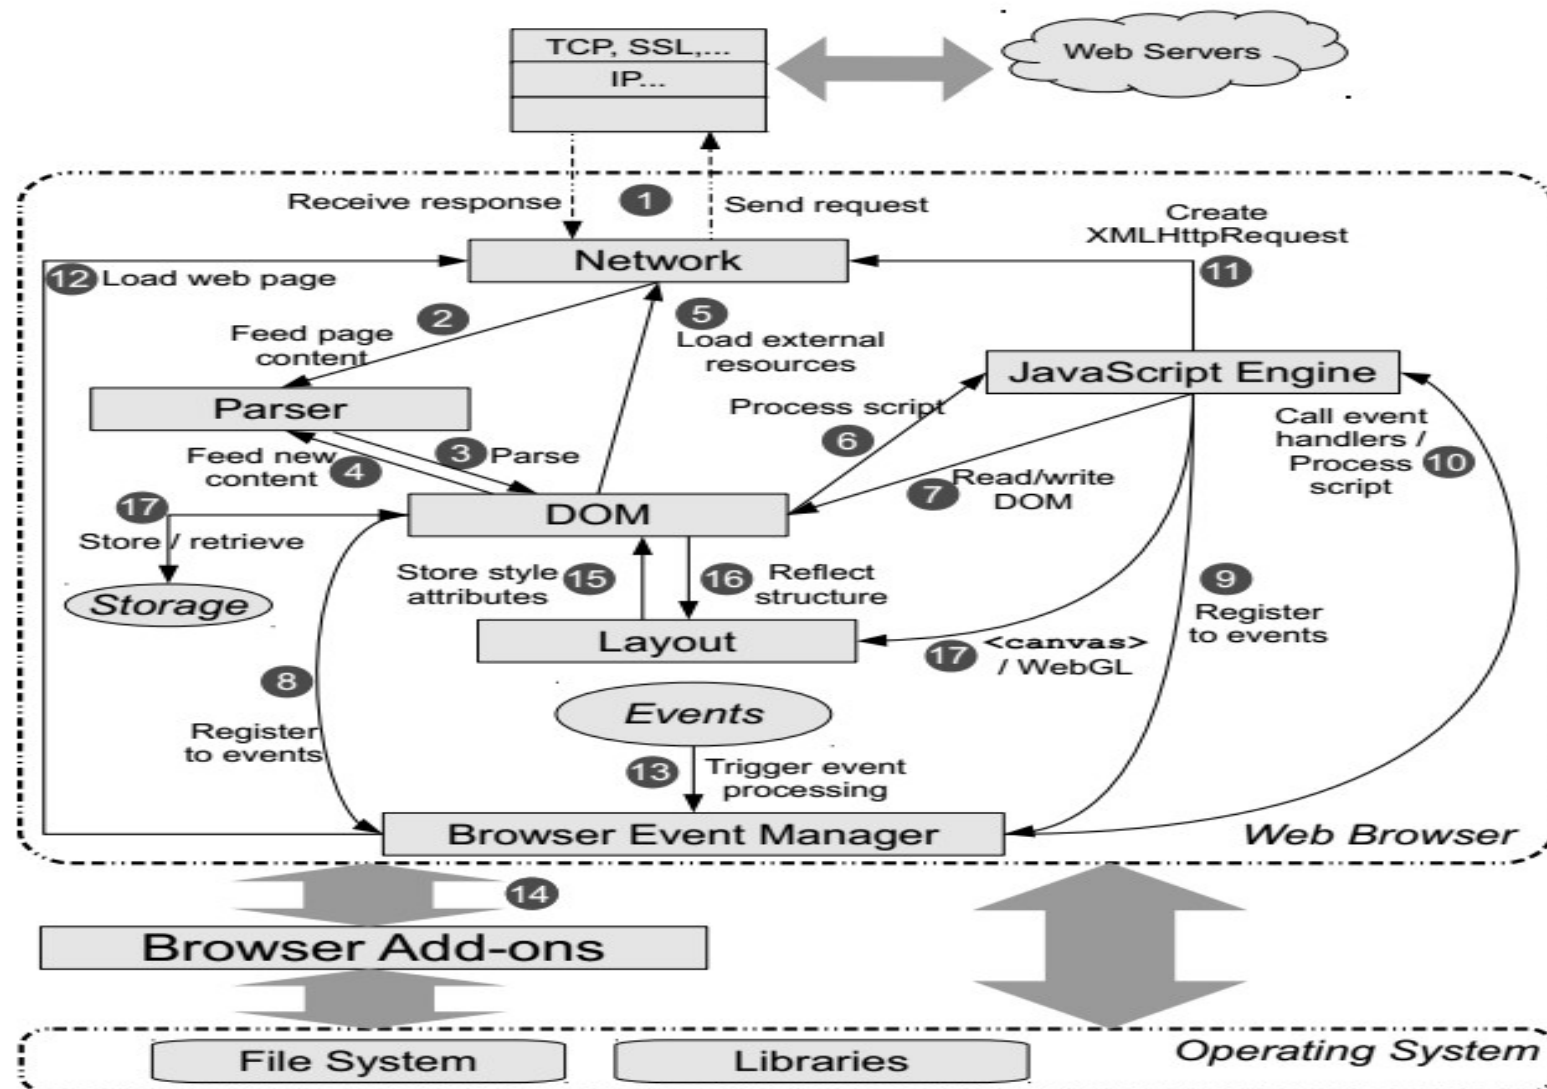
- Each compartment gets the least set of privileges it needs for its function



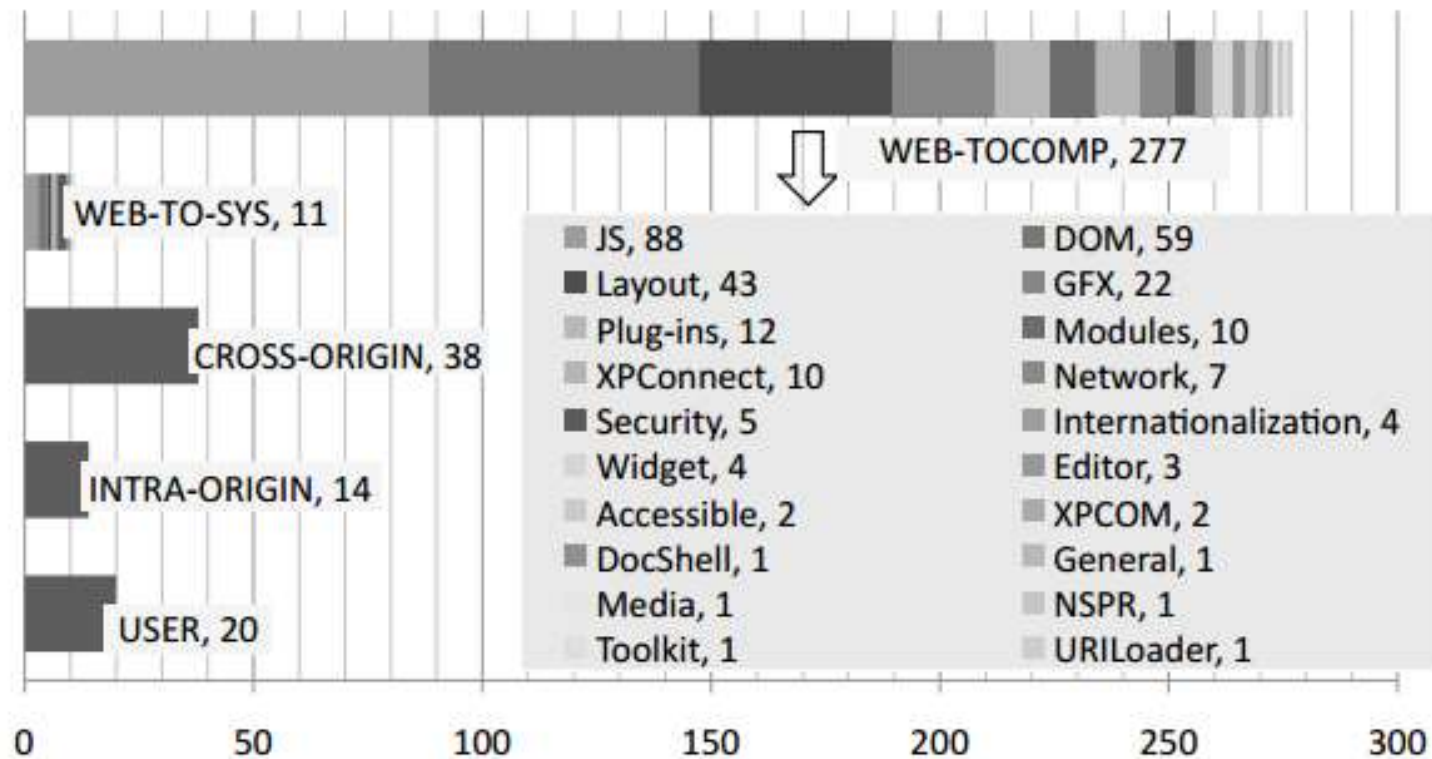
# **Privilege Separation Case Study (I): Design of the Web Browsers**

# Security Issues for Implementation

- A huge codebase (2 MLOC)
- Many languages (JS, CSS, HTML, URL, ...)



# Distribution of Browser Implementation Bugs




(a) Number of Historical Security Vulnerabilities in Firefox, Categorized by Severity and Firefox Components

# So, what should we do?

- Auto-patching
  - E.g. Google Chrome
- Consider Firefox: Single- process
  - 1 Vulnerability leads to accessing all origins
- Solution: Privilege Separation
  - Compartmentalize & assign least privilege
- Google Chrome
  - Goal: Separate Filesystem from web code

# Google Chrome Design

- Goal: Prevent web & network attacker from compromising OS resources (e.g. Filesystem)



Rendering Engine	Browser Kernel
HTML parsing	Cookie database
CSS parsing	History database
Image decoding	Password database
JavaScript interpreter	Window management
Regular expressions	Location bar
Layout	Safe Browsing blacklist
Document Object Model	Network stack
Rendering	SSL/TLS
SVG	Disk cache
XML parsing	Download manager
XSLT	Clipboard
<b>Both</b>	
URL parsing	
Unicode parsing	

# 3 Security Principles

- Separation of Concerns:
  - Separate the **policy** from its **enforcement**
- Minimize Trusted Code Base (TCB)
  - Reduce what one needs to trust
  - Separate **verifier** from the **enforcement**
- Least Privilege
  - Give each component only the privileges necessary

# Wrap Up: The Key Takeaway

- Concept: A Threat Model defines:
  - Desired Security Property / Goal
  - Attacker Capabilities
  - Assumptions about the setup
- **This Module at a glance:**
  - Stack of Threats
  - Real attacks
  - Defenses
  - How to argue security!





# Wrap Up...

- **The Attacker Perspective: A New Viewpoint**
- We aren't teaching you an artefact (e.g. OS, DB)
  - We're teaching you how to build it securely!
- As much about "new knowledge", as about
  - **Making sound security arguments (with threat models)**
  - **Principles and abstractions** of "secure" construction
- You learnt:
  - Basics of Crypto, web, OS, systems security
  - With 2 hands-on coding assignments!
  - How theory meets practice in computer security

**Thanks!**  
**See you next semester...**

**(Good luck!)**