

Observing Function Call and Return using GDB

The goal of this group assignment is to get familiar with the GDB debugger, and use it to understand the low-level function call and return mechanism used by Intel CPUs.

Lab setup

You will work on your Ubuntu machine or VM.

Practice steps

1. To mitigate memory exploits, including buffer overflow, Ubuntu has its address-space randomization turned on by default. We need to turn it off for easily observing the low-level mechanisms for call and return. Use the following command to disable address-space randomization:

```
sudo sysctl -w kernel.randomize_va_space=0
```

Note: also look in `/proc/sys/kernel`

GDB by default disables randomization. To turn it off:

```
set disable-randomization off
```

2. Download the sample program from the following address:

<http://www.comp.nus.edu.sg/~liangzk/cs4238/overflowsample.tar.gz>

Find the downloaded file in Downloads directory in your home folder, and extract the file `sample.c` from the compressed file.

3. Compile the source file `sample.c` with stack-protector disabled (`-fno-stack-protector`), debugging information (`-g`), and generate an executable file named `sample` (`-o sample`):

```
gcc -fno-stack-protector -g -o sample sample.c
```

Ignore any warnings.

4. Start the GDB debugger:

```
gdb ./sample
```

5. Set a breakpoint at the beginning of the `main()` function:

(Under the gdb prompt) `break main`

```
(gdb) break main
Breakpoint 1 at 0x11ec: file sample.c, line 18.
(gdb) █
```

Note: The command can be shortened to “`b main`”.

6. Now we can start to execute the program:

(Under the gdb prompt) `run ./sample`

Or simply `run`

```
(gdb) run
Starting program: /home/kaihang/CS4238/overflowsample/sample

Breakpoint 1, main () at sample.c:18
18      {
(gdb) █
```

Now the program stops at `main()`.

7. Before we continue the program execution, disassemble the main function to note down an important value from the program:

```
(gdb) disassemble main
Dump of assembler code for function main:
=> 0x0000555555551ec <+0>:    endbr64
    0x0000555555551f0 <+4>:    push   %rbp
    0x0000555555551f1 <+5>:    mov    %rsp,%rbp
    0x0000555555551f4 <+8>:    sub    $0x10,%rsp
    0x0000555555551f8 <+12>:   lea     -0x4(%rbp),%rax
    0x0000555555551fc <+16>:   mov    %rax,%rsi
    0x0000555555551ff <+19>:   lea     0xeca(%rip),%rdi        # 0x5555555560d0
    0x000055555555206 <+26>:   mov    $0x0,%eax
    0x00005555555520b <+31>:   callq  0x55555555060 <printf@plt>
    0x000055555555210 <+36>:   mov    $0x0,%eax
    0x000055555555215 <+41>:   callq  0x55555555169 <sample_function>
    0x00005555555521a <+46>:   mov    $0x0,%eax
    0x00005555555521f <+51>:   leaveq
    0x000055555555220 <+52>:   retq
End of assembler dump.
(gdb) █
```

(Under the gdb prompt) `disassemble main`

If you are more comfortable with intel style of disassembly, you can change it via command `set disassembly-flavor intel`

```
(gdb) set disassembly-flavor intel
(gdb) disassemble main
Dump of assembler code for function main:
=> 0x0000555555551ec <+0>:      endbr64
    0x0000555555551f0 <+4>:      push   rbp
    0x0000555555551f1 <+5>:      mov    rbp, rsp
    0x0000555555551f4 <+8>:      sub    rsp, 0x10
    0x0000555555551f8 <+12>:     lea    rax, [rbp-0x4]
    0x0000555555551fc <+16>:     mov    rsi, rax
    0x0000555555551ff <+19>:     lea    rdi, [rip+0xec2]      # 0x5555555560c8
    0x000055555555206 <+26>:     mov    eax, 0x0
    0x00005555555520b <+31>:     call   0x55555555060 <printf@plt>
    0x000055555555210 <+36>:     mov    eax, 0x0
    0x000055555555215 <+41>:     call   0x55555555169 <sample_function>
    0x00005555555521a <+46>:     mov    eax, 0x0
    0x00005555555521f <+51>:     leave
    0x000055555555220 <+52>:     ret
End of assembler dump.
```

This is the assembly code of the `main()` function. Each instruction line starts with the memory address of that instruction, followed by the disassembled instruction. Note that the instruction at the address `0x 55555555215` (the instruction above the red line) is the call to `sample_function()`. Therefore, when the function returns, it should continue to execute the next instruction, whose address is `0x5555555521a`. Note down this address for a comparison later.

Note: the listings above are illustrative. It is possible for the exact details to vary with different compiler flags, compiler versions, and other compilers.

You can also use `gcc` with `-S` to only compile `sample.c` into the assembly code in `sample.s`.

8. Do “next” for twice to step over executing the `printf()` functions. From the output, you can see the memory address of the variable `x`.

(Under the gdb prompt) `next`

```
(gdb) next
21          printf("In main(), x is stored at %p.\n", &x);
(gdb) next
In main(), x is stored at 0x7fffffffdebc.
22          sample_function();
```

Now the program is about to call the function `sample_function()`.

9. Let's inspect the register values:

(Under the gdb prompt) `info registers`

```
(gdb) info register
rax          0x2a          42
rbx          0x55555555230   93824992236080
rcx          0x0           0
rdx          0x0           0
rsi          0x5555555592a0  93824992252576
rdi          0x7fffffff7fac4c0 140737353794752
rbp          0x7fffffffdec0  0x7fffffffdec0
rsp          0x7fffffffdeb0  0x7fffffffdeb0
r8           0x0           0
r9           0x2a          42
r10          0x5555555560e4  93824992239844
r11          0x246         582
r12          0x555555555080  93824992235648
r13          0x7fffffffdfb0  140737488347056
r14          0x0           0
r15          0x0           0
rip          0x555555555210  0x555555555210 <main+36>
eflags      0x202         [ IF ]
cs          0x33          51
ss          0x2b          43
ds          0x0           0
es          0x0           0
fs          0x0           0
gs          0x0           0
(gdb) |
```

This command shows the value of integer registers and the decoded value. Here we just need to use the first number (hexadecimal value of the register).

We can see that the stack pointer RSP is at `0x7fffffffdeb0`. The frame (base) pointer RBP is at `0x7fffffffdec0`. The instruction pointer RIP is at

0x55555555210. Can you check, from the disassembly of `main()`, which instruction will be executed next?

10. Before we enter the `sample_function()`, do a disassemble of the sample function.

```
(gdb) disassemble sample_function
Dump of assembler code for function sample_function:
0x000055555555169 <+0>:    endbr64
0x00005555555516d <+4>:    push    rbp
0x00005555555516e <+5>:    mov     rbp, rsp
0x000055555555171 <+8>:    sub     rsp, 0x10
0x000055555555175 <+12>:   mov     DWORD PTR [rbp-0x4], 0x0
0x00005555555517c <+19>:   lea     rax, [rbp-0x4]
0x000055555555180 <+23>:   mov     rsi, rax
0x000055555555183 <+26>:   lea     rdi, [rip+0xe7e]          # 0x555555556008
0x00005555555518a <+33>:   mov     eax, 0x0
0x00005555555518f <+38>:   call    0x55555555060 <printf@plt>
0x000055555555194 <+43>:   lea     rax, [rbp-0xe]
0x000055555555198 <+47>:   mov     rsi, rax
0x00005555555519b <+50>:   lea     rdi, [rip+0xe96]          # 0x555555556038
0x0000555555551a2 <+57>:   mov     eax, 0x0
0x0000555555551a7 <+62>:   call    0x55555555060 <printf@plt>
0x0000555555551ac <+67>:   mov     eax, DWORD PTR [rbp-0x4]
0x0000555555551af <+70>:   mov     esi, eax
0x0000555555551b1 <+72>:   lea     rdi, [rip+0xeb0]          # 0x555555556068
0x0000555555551b8 <+79>:   mov     eax, 0x0
0x0000555555551bd <+84>:   call    0x55555555060 <printf@plt>
0x0000555555551c2 <+89>:   lea     rax, [rbp-0xe]
0x0000555555551c6 <+93>:   mov     rdi, rax
0x0000555555551c9 <+96>:   mov     eax, 0x0
0x0000555555551ce <+101>:  call    0x55555555070 <gets@plt>
0x0000555555551d3 <+106>:  mov     eax, DWORD PTR [rbp-0x4]
0x0000555555551d6 <+109>:  mov     esi, eax
0x0000555555551d8 <+111>:  lea     rdi, [rip+0xeb9]          # 0x555555556098
0x0000555555551df <+118>:  mov     eax, 0x0
0x0000555555551e4 <+123>:  call    0x55555555060 <printf@plt>
0x0000555555551e9 <+128>:  nop
0x0000555555551ea <+129>:  leave
0x0000555555551eb <+130>:  ret
End of assembler dump.
(gdb) |
```

The first instruction is an indirect branch terminating instruction. You may think of it as a `nop` instruction which does nothing.

The following three instructions of this function is common across most of the functions generated by the gcc compiler without optimization. It saves the base pointer on the stack (`push rbp`), point the base pointer to the current stack top (`mov rbp, rsp`), and move down the stack pointer to

allocate space for local variables (`sub rsp, 0x10`). The rest of the instructions is generated from the C code of `sample_function()`.

11. Let's see what will happen to the stack when the program enters `sample_function()`. The stack pointer is originally at `0x7fffffffdeb0`, shown in the previous “`info registers`” command.

Turn on disassemble by “`set disassemble-next-line on`” and use “`stepi`” to step into `sample_function`.

```
(gdb) set disassemble-next-line on
(gdb) stepi
0x000055555555215      22      sample_function();
0x000055555555210 <main+36>:    b8 00 00 00 00  mov    eax,0x0
=> 0x000055555555215 <main+41>:    e8 4f ff ff ff  call   0x55555555169 <sample_function>
0x00005555555521a <main+46>:    b8 00 00 00 00  mov    eax,0x0
(gdb) si
sample_function () at sample.c:4
4      {
=> 0x000055555555169 <sample_function+0>:    f3 0f 1e fa      endbr64
0x00005555555516d <sample_function+4>:    55              push   rbp
0x00005555555516e <sample_function+5>:    48 89 e5         mov    rbp, rsp
0x000055555555171 <sample_function+8>:    48 83 ec 10      sub    rsp, 0x10
(gdb)
```

First, a return address will be pushed on the stack by the `call` instruction. A return address is 8 bytes on a 64-bit computer. Therefore, the stack pointer should be at `0x7fffffffdeb0 - 0x8 = 0x7fffffffdea8`.

We can check by “`info registers`” command.

```
(gdb) info registers
rax          0x0          0
rbx          0x55555555230  93824992236080
rcx          0x0          0
rdx          0x0          0
rsi          0x5555555592a0  93824992252576
rdi          0x7fffffff7fac4c0  140737353794752
rbp          0x7fffffffdec0  0x7fffffffdec0
rsp          0x7fffffffdea8  0x7fffffffdea8
r8           0x0          0
r9           0x2a         42
r10          0x5555555560e4  93824992239844
r11          0x246        582
r12          0x555555555080  93824992235648
r13          0x7fffffffdfb0  140737488347056
r14          0x0          0
r15          0x0          0
rip          0x55555555169  0x55555555169 <sample_function>
eflags       0x202        [ IF ]
cs           0x33         51
ss           0x2b         43
ds           0x0          0
es           0x0          0
fs           0x0          0
gs           0x0          0
(gdb) |
```

And check the saved return address on the stack, and it should be the value we recorded in STEP 7:

```
(gdb) x/xg $rsp
0x7fffffffdea8: 0x000055555555521a
(gdb) |
```

(Please note the endianness)

12. Next, the `push rbp` instruction will push an 8-byte RBP on to the stack.

The stack pointer will be moved down by 8, resulting in a new value $0x7fffffffdea8 - 0x8 = 0x7fffffffdea0$.

Then, the `mov rbp, rsp` instruction will set RBP to the value of RSP, $0x7fffffffdea0$.

Finally, the stack pointer is moved down by $0x10$ to make space for local variables. The new stack pointer ESP is $0x7fffffffdea0 - 0x10 =$

$0x7fffffffde90$. Therefore, the local variables of `sample_function()`

should be in the range of 0x7fffffffde90 to 0x7fffffffdea0.

We can check whether the registers match our analysis or not by stepping over the `sub rsp, 0x10` instruction and print the registers.

```
(gdb) ni
0x000055555555171    4    {
    0x000055555555169 <sample_function+0>:    f3 0f 1e fa    endbr64
    0x00005555555516d <sample_function+4>:    55          push    rbp
    0x00005555555516e <sample_function+5>:    48 89 e5      mov     rbp, rsp
=> 0x000055555555171 <sample_function+8>:    48 83 ec 10    sub     rsp, 0x10
(gdb) ni
5                int i = 0;
=> 0x000055555555175 <sample_function+12>:    c7 45 fc 00 00 00 00    mov     DWORD PTR [rbp-0x4], 0x0
(gdb) info registers
rax                0x0
rbx                0x55555555230    93824992236080
rcx                0x0
rdx                0x0
rsi                0x5555555592a0    93824992252576
rdi                0x7fffffff7fac4c0    140737353794752
rbp                0x7fffffffdea0    0x7fffffffdea0
rsp                0x7fffffffde90    0x7fffffffde90
r8                 0x0
r9                 0x2a            42
r10                0x5555555560e4    93824992239844
r11                0x246            582
r12                0x555555555080    93824992235648
r13                0x7fffffffdfb0    140737488347056
r14                0x0
r15                0x0
rip                0x55555555175    0x55555555175 <sample_function+12>
eflags             0x206            [ PF IF ]
cs                 0x33            51
ss                 0x2b            43
ds                 0x0
es                 0x0
fs                 0x0
gs                 0x0
(gdb) |
```

13. Now where is the saved return address?

```
(gdb) x/xg $rbp+0x8
0x7fffffffdea8: 0x000055555555521a
(gdb) x/xg $rsp+0x18
0x7fffffffdea8: 0x000055555555521a
(gdb) |
```

Task:

Use a figure to illustrate the stack layout when the program is:

- (1) right before `sample_function()` is called;
- (2) inside `sample_function()` after stack memory allocation;
- (3) right after `sample_function()` returns.

Mark the location of the stack pointer, the base pointer, and return address.

Work out the role of the stack pointer (RSP), the base pointer (RBP), and the instruction pointer (RIP) in the program.

You can also modify the code to see how things change, e.g. adding more variables to see how the stack frame changes, change the compiler options (e.g. `-O2`), etc.

Take note that there are many GDB extensions which provide enhanced feature on GDB, you can explore them on your own. Just to list out the most common three:

peda: <https://github.com/longld/peda>

gef: <https://github.com/hugsy/gef>

pwndbg: <https://github.com/pwndbg/pwndbg>