


# CS5231: Systems Security

---

## Lecture 6: Kernel, Auditing, and System Provenance

# Kernel Introduction

- Core of Operating System
- A software that interface between applications & hardware
- Runs in a different CPU privilege level than normal software



# Under the Hood: Kernel Components

- Process Scheduler
  - Fair distribution of CPU time across processes
- Memory Management Unit (MMU)
  - Fair distribution of memory resources among processes
- Virtual File System (VFS)
  - Provide interface for accessing storage devices
- Networking
- Inter Process Communication (IPC)
- and others...
  - I/O scheduling & protection, buffering, caching, error handling, interrupt & event handling ...


# Types of Kernels

- Monolithic Kernel
  - Simple, common, and **same address space**
- Microkernel
  - Contain essential functionalities, hence smaller kernel
  - User & kernel have **different address space**
- Hybrid Kernel
  - Speed from monolithic & modularity from microkernel
- Nano Kernel
  - Small & bare minimum functionalities
- Exo Kernel
  - Distinct resource protection & management

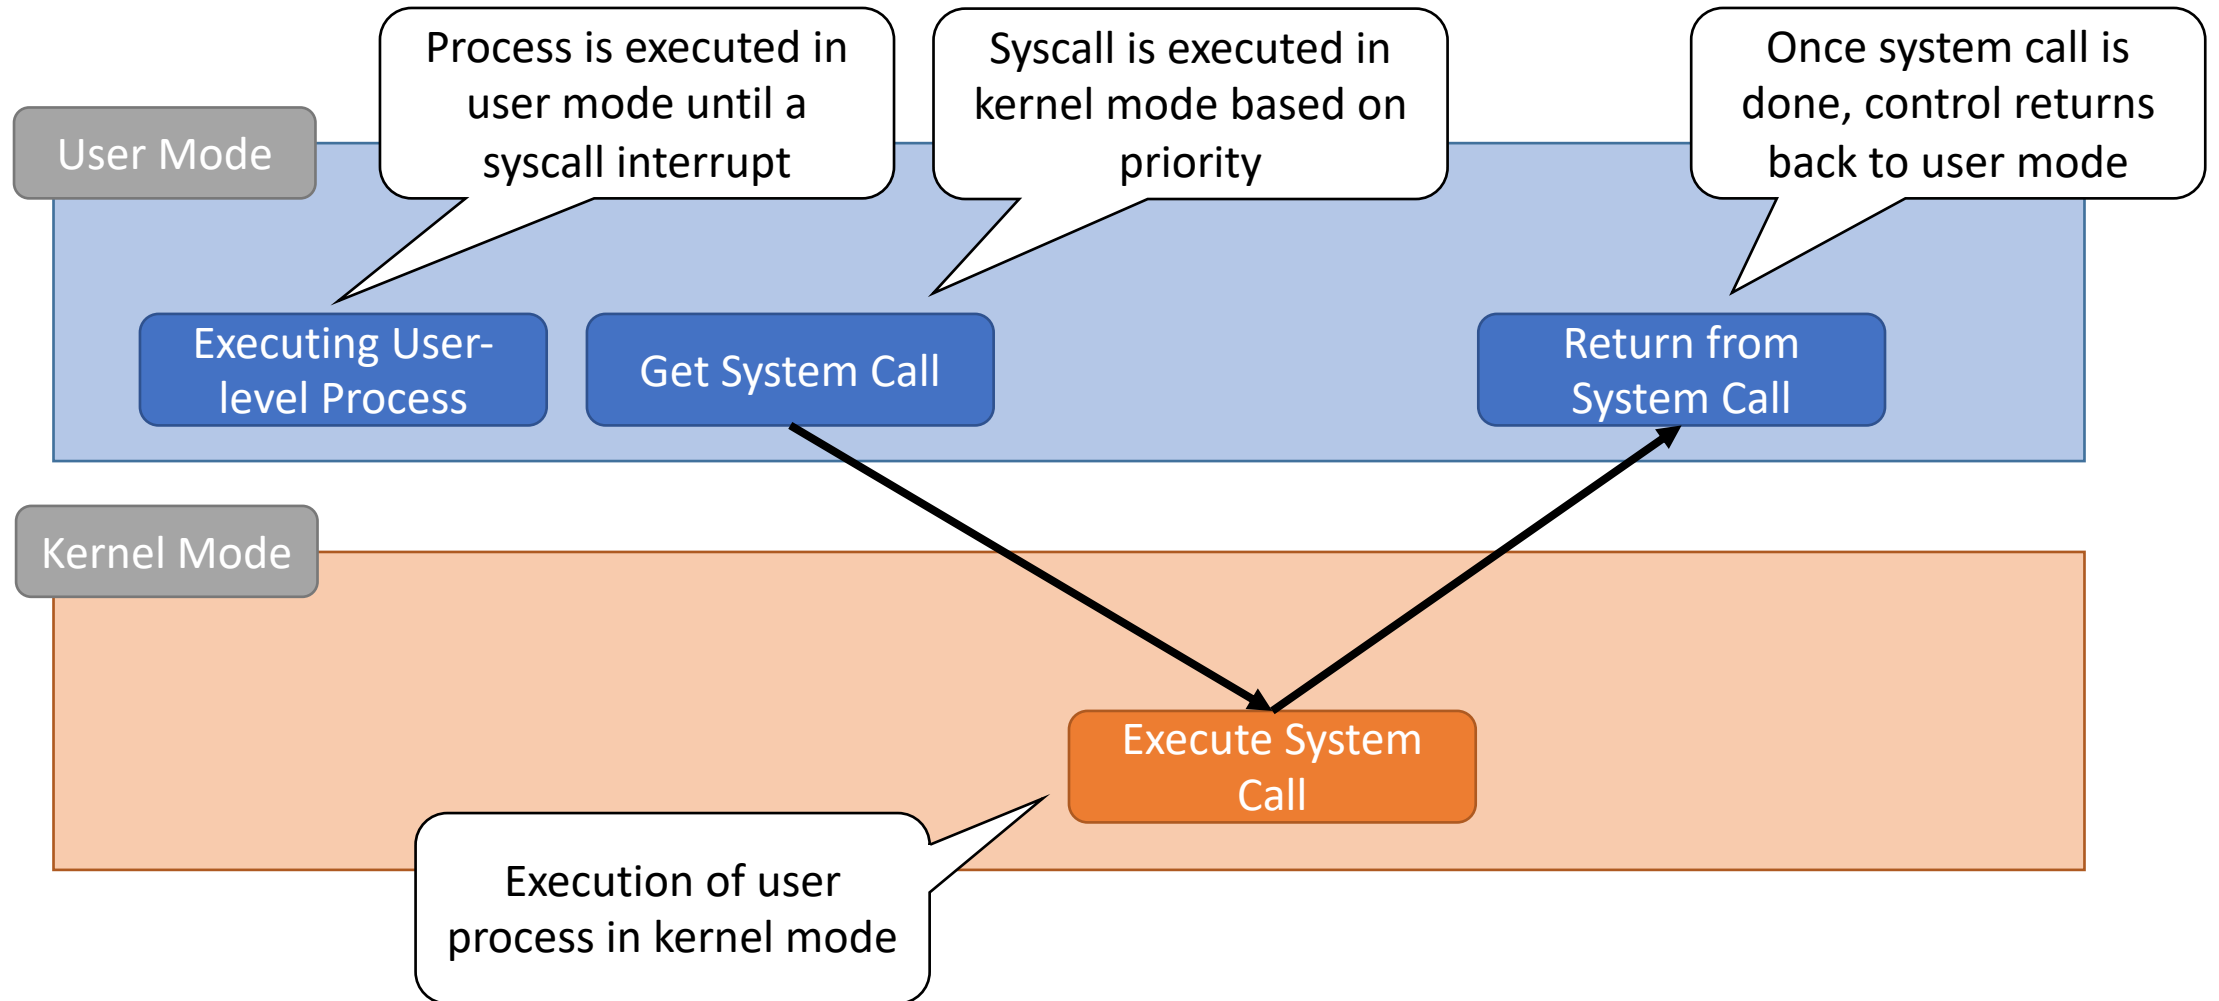
# Entering the Kernel

## System Call


- Entry point into kernel
- Application can call *syscall* to perform task
  - Create process
  - Networking
  - File I/O
  - ...
- All syscall in Linux manpage
  - <https://man7.org/linux/man-pages/man2/syscalls.2.html>



# Syscall Process



# Kernel Modules & Drivers



- Kernel module
  - Codes that can be loaded & unloaded into the kernel on-demand
  - Also known as *loadable kernel module*
- Kernel driver
  - Can be built statically in the kernel, or built as module for dynamic loading
  - Example of basic USB driver
    - [https://github.com/muratdemirtas/Linux-Kernel-Examples/blob/master/Basic\\_USB\\_Driver/Basic\\_USB\\_Driver/Basic\\_USB\\_Driver\\_main.c](https://github.com/muratdemirtas/Linux-Kernel-Examples/blob/master/Basic_USB_Driver/Basic_USB_Driver/Basic_USB_Driver_main.c)

# Kernel Module Example

```
1 //Hello World Example
2
3 //include needed libraries for building and printk() function
4 #include <linux/module.h>
5 #include <linux/init.h>
6
7 //start function for loading our module
8 int init_module(void)
9 {
10     //kernel message, you can see with dmesg command
11     printk(KERN_INFO "MOD: HELLO WORLD EXAMPLE LOADED.\n");
12     //return 0 (success) if module was loaded correctly, a non 0 return means module failed.
13     return 0;
14 }
15
16 //this function will execute when you remove this module from kernel.
17 void cleanup_module(void)
18 {
19     //print debug message to kernel.
20     printk(KERN_INFO "MOD: WORLD EXAMPLE LEAVING.\n");
21 }
```

*printk* is a C function that prints information to kernel log

[https://github.com/muratdemirtas/Linux-Kernel-Examples/blob/master/Hello\\_World/Hello\\_World/Hello\\_World\\_main.c](https://github.com/muratdemirtas/Linux-Kernel-Examples/blob/master/Hello_World/Hello_World/Hello_World_main.c)



# Kernel Module Example

- Compile & Run

- Need to have `Kbuild`, `Makefile`, and `module.c` before `make`
- Use `insmod module.ko` & `rmmod module.ko` & `lsmod`

- View Log

```
# cat /var/log/syslog | tail -2
Oct 06 21:38:41 mysystem kernel: MOD: HELLO WORLD EXAMPLE LOADED.
Oct 06 21:38:47 mysystem kernel: MOD: WORLD EXAMPLE LEAVING.

# dmesg | tail -2
MOD: HELLO WORLD EXAMPLE LOADED.
MOD: WORLD EXAMPLE LEAVING.
```

# System Auditing

Recording key events in the kernel

# Endpoint Monitoring Solutions


Endpoint monitoring solutions record **audit logs** for attack investigation



## Audit logs:

- A history of events representing OS-level activities
- Provide visibility into security incidents with data provenance

```
type=SYSCALL msg=audit(30/09/19 20:34:53.383:98866813) : arch=x86_64  
syscall=read exit=25 a0=0x3 ppid=15757 pid=30204 audit=junzeng sess=6309
```



# Investigation Using Audit Logs

Researchers use a **provenance graph** to navigate through audit logs:



- Nodes: system entities (e.g., process, file, and socket) & Edges: system calls
- **Backward/forward tracking** to find root cause of an attack and its ramifications

Real-world audit logs are always **large-Scale**, and provenance graphs are **Sophisticated!**

```
{ "@timestamp": "2020-10-31T14:14:47.777Z", "@metadata": { "beat": "auditbeat", "type": "doc", "version": "6.8.12", "user": { "suid": "0", "fsgid": "0", "gid": "0", "name_map": { "egid": "root", "gid": "root", "uid": "root", "sgid": "root", "suid": "root", "audit": "yinfang", "euid": "root", "fsgid": "root", "fsuid": "root", "sgid": "0", "uid": "0", "euid": "0", "egid": "0", "fsuid": "0", "audit": "1000", "process": { "name": "sshd", "exe": "/usr/sbin/sshd", "pid": "18104", "ppid": "1689", "auditd": { "sequence": 166719, "result": "success", "session": "705", "data": { "tty": "(none)", "a3": "8", "a0": "4", "exit": "384", "arch": "x86_64", "syscall": "read", "a2": "180", "a1": "7ff97aeba120" } } } } } }

{ "@timestamp": "2020-10-31T14:14:47.777Z", "@metadata": { "beat": "auditbeat", "type": "doc", "version": "6.8.12", "user": { "audit": "1000", "fsuid": "0", "fsgid": "0", "egid": "0", "sgid": "0", "suid": "0", "uid": "0", "euid": "0", "name_map": { "fsgid": "root", "sgid": "root", "suid": "root", "uid": "root", "audit": "yinfang", "egid": "root", "euid": "root", "fsuid": "root", "gid": "root", "gid": "0", "process": { "pid": "18104", "ppid": "1689", "name": "sshd", "exe": "/usr/sbin/sshd", "auditd": { "data": { "a3": "8", "a2": "180", "arch": "x86_64", "tty": "(none)", "a0": "4", "exit": "384", "a1": "7ff97aeba120", "syscall": "read", "sequence": 166720, "result": "success", "session": "705" } } } } } }


{ "@timestamp": "2020-10-31T14:14:47.777Z", "@metadata": { "beat": "auditbeat", "type": "doc", "version": "6.8.12", "user": { "egid": "0", "audit": "1000", "fsgid": "0", "name_map": { "euid": "root", "fsuid": "root", "gid": "root", "sgid": "root", "audit": "yinfang", "fsgid": "root", "suid": "root", "uid": "root", "egid": "root", "suid": "0", "euid": "0", "uid": "0", "sgid": "0", "fsuid": "0", "gid": "0", "process": { "name": "sshd", "exe": "/usr/sbin/sshd", "pid": "18104", "ppid": "1689", "auditd": { "sequence": 166721, "result": "success", "session": "705", "data": { "a1": "7ff97aeba120", "arch": "x86_64", "a3": "8", "exit": "384", "syscall": "read", "a2": "180", "a0": "4", "tty": "(none)" } } } } } }
```



# Audit Log Analysis


- Starting from a detection point, *Backtracker* does:
  - Events & objects identification related detection point
  - Generate dependency graph
  - Use rules to prune unrelated nodes in the dependency graph

Dependency explosion!



# Audit Log Analysis

- Resolve *dependency explosion* problem in a long running application
  - Fine-grained provenance tracing technique
  - Identifying unit boundaries & dependences
  - Partition into individual *unit*
  - Code instrumentation




# Audit Log Analysis

turn a security problem into a graph analysis problem


- Address *threat alert fatigue* during threat investigation
  - Assign anomaly scores to every edge in dependency graph
  - Based on frequency of events that have occurred (historical & contextual information)
  - Propagated score through edges in the graph
  - Generate aggregated anomaly score for triaging

NoDoze (Hassan et al., 2019)



# Audit Log Analysis

- Generate high-level graph during threat investigation
  - Develop robust & reliable detection signal
  - Correlate between suspicious information flow





# What to be logged?

- It depends!
  - Application scenarios
  - Business purposes
- Questions
  - Universities may keep different kinds of logs
    - Student/staff record
    - Student/staff emails?

# Why logging?

- Main goals
  - Record valuable data about applications, system and network activities
  - Provide clues for later investigation
- It must be performed by a trusted/ authorized entity ...

# A Taxonomy of Logs

- **Application logs**
  - Contain application-level events
  - e.g., access log / error log
- **System logs**
  - Contain system-related events
  - e.g., device driver loading/unloading
- **Security logs**
  - Contain security-related events
  - e.g., valid/invalid logon attempts

# Application Log

- Mail server log
  - Connection status; SMTP queues;...
- FTP server log
  - Current logins; file uploaded/downloaded;...
- Database server log
  - Objects accessed; creation of new tables; ...
- Web browser log
  - URLs visited; file downloaded; ...

# System Log

- System startup/initialization log
  - *dmesg*
- System running status
  - */proc*
- When a system goes wrong (e.g., crash)
  - What should be logged?
  - How to collect log?
    - Kernel dumps (egg vs. chicken)

# Security Log

- Firewall log
  - Inbound/outbound packets
  - Packets dropped
  - TCP connections rejected
- IDS log
  - Unauthorized file modifications
  - Intrusion alerts
  - Suspicious system activities
  - Attack statistics

# General Logging Mechanisms

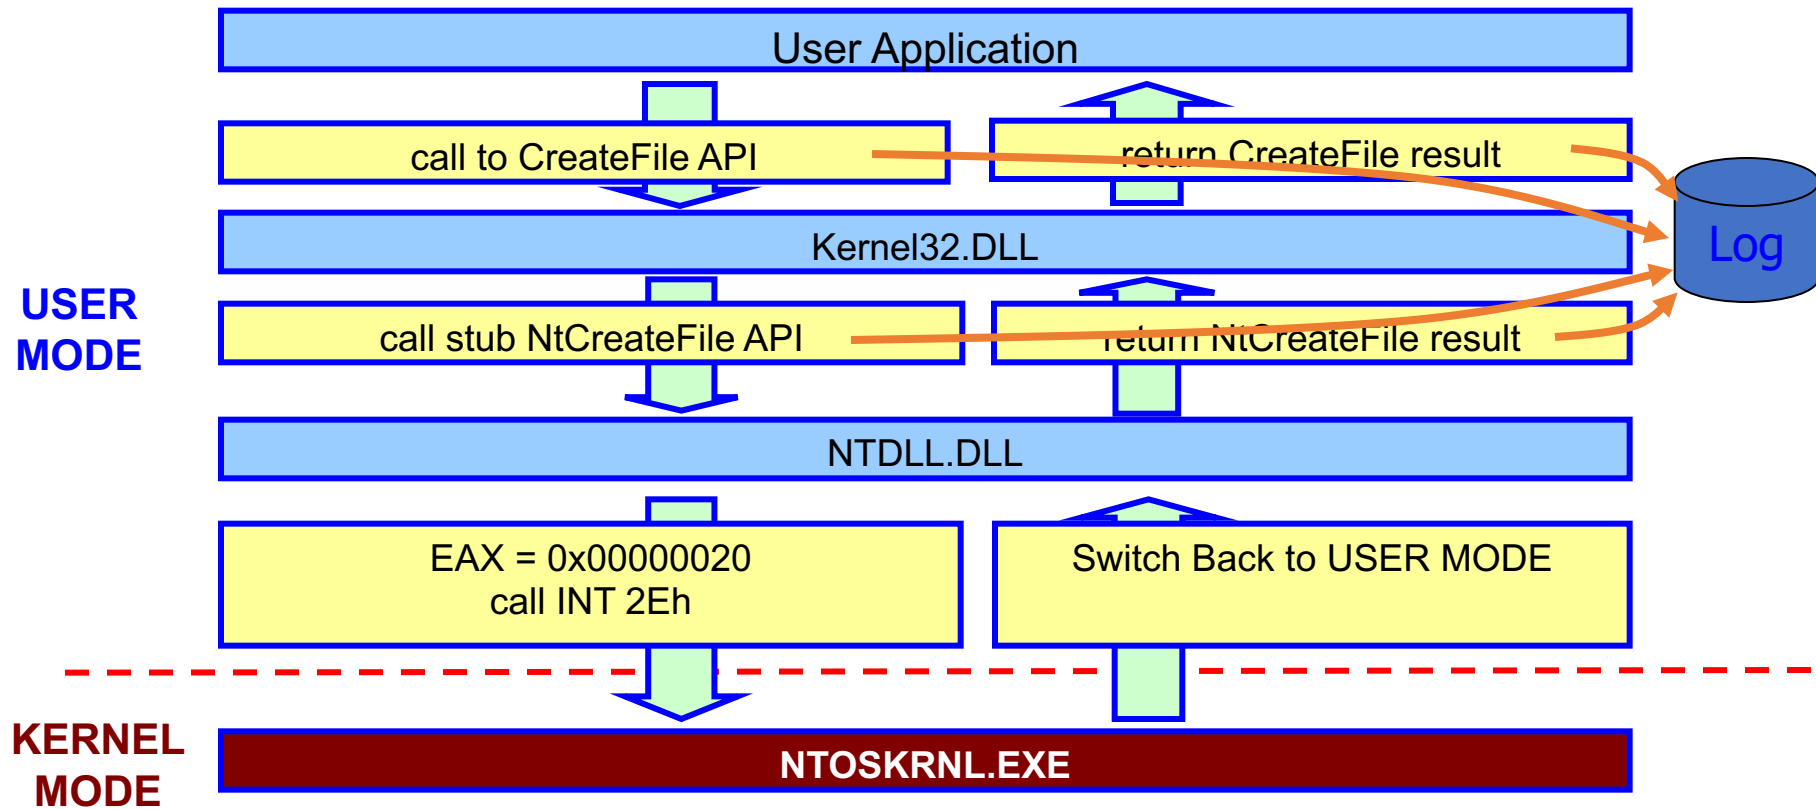
- Application-level
  - Library wrapping / API hooking
- Kernel-level
  - Syslogd/klogd
  - System call interception
  - Linux security module
- Virtual Machine Monitor-level
  - System call interception

# API Hooking

- Similar to **functional overloading** in programming language
- Commonly used for debugging purposes
- Platform-specific
  - Windows:
    - DLL injection/binary rewriting
  - Unix:
    - LD\_PRELOAD trick



# API Hooking



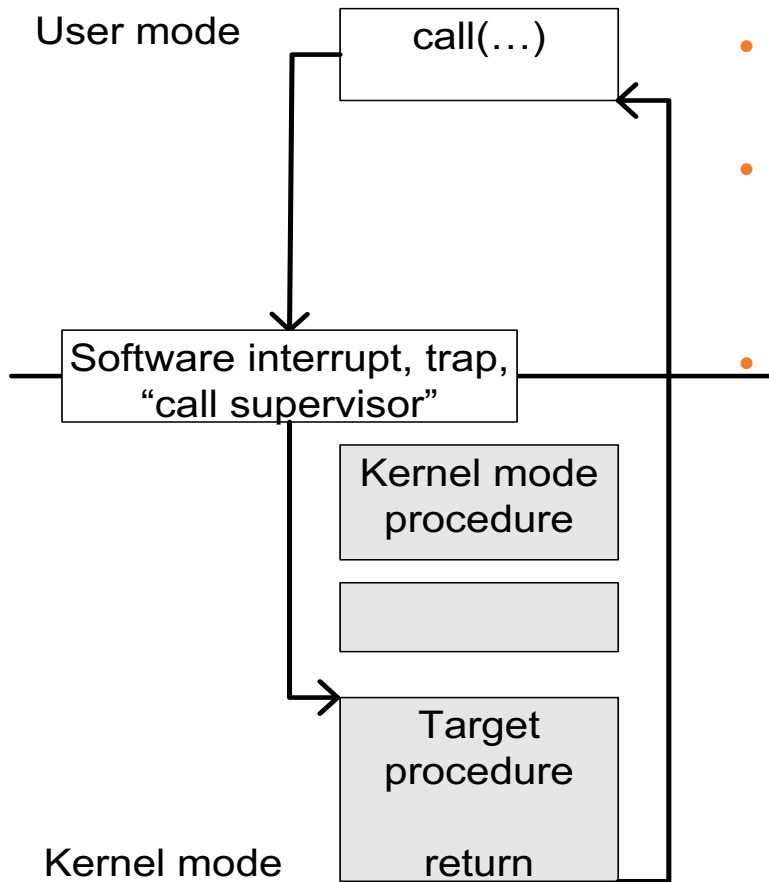
# Example API Hooking Log

- iexplore.exe (268) : File::Write (C:\WINDOWS\Downloaded Program Files\ieloader.exe)
- iexplore.exe (268) : Sys::Execute (C:\WINDOWS\Downloaded Program Files\ieloader.exe)
- ieloader.exe (1728) : Reg::SetValue (HKLM\Software\Microsoft\Windows\CurrentVersion\Run, FX)
- iexplore.exe (268) : Net::Connect (205.205.86.51, 80)
- ieloader.exe (1728) : File::Write (C:\WINDOWS\System32\uvbdcgrtjce.dll)
- ieloader.exe (1728) : Sys::Execute (C:\Program.exe)
- iexplore.exe (268) : File::Write (C:\Documents and Settings\HS\Local Settings\Temporary Internet Files\Content.IE5\EJ89IPOV\init[1].js)
- iexplore.exe (1812) : Reg::SetValue (HKCU\Software\Microsoft\Windows\CurrentVersion\Explorer\Shell Folders, Desktop)
- iexplore.exe (1812) : Reg::SetValue (HKCU\Software\Microsoft\Windows\ShellNoRoam\BagMRU, NodeSlots)
- iexplore.exe (1812) : Reg::SetValue (HKCU\Software\Microsoft\Windows\ShellNoRoam\BagMRU, MRUListEx)
- iexplore.exe (268) : File::Write (C:\WINDOWS\Downloaded Program Files\SET4A.tmp)
- iexplore.exe (268) : File::Write (C:\WINDOWS\Downloaded Program Files\ieloader.exe)
- iexplore.exe (268) : Sys::Execute (C:\WINDOWS\Downloaded.exe)
- ieloader.exe (1728) : Reg::SetValue (HKLM\Software\Microsoft\Windows\CurrentVersion\Run, FX)
- iexplore.exe (268) : Net::Connect (205.205.86.51, 80)
- ieloader.exe (1728) : File::Write (C:\WINDOWS\System32\uvbdcgrtjce.dll)
- iexplore.exe (268) : File::Write (C:\Documents and Settings\HS\Local Settings\Temporary Internet Files\Content.IE5\EJ89IPOV\init[1].js)

# General Logging Mechanisms

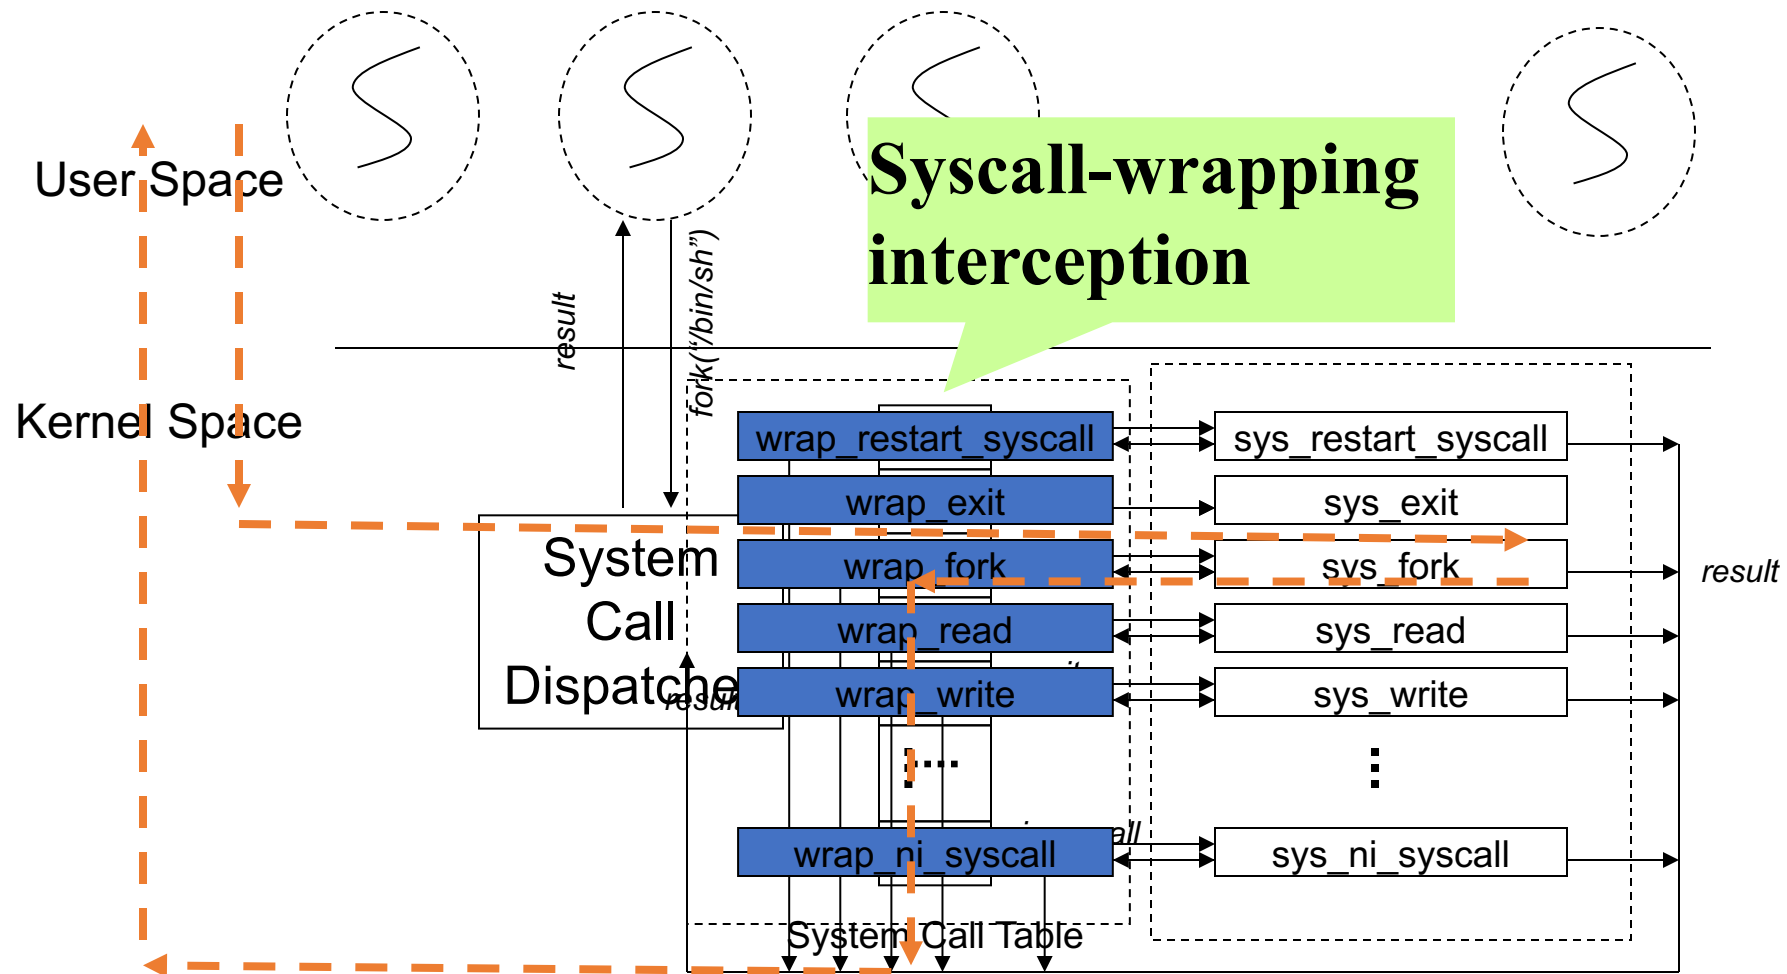
- Application-level
  - Library wrapping / API hooking
- Kernel-level
  - Syslogd/klogd
  - System call interception
  - Linux security module
- Virtual Machine Monitor-level
  - System call interception

# System Call



- The parameters of the call are passed according to certain OS/hardware specific convention
- Switch to protected mode
  - E.g., int \$0x80
- A special module takes over, that will analyze the parameters and the access rights; this module can reject the system call
- If accepted, then the corresponding routine from the operating system is executed and the result is returned to the user; upon return, the user mode is restored

# System Call Interception



# Example System Call Log

```
673["sendmail"]: 5_open("/proc/loadavg", 0, 438) = 5
673["sendmail"]: 192_mmap2(0, 4096, 3, 34, 4294967295, 0) = 1073868800
673["sendmail"]: 3_read(5, "0.26 0.10 0.03 2...", 4096) = 25
673["sendmail"]: 6_close(5) = 0
673["sendmail"]: 91_munmap(1073868800, 4096) = 0
...
2568["httpd"]: 102_accept(16, sockaddr{2, cbbdff3a}, cbbdff38) = 5
2568["httpd"]: 3_read(5, "\1281\1\0\2\0\24...", 11) = 11
2568["httpd"]: 3_read(5, "\7\0À\5\0\128\3\...", 40) = 40
2568["httpd"]: 4_write(5, "\132@\4\0\1\0\2\...", 1090) = 1090
...
2568["httpd"]: 4_write(5, "\128\19Ê\136\18\...", 21) = 2
```

# Applications of Syscall Interception

- Logging
  - Syscalltrack, sebek, ...
    - *<http://syscalltrack.sf.net>, ...*
- Sandboxing
  - Systrace, Janus, ...
    - *<http://www.citi.umich.edu/u/provos/systrace/>, ...*
- Virtual Machine
  - UML, UMLinux
    - *<http://user-mode-linux.sf.net>, ...*

# General Logging Mechanisms

- Application-level
  - Library wrapping / API hooking
- Kernel-level
  - Syslogd/klogd
  - System call interception
  - Linux security module
- Virtual Machine Monitor-level
  - System call interception



# Linux Security Module: Overview

- Motivation: Separate kernel from security features
  - Minimize the impact to kernel
  - SELinux motivated the creation of LSM
- Note LSM doesn't provide any security
  - Adds security fields to kernel
  - Provides interface to manage these fields

# Linux Security Module (LSM)

- Security extension for kernel
- Reduce attack surface by enforcing access policies
- Mandatory Access Control
  - Subjects, objects, and operations



Unlike DAC in Linux's access control..!!

Introduce *LSM hooks*

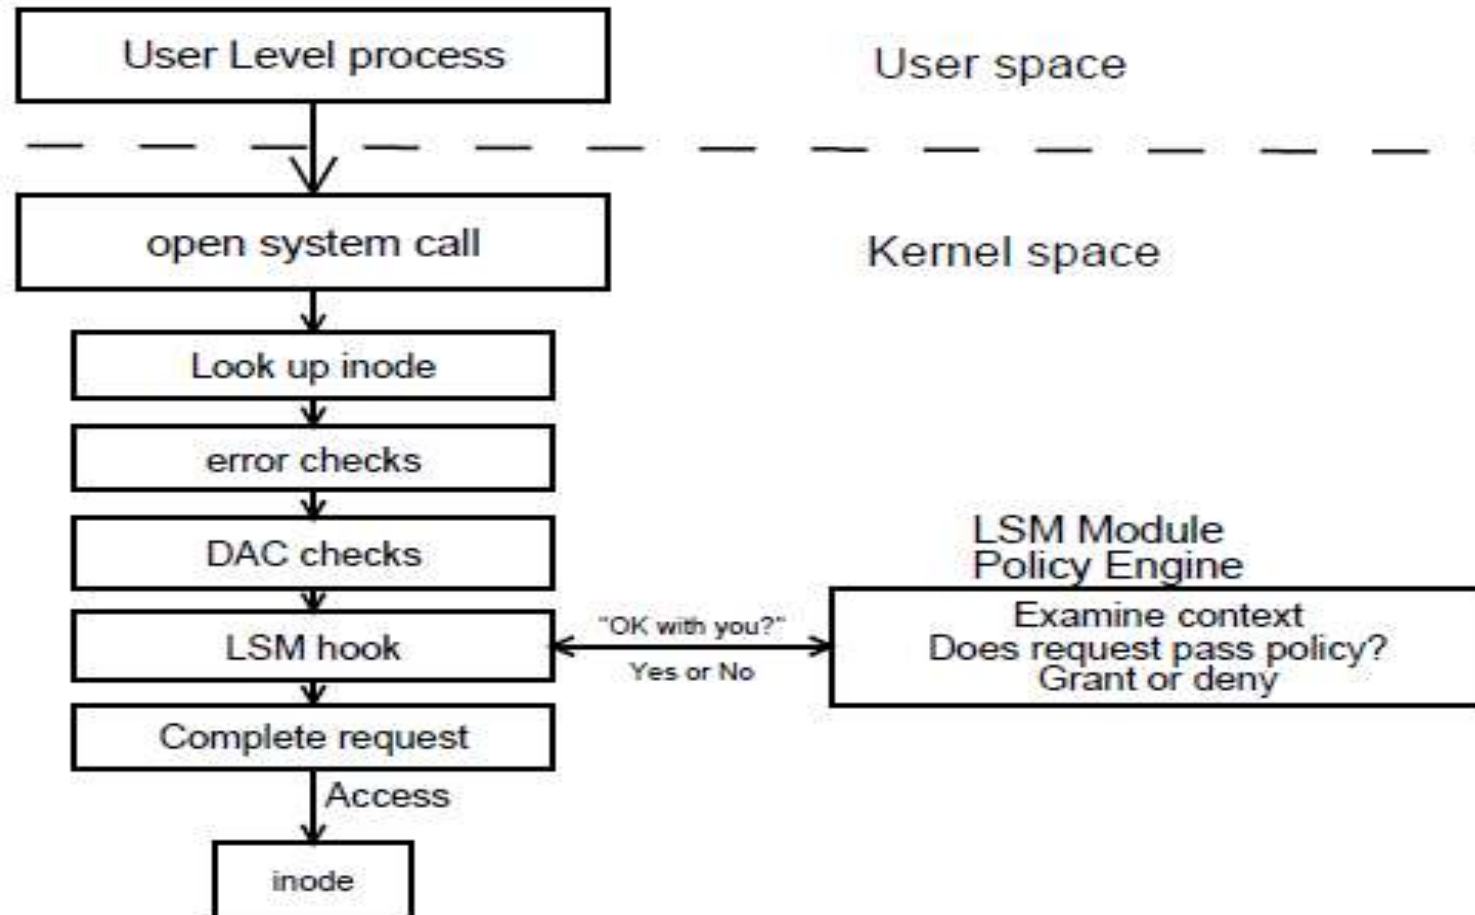
- Control access to kernel objects
  - E.g. inodes, files, credentials, devices, tasks, IPCs

*“Can a subject  $S$  perform kernel Operation  $O$  on a kernel object  $A$ ?”*

# Linux Security Module: Hooks

- Hooks
  - A set of functions to control operations on kernel objects and security fields in kernel data structures.
  - Management Hooks
    - used to manage security fields (e.g., `file_alloc_security`)
  - Control Hooks
    - used to perform access controls (e.g., `selinux_inode_permission`)

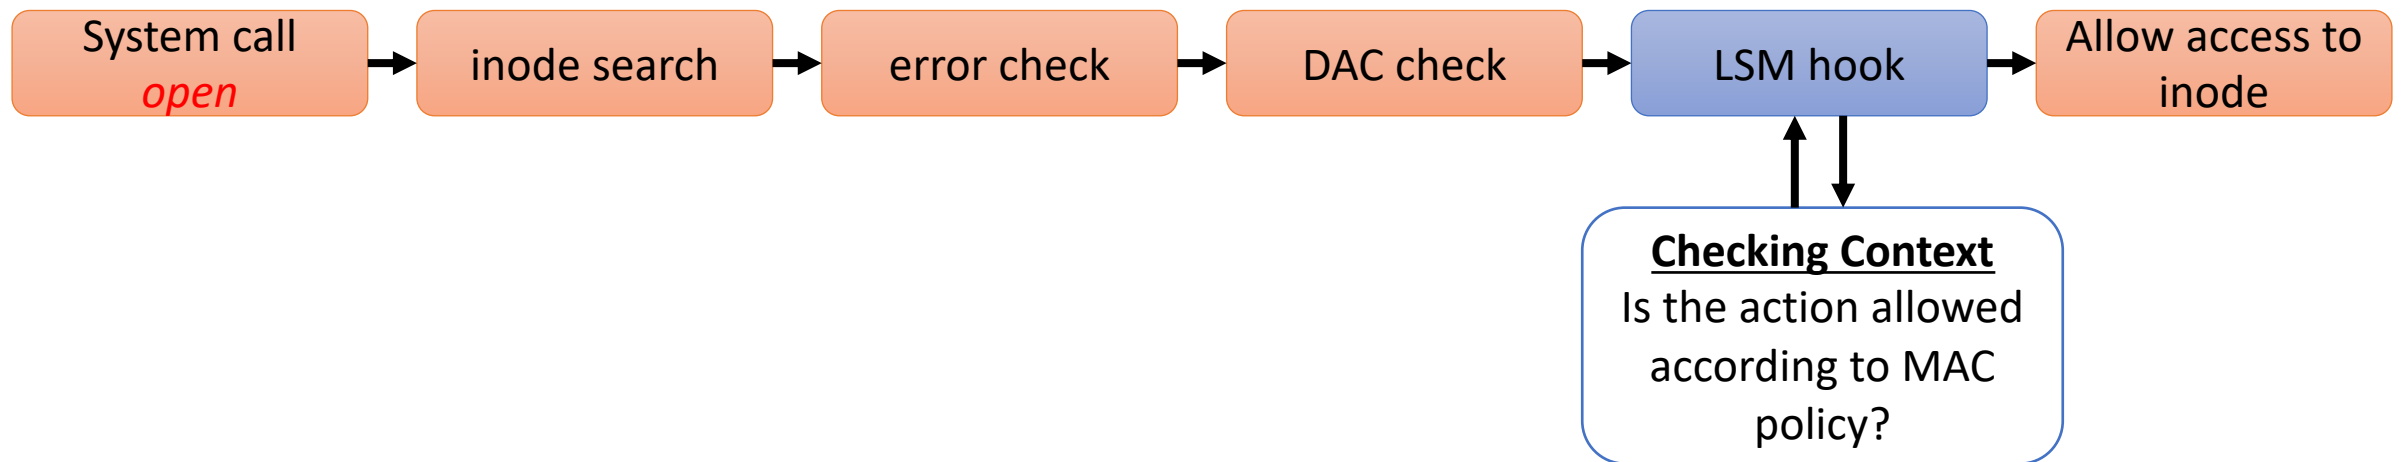
# LSM Hook Architecture



# Linux Security Module (LSM)

- LSM's MAC is checked only after DAC & other checks are performed
- LSM hooks are applied inline during kernel code execution
- Examples – AppArmor, SELinux, Smack, TOMOYO

## LSM workflow during open() syscall in kernel space




# General Logging Mechanisms

- Application-level
  - Library wrapping / API hooking
- Kernel-level
  - Syslogd/klogd
  - System call interception
  - Linux security module
- Virtual Machine Monitor-level
  - System call interception

# Virtual Machine Monitor-level

- Strength?
- Weakness?



# Logging Storage Policies

- Reset log files at periodic intervals
- Rotate log files, keeping data for a fixed time
- Compress and archive to a tape or other permanent media



# Main Applications

- Logging-based Applications
  - Intrusion Detection
  - Intrusion Recovery
  - Software Debugging

# Intrusion Detection

- Intrusion Detection is the process of **identifying** and **responding** to malicious activity targeted at computing and networking resources
- Resources:
  - One computer, or
  - A local/wide area network

# Models of Intrusion Detection

- Anomaly detection
  - What is usual, is known
  - What is unusual, is bad
- Misuse detection
  - What is bad, is known
  - What is not bad, is good
- Specification-based detection
  - What is good, is known
  - What is not good, is bad
- Goal → generating a **Detection Point**

# Detection Point

- Suggests a possible intrusion
- Examples:
  - An anomaly log entry
    - e.g., a shell process launched
  - A suspicious system activity
    - e.g., an outbound TCP connection to a remote IRC server
  - An unauthorized modification to a critical configuration file
    - e.g., /etc/inetd.conf

# After an Intrusion Is Identified

- For each intrusion, it is desirable to find out:
  - **Break-in Point:**
    - How did the attacker gain access to the system?
  - **Contaminations:**
    - What did the attacker do after the break-in?



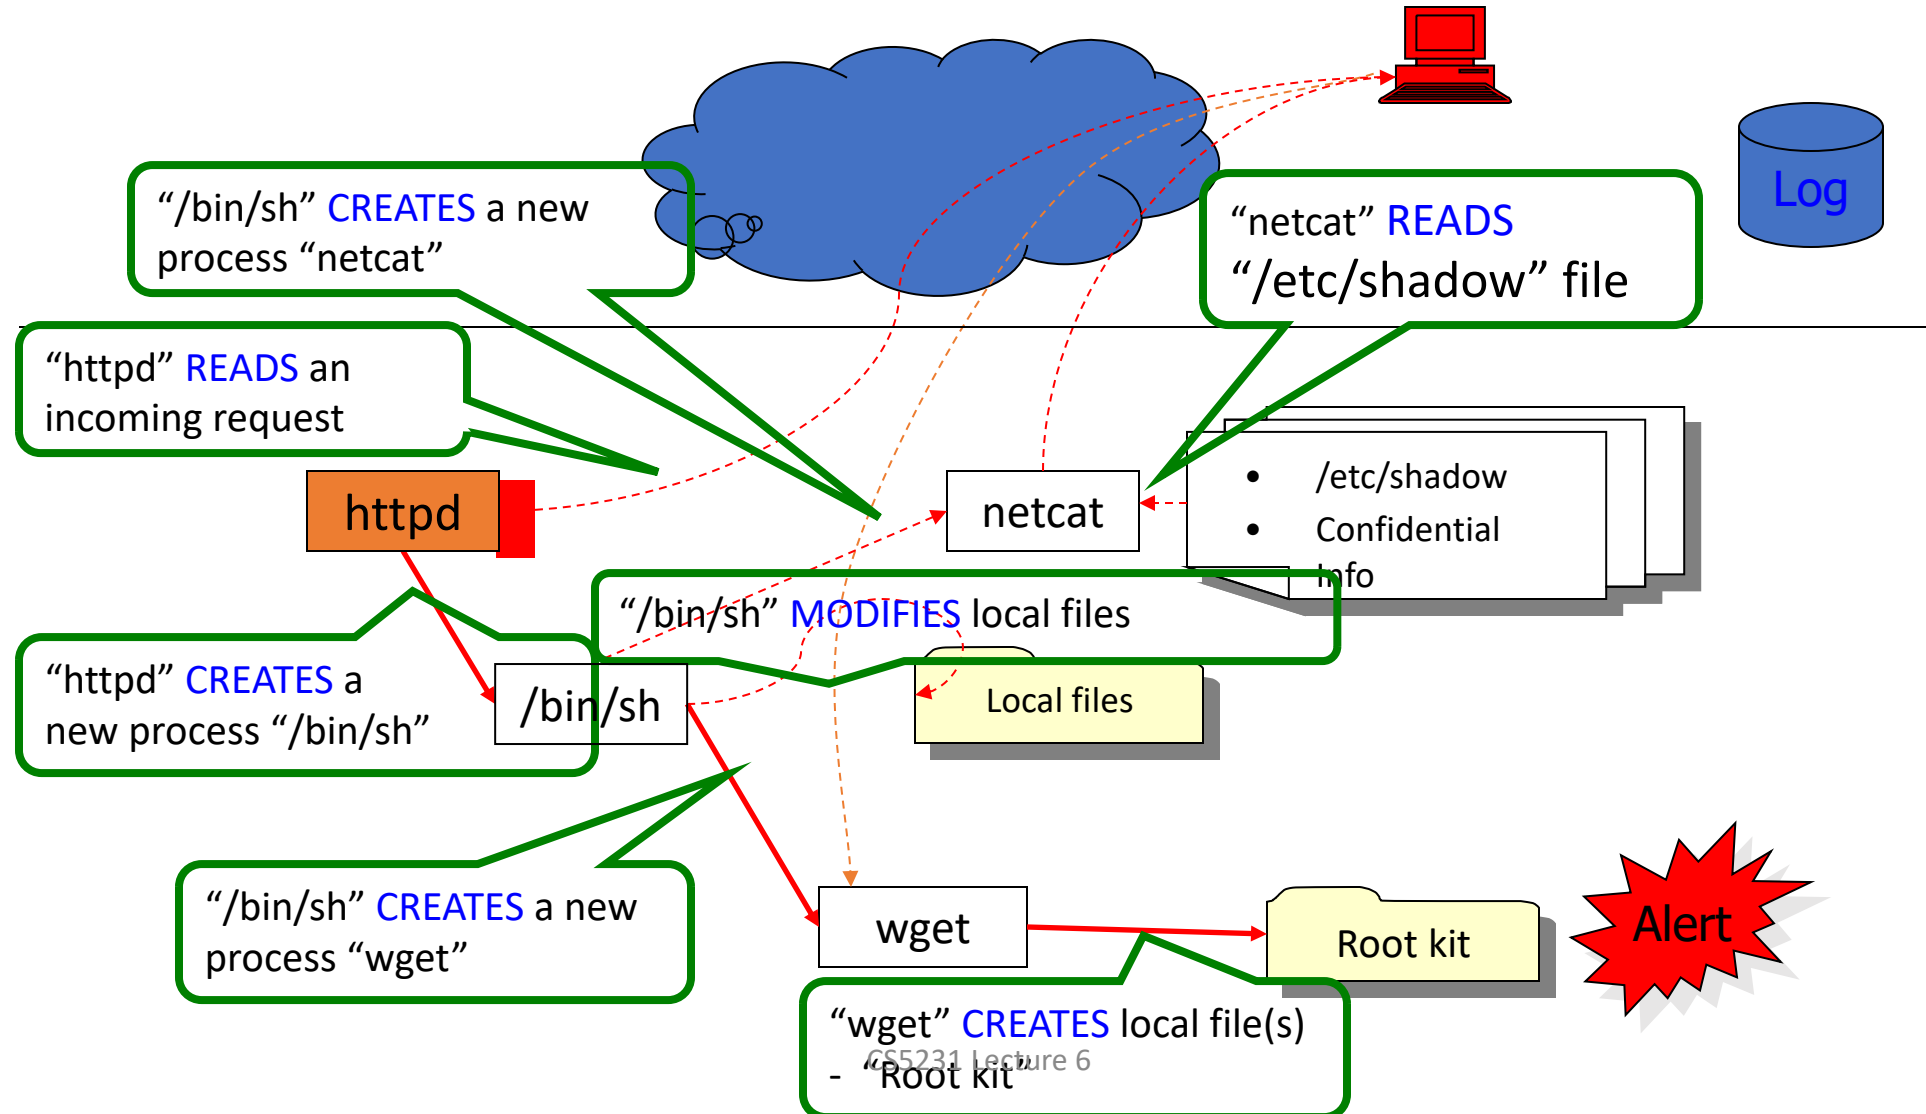
The diagram consists of two green curved arrows. One arrow starts from the 'Detection Point' box and points to the 'Break-in Point' bullet point. The other arrow starts from the 'Detection Point' box and points to the 'Contaminations' bullet point.

**Detection Point**

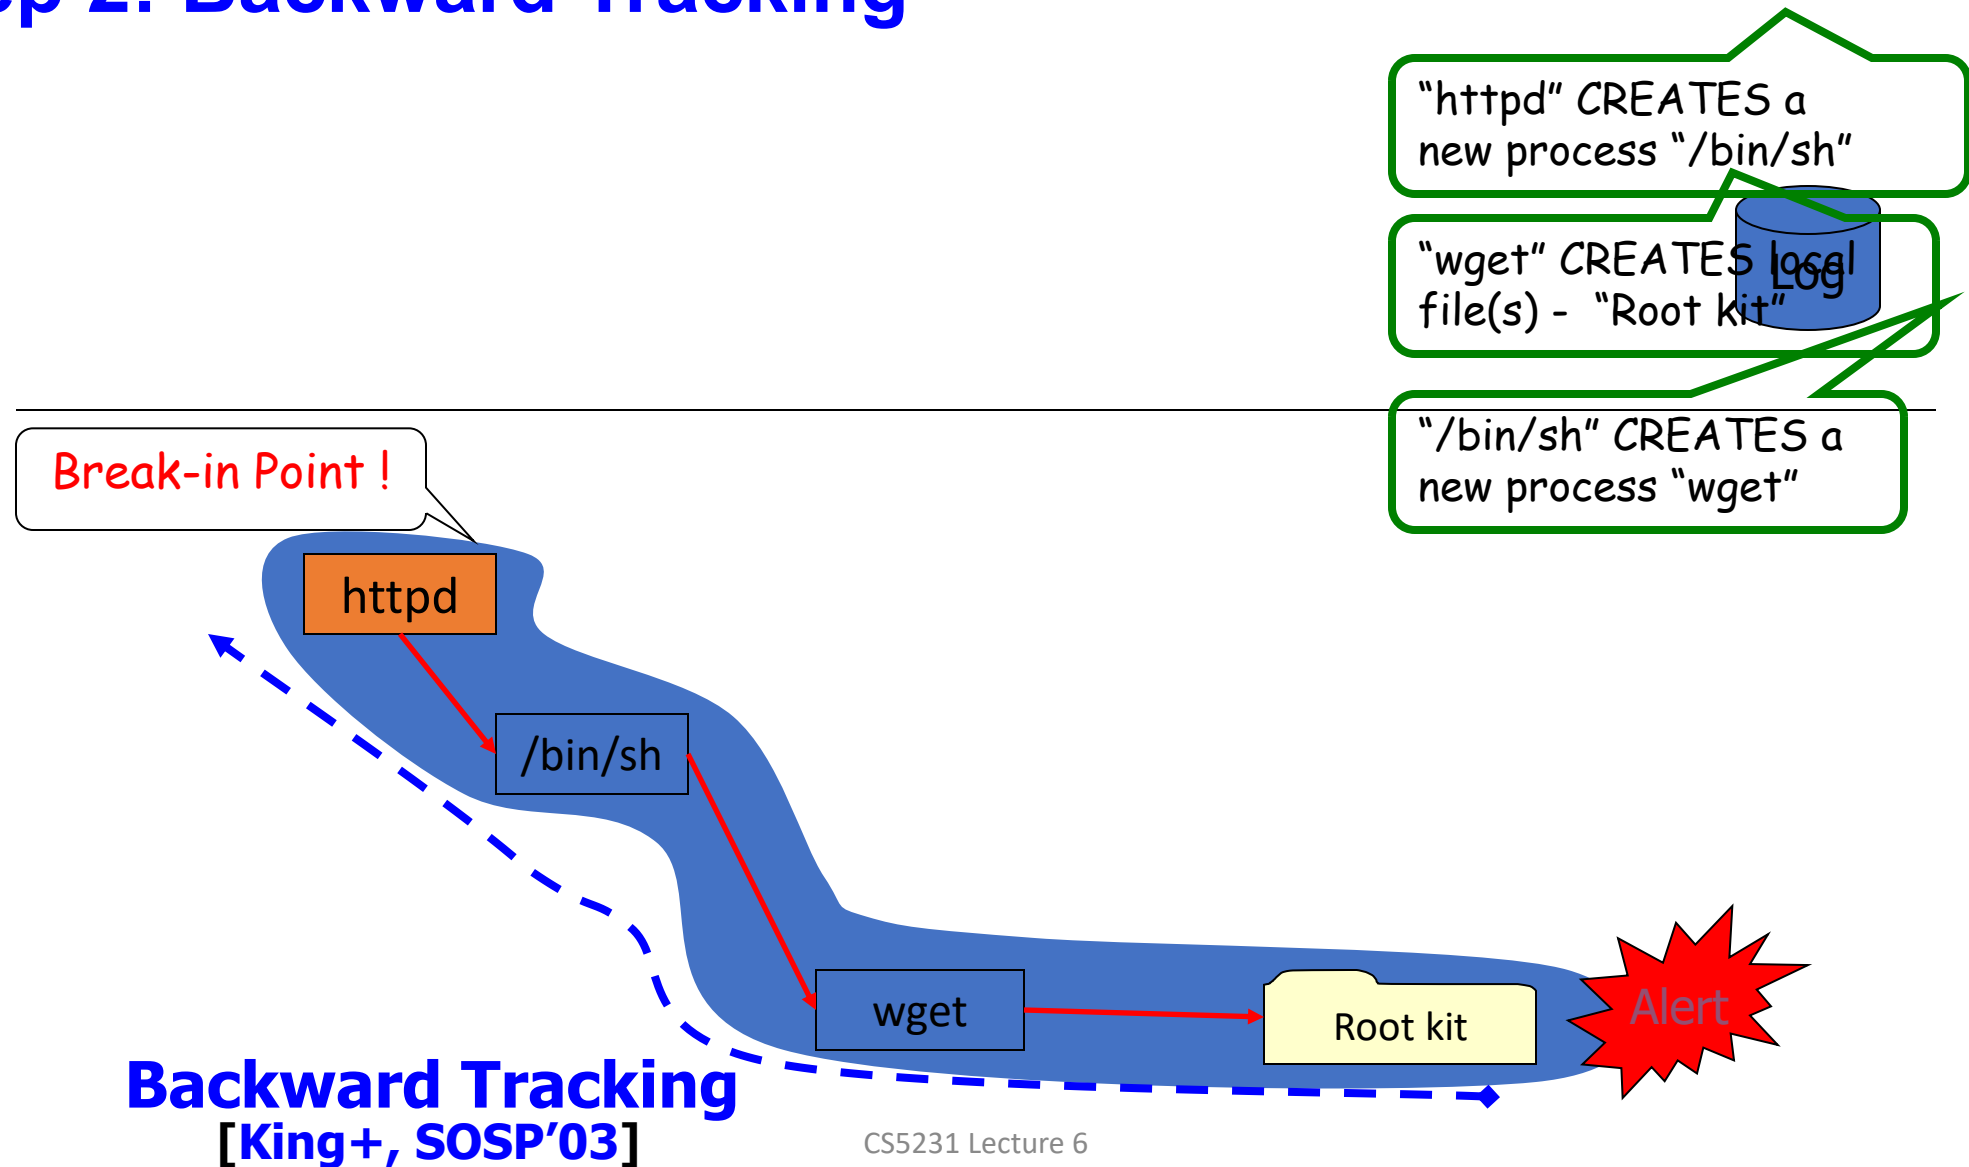
# Intrusion Investigation

- Three Main Steps
  - Step 1: Online Log Collection
  - Step 2: Backward Tracking
  - Step 3: Forward Tracking

# Step 1: Online Log Collection

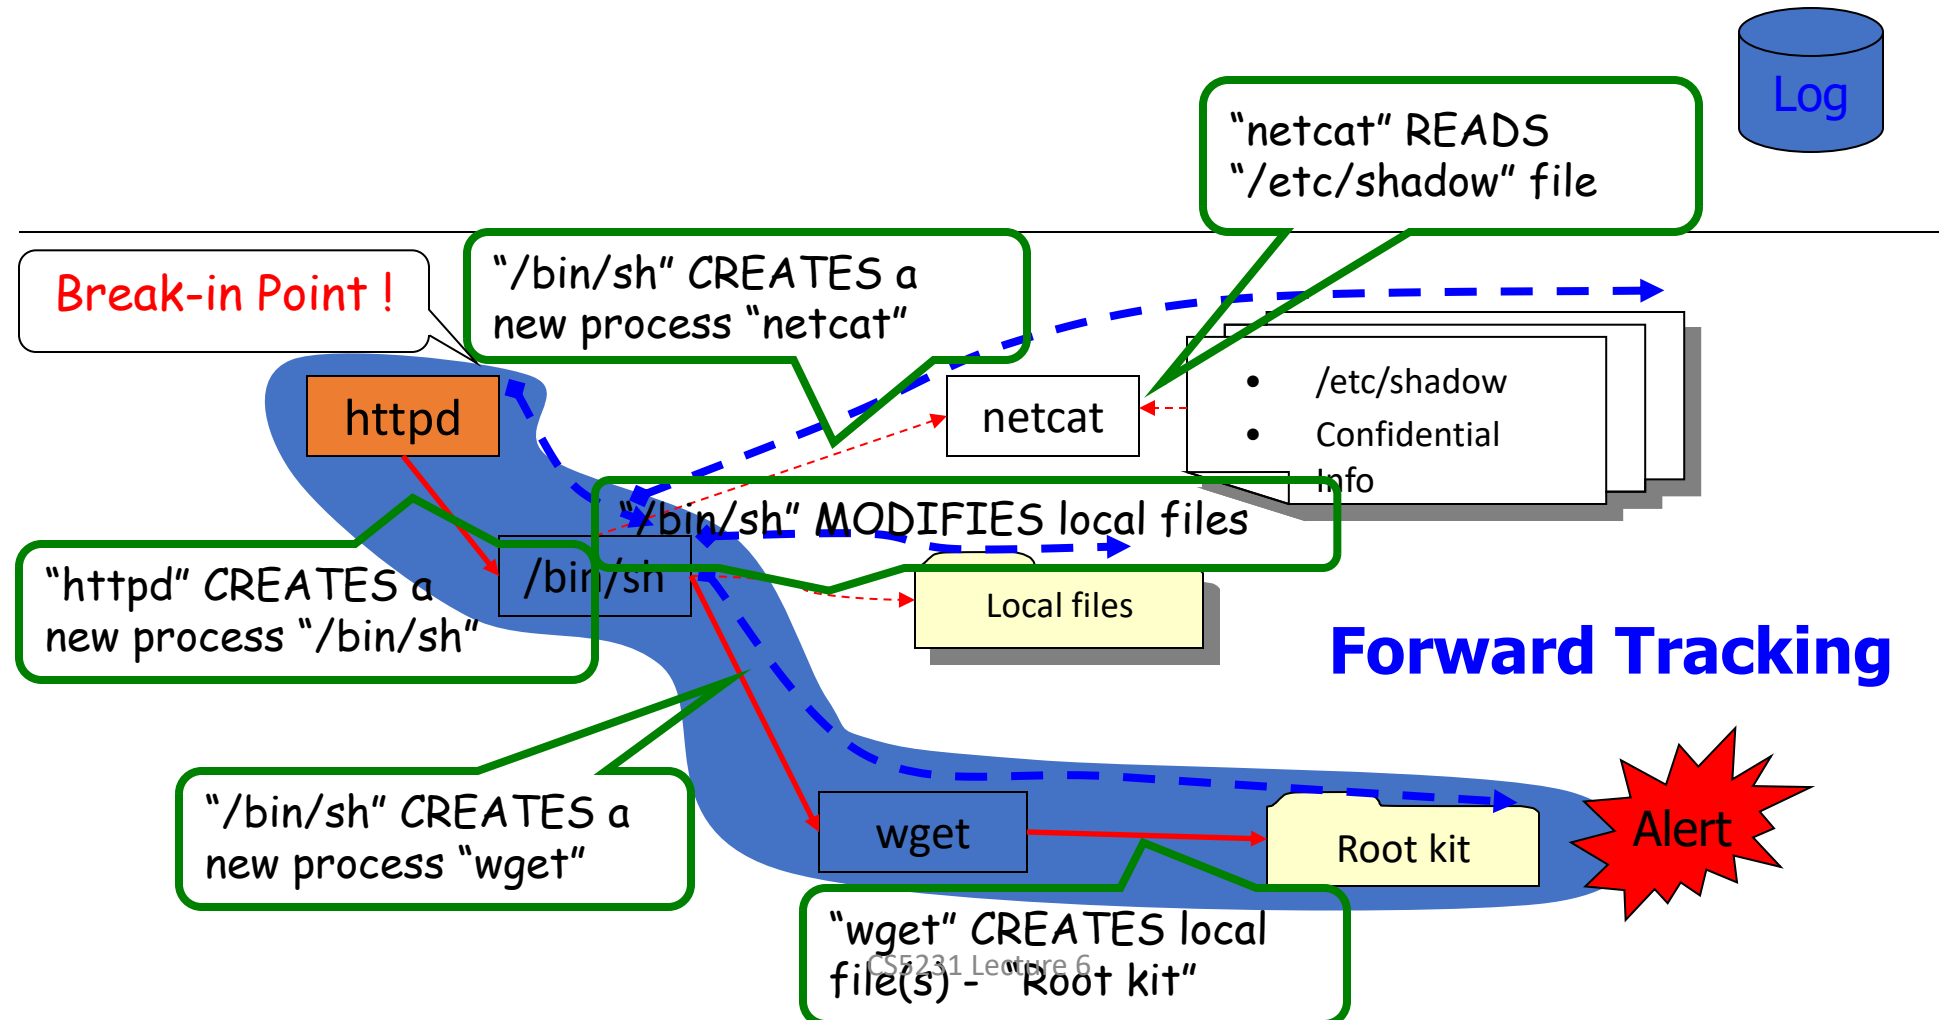


## Step 2: Backward Tracking

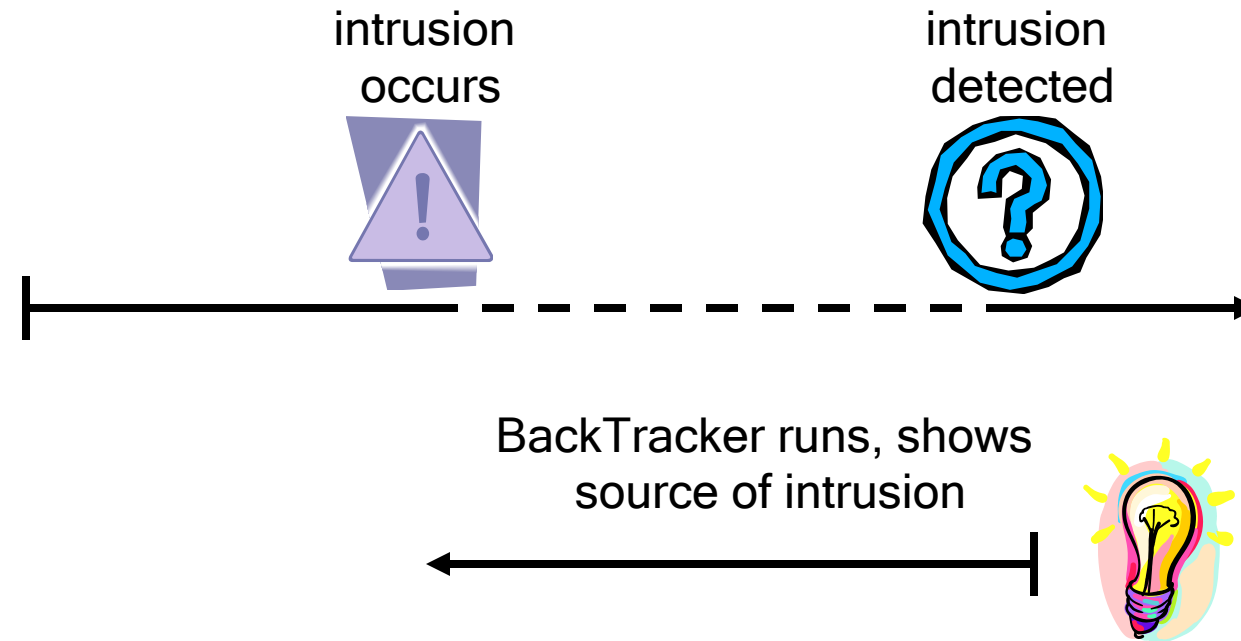




## Step 3: Forward Tracking



# BackTracker



- ❑ Online component logs objects and events
- ❑ Offline component generates graphs

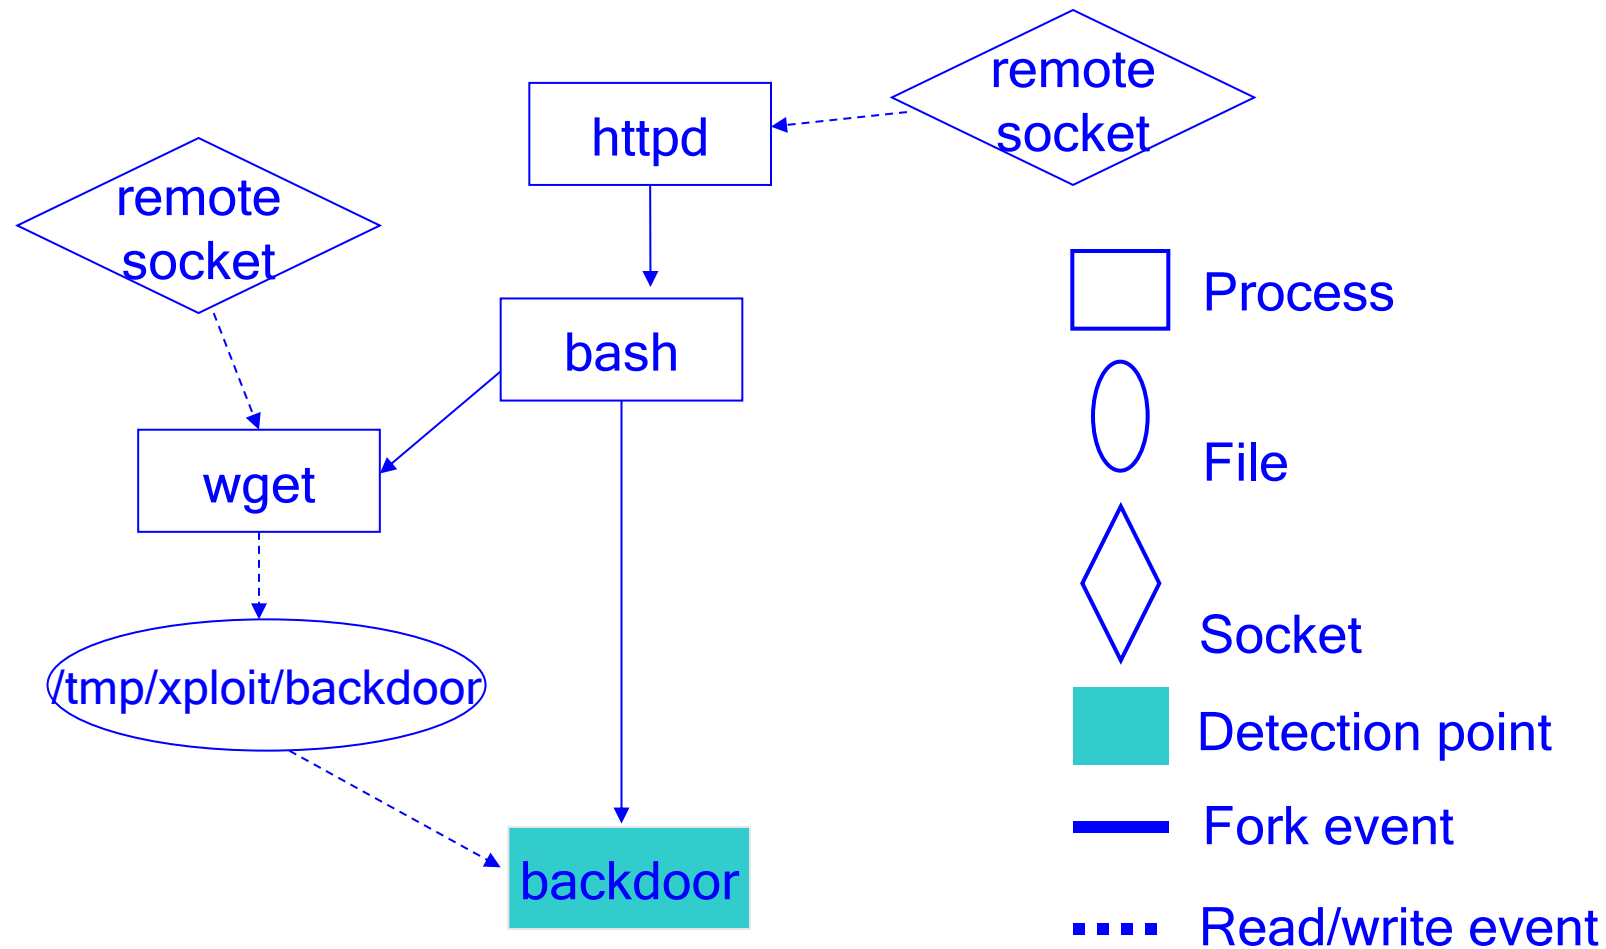
# BackTracker Objects

- Process
- File
- Filename

# Dependency-Forming Events

- Process / Process
  - fork, clone, vfork
- Process / File
  - read, write, mmap, exec
- Process / Filename
  - open, creat, link, unlink, mkdir, rmdir, stat, chmod, ...
- Dependency-tracking is an effective technique for highlighting actions of attacker


# BackTracker Example



# Challenge in Scalability

- Backward Tracking → Break-in Point
  - Inputs: Detection Point, **the Whole Log**
- Forward Tracking → Contaminations
  - Inputs: Break-in Point, **the Whole Log**

Analyze the whole log !



# Related Work

- Scale up provenance analysis:
  - Data reduction [NDSS'16, 18 ...] & Query system [Security'18, ATC'18 ...]
  - Recognizing behaviors of interest requires intensive manual efforts

A semantic gap between low-level events and high-level behaviors

- Apply expert-defined specifications to bridge the gap
  - Match audit events against domain rules that describe behaviors
  - Query graph [VLDB'15, CCS'19], Tactics Techniques Procedures (TTPs) specification [SP'19,20], and Tag policy [Security'17,18]

Behavior-specific rules heavily rely on domain knowledge (time-consuming)

# Related Work

- Scale up provenance analysis:
  - Data reduction [NDSS'16, 18 ...] & Query system [Security'18, ATC'18 ...]

Can we automatically **abstract** high-level behaviors from low-level audit logs and **cluster** semantically similar behaviors before human inspection?


- Query graph [VLDB'15, CCS'19], Tactics Techniques Procedures (TTPs) specification [SP'19,20], and Tag policy [Security'17,18]

Behavior-specific rules heavily rely on domain knowledge (**time-consuming**)




# Motivating Example

Attack Scenario: A software tester **exfiltrates sensitive data** that he has access to



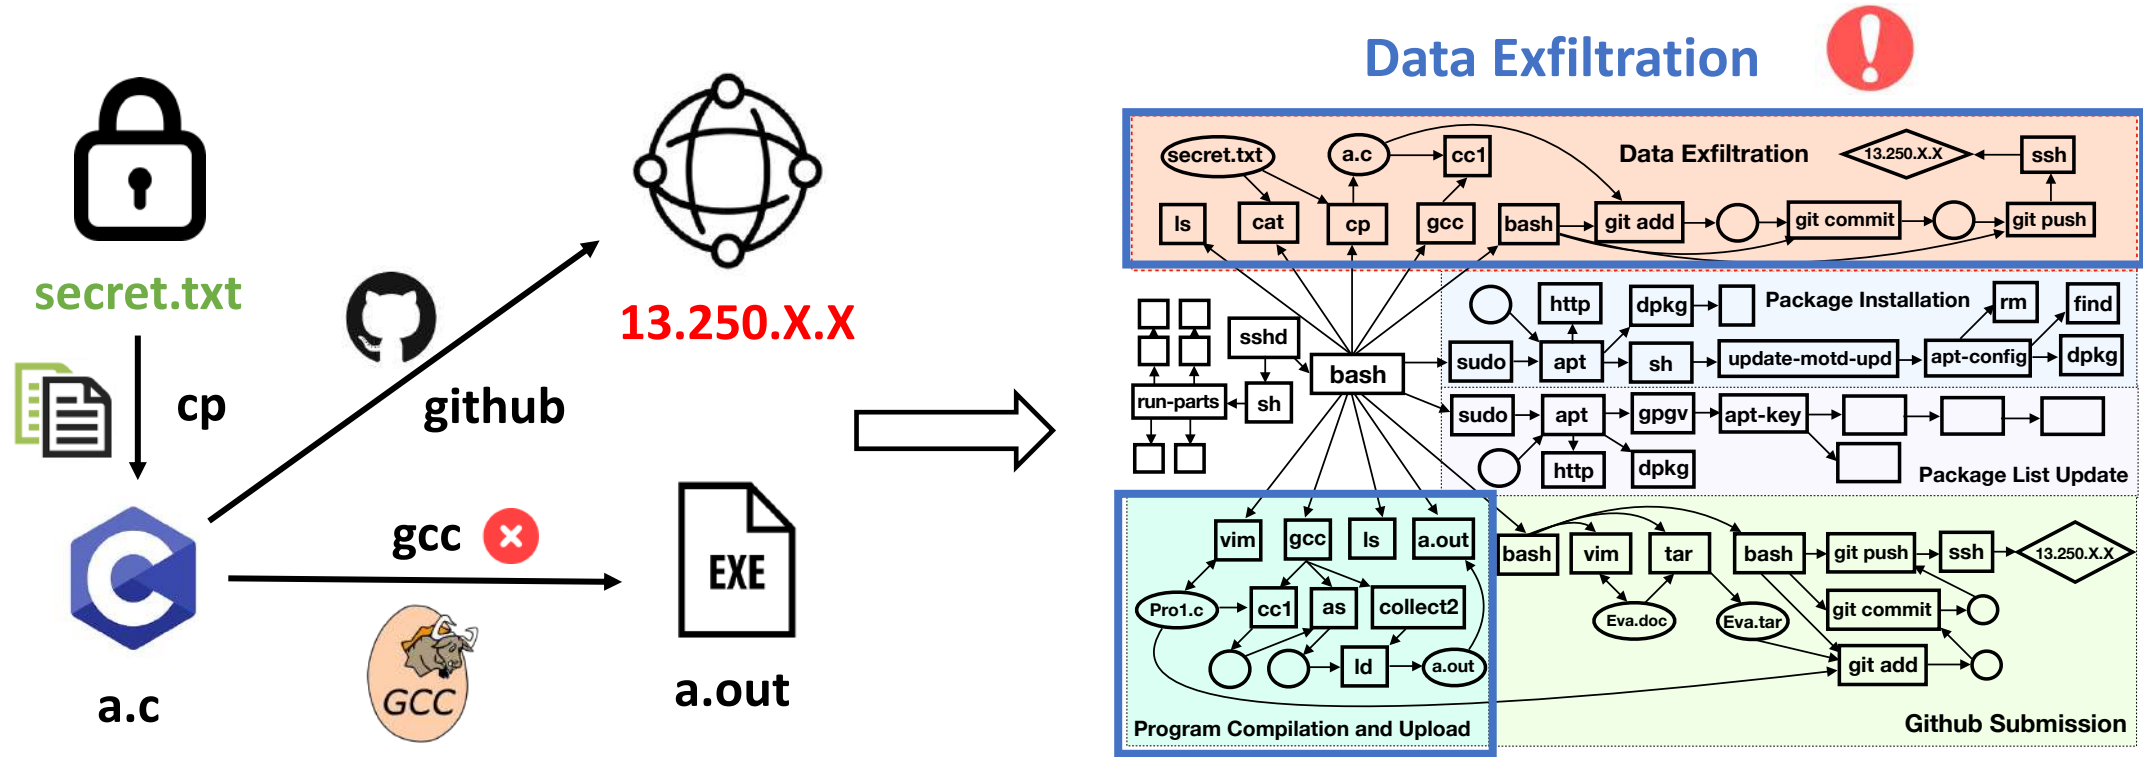
Data Exfiltration Steps



Motivating Example Logs

# Motivating Example

Attack Scenario: A software tester **exfiltrates sensitive data** that he has access to



Data Exfiltration Steps

Program Compiling and Upload (cluster)

Motivating Example Logs

# Challenges for Behavior Abstraction

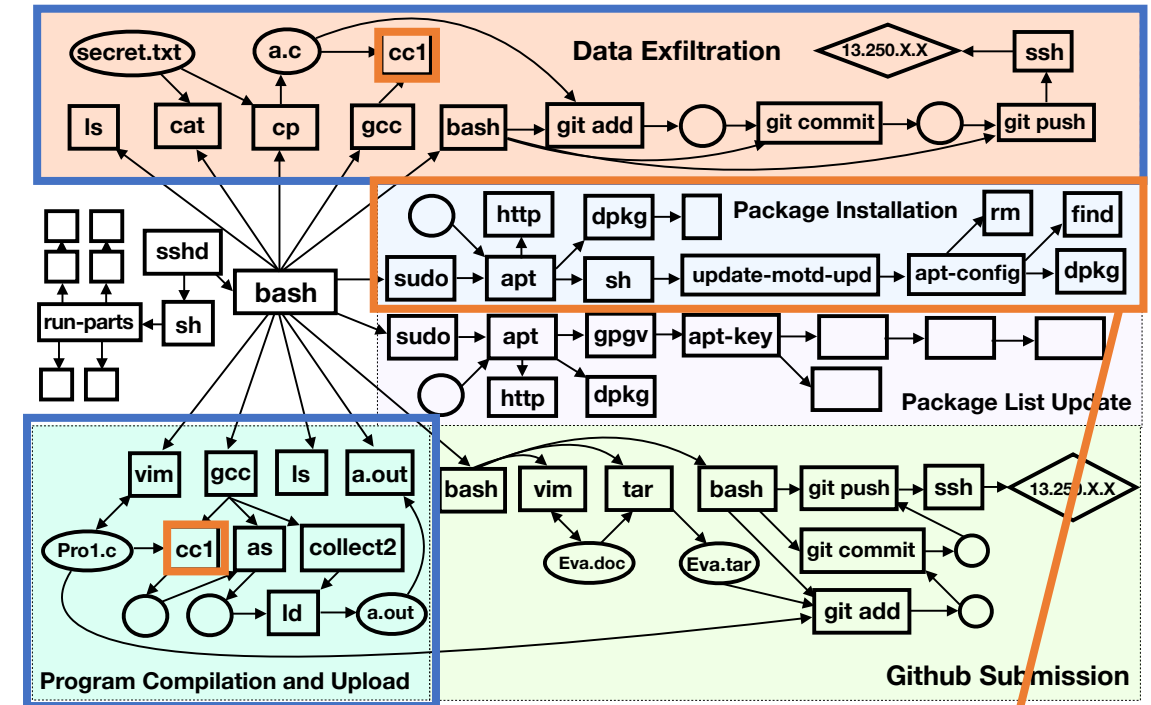
## Data Exfiltration

### Event Semantics Inference:

- Logs record **general-purpose** system activities but lack knowledge of **high-level semantics**

### Individual Behavior Identification:

- The volume of audit logs is **overwhelming**
- Audit events are **highly interleaving**




## Program Compiling and Upload

**Package Installation Events > 50,000**

# Our Insights

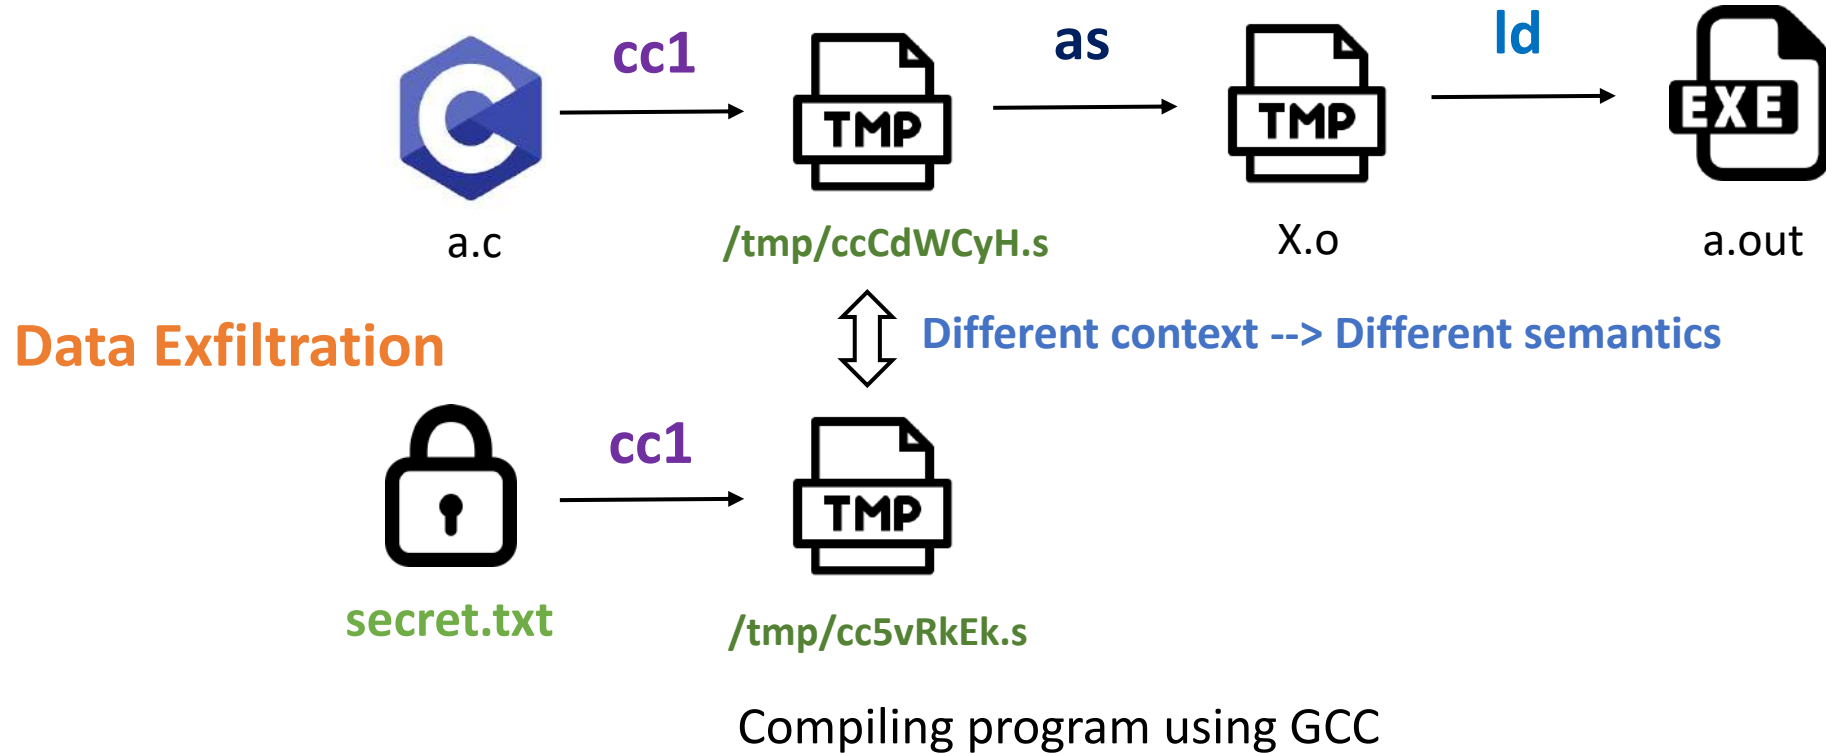
How do analysts manually interpret the semantics of audit events?



Compiling program using GCC

# Our Insights

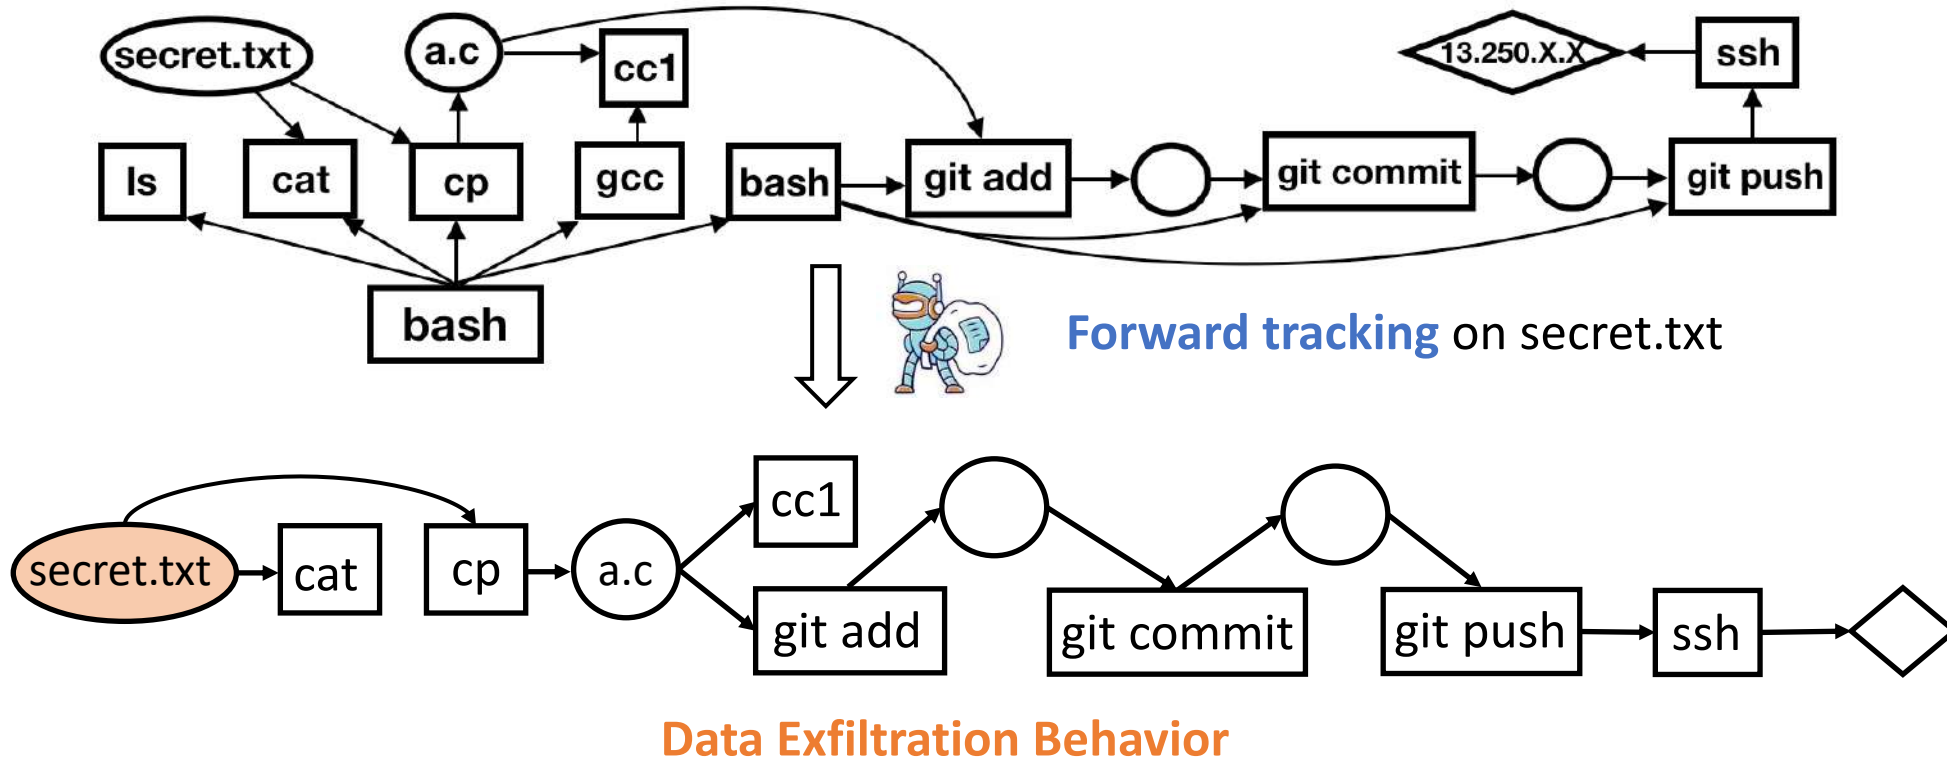
How do analysts manually interpret the semantics of audit events?



Reveal the semantics of audit events from their usage **contexts** in logs

# Our Insights

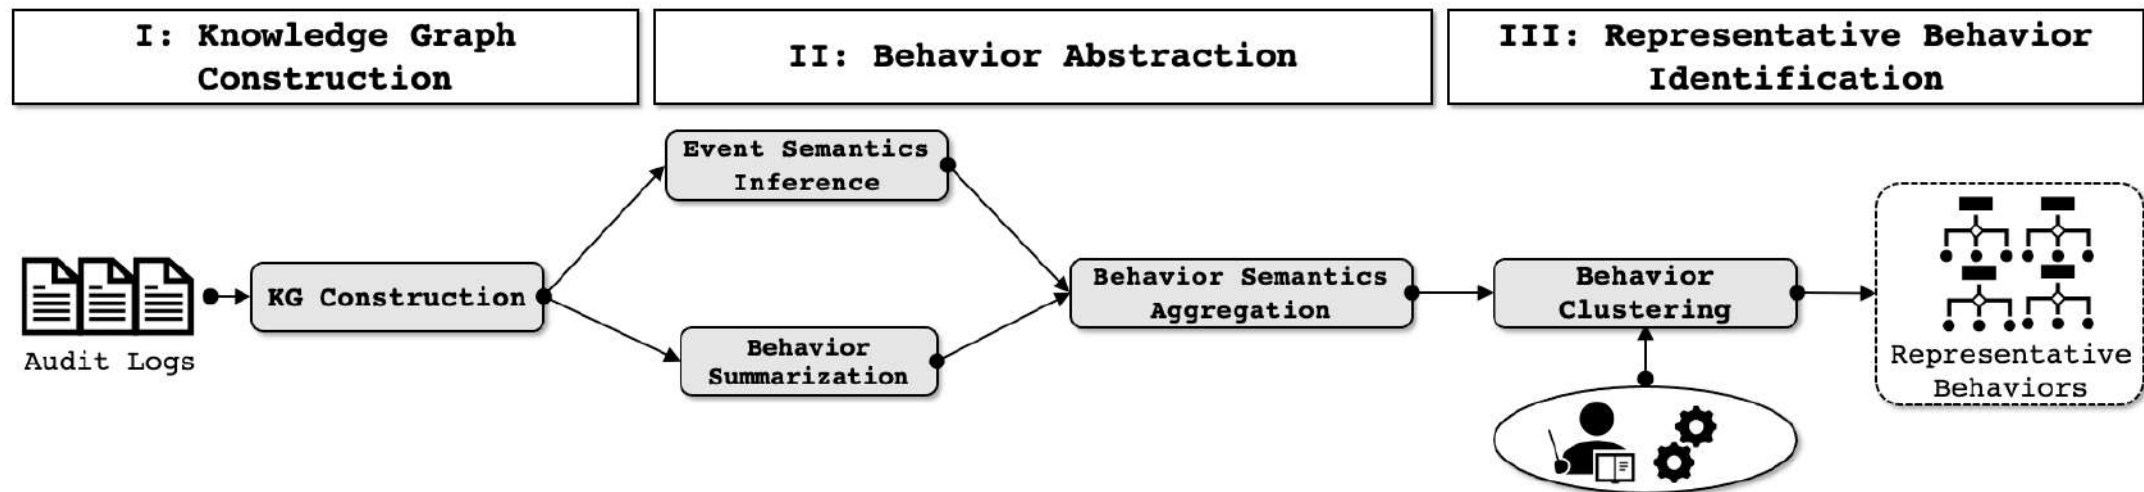
How do analysts manually identify behaviors from audit events?



Summarize behaviors by tracking **information flows** rooted at **data objects**

An automated behavior abstraction approach that **aggregates the semantics of audit logs to model behavioral patterns**

- Input: audit logs (e.g., Linux Audit<sup>[1]</sup>)
- Output: representative behaviors



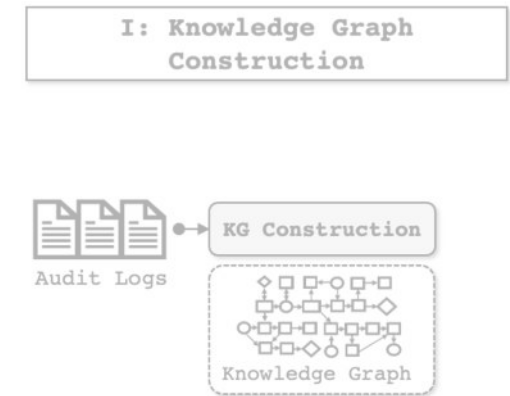
# Knowledge Graph Construction

We propose to use a **knowledge graph** (KG) to represent audit logs:

- KG is a directed acyclic graph built upon triples
- Each triple, corresponding to an audit event, consists of three elements (head, relation, and tail):

$$\mathcal{KG} = \{(h, r, t) | h, t \in \{Process, File, Socket\}, r \in \{Syscall\}\}$$


- KG unifies **heterogeneous** events in a **homogeneous** manner





# Event Semantics Inference


- Suitable **granularity** to capture contextual semantics
  - Prior work [CCS'17] studies log semantics using events as basic units.
  - Lose contextual information within events
  - Working on **Elements** (head, relation, and tail) preserves more contexts
- Employ an embedding model to extract contexts
  - Map elements into a vector space
  - Spatial distance represents semantic similarities
  - **TransE**: a translation-based embedding model
  - **Head + Relation  $\approx$  Tail  $\rightarrow$  Context decides semantics**




# Behavior Summarization

Individual behavior identification: Apply an adapted depth-first search (DFS) to track information flows rooted at a data object:


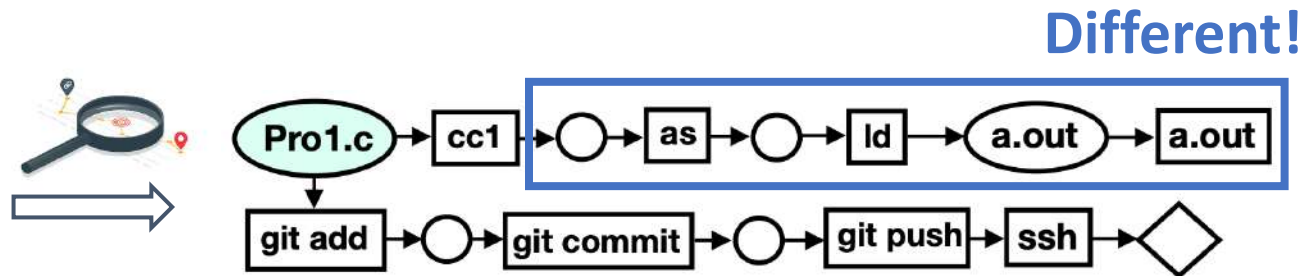
- Perform the DFS on every data object except libraries
- Two behaviors are merged if one is the subset of another




## Program Compiling and Upload

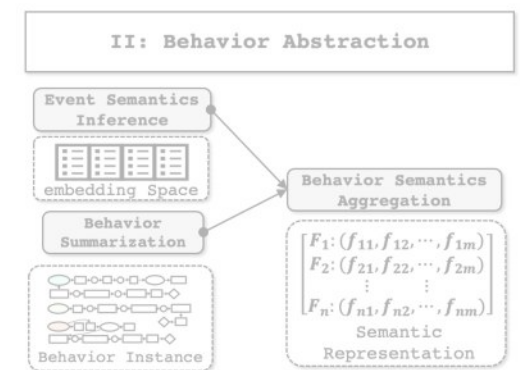


## Data Exfiltration



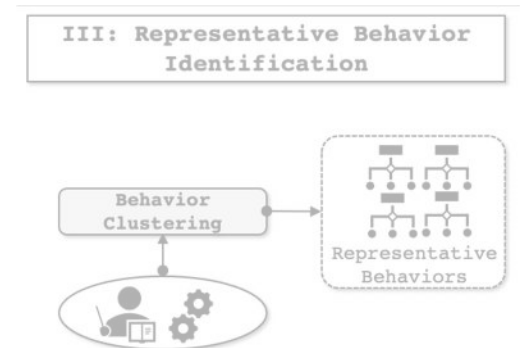
# Behavior Semantics Aggregation

- How to aggregate event semantics to represent behavior semantics?
  - Naïve approach: Add up the semantics of a behavior's constituent events
  - Assumption: audit events equally contribute to behavior semantics 
- **Relative event importance**
  - Observation: behavior-related events are common across behaviors, while behavior-unrelated events the opposite
  - Apply frequency as a metric to define event importance
  - Quantify the frequency: **Inverse Document Frequency (IDF)**
- The presence of **noisy events**
  - Redundant events [CCS'16] & Mundane events



# Representative Behavior Identification

- Cluster semantically similar behaviors: **Agglomerative Hierarchical Clustering analysis (HCA)**
- Extract the most representative behaviors
  - Representativeness: Behavior's average similarity with other behaviors in a cluster
  - **Analysis workload reduction**: Do not go through the whole behavior space



# Summary

- Logging mechanisms
  - Application-level: Library wrapping / API hooking
  - Kernel-level: Syslogd/klogd, System call interception, Linux security module
  - Virtual Machine Monitor-level: System call interception
- Applications for auditing
  - Intrusion detection, recovery and investigation