

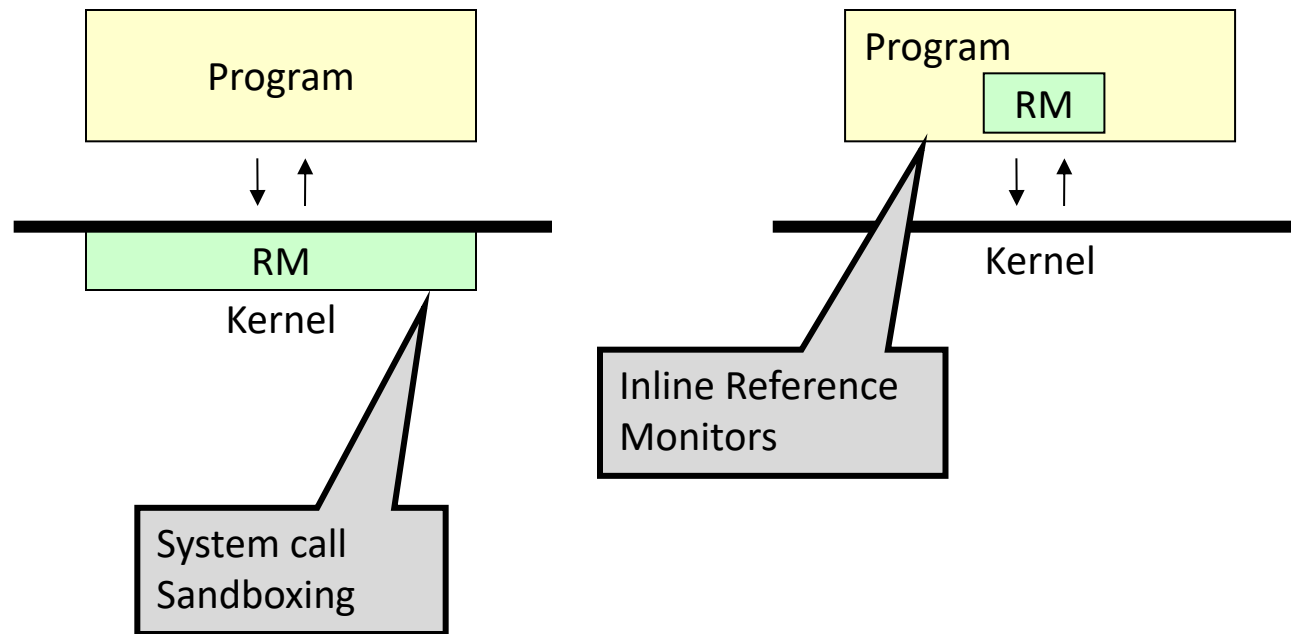
CS5231: Systems Security

Lecture 09a: Virtualization and Trusted Computing Environment

Recap: Reference Monitors

Reference Monitor: A piece of code that *checks all references* to an **object**

Syscall Sandbox: A reference monitor for protecting OS resource objects from an app



Recap: Policy vs. Enforcement Mechanism


- Access Control Policies
- Enforcement:
 - Process sandboxing
 - Inline Reference Monitors
 - *Virtualization*
 - *Hardware-based isolation / Trusted Execution Env.*

Isolation: Virtualization

Problem: Isolated Computation on Shared CPU

Operating
System

Operating
System



sharing of hardware resources allow them to bypass isolation



the normal problem with this being on the same system is that an attack on the game can affect the banking application as well

Defense(I): Virtualization

to escape vm

Game
VM

Banking
VM

Virtual Machine Monitor (VMM)



A Bit of History...

- Virtual Machines – 1960's
 - **Motivation**: Sharing of machines between users
 - Many implementations by IBM
- Virtual Machines on RISC / CISC – Late 90's
 - **Motivation**: Unify under-utilized machines, ease-of-maintenance, security
 - E.g. VMWare
- Heavy utilization in cloud computing...

Public Clouds: EC2

- Virtual Machine Monitor (VMM): **Xen**
- Instance: **A running OS image of virtual machine**
- ECU: **EC2 Compute Unit** \approx 1.2GHz Opteron/Xeon CPU

Amazon CloudWatch

Region:

Detailed Monitoring for Amazon EC2 Instances

- \$3.50 per instance per month, provided at 1-minute frequency

Basic Monitoring for Amazon EC2 instances

- \$0.00 (free of charge) per instance per month, provided at 5-minute frequency

Monitoring for Custom Metrics

- \$0.50 per metric per month

Assumptions

- **Goal:** Isolation of Code, data, resources between:
 - Guest VM and Host VMM
 - Between VMs
- **Assumptions:**
 - Bug-free TCB: Host OS, VMM
 - Malware can affect the guest OS & apps.

Security Applications of Virtualization

- Virtual Machine Isolation
 - Red-Green Systems
 - E.g. Banking VM vs. Normal VM
 - Dynamic Analysis / Containment of Malware
- Virtual Machine Introspection
 - E.g. Run an anti-virus in the VMM

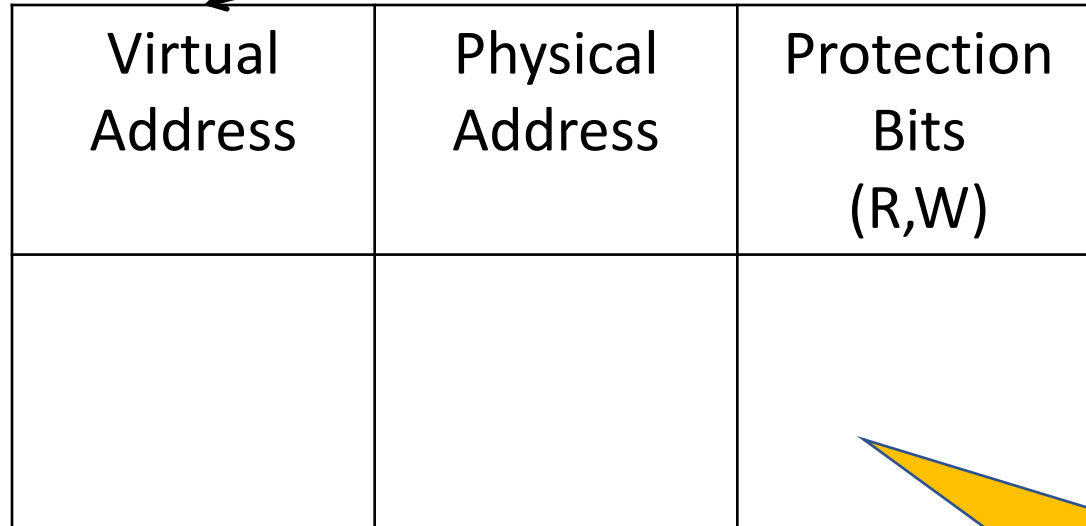
Enforcement Goals for a VMM

- Security VMM Goals:
 - Complete Mediation
 - Trap on all MMU, DMA, I/O accesses
 - Transparency
- Commercial VMM Goals:
 - Performance
 - Compatibility: Run on commodity OSes

Compatibility Challenges: An Example

```
mov eax, (ebp)
```

One of the core vm concept is to hack around with page tables, having their own page tables



A diagram of a page table structure. It consists of a table with two rows and three columns. The columns are labeled 'Virtual Address', 'Physical Address', and 'Protection Bits (R,W)'. An arrow points from the assembly instruction 'mov eax, (ebp)' to the 'Virtual Address' column. The table is currently empty of data.


Virtual Address	Physical Address	Protection Bits (R,W)

Page Table

OS uses to isolate kernel code / data

Compatibility Challenges


Non-virtualized execution



Privileged Instructions

- that trap: e.g. `cli`

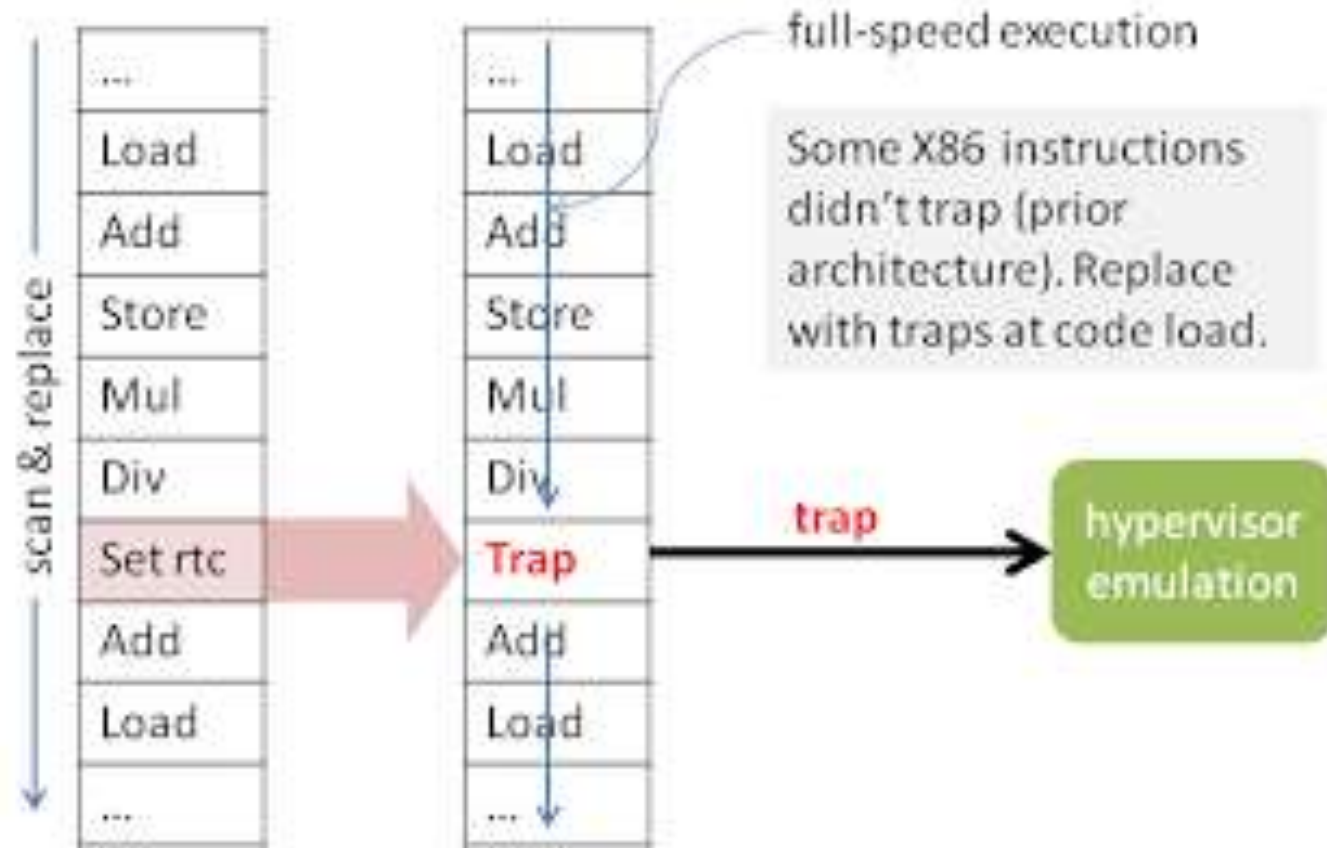
Virtualized execution



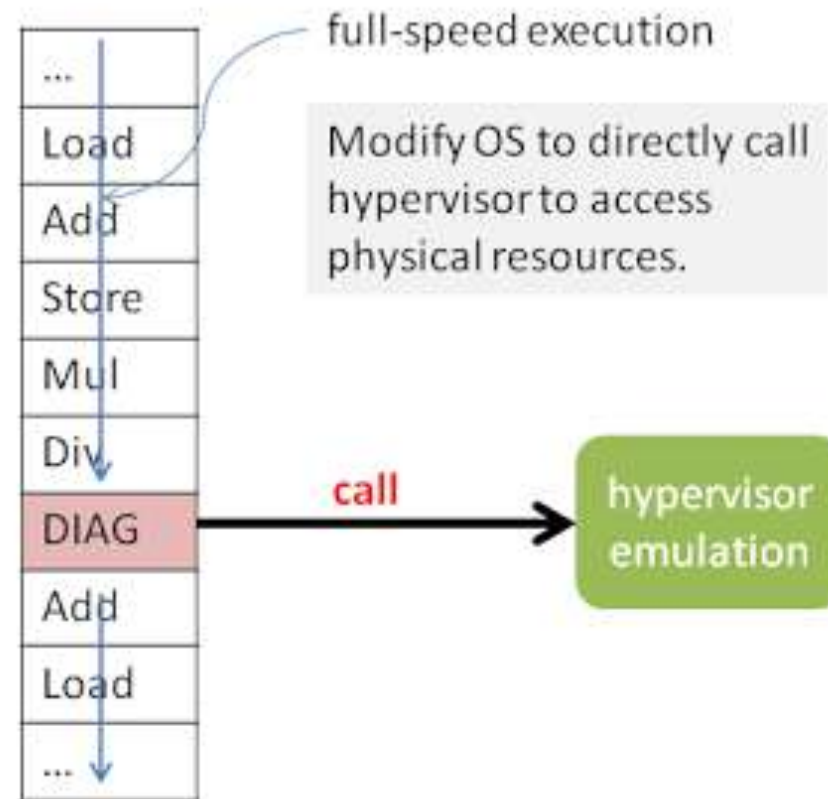
Privileged Instructions

- That don't trap: e.g. `popf`

Virtualization Techniques: Binary Translation (VMware)



Virtualization: Paravirtualization (Xen)

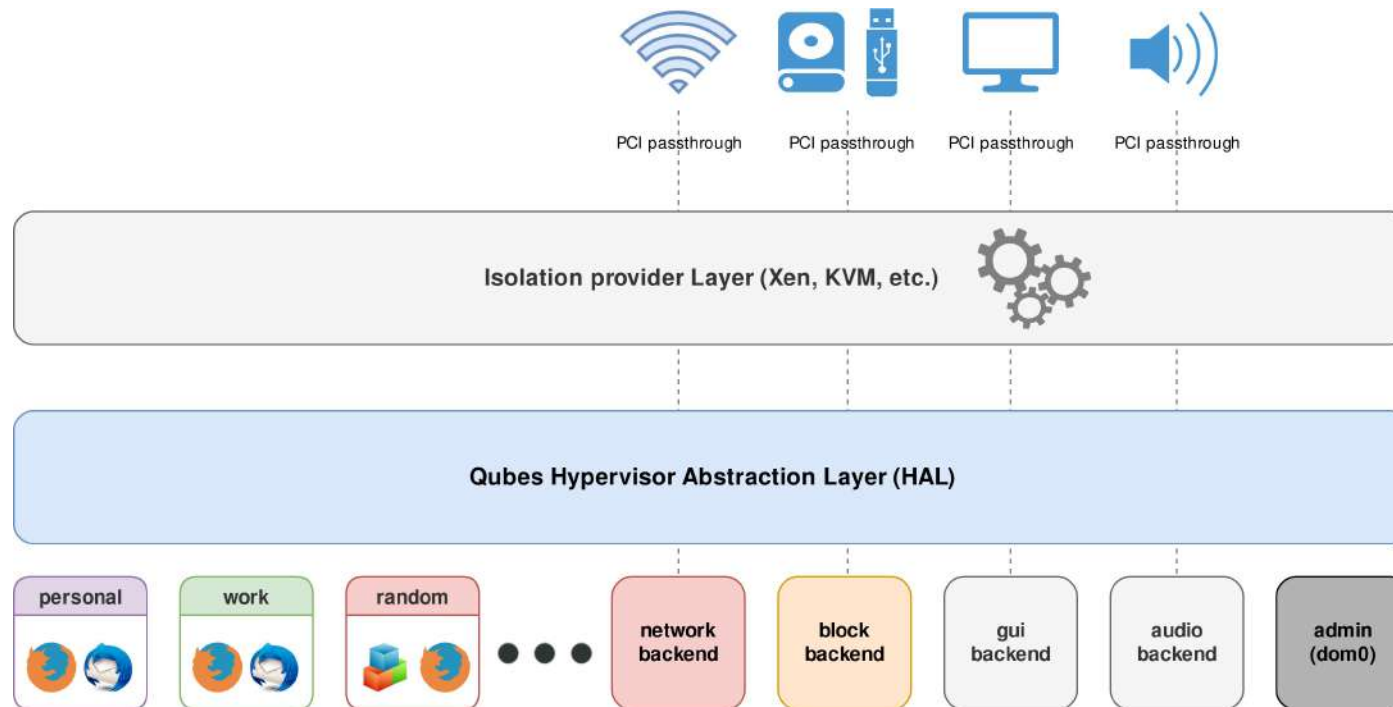


Hardware Assisted Virtualization

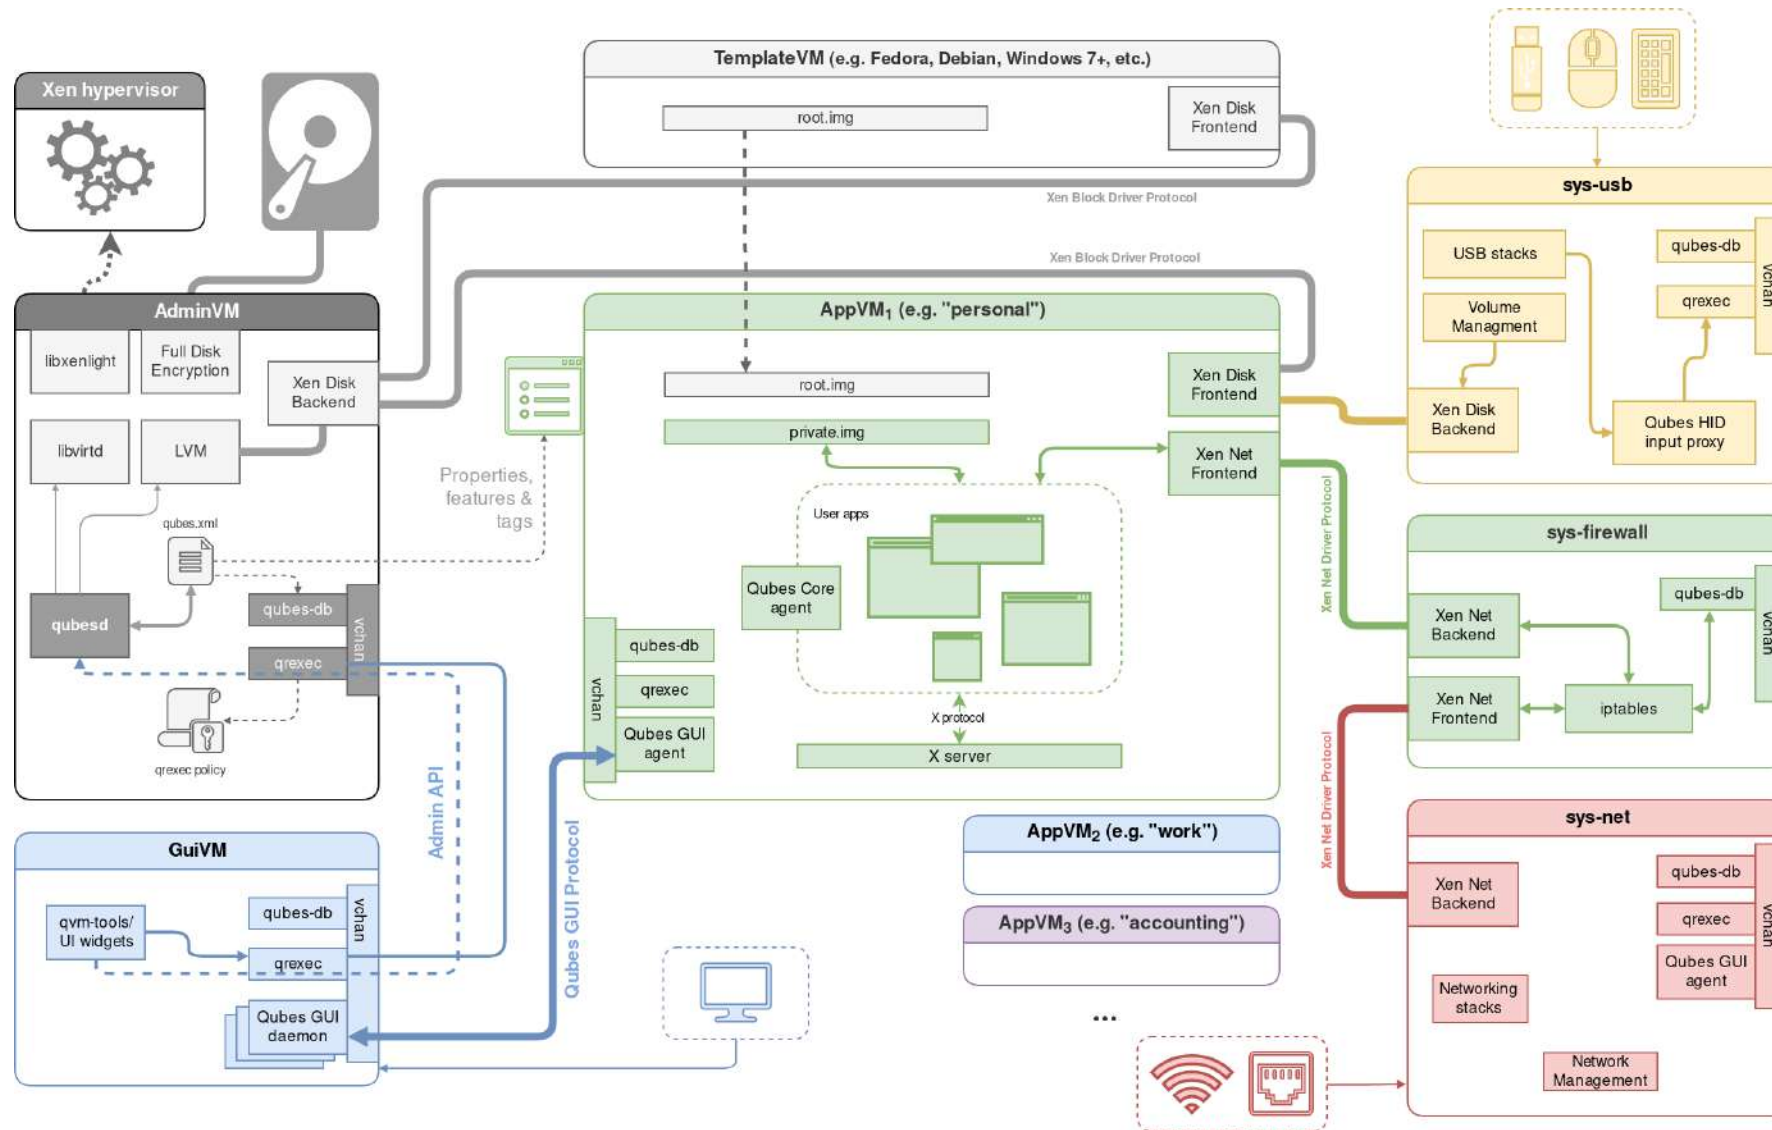
- CPUs adding support over the years
- Goal: Better Performance, Security
 - Intel VT-x
 - MMU virtualization using EPT (2009),
 - Nested virtualization - VMCS (2012)
 - I/O virtualization – IOMMU (2009)
 - Intel VT-d
 - DMA remapping (2009)

Qubes OS

- A reasonably secure operating system
 - A network of virtual machines in a computer




Qubes OS Architecture




Limitations of Virtualization

Virtual Machine Based Rootkits



The Problem of Secure “Root of Trust”: Is highest layer of privilege malicious?



Implication on Malware Containment

- In principle, is some containment possible?
 - Yes, When the highest layer of privilege is trusted
 - E.g. the VMM is trustworthy

Virtual Machine Based Rootkits: Can the software know its virtualized?

- Blue Pill: “blissful ignorance of illusion”
- Red Pill: “Detects you are virtualized”




VMM Detection: The Red Pill

- Red Pill: “Detects you are virtualized”
- Ways to achieve a “red pill” attack:
 - Commercial VMMs aren’t fully transparent
 - E.g. VmWare emulates i440bx chipset (old)
 - Virtualization Timing latencies Measurements
 - Many other measurement channels [[HotOS’07](#)]
- Applications of VM Detection:
 - Malware can detect introspection software (e.g. AV)
 - Can utilize “Anti-VM” techniques
 - Benign use: Copy protection by VM duplication

Implication on Malware Containment

- In principle, is some containment possible?
 - Yes, When the highest layer of privilege is trusted
 - E.g. the VMM is trustworthy
- Detecting virtualization is easy
 - It can thwart malware analysis (introspection)



Can attacker leak secret document to evil.com?

The Problem of Covert Channels

- Definition: “An unintended channel of communication between 2 untrusted programs”
- E.g. Shared Cache Latency
 - Sender
 - Send bitval 1: Perform random memory access
 - Send bitval 0: Do nothing
 - Receiver
 - Rcv bitval 1: If long read time for a fixed memory loc.
 - Rcv bitval 0: If short read time for fixed memory loc.
- Can get 0.02 bits/sec on Amazon EC2 [CCS'09]
- Many channels: Disk, I/O, Virtualization latency, ...

Implication on Malware Containment


- In principle, is some containment possible?
 - Yes, When the highest layer of privilege is trusted
 - E.g. the VMM is trustworthy
- Detecting virtualization is easy
 - It can thwart malware analysis (introspection)
- Which containment using VMs is possible: integrity vs. confidentiality?
 - Yes, for Integrity policy – I.e., protecting contained malware from corrupting benign data outside the VM
 - No, for confidentiality, covert channels are a problem

Optional Reading Material for the curious...

- [SubVirt: Implementing malware with virtual machines](#)
- [Overshadow: A Virtualization-Based Approach to Retrofitting Protection in Commodity Operating Systems](#)
- [HyperSafe: A Lightweight Approach to Provide Lifetime Hypervisor Control-Flow Integrity](#)
- [Hey, You, Get Off of My Cloud: Exploring Information Leakage in Third-Party Compute Clouds](#)

Trusted Execution Environments: The Basic Idea

The Problem of “Root of Trust”: Is highest layer of privilege malicious?

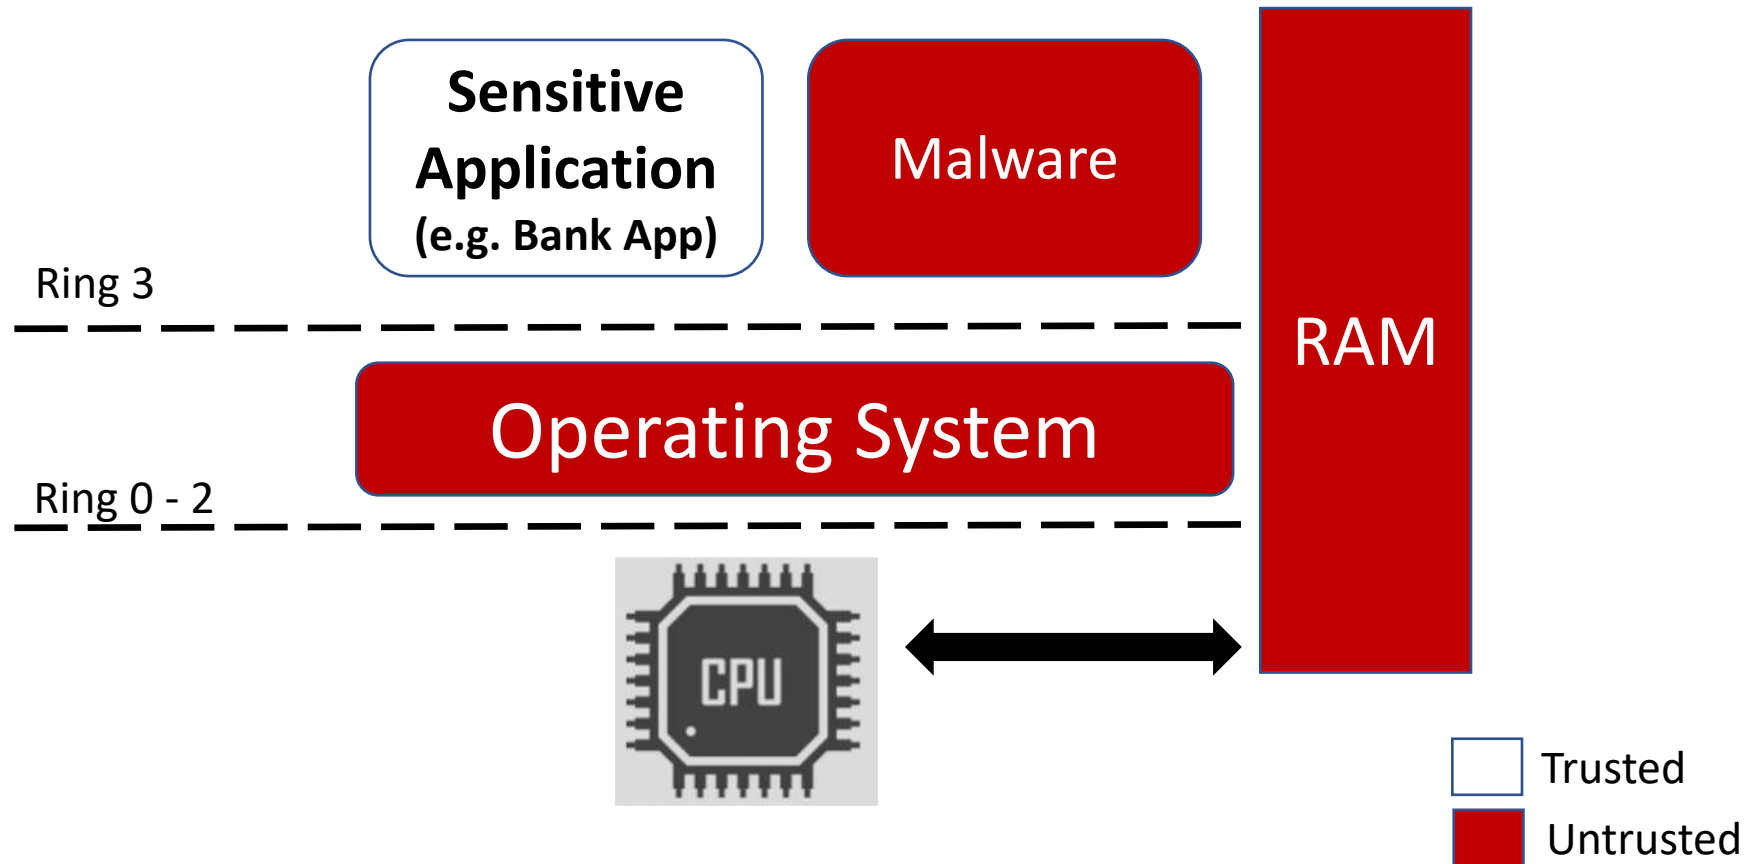


Solution: TEEs ensure Secure “Root of Trust”

- Trusted Execution Environments (TEEs)
 - A **hardware root-of-trust**
 - Can assume that all software is malicious
- Why trust hardware?
 - Tamper-resistant from all software malware
 - Perhaps less complex, easier to verify?

Trusted Execution Environment (TEE): Security Model

- Trust the hardware
- Don't trust other software on the system



To Continue

- Latest trusted computing technology