

CS5231: Systems Security

Lecture 4: Advanced Memory Error Defenses

Control-flow Hijacking: Code Injection

- Control-oriented a.k.a control-flow hijacking
- Outcome 1: Code Injection
 - **Definition:** *A memory exploit that hijacks control to jump to attacker's data payload*
- Req 1: Write Attack Payload in memory
- Req 2: Have Attack Payload Be Executable
- Req 3: Divert control-flow to payload

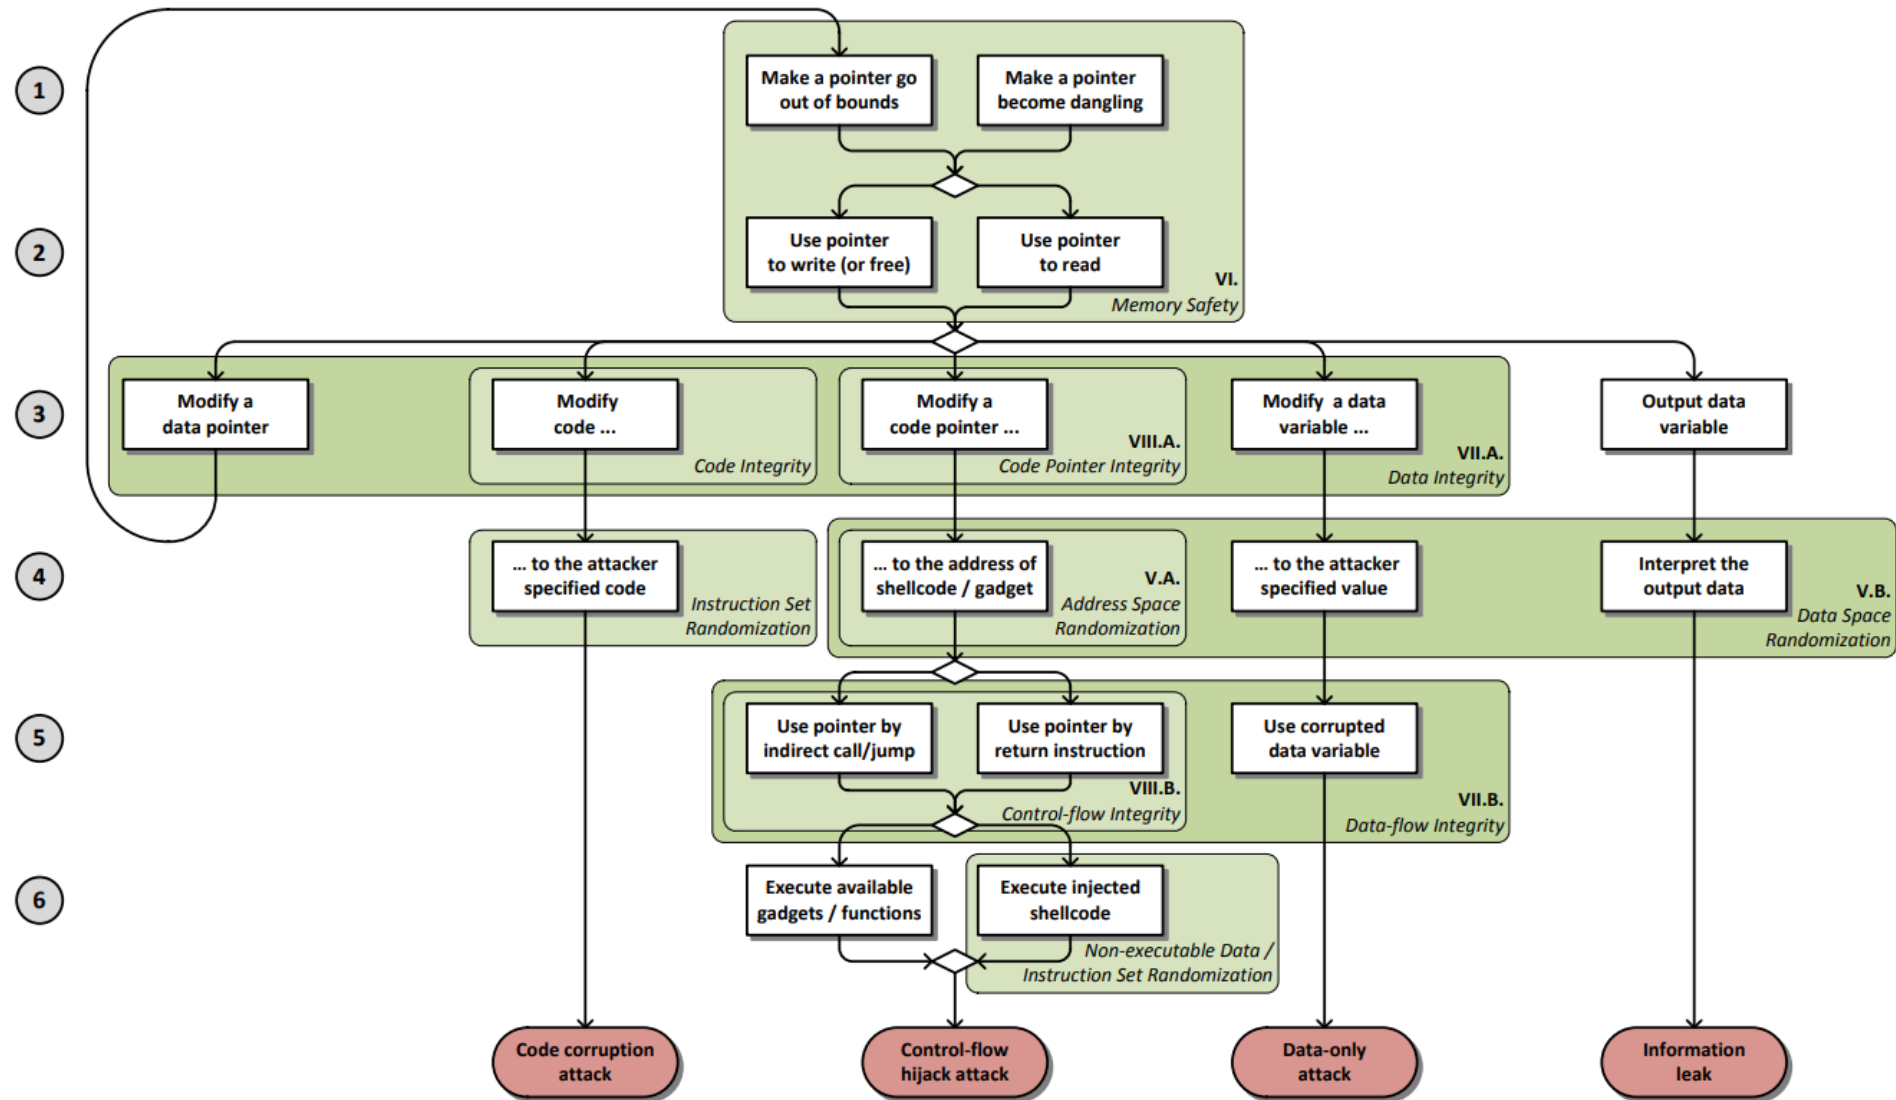
Control-oriented Exploits (II): Code Reuse

- Outcome 2: Code Reuse
 - **Definition:** *A memory exploit that hijacks control to jump to attacker's controlled code address*
- Requirements for Code Reuse
 - ~~Req 1: Write Attack Payload in memory~~
 - Req 2: Have Attack Payload Be Executable
 - Req 3: Divert control-flow to payload
- Insight: Re-use the existing code as payload

Data-oriented Exploits

- Don't need any execution of illegitimate code
- Requirements for Code Reuse
 - ~~Req 1: Write Attack Payload in memory~~
 - ~~Req 2: Have Attack Payload Be Executable~~
 - ~~Req 3: Divert control-flow to payload~~
- Insight: changing data to affect the computation done by a program

Taxonomy of Safety Properties



Memory Error Defense Summary

- Safe coding practice
- Randomization
 - Address-space randomization, data-space randomization, instruction set randomization
- Partial memory safety
 - StackGuard, stack canaries
 - Non-executable data/DEP
- Full memory safety

e.g. XOR / encrypt with a key


Full Memory Safety

Definition: Memory Safety

- Goals:
 - Create memory pointers via permitted operations
 - E.g. malloc(), p = &q;
 - Only access memory allocated to the pointer
 - Spatially → within the allocated range
 - Temporally → while the memory is in scope
 - All “objects” are spatially disjoint at all times
- Enforcement:
 - Can be done by ^(new compiler) compilation or ^(after compilation) binary rewriting
 - Insert metadata & inline reference monitors

Spatial Safety

1. Distinguish pointers from non-pointers
2. Check object allocation
3. Check each pointer access
 - Recall: Pointers can be incremented, type cast, etc.

- Proposals: *alternative: low-fat pointers*
 - Fat pointers 
 - Shadow-memory data structure [e.g. [JK-Tree](#)]
 - Encode the size information in pointer value [[BB](#)]
- Overheads: About 30% or more (SPEC)
- Hardware support: Intel MPX

if enforced, hopefully the memory will not be overwritten.

Temporal Safety

1. Track creation and destruction of pointers
2. Ensure: De-allocated pointers are not accessed

- A Proposal:

- Lock-and-key [[CETS](#)]



- The key K and lock L will match only if P is live
 - When de-allocated, change the key K
 - Where to store the lock & key?

- Fat pointers, shadow memory data structure...

- Canary-based defenses: Set to NULL on de-allocate [[DN](#)]


- Overheads:

- About 50% or more for lock-n-key
 - Unclear, but could be about 10% for canary-based

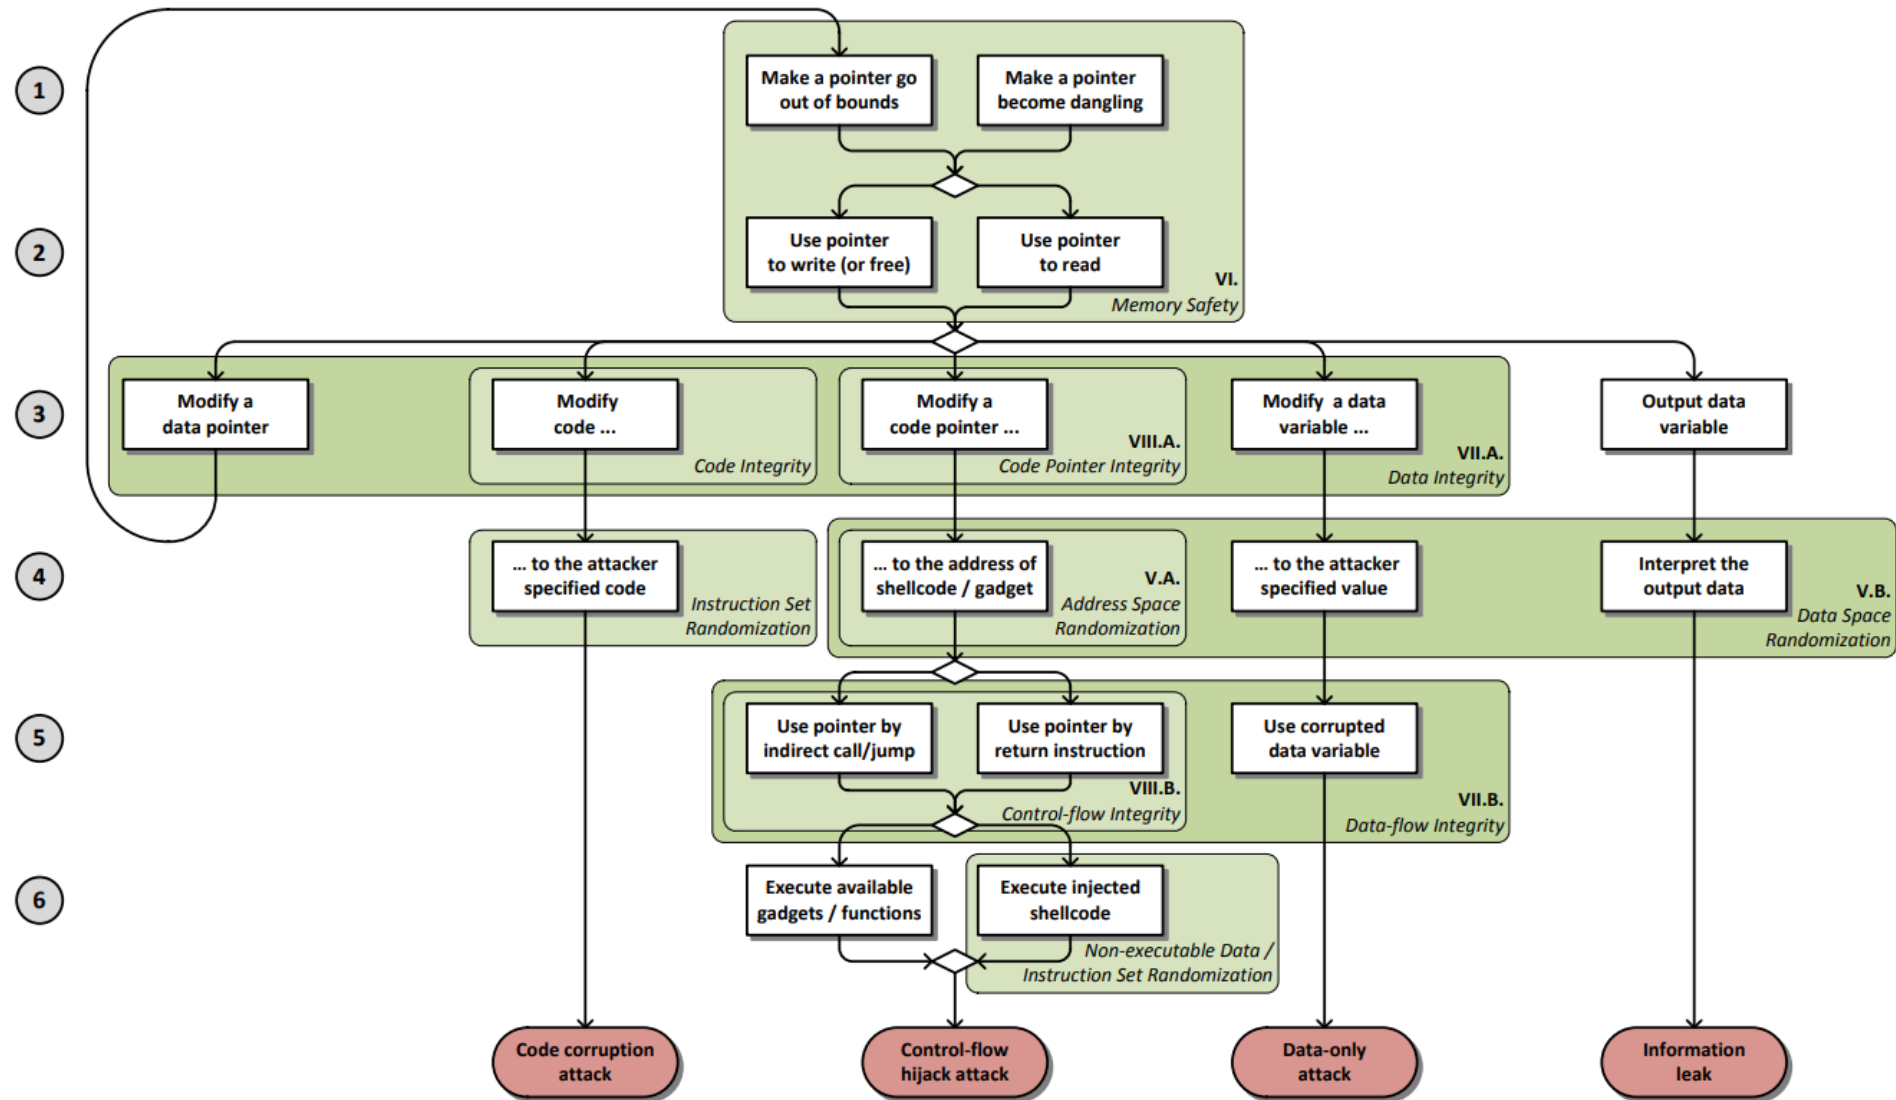
Approaches to Enforce Memory Safety: Static vs. Dynamic

- Statically disallow
 - Type casts
 - Unchecked buffer accesses
 - Pointer Arithmetic
 - Explicit Alloc / Dealloc
 - Examples:
 - Memory-safe languages
 - Safe C subsets (e.g. Cyclone)
- Dynamically check
 - Spatial Errors
 - Temporal Errors
 - Check Type casts?
 - Need to track type-info at runtime
 - Examples:
 - See previous slides

Approaches to Enforce Memory safety: Static + Dynamic



Summary of Memory Defenses



Control Flow Integrity

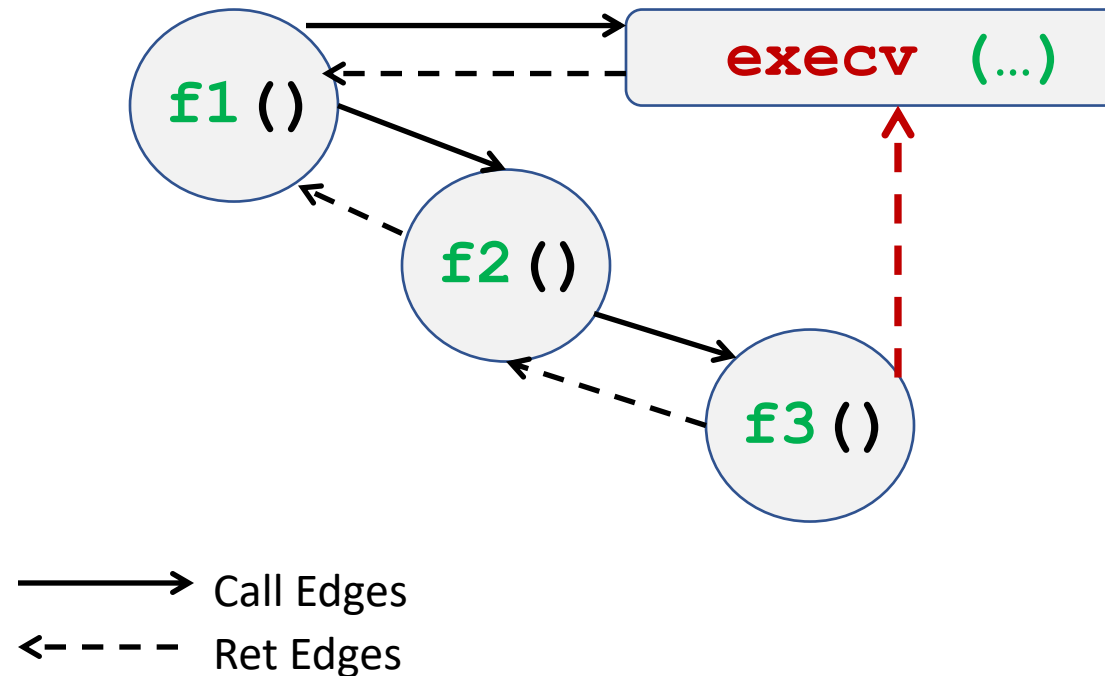
Control-flow Integrity

- Goal of CFI enforcement:

- **Control Flow Integrity**


"Follow the statically determine CFG at runtime"

Control
Flow
Graph



Control Flow Integrity

- **Definition of Control Flow Integrity**
 - Each control transfer jumps to a **statically-known set of locations** *Check if return address is a known static address before jumping*
 - E.g. Returns -> Return points, Call Instructions -> calls
- CFI blocks all control-flow hijacking exploits



Control Flow Graph



CFI: Theoretical power


- Goal of CFI enforcement:
 - **Control Flow Integrity**
"Follow the statically determined CFG at runtime"
- Can block all control-flow hijacking attacks!

Legitimate Call Sequences:

A -> B -> D

A -> C -> D

A -> B -> C



Illegitimate Call Sequences:

A -> C -> B -> D *(maybe missing condition)*

A->D

CFI can protect against all such bad call sequences

Inline Reference Monitors (II): CFI – Implementation 1

- Goal of CFI enforcement:
 - **Control Flow Integrity**
“Follow the statically determined CFG at runtime”

```
jmp ecx ; computed jump
```

```
cmp ecx, 0x80480aa ;  
jne error_label ;  
..... ;  
jmp ecx ; jump to dst
```

Control-flow Integrity: Return Edges?

- A function can have several callers....
- If a small set of return points
 - Instrument code to enforce return target
- If a large possible set of return points
 - Use a shadow stack!
 - Shadow stack can be protected by SFI

Software
Fault
Isolation.

CFI Implementations

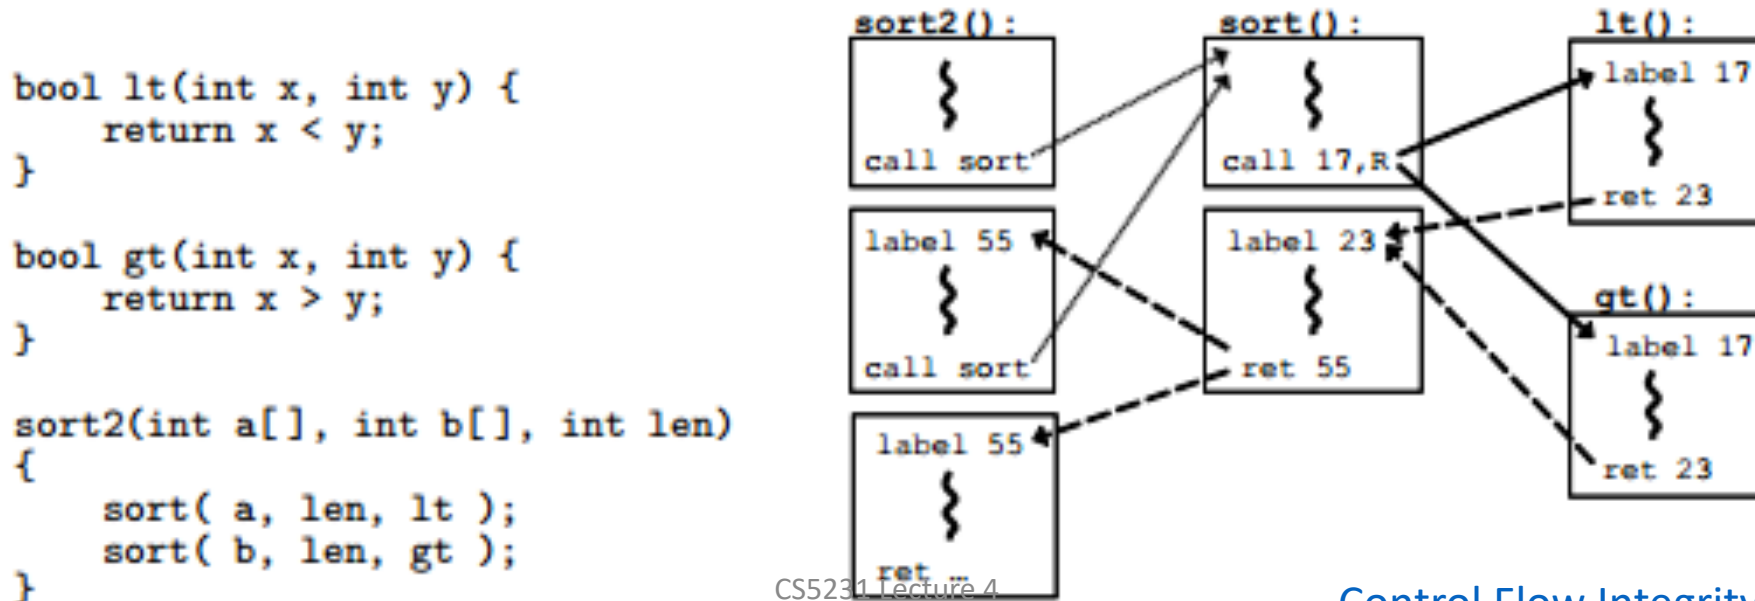
- Can we do faster than CFI-1?

```
jmp ecx ; computed jump
```

```
cmp ecx, 0x80480aa ;  
jne error_label    ;  
.....            ;  
jmp ecx            ; jump to dst
```

CFI Implementation With Randomized Tags

- Each code block must start with a tag
 - The tag should be a random, secret value
 - If f can jump to block g, h, \dots
 - Then blocks g, h, \dots should have the same tag




CFI Implementation With Randomized Tags

- Each code block must start with a tag
 - The tag should be a random, secret value
 - If f can jump to block g,h,...
 - Then blocks g,h,... should have the same tag

check tag without storing
known locations in memory.
No checking list at every return.

code block




```
jmp ecx ; computed jump
```

```
cmp [ecx], 12345678h ; comp ID & dst
jne error_label        ; if != fail
lea ecx, [ecx+4]        ; skip ID at dst
jmp ecx                ; jump to dst
```

CFI, In Practice

- Powerful in theory, but...
 - It is challenging to recover the precise control flow graph at compile time
- Why?
 - Do we know what will a function pointer point to?
 - Pointer analysis:
 - Theoretically is undecidable
 - Practically is a difficult problem for real-world programs
- Implemented in [LLVM – Clang](#)
- Implemented in [Microsoft V. Studio compiler](#)


Taxonomy of Safety Properties



Pointer Integrity

Pointer Integrity

- Runtime Property:
 - Pointers should point to valid addresses only
- Code Pointer (Definition):
 - Rule 1: A pointer that can be legally point to code
 - Rule 2: Pointers that legally can point to pointers of Type 1, by transitively dereferencing and legal arithmetic operations



Code Pointer Integrity

Code Pointer Integrity (CPI) Defeats CI

- **Definition** of Code Pointer Integrity :
 1. *Enforce that code pointers point to code-segment only!*
 2. *Enforce that control transfers use code pointers.*
- Recall, the requirements for code injection:
 - Req 1: Write Attack Payload in memory
 - Req 2: Have Attack Payload Be Executable
 - Req 3: Divert control-flow to payload
- Rule 1 of CPI defeats requirement 3
 - Code segment is not writable
 - Enforcement Details: [CPI Paper \(OSDI'14\)](#)

Protecting Code Pointers

- Examples of Code pointers:
 - Return Address Storage
 - Jump Tables / Global Offset Tables
 - Function Pointers
 - Virtual Method Tables (e.g. C++ classes)

```
mov ecx, 0x4[esp]  
call [ecx]
```

```
mov edx, 0x14[esp]  
jmp [edx]
```

```
ret
```

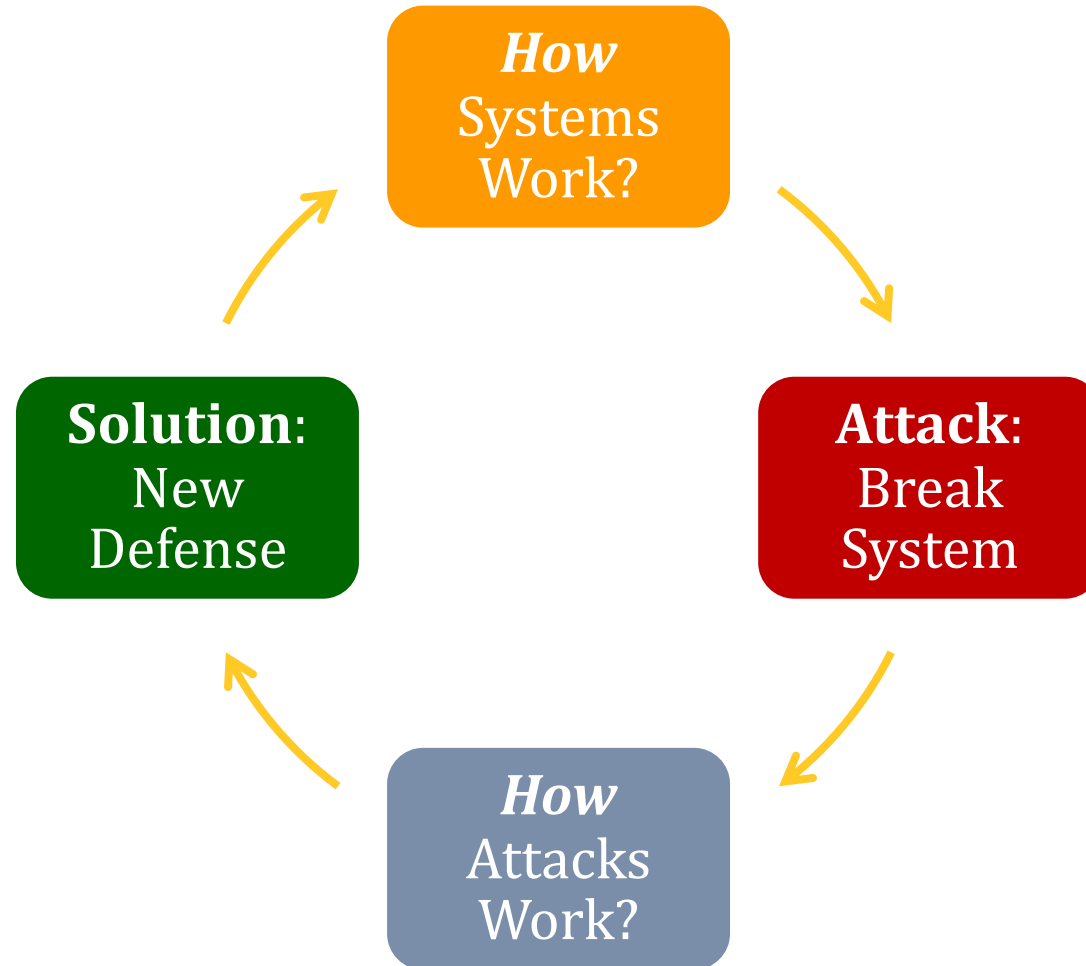
- *Code Pointer Corruption:*
 - *Forging the runtime value of a code pointer to an “invalid” one!*
 - Valid value: A value that is possible under a memory safe execution of the program

Data & Code Pointer Integrity

Is a research topic...

- One approach: **Pointer authentication**
 - Available in ARM processors as hardware primitive
- The basic idea:
 - Cryptographically bind a pointer address to its legitimate value when it is created
 - When legitimate instructions use this value, they can check whether the value has been tampered

Arms Race between Attackers and Defenders



A New Round in Arms Race: Data-Oriented Programming

Non-Control Data Attacks

- Corrupt/leak several bytes of **security-critical data**

```
//set root privilege *  
setuid(0);  
.....  
//set normal user privilege  
setuid(pw->pw_uid);  
//execute user's command
```

```
//offset depends on IE version +  
safemode = *(DWORD *)  
                (jsobj + offset);  
if(safemode & 0xB == 0) {  
    Turn_on_God_Mode();  
}
```

- Special cases relying on particular data/functions
 - user id, safemode, private key, etc
 - interpreter – printf(), etc
- *What is the expressiveness of general non-control data attacks?*

specific

trivial-to-prevent

* Shuo Chen, Jun Xu, Emre C. Sezer, Prachi Gauriar, and Ravishankar K. Iyer. Non-Control-Data Attacks Are Realistic Threats. In USENIX 2005.

+ Yang Yu. Write Once, Pwn Anywhere. In Black Hat USA 2014


Motivating Example

```
1 struct server{int *cur_max, total, typ;} *srv;
2 int quota = MAXCONN; int *size, *type;
3 char buf[MAXLEN];
4 size = &buf[8]; type = &buf[12]
5 ...
6 while (quota-->0) {
7     readData(sockfd, buf); // stack bof
8     if(*type == NONE ) break;
9     if(*type == STREAM)
10         *size = *(srv->cur_max);
11     else {
12         srv->typ = *type;
13         srv->total += *size;
14     } //...(following code skipped)...
15 }
```



```
1 struct Obj {struct Obj *next; int prop;}
2
3 void updateList(struct Obj *list, int addend){
4     for(; list != NULL; list = list->next)
5         list->prop += addend;
6 }
```

Vulnerable
Program

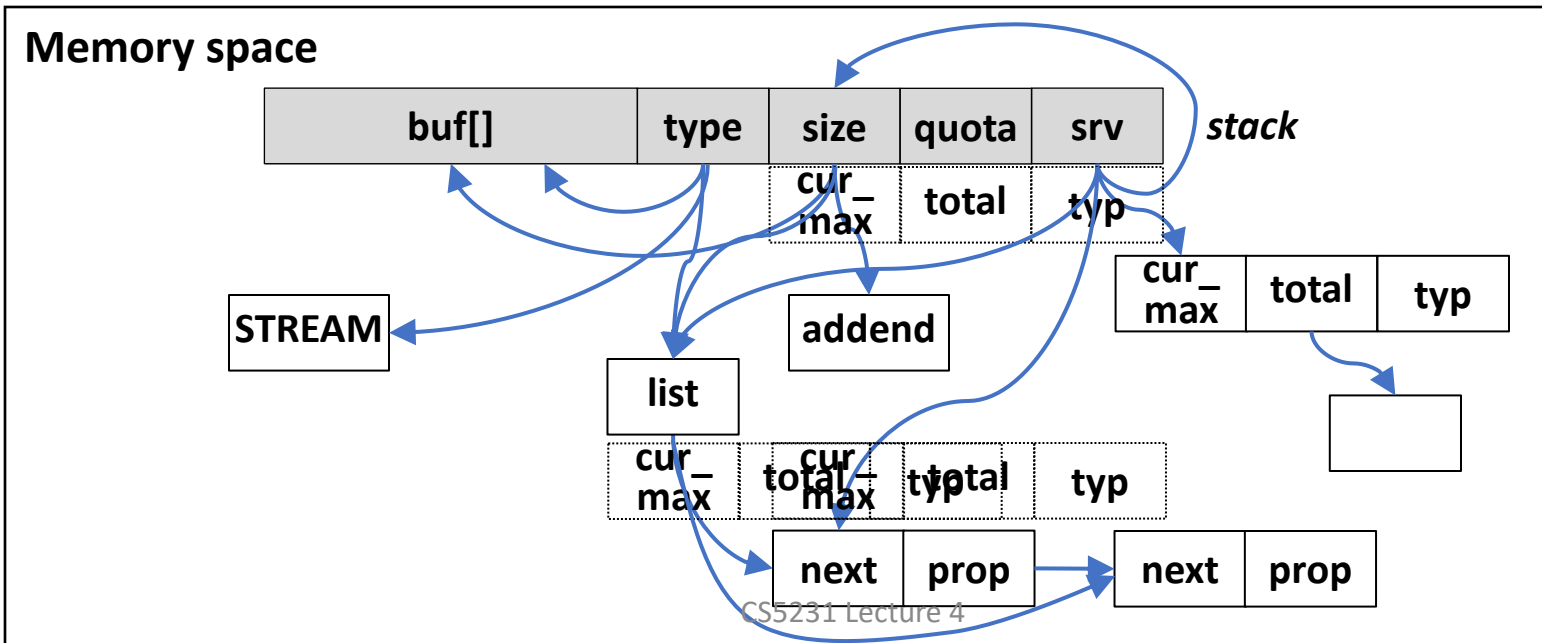


Expected
Computation

Motivating Example (cont.)

```
6 while (quota-- ) {  
7   readData(sockfd, buf);  
8   if(*type == NONE ) break;  
9   if(*type == STREAM)  
10    *size = *(srv->cur_max);  
11   else {  
12     srv->typ = *type;  
13     srv->total += *size;  
14   }  
15 }
```

```
4 for(; list != NULL; list = list->next)  
5   list->prop += addend;
```




Data-Oriented Programming (DOP)

- General construction
 - w/o dependency on security-critical data / functions
- Expressive attacks
 - towards Turing-complete computation
- Rely on data-oriented gadgets & dispatchers

```
6  while (quota--) {
7      readData(sockfd, buf);    //stack bof
8      if(*type == NONE ) break;
9      if(*type == STREAM)
10         *size = *(srv->cur_max);
11     else {
12         srv->typ = *type;
13         srv->total += *size;
14     } //...(following code skipped)...
15 }
```

Data-Oriented Gadgets

- x86 instruction sequence
 - Shown in normal execution
 - Simulating registers with memory
 - **Load** *micro-op* --> **Semantics** *micro-op*
op --> **Store** *micro-op*




Addition: `srv->total += *size;`

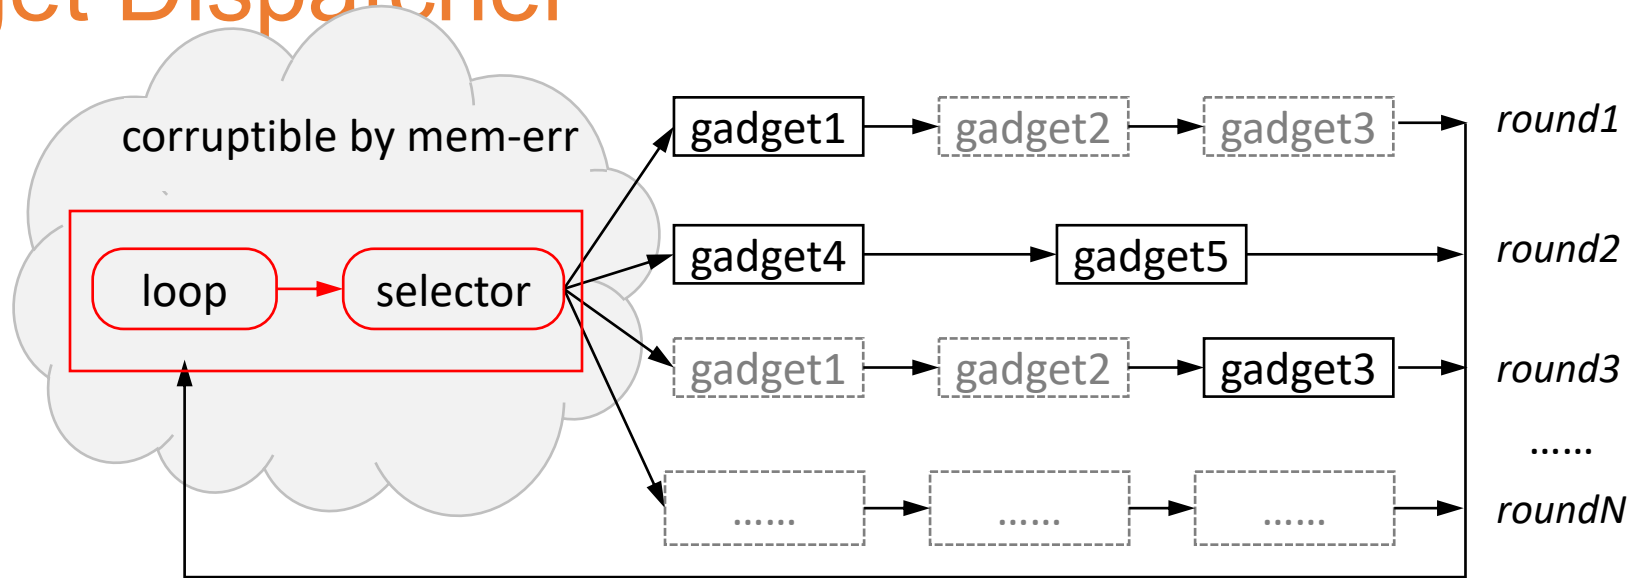
```
1  mov (%esi), %ebx    //load micro-op
2  mov (%edi), %eax    //load micro-op
3  add %ebx, %eax      //addition
4  mov %eax, (%edi)    //store micro-op
```

Load: `*size = *(srv ->cur_max);`

```
1  mov (%esi), %ebx    //load micro-op
2  mov (%edi), %eax    //load micro-op
3  mov 0xb(%ebx), %eax //load
4  mov %eax, (%edx)    //store micro-op
```



Gadget Dispatcher



- Chaining data-oriented gadgets
 - **Loop** ---> repeatedly invoke gadgets
 - **Selector** ---> selectively active gadgets

```
6  while (quota--) {                               // loop
7      readData(sockfd, buf);                       // selector
8      if(*type == NONE ) break;
9      if(*type == STREAM) *size = *(srv->cur_max);
10     else{ srv->typ = *type;  srv->total += *size; }
14 }
```

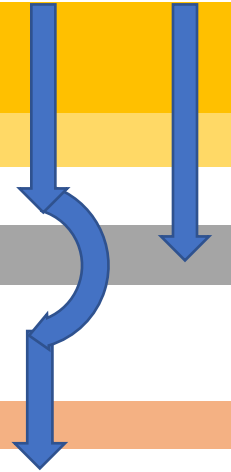
Turing Completeness

- DOP emulates a minimal language *MINDOP*
 - *MINDOP* is Turing-complete

Semantics	Statements In C	Data-Oriented Gadgets in DOP
arithmetic / logical	a op b	*p op *q
assignment	a = b	*p = *q
load	a = *b	*p = **q
store	*a = b	**p = *q
jump	goto L	vpc = &input
conditional jump	if (a) goto L	vpc = &input if *p
p – &a; q – &b; op – any arithmetic / logical operation		

Attack Construction

```
6  while (quota-- ) {  
7      readData(sockfd, buf);  
8      if(*type == NONE ) break;  
9      if(*type == STREAM)  
10         *size = *(srv->cur_max);  
11     else {  
12         srv->typ = *type;  
13         srv->total += *size;  
14     } //...(code skipped)...  
15 }
```



- Gadget identification
 - statically identify load-semantics-store chain from LLVM IR
- Dispatcher identification
 - static identify loops with gadgets from LLVM IR
- Gadget stitching
 - select gadgets and dispatchers (manual)
 - check stitchability (manual)


Evaluation – Feasibility

Nine x86 programs with nine vulnerabilities

- x86 Gadgets
 - 7518 in total
 - 1273 reachable via selected CVEs
 - 8 programs can simulate all MINDOP operations
 - manually confirmed 2 can build Turing-complete attacks
 - x86 Dispatchers
 - 5052 in total, 1443 contains gadgets
- > DOP elements are abundant

Case Study: Exploiting dlopen

- *dlopen* allows arbitrary computation
 - send malicious payload
 - corrupt link list & call dlopen
- invalid input
no call to dlopen



Case Study: Exploiting dlopen

- DOP attack addresses the problems
 - construct payload in memory
 - force call to dlopen
- invalid input
no call to dlopen

