

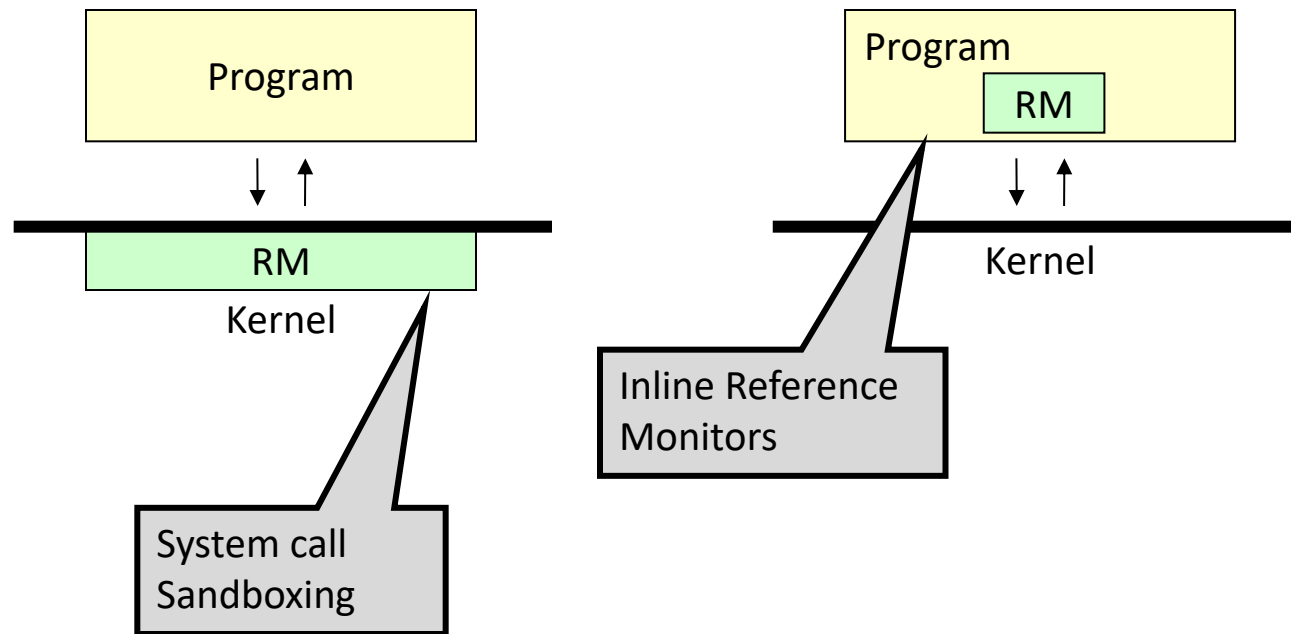
CS5231: Systems Security

Lecture 11: Virtualization and TEE

Recap: Reference Monitors

Reference Monitor: A piece of code that *checks all references* to an **object**

Syscall Sandbox: A reference monitor for protecting OS resource objects from an app



Recap: Policy vs. Enforcement Mechanism

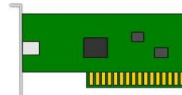
- Access Control Policies
- Enforcement:
 - Process sandboxing
 - Inline Reference Monitors
 - *Virtualization*
 - *Hardware-based isolation / Trusted Execution Env.*

Isolation: Virtualization

Problem: Isolated Computation on Shared CPU

Operating
System

Operating
System

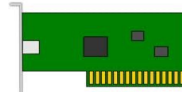
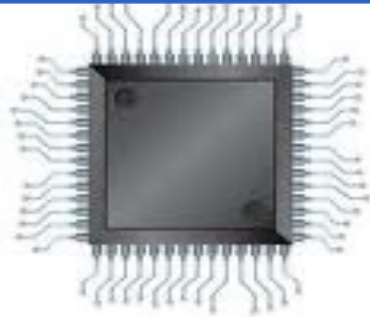


Defense(I): Virtualization

Game
VM

Banking
VM

Virtual Machine Monitor (VMM)



A Bit of History...

- Virtual Machines – 1960's
 - **Motivation**: Sharing of machines between users
 - Many implementations by IBM
- Virtual Machines on RISC / CISC – Late 90's
 - **Motivation**: Unify under-utilized machines, ease-of-maintenance, security
 - E.g. VMWare
- Heavy utilization in cloud computing...

Public Clouds: EC2

- Virtual Machine Monitor (VMM): **Xen**
- Instance: **A running OS image of virtual machine**
- ECU: **EC2 Compute Unit** \approx 1.2GHz Opteron/Xeon CPU

Amazon CloudWatch

Region:

Detailed Monitoring for Amazon EC2 Instances

- \$3.50 per instance per month, provided at 1-minute frequency

Basic Monitoring for Amazon EC2 instances

- \$0.00 (free of charge) per instance per month, provided at 5-minute frequency

Monitoring for Custom Metrics

- \$0.50 per metric per month

Assumptions

- **Goal:** Isolation of Code, data, resources between:
 - Guest VM and Host VMM
 - Between VMs
- **Assumptions:**
 - Bug-free TCB: Host OS, VMM
 - Malware can affect the guest OS & apps.

Security Applications of Virtualization


- Virtual Machine Isolation
 - Red-Green Systems
 - E.g. Banking VM vs. Normal VM
 - Dynamic Analysis / Containment of Malware
- Virtual Machine Introspection
 - E.g. Run an anti-virus in the VMM

Enforcement Goals for a VMM

- Security VMM Goals:
 - Complete Mediation
 - Trap on all MMU, DMA, I/O accesses
 - Transparency
- Commercial VMM Goals:
 - Performance
 - Compatibility: Run on commodity OSes

Compatibility Challenges: An Example

```
mov eax, (ebp)
```



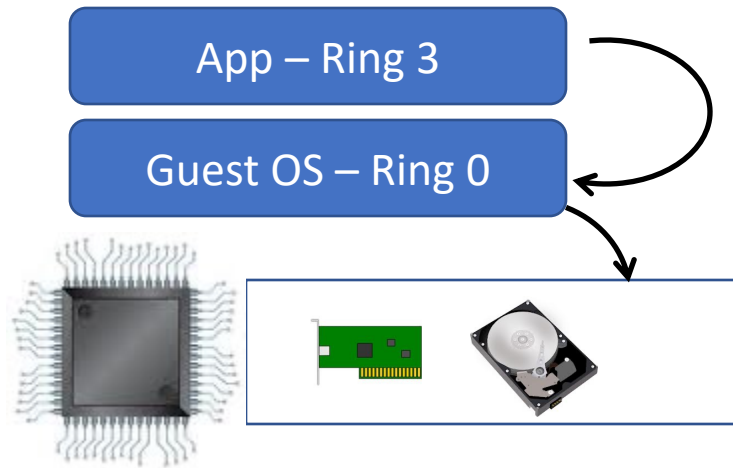
Virtual Address	Physical Address	Protection Bits (R,W)

Page Table

OS uses to isolate kernel code / data

Compatibility Challenges

Non-virtualized execution



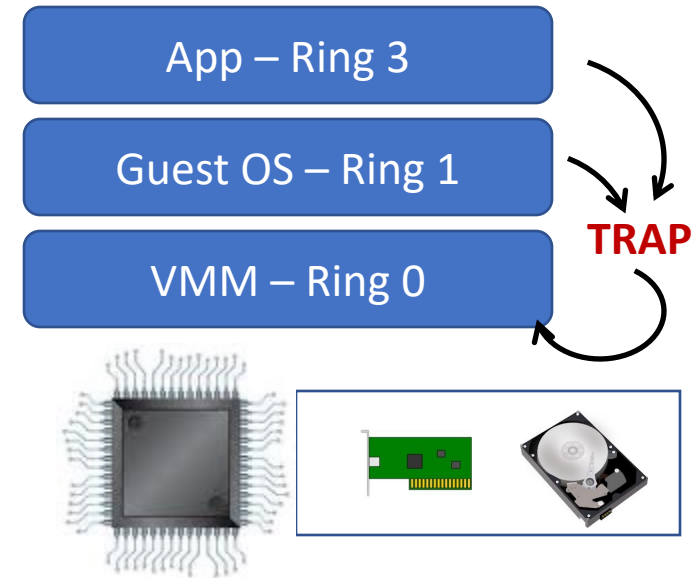
Privileged Instructions

- that trap: e.g. `cli`

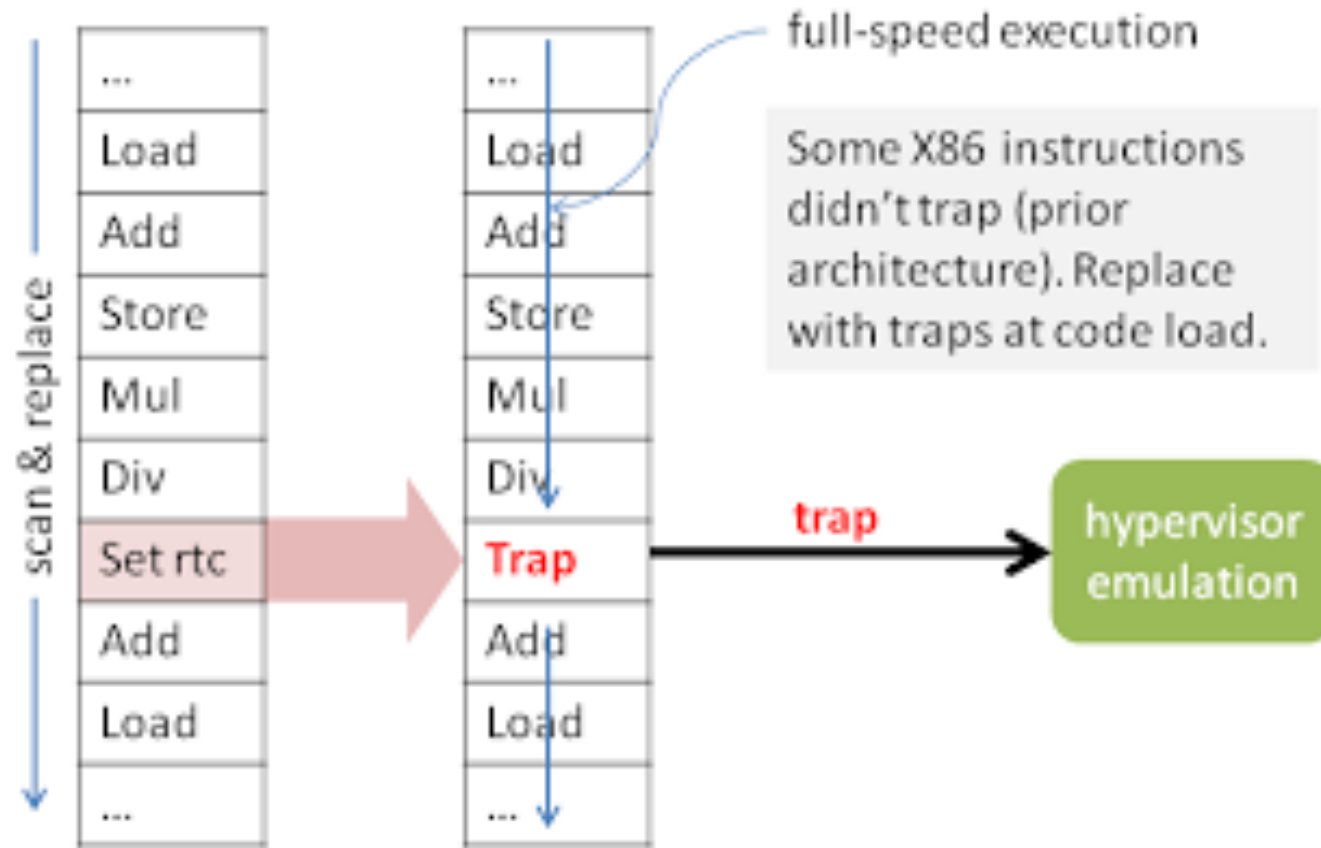
Privileged Instructions

- That don't trap: e.g. `popf`

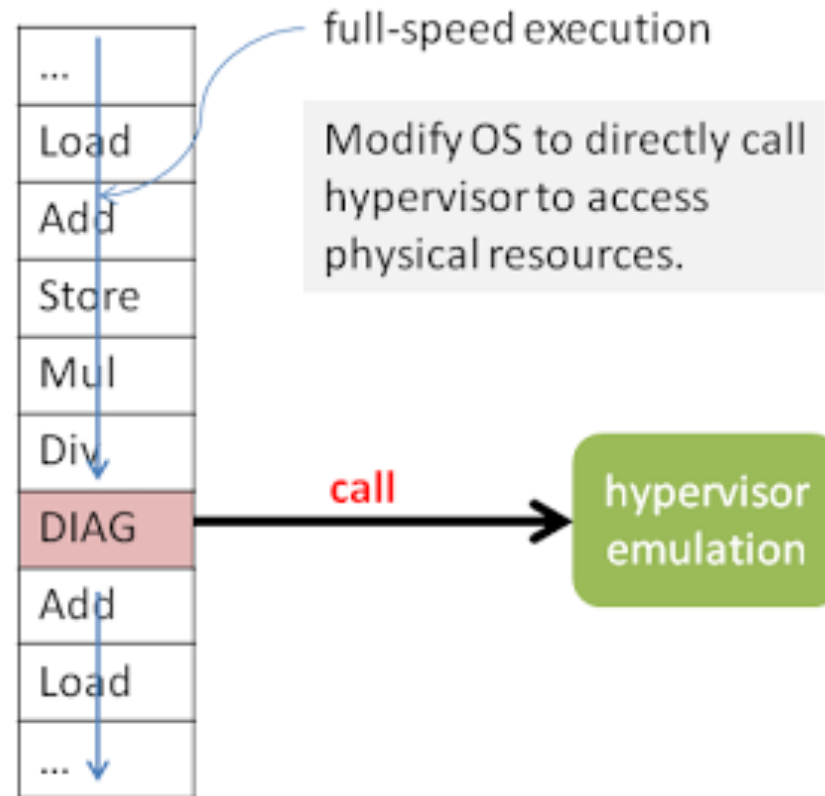
Virtualized execution



Virtualization Techniques: Binary Translation (VMware)



Virtualization: Paravirtualization (Xen)

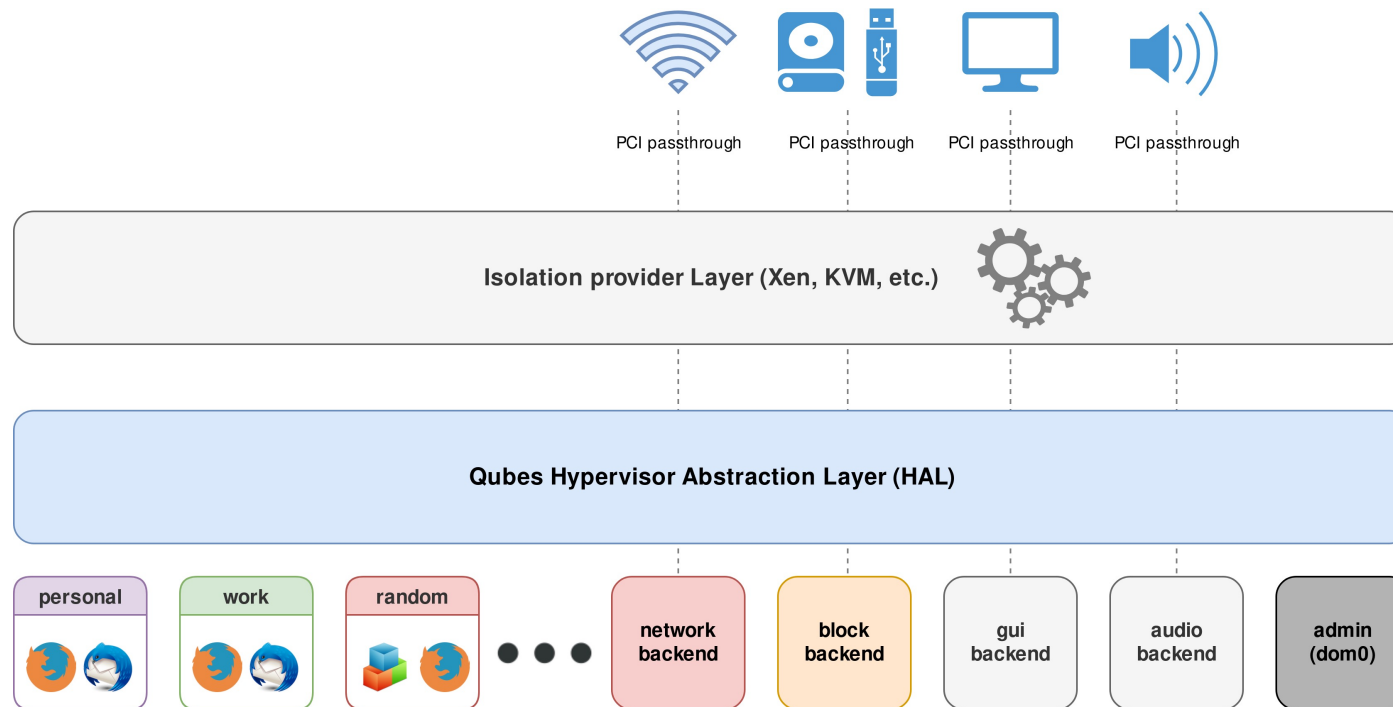


Hardware Assisted Virtualization

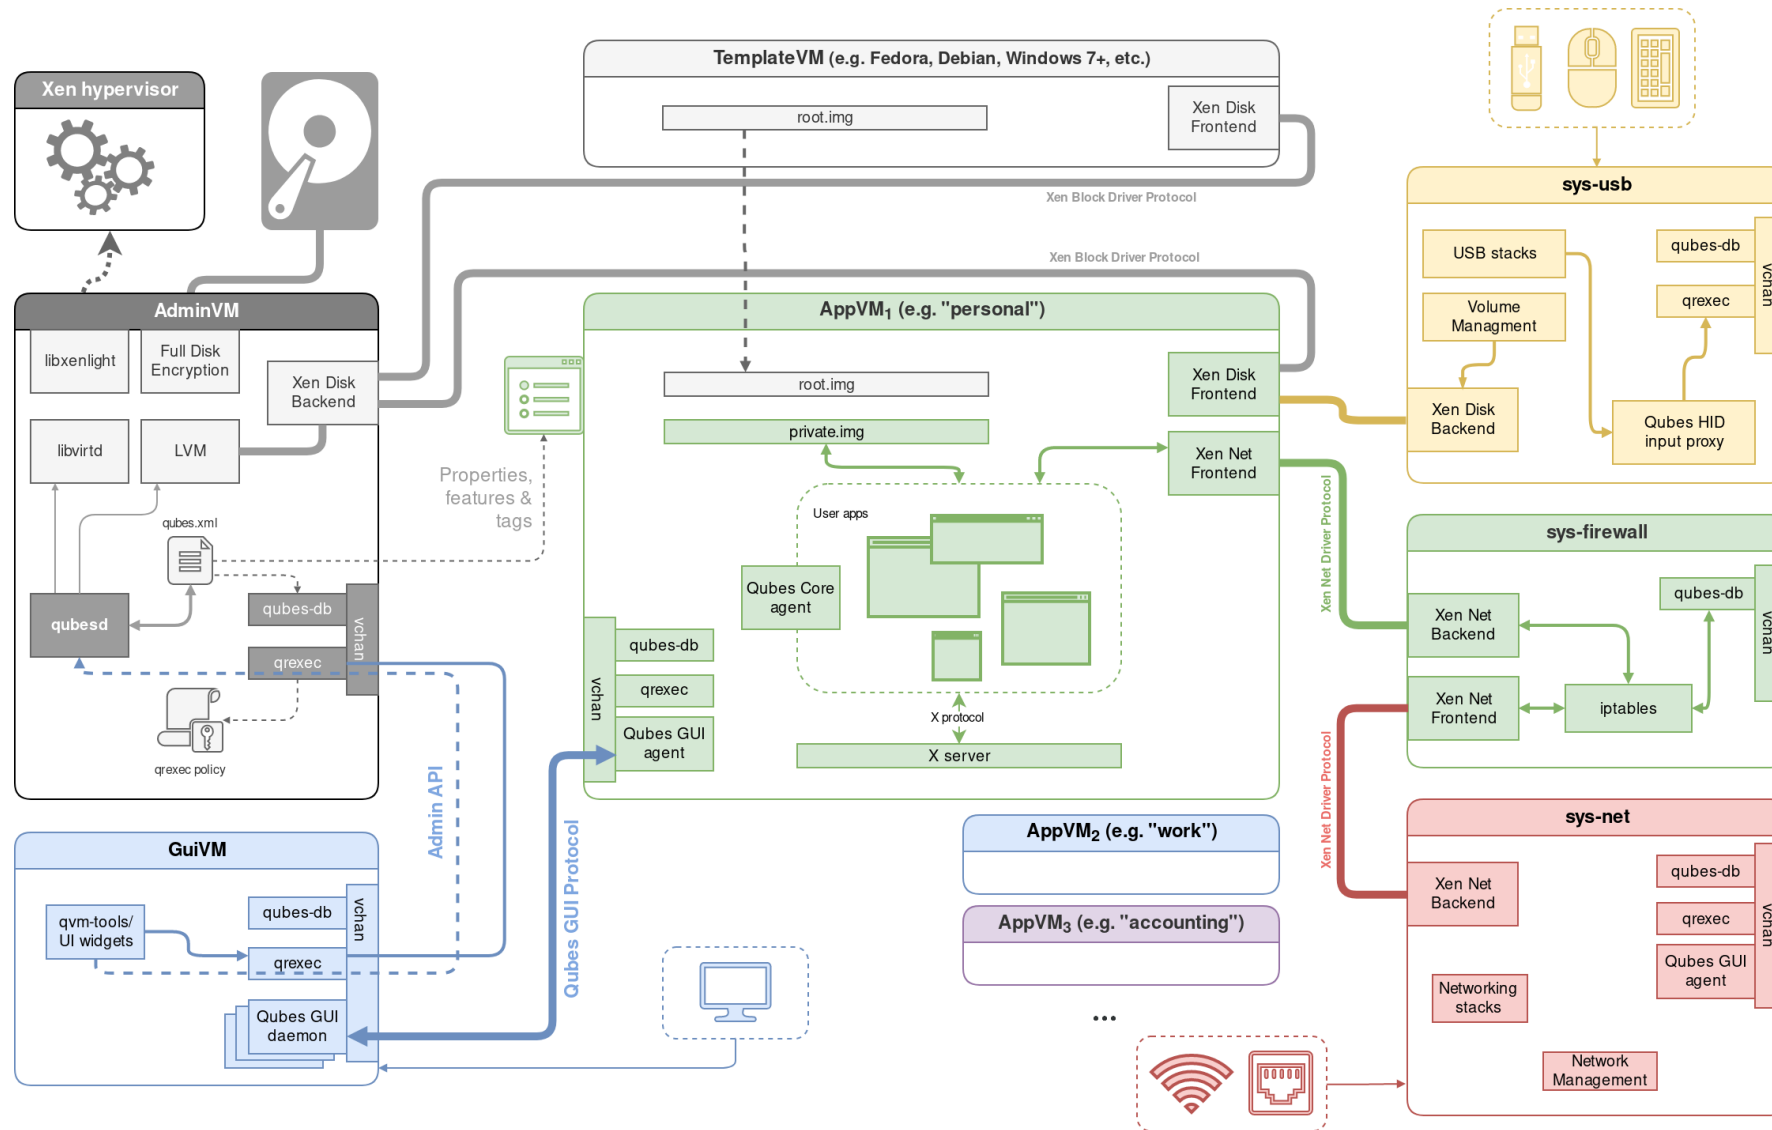
- CPUs adding support over the years
- Goal: Better Performance, Security
 - Intel VT-x
 - MMU virtualization using EPT (2009),
 - Nested virtualization - VMCS (2012)
 - I/O virtualization – IOMMU (2009)
 - Intel VT-d
 - DMA remapping (2009)

Qubes OS

- A reasonably secure operating system
 - A network of virtual machines in a computer

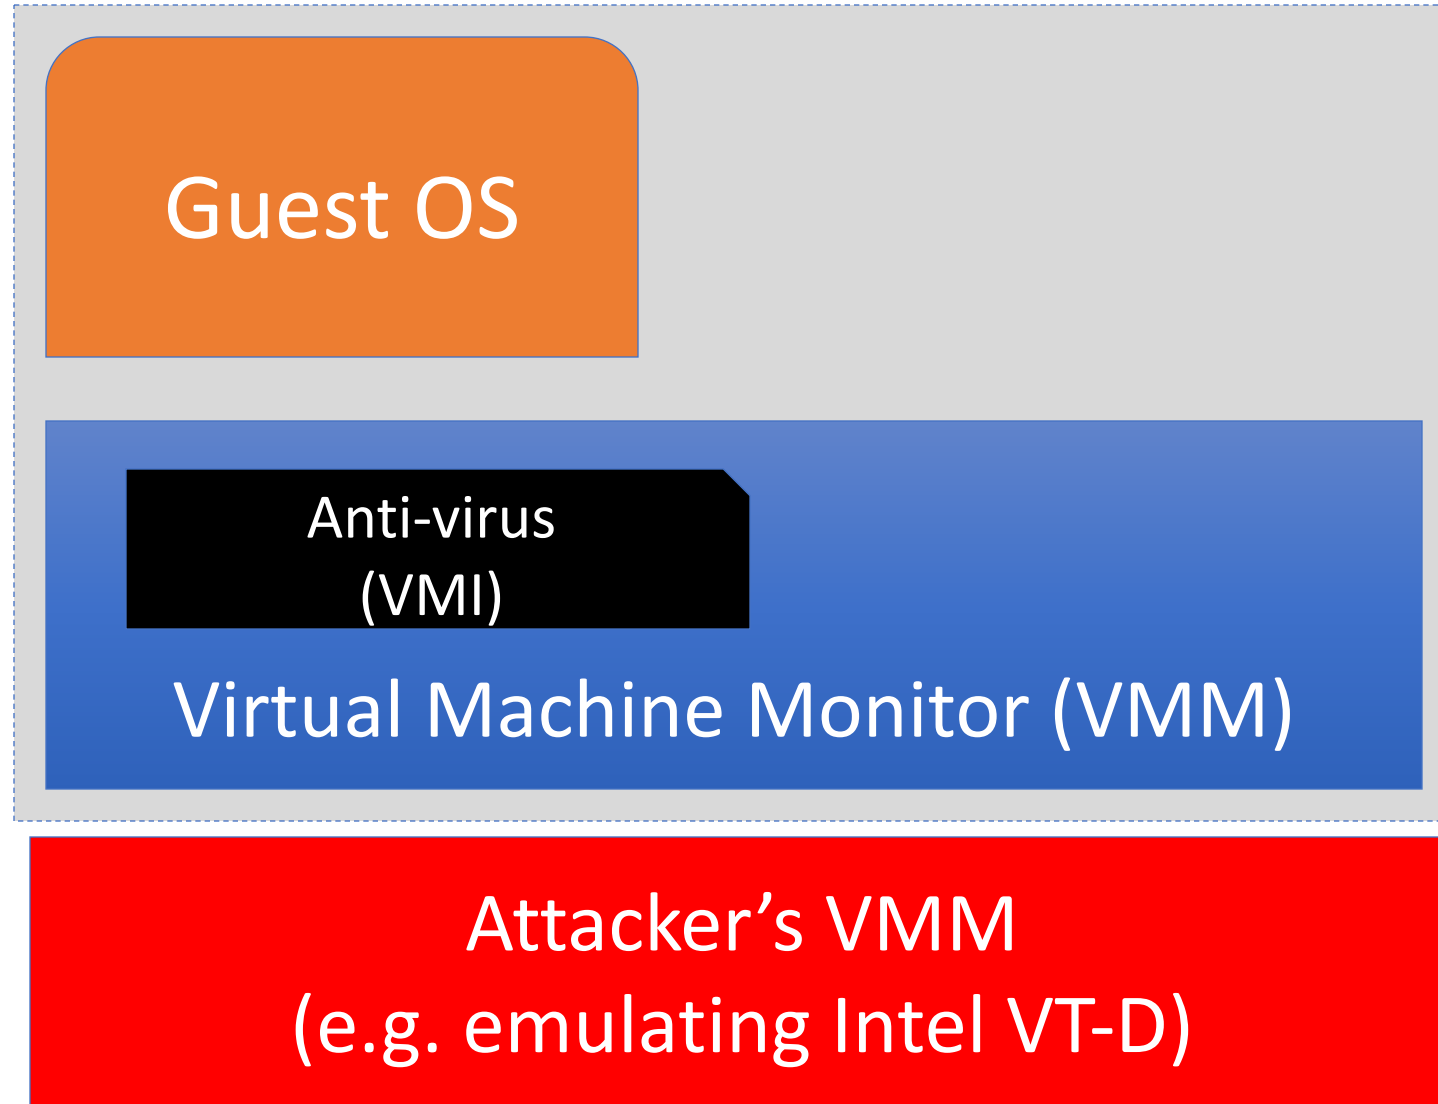


Qubes OS Architecture

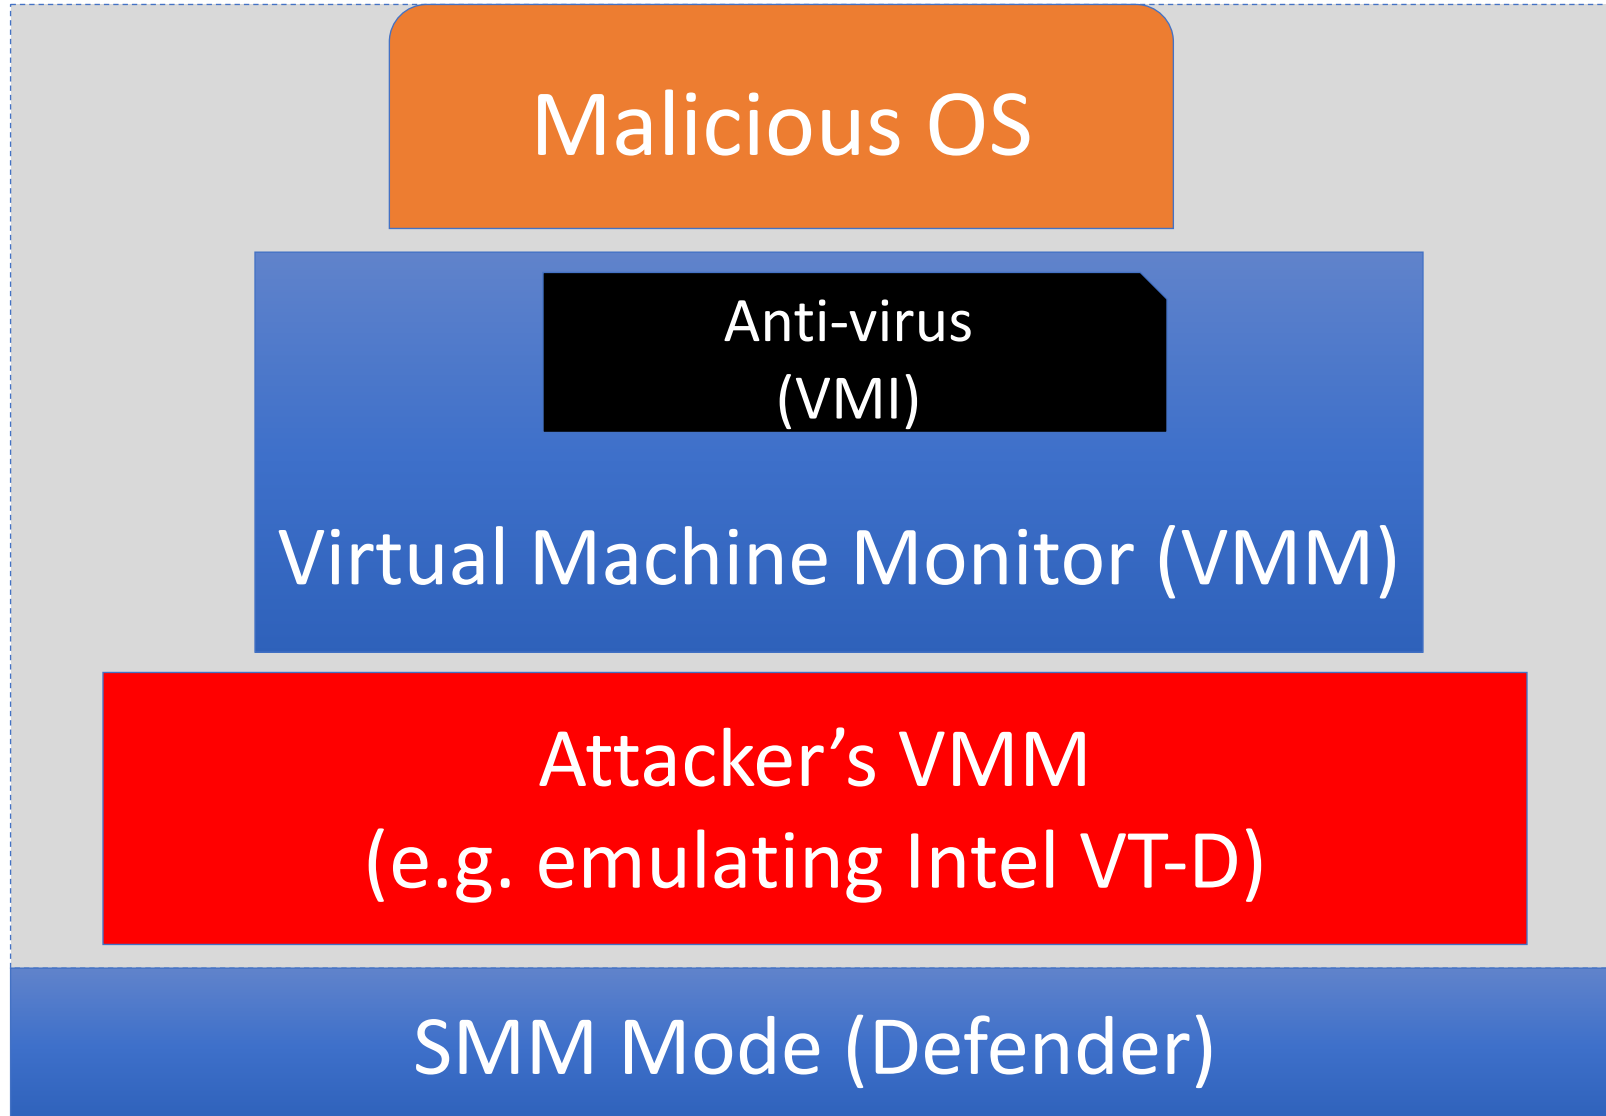


Limitations of Virtualization

Virtual Machine Based Rootkits



The Problem of Secure “Root of Trust”: Is highest layer of privilege malicious?

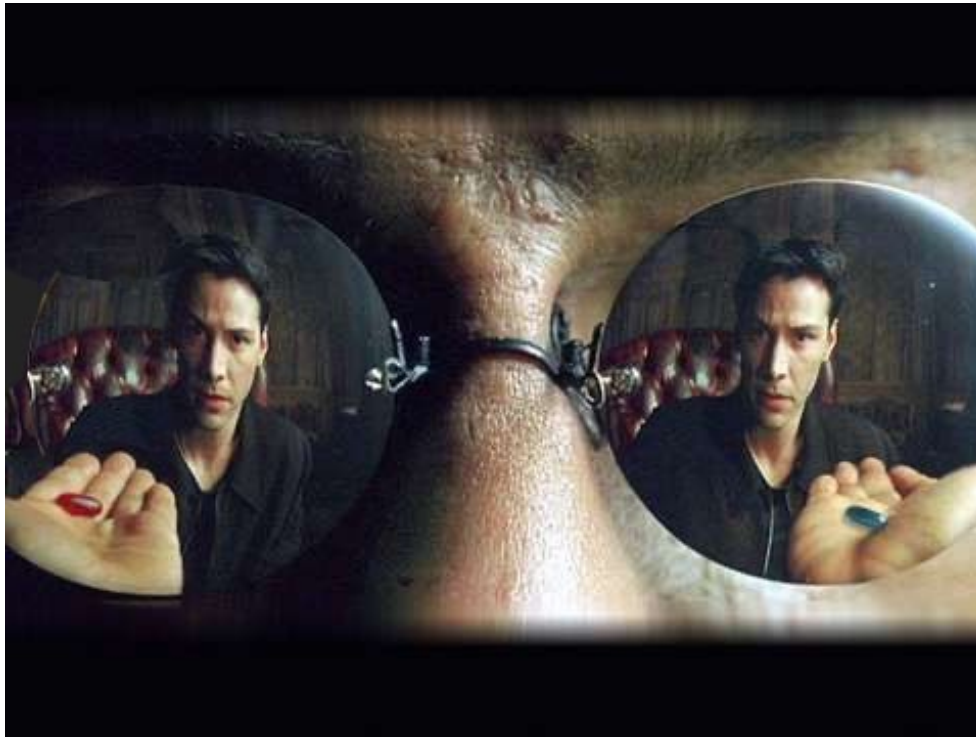


Implication on Malware Containment

- In principle, is some containment possible?
 - Yes, When the highest layer of privilege is trusted
 - E.g. the VMM is trustworthy

Virtual Machine Based Rootkits: Can the software know its virtualized?

- Blue Pill: “blissful ignorance of illusion”
- Red Pill: “Detects you are virtualized”

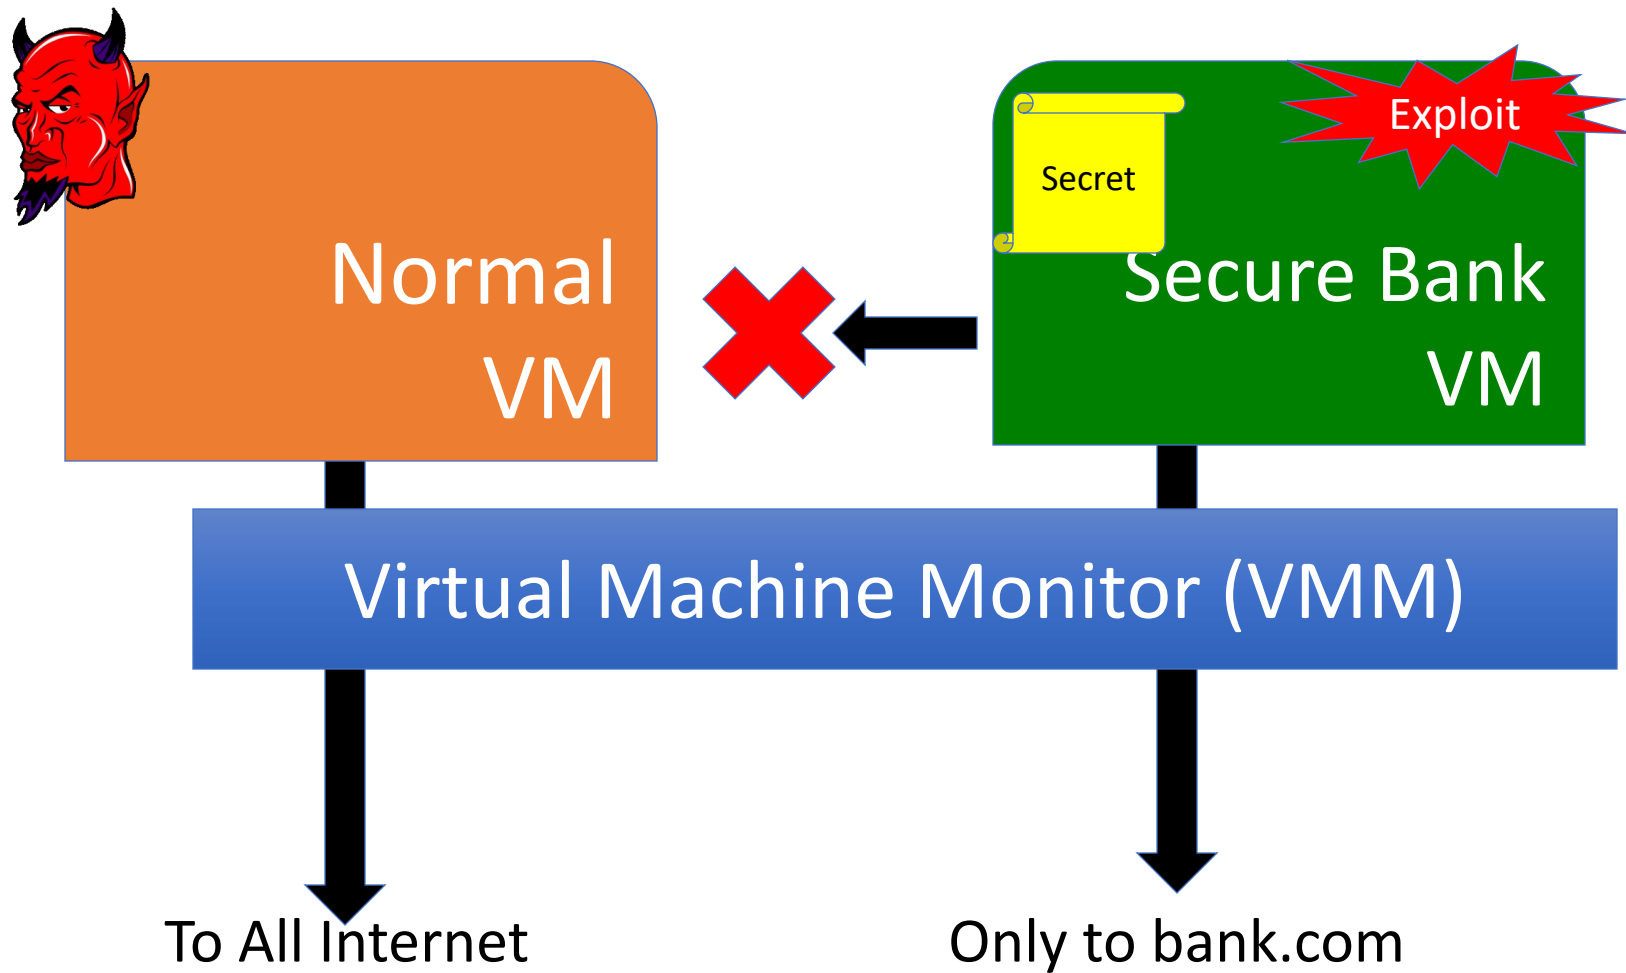


VMM Detection: The Red Pill

- Red Pill: “Detects you are virtualized”
- Ways to achieve a “red pill” attack:
 - Commercial VMMs aren’t fully transparent
 - E.g. VmWare emulates i440bx chipset (old)
 - Virtualization Timing latencies Measurements
 - Many other measurement channels [[HotOS’07](#)]
- Applications of VM Detection:
 - Malware can detect introspection software (e.g. AV)
 - Can utilize “Anti-VM” techniques
 - Benign use: Copy protection by VM duplication

Implication on Malware Containment

- In principle, is some containment possible?
 - Yes, When the highest layer of privilege is trusted
 - E.g. the VMM is trustworthy
- Detecting virtualization is easy
 - It can thwart malware analysis (introspection)



Can attacker leak secret document to evil.com?

The Problem of Covert Channels

- Definition: "An unintended channel of communication between 2 untrusted programs"
- E.g. Shared Cache Latency
 - Sender
 - Send bitval 1: Perform random memory access
 - Send bitval 0: Do nothing
 - Receiver
 - Rcv bitval 1: If long read time for a fixed memory loc.
 - Rcv bitval 0: If short read time for fixed memory loc.
- Can get 0.02 bits/sec on Amazon EC2 [CCS'09]
- Many channels: Disk, I/O, Virtualization latency, ...

Implication on Malware Containment

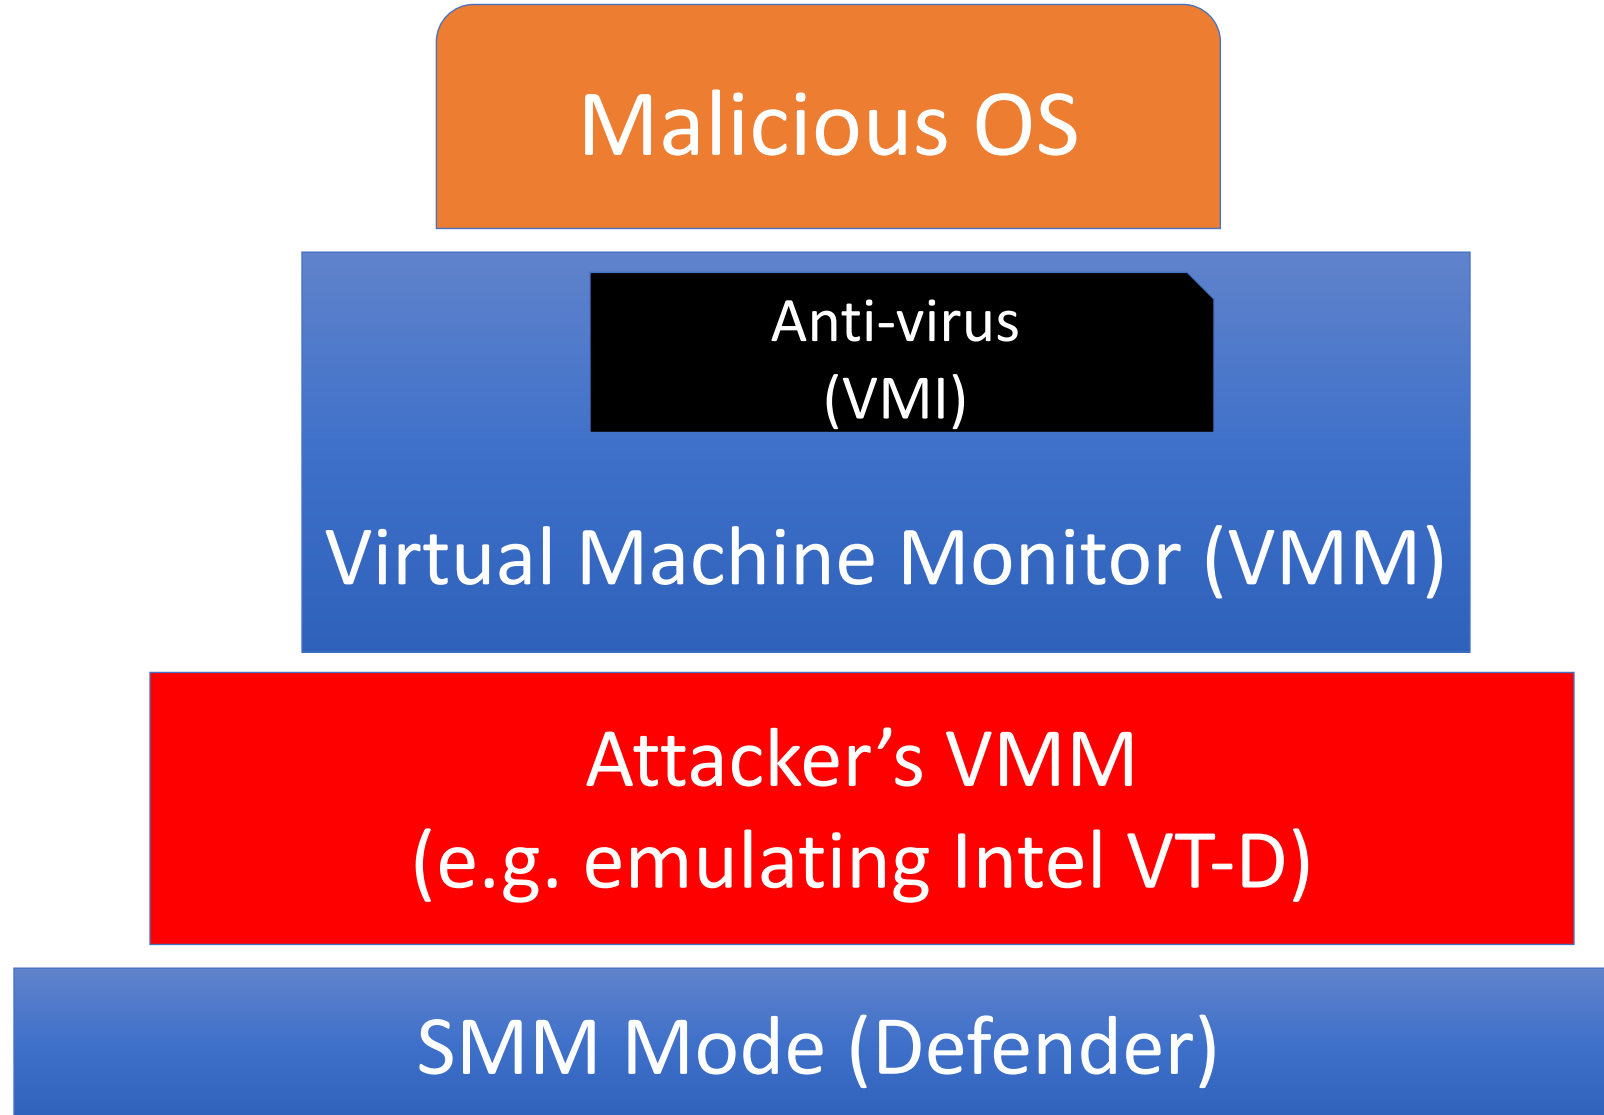
- In principle, is some containment possible?
 - Yes, When the highest layer of privilege is trusted
 - E.g. the VMM is trustworthy
- Detecting virtualization is easy
 - It can thwart malware analysis (introspection)
- Which containment using VMs is possible: integrity vs. confidentiality?
 - Yes, for Integrity policy – I.e., protecting contained malware from corrupting benign data outside the VM
 - No, for confidentiality, covert channels are a problem

Optional Reading Material for the curious...

- [SubVirt: Implementing malware with virtual machines](#)
- [Overshadow: A Virtualization-Based Approach to Retrofitting Protection in Commodity Operating Systems](#)
- [HyperSafe: A Lightweight Approach to Provide Lifetime Hypervisor Control-Flow Integrity](#)
- [Hey, You, Get Off of My Cloud: Exploring Information Leakage in Third-Party Compute Clouds](#)

Trusted Execution Environments: The Basic Idea

The Problem of “Root of Trust”: Is highest layer of privilege malicious?

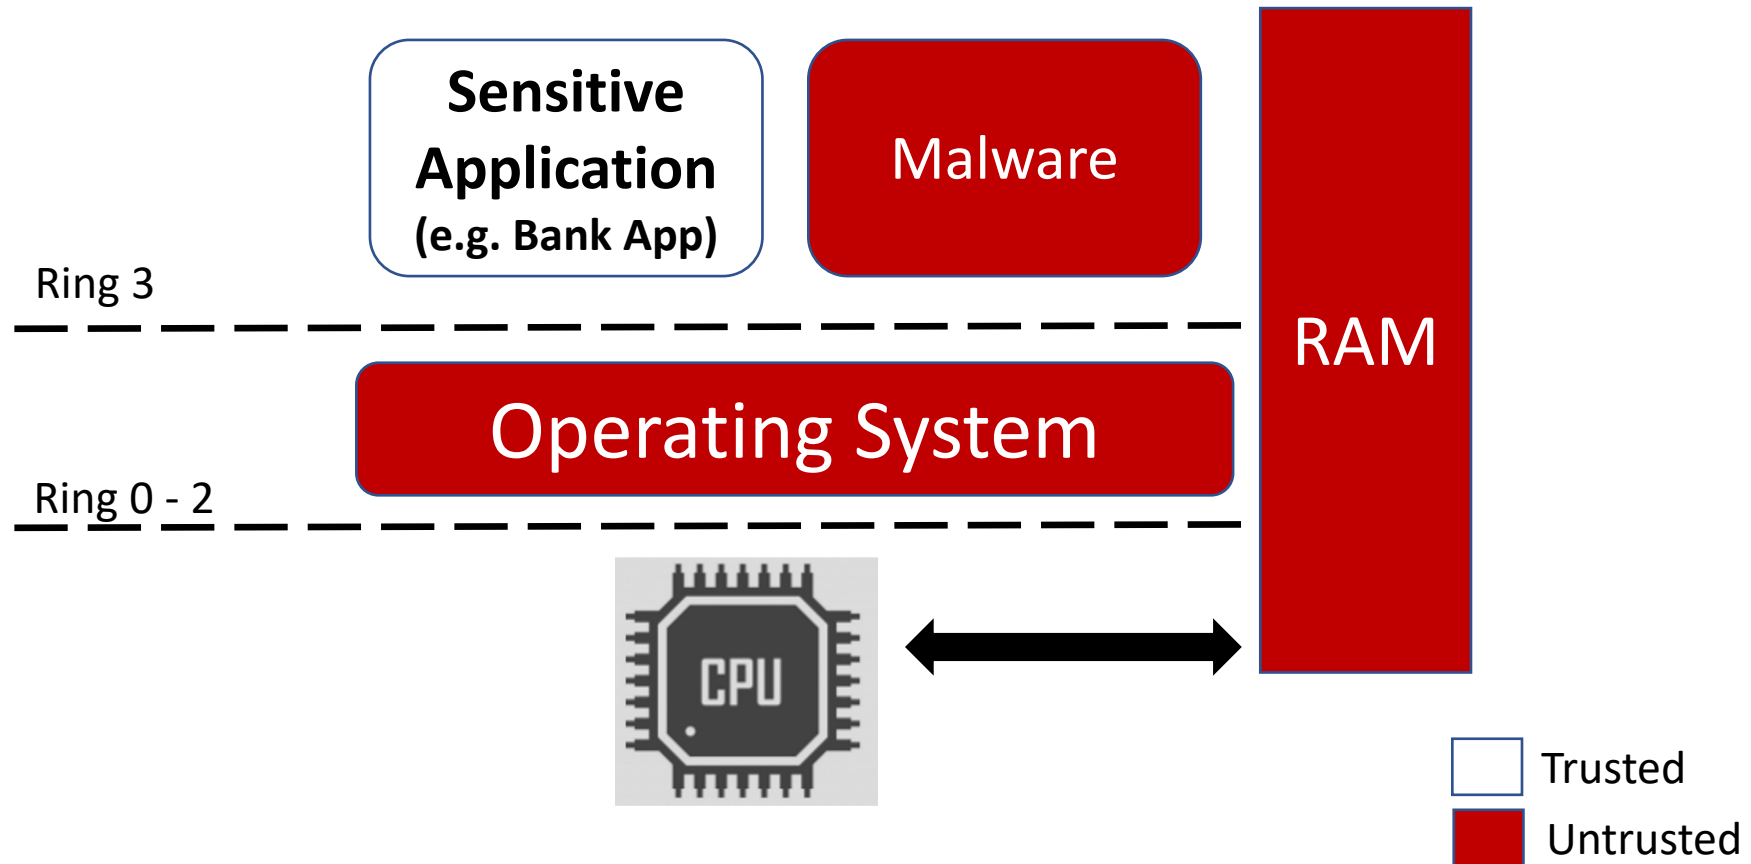


Solution: TEEs ensure Secure “Root of Trust”

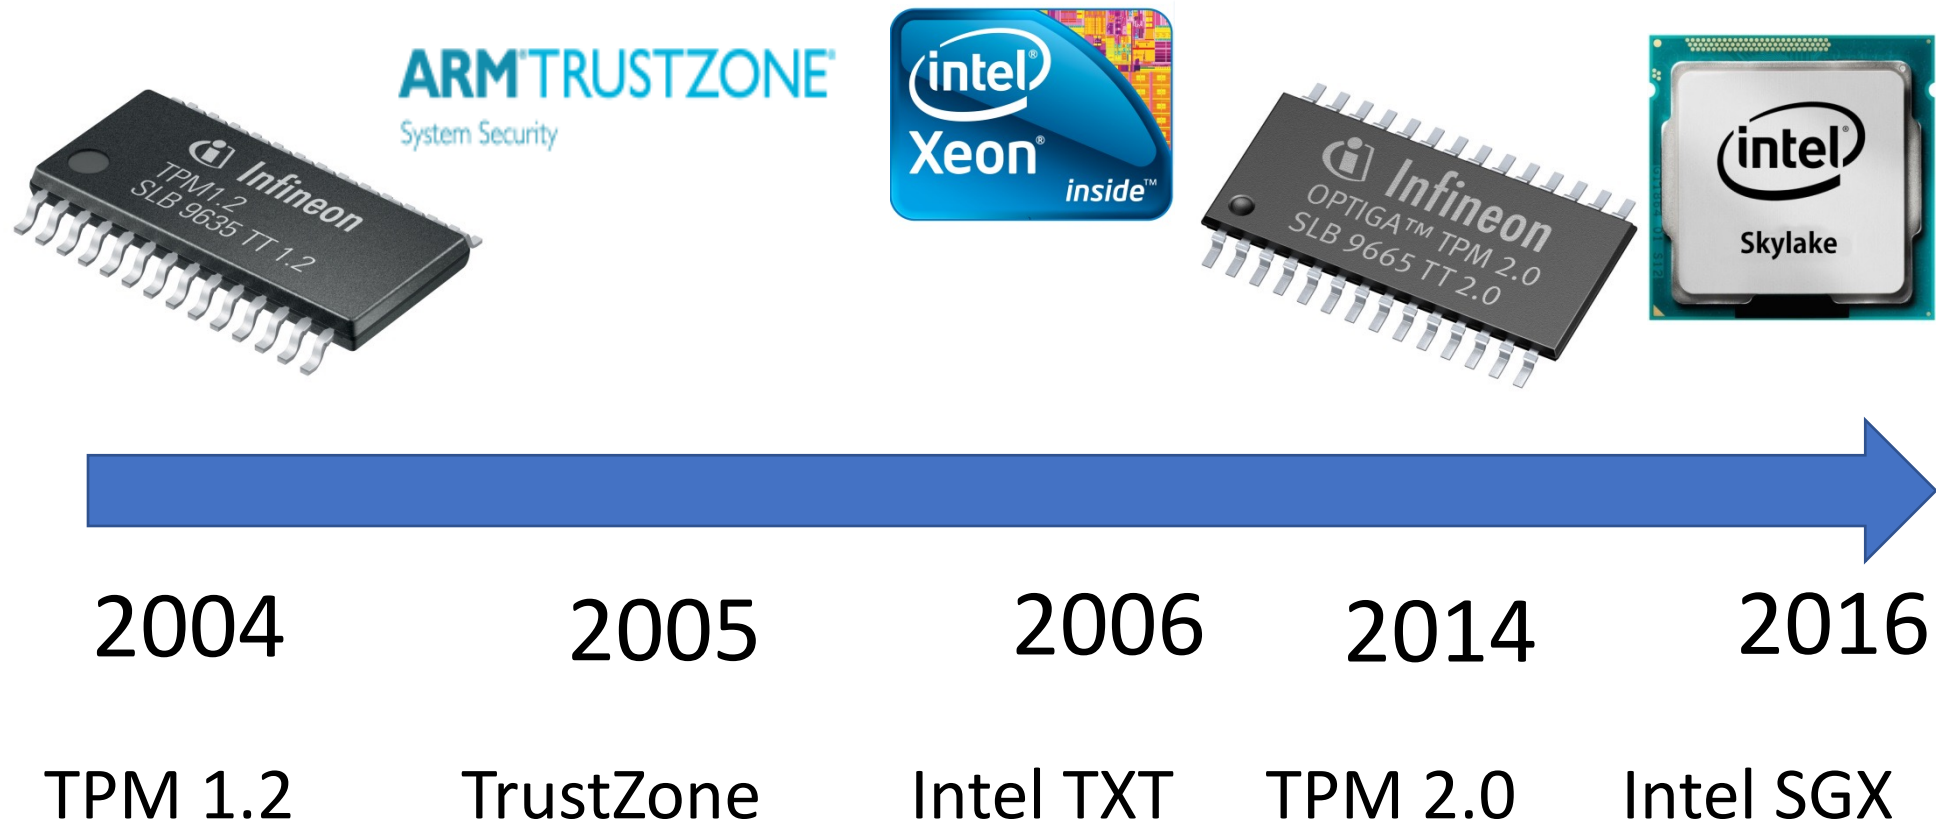
- Trusted Execution Environments (TEEs)
 - A **hardware root-of-trust**
 - Can assume that all software is malicious
- Why trust hardware?
 - Tamper-resistant from all software malware
 - Perhaps less complex, easier to verify?

Trusted Execution Environment (TEE): Security Model

- Trust the hardware
- Don't trust other software on the system

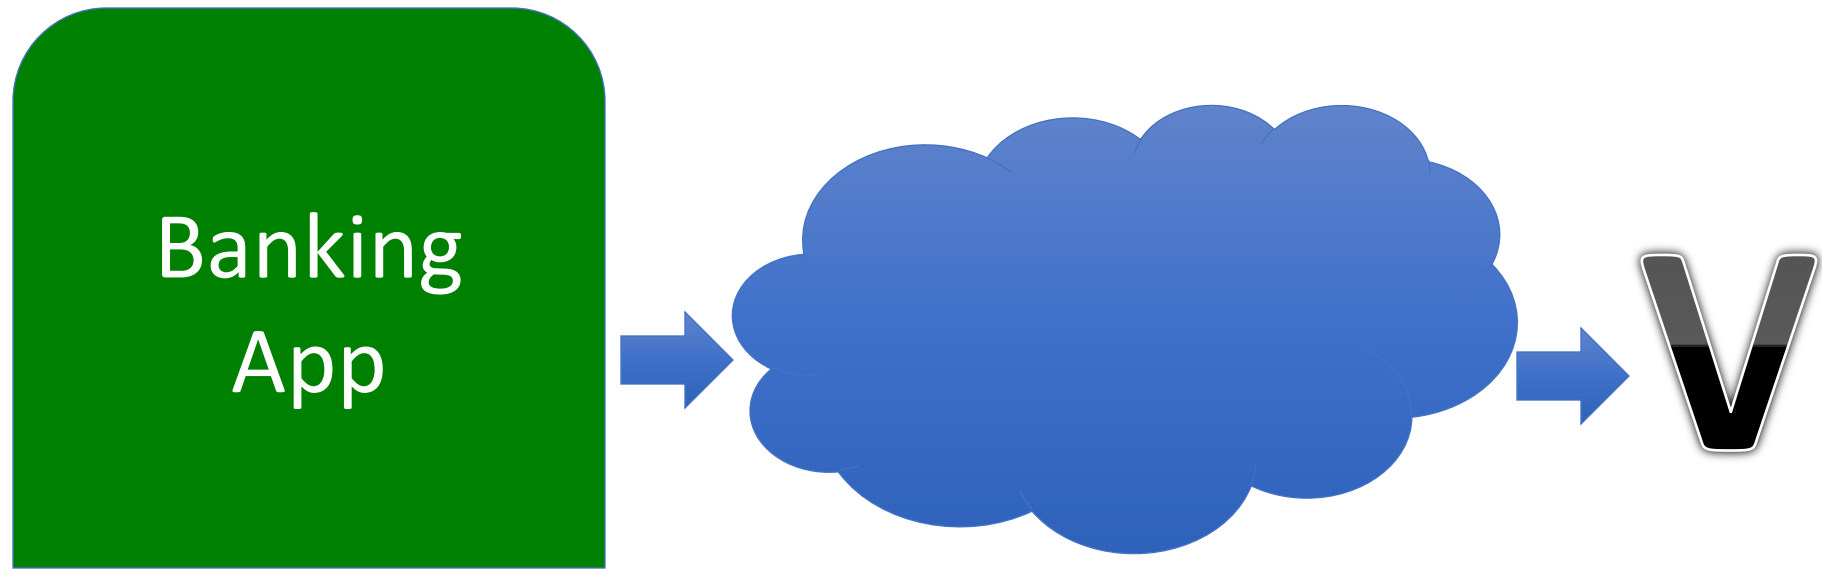


Evolution of TEEs

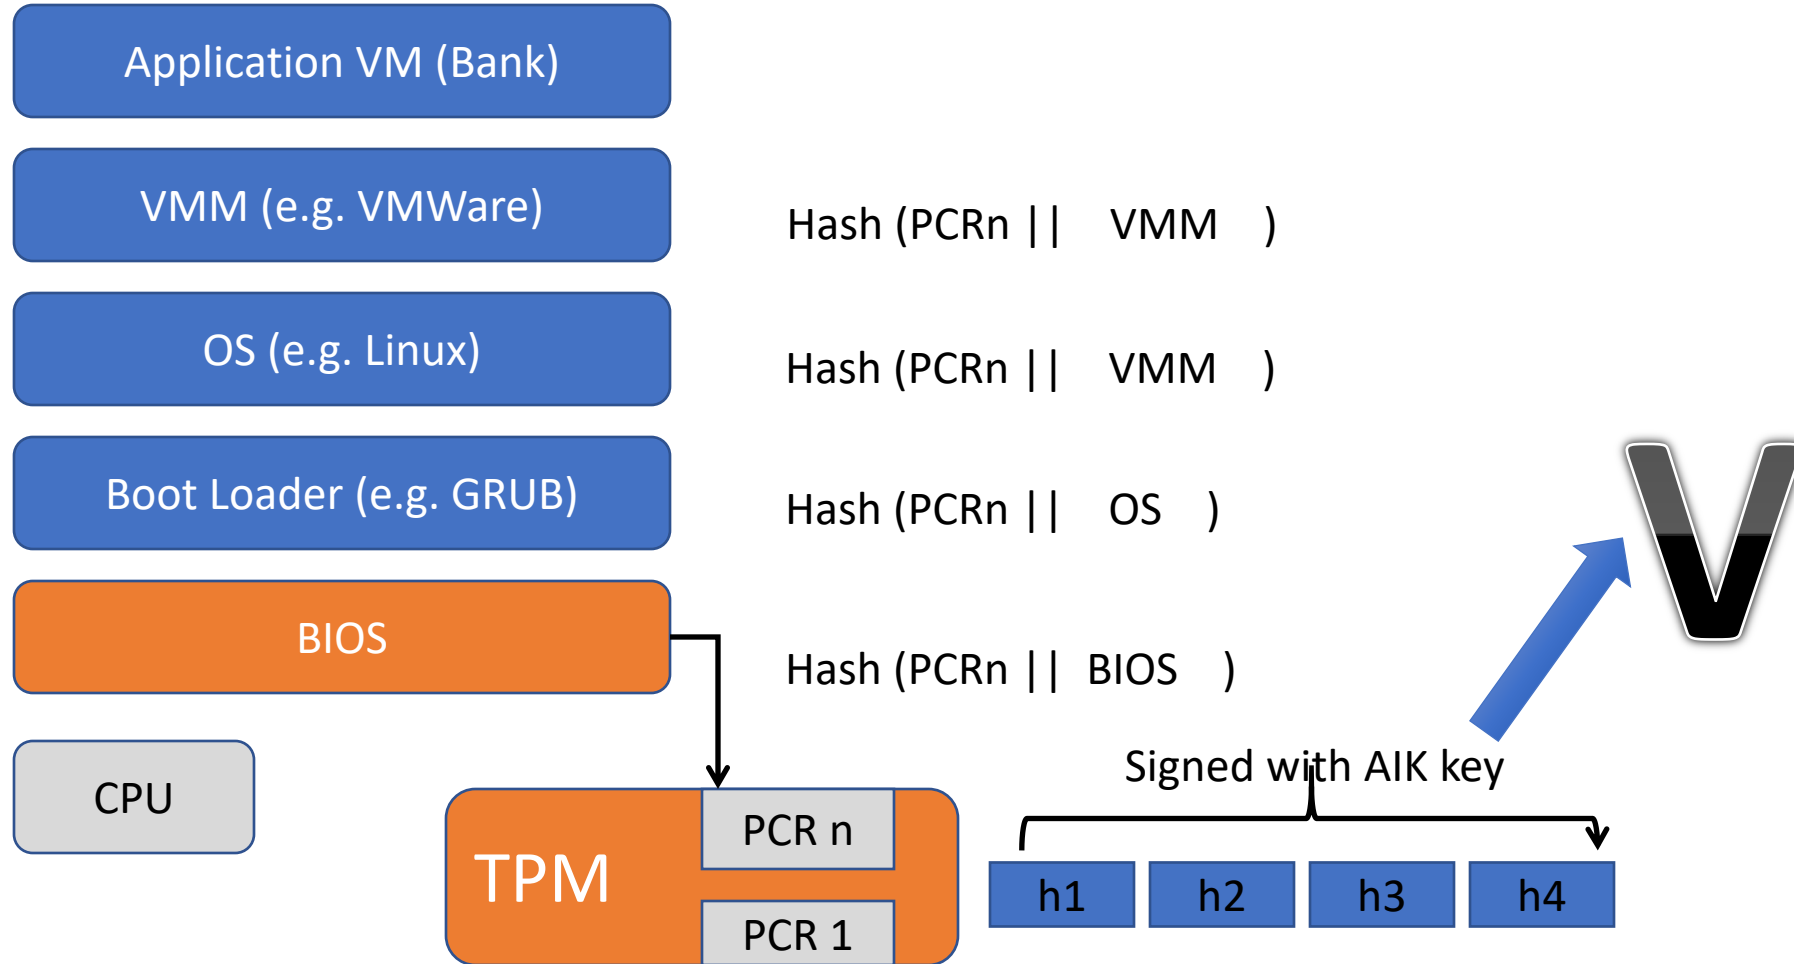


Trusted Execution Environment

Trusted Execution Primitives (I): Remote Attestation

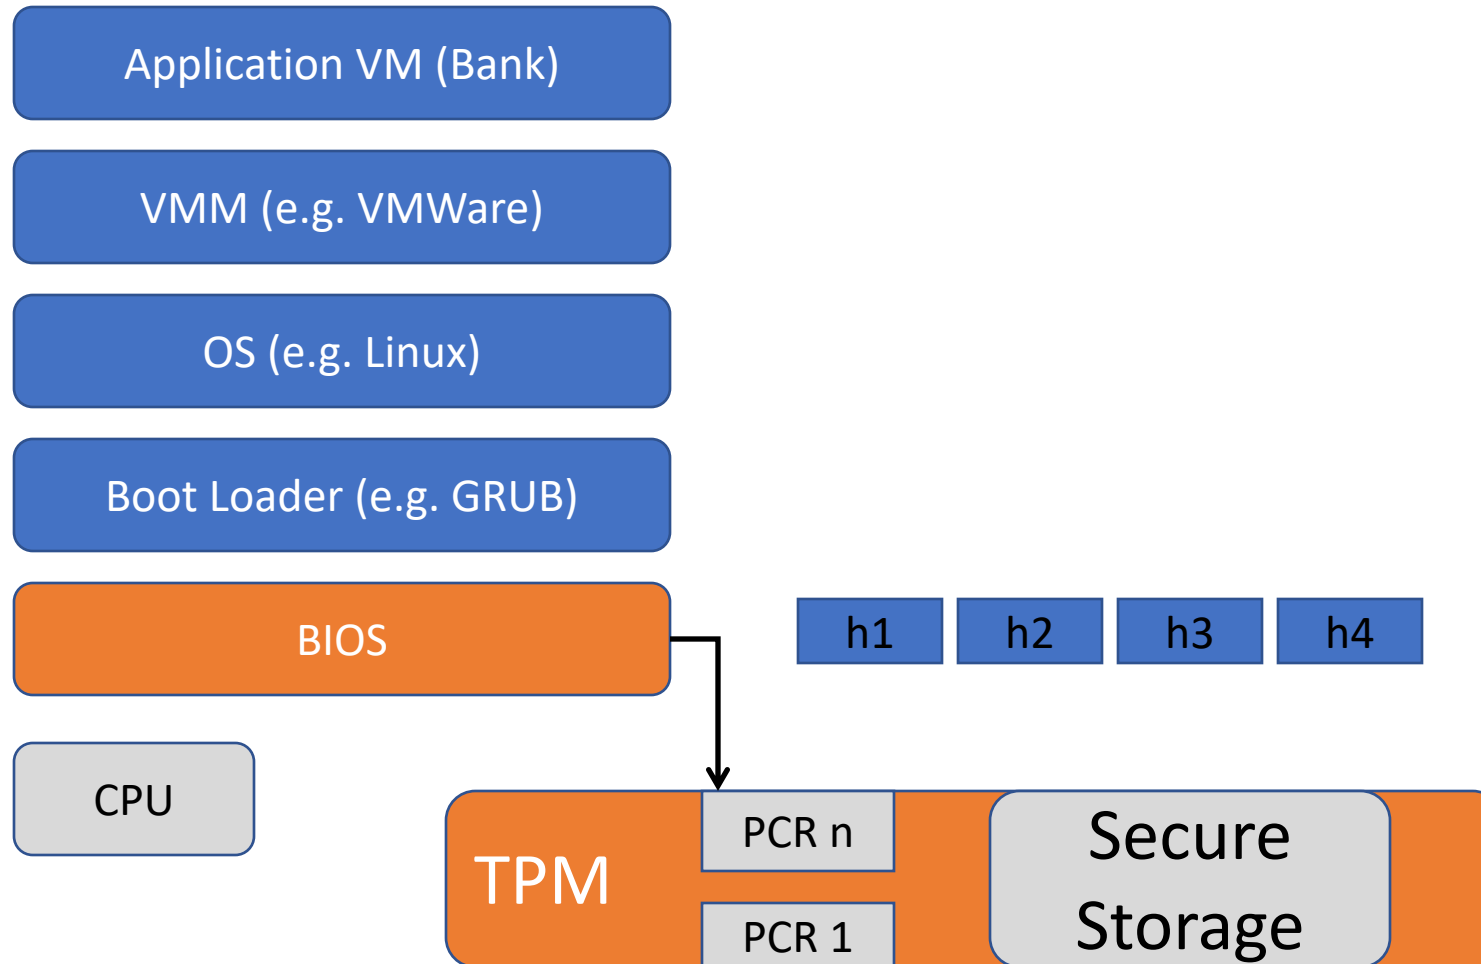


Trusted Execution Primitives (I): Remote Attestation



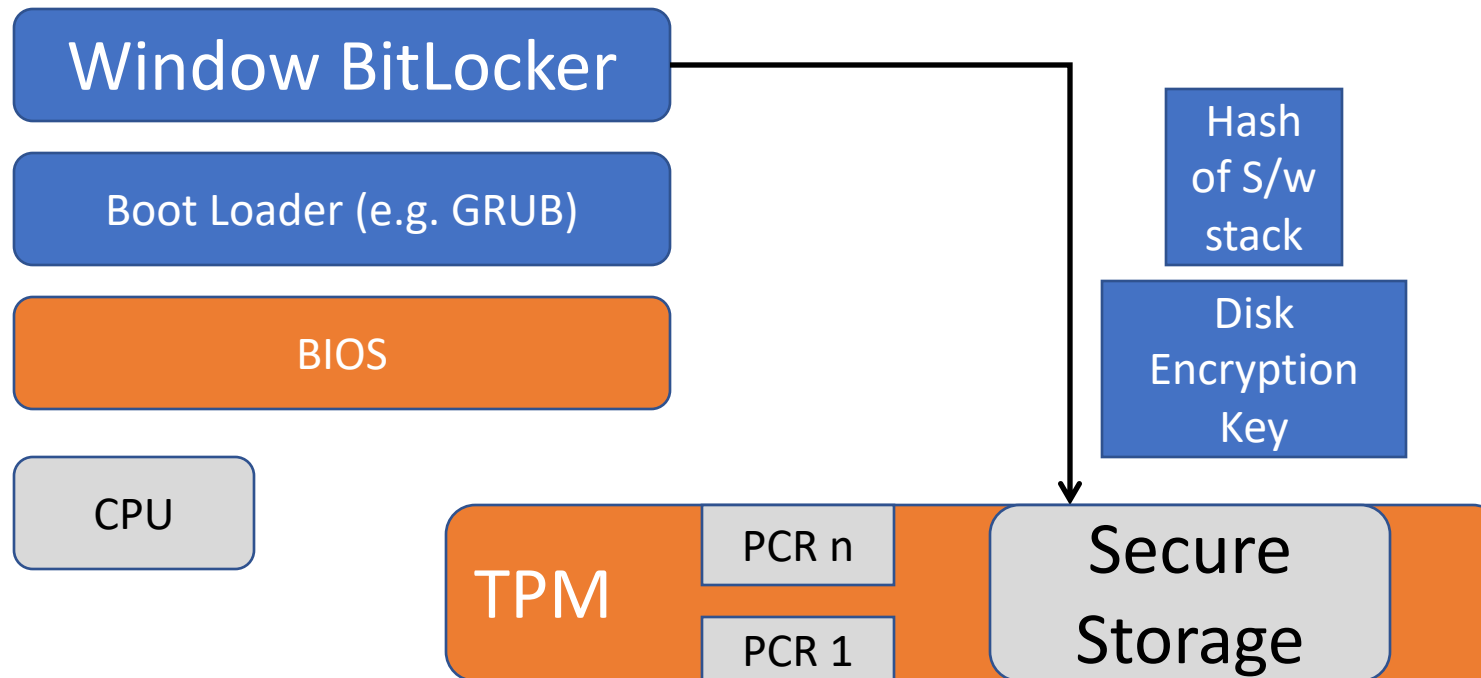
Each TPM has a AIK signing key

Trusted Execution Primitives (II): Trusted Boot

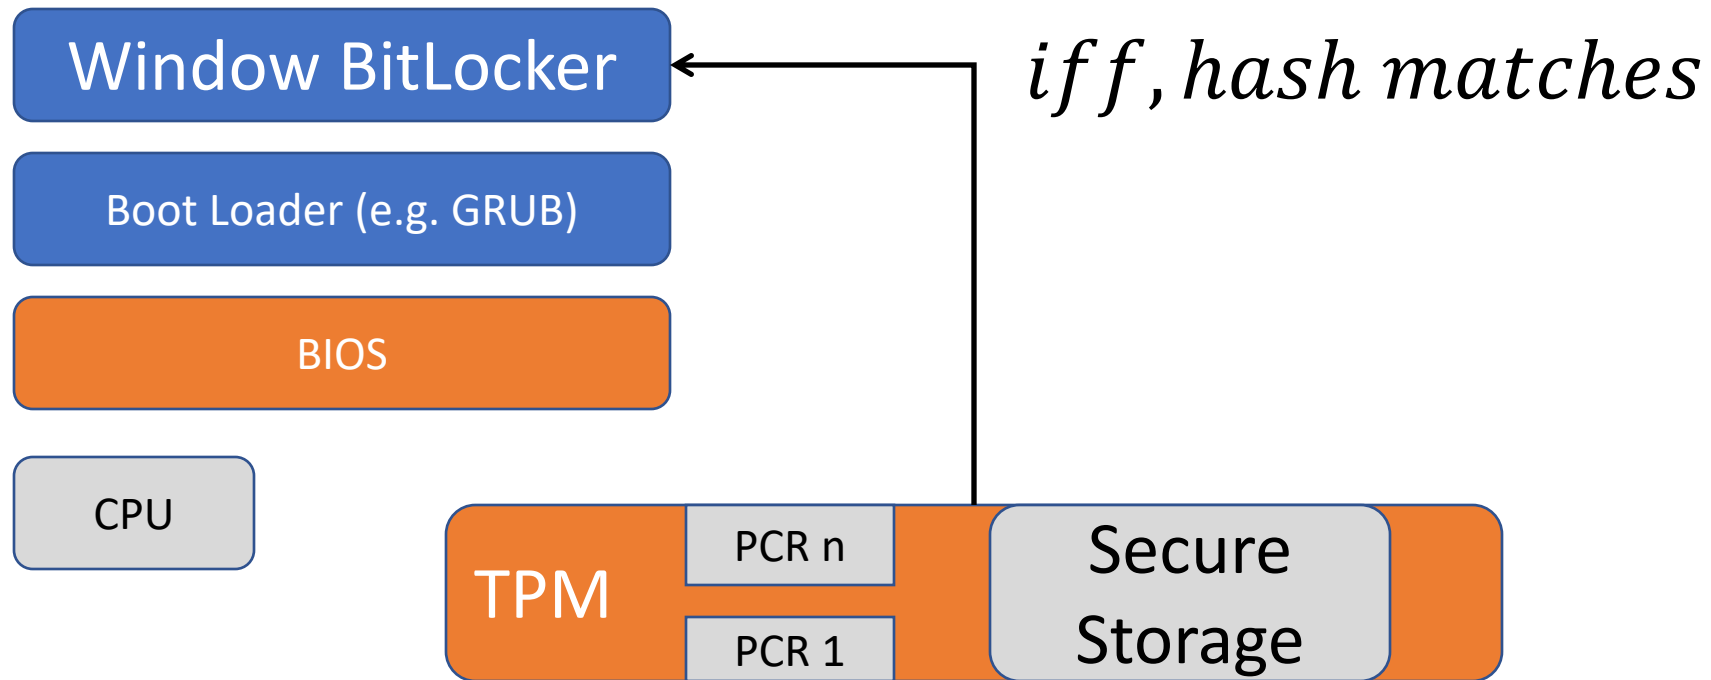


Trusted Execution Primitives(III): Sealing Data

Use TPM Measurements & Secure Storage for
Disk Encryption Systems?



Trusted Execution Primitives (III): Unsealing Data



Trusted Execution Primitives (III): Sealed Storage

```
# echo 'Secret!!!' | tpm_sealdata -z -i/proc/self/fd/0  
-o./mysecret.blob -p17 -p18 -p19
```

```
// assuming PCR's are the same  
# tpm_unsealdata ./mysecret.blob  
Secret!!!
```

```
// assuming PCR's are different  
# tpm_unsealdata ./mysecret.blob  
error 24: Tspi_Data_Unseal: 0x00000018 - layer=tpm,  
code=0018 (24), Wrong PCR value
```

Trusted Computing Primitives (I – II): Static Root-of-Trust

- So far, we've seen SRTM systems
 - Checks / Verification at load time
- Many Applications
- Windows BitLocker
- Linux TrustedGrub (because TXT is too slow!)
- Build your own “secure apps”
 - [eXtensible, Modular Hypervisor Framework](#)
(from CMU)

Use Case: Full Volume Encryption

- Encryption at the block level underneath file system
- Everything in the volume is encrypted.
- BitLocker is used by Microsoft since Windows Vista
- BitLocker takes advantage of TPM (TEEs)
 - Top level root key sealed in hardware
 - Root key encrypts disk encryption key, which encrypts sector data
- CPU protects disk encryption key by encrypting it
- CPU releases key only after comparing hash of early (unencrypted) boot files with previous hash

To Continue

- Other trusted computing technology
 - ARM TrustZone
 - Intel SGX