

CS5321 Network Security

Week2: Crypto Basics (2)

Daisuke MASHIMA

<http://www.mashima.us/daisuke/index.html>

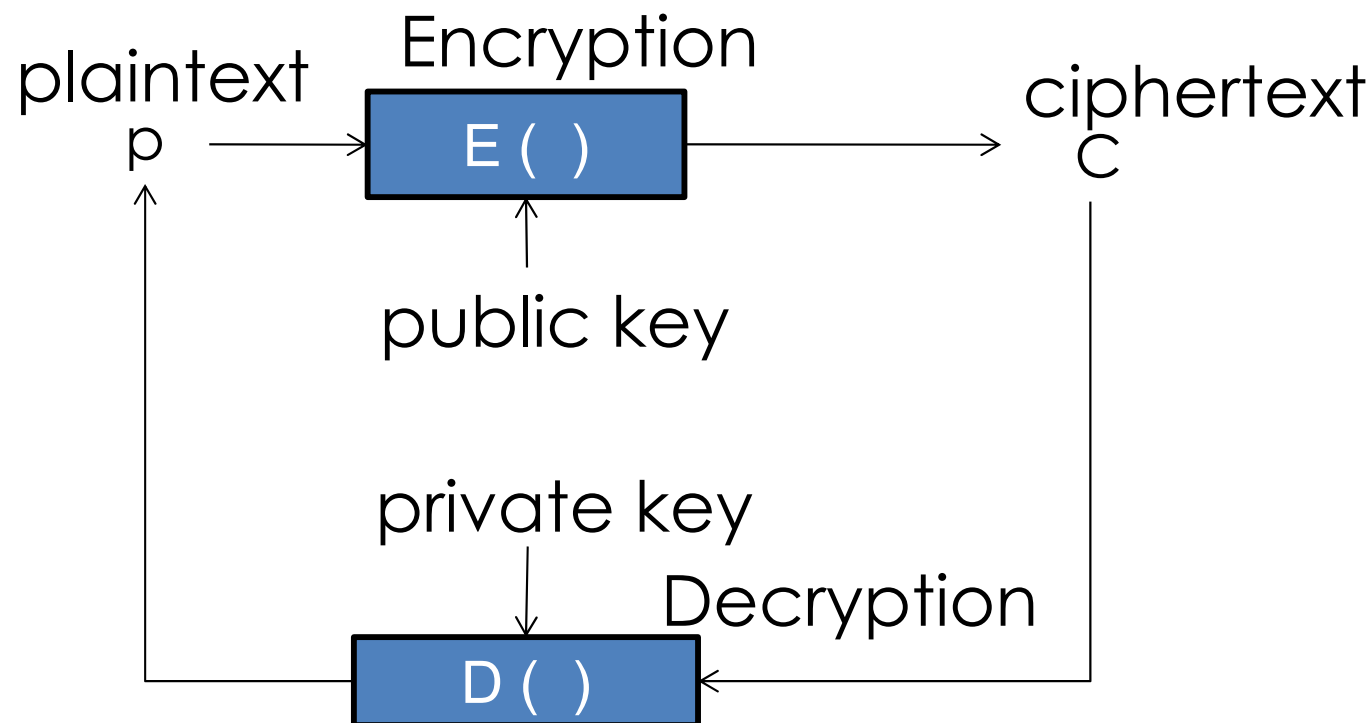
2022/23 Sem 2

Basic Cryptographic Primitives


- Symmetric (shared-key, same-key)
 - Stream cipher (one-time pad)
 - Block cipher (pseudo-random permutation PRP)
 - Modes of encryption
 - Message authentication code (MAC)
 - Authenticated Encryption (Encryption + MAC)
- **Asymmetric (public-private key)**
 - **Diffie-Hellman key agreement**
 - **Public-key encryption**
 - **Digital signature**
- **Others (unkeyed symmetric)**
 - **Cryptographic hash function and more**

Asymmetric Encryption Primitives

- Also known as public-key cryptography
- The encryption key is different from the decryption key
- A pair of public key and private key
 - Only private key must be secret.




Public-key encryption



- Receiver (Alice) generates (pk, sk) and sends pk to Bob
 - or publicize her pk (e.g., her webpage, central database)
 - assume that intended parties receive pk via *authenticated channels* (what if there's no such channel? We'll learn **PKI**)
- Anyone with pk can encrypt message
- Only the one with sk can decrypt message
 - cannot decrypt with pk (Eve may encrypt but can't decrypt)

Digital signature



- Only the one with sk can generate signature σ for message m
- Anyone with pk can verify the signature
 - assume that intended parties receive pk via *authenticated channels*
- **Non-repudiation:** signed document becomes proof that Alice indeed signed the document

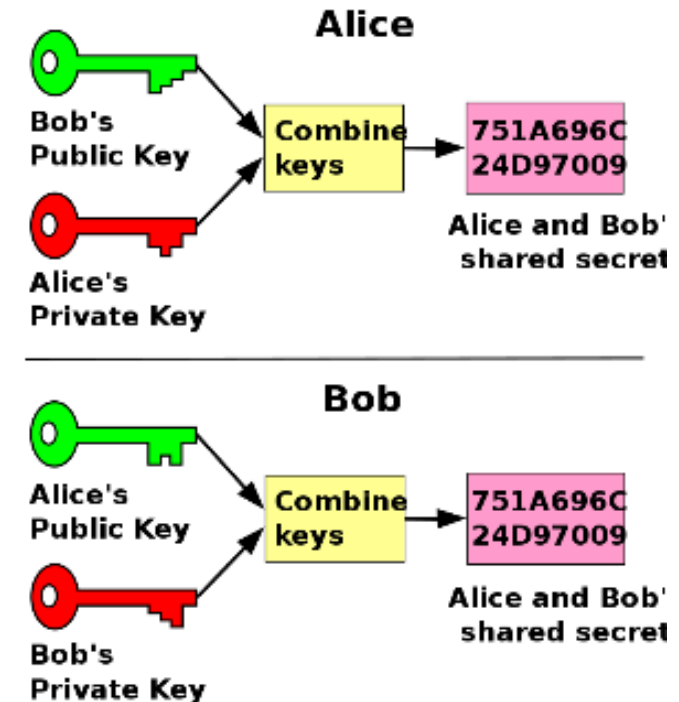
Asymmetric Encryption Primitive: DH

- **Diffie-Hellman (DH) key agreement**

- Public values: large prime p , generator g
- Alice has **secret value a** , Bob has **secret b**
- $A \rightarrow B$: **$g^a \pmod{p}$**
- $B \rightarrow A$: **$g^b \pmod{p}$**
- Bob computes $(g^a)^b = g^{ab} \pmod{p}$
- Alice computes $(g^b)^a = g^{ab} \pmod{p}$
- Can Eve compute $g^{ab} \pmod{p}$?

=> Secure against *eavesdropper*

- Why? **Discrete Logarithm Problem** is conjectured to be hard!



(Wikipedia)

Discrete Logarithm Problem

- Public values: **large prime p**, **generator g** of cyclic group G
 - A group G consisting of numbers relatively prime to p is a cyclic group. Generator g is an element of G that satisfies $g^{p-1} = 1 \pmod p$ (only p-1 satisfies this)
 - <https://eli.thegreenplace.net/2019/diffie-hellman-key-exchange/>
- $g^a \pmod p = x$
- Discrete logarithm problem: given x, g, and p, find a
- Table g=2, p=11

a	1	2	3	4	5	6	7	8	9	10	11
x	2	4	8	5	10	9	7	3	6	1	2


- Number field sieve** is fastest algorithm known today to solve discrete logarithm problem
 - Running time: $O(e^{(1.923+o(1))(\ln(p))^{1/3} (\ln(\ln(p)))^{2/3}})$ <- **Subexponential**
- Going back to DH Protocol:
 - Eve cannot find a (or b), and thus cannot compute $g^{ab} \pmod p$

Example

- $a=3, b=6, g=2, p=11$
 - $A \rightarrow B: g^a \pmod{p} = 2^3 \pmod{11} = 8$
 - $B \rightarrow A: g^b \pmod{p} = 2^6 \pmod{11} = 64 \pmod{11} = 9$
 - Bob computes $(g^a)^b \pmod{p} = 8^6 \pmod{11}$
 $= 262144 \pmod{11} = 3$
 - Alice computes $(g^b)^a \pmod{p} = 9^3 \pmod{11}$
 $= 729 \pmod{11} = 3$

Problem: Man-in-the-Middle Attack

- Public values: large prime p , generator g
- Problem: in Man-in-the-Middle attack, Mallory impersonates Alice to Bob and Bob to Alice



Asymmetric Primitive: RSA

- **RSA algorithm**

- Invented by Rivest, Shamir, Adleman in 1978

- Let p, q be large secret primes

- Pick e , compute d so $ed \equiv 1 \pmod{\phi(pq)}$

Relatively prime to $\phi(pq)$


$$= 1 + v \phi(pq)$$

$\Phi(n)$ counts the number of positive integers up to n that are relatively prime to n .
when p, q are prime numbers, $\Phi(pq) = (p-1)(q-1)$

- Public key $N=pq$, e

- Private key p, q, d

- Only I can decrypt
- Only I can sign



- Encryption/decryption of message M

- Encryption: $C = M^e \pmod{N}$

- Decryption: $M = C^d = M^{ed} \pmod{N} (= M^{1+v\phi(N)} \pmod{N} = M^{v\phi(N)} M \pmod{N})$

$M^{\phi(N)} \pmod{N} = 1$
(Fermat-Euler generalization)

- Signing message M

- Signature generation: $\Sigma = M^d \pmod{N}$

- Signature verification: $\Sigma^e = M^{ed} = M^{1+v\phi(N)} = M \pmod{N}$

- Secure? **Large prime factorization** is conjectured to be hard!

Example

- $p=3, q=7, N=21, \phi(pq)=12, e=5, d=5$
 - $\phi(pq)=(3-1)*(7-1)=2*6=12$
 - $e=5$
 - Want d such that $ed = 1 \bmod \phi(pq) = 1 \bmod 12$
 - Pick $d=5$
 - $M=2$
 - Signature $\Sigma = M^d \bmod N = 2^5 \bmod 21 = 11$
 - Signature verification:
 $\Sigma^e = (M^d)^e \bmod 21 = 11^5 \bmod 21 = 161051 \bmod 21 = 2$
 - Could also pick $e=7, d=7 \dots$

PKI (Public-key Infrastructure) Overview



- We need an **infrastructure to distribute keys**, in particular, the public key. This is generally known as key distribution or key management.
 - “RFC 4949 defines public-key infrastructure (PKI) as the set of hardware, software, people, policies, and procedures needed to **create, manage, store, distribute, and revoke digital certificates** based on asymmetric cryptography”
- Three ways of key distribution:
 - Public announcement (e.g., Email, web site)
 - Publicly-available directory (e.g., PGP public key server)
 - **Certificate**

Certificate and Certification Authority

- A **(digital) certificate** is a digital document that contains at least the following 4 main items
 - 1) The identity of an entity
 - 2) The associated public key
 - 3) The time window that this certificate is valid.
 - 4) The **signature of the CA**
 - Typically, it has other meta-information. Including (serial number, purpose, crypto algorithms)
- The CA keeps a **directory of public keys**.
- The CA has its own public-private key. The CA's public key has been securely distributed to all entities involved.
- CA is also responsible for **revocation** of certificates by maintaining CRL (certificate revocation list)

Certificate Example

Certificate	
www.google.com	
	GTS CA 1C3
	GTS Root R1
	GlobalSign Root CA
Subject Name	
Common Name	www.google.com
Issuer Name	
Country	US
Organization	Google Trust Services LLC
Common Name	GTS CA 1C3
Validity	
Not Before	Mon, 12 Dec 2022 08:19:43 GMT
Not After	Mon, 06 Mar 2023 08:19:42 GMT
Subject Alt Names	
DNS Name	www.google.com
Public Key Info	
Algorithm	Elliptic Curve
Key Size	256
Curve	P-256
Public Value	04:6E:C8:BC:8E:97:E3:66:D1:1E:3D:04:D8:F0:1D:69:E0:05:8A:16:19:72:D5:0D:B8:2D:0C:79:F4:F8:1E:99:65:2F...

Comparison **Sym** vs **Asym** Crypto

Symmetric crypto

- Need shared secret key
- 86 bit key for high security (year 2020)
- ~1,000,000 ops/s on 1GHz processor
- 10x speedup in HW

Asymmetric crypto

- Need authentic public key
- 2048 bit key (RSA) for high security
- ~100 signatures/s, ~1000 verify/s (RSA) on 1GHz processor
- Limited speedup in HW

Recommendation of Key size [Lenstra and Verheul]

	1982	1995	2002	2010	2020	2030	2040
Sym	56	66	72	78	86	93	101
RSA	417	777	1028	1369	1881	2493	3214

Basic Cryptographic Primitives

- Symmetric (shared-key, same-key)
 - Stream cipher (one-time pad)
 - Block cipher (pseudo-random permutation PRP)
 - Modes of encryption
 - Message authentication code (MAC)
 - Authenticated Encryption (Follow-up for Week 1)
- Asymmetric (public-private key)
 - Diffie-Hellman key agreement
 - Public-key encryption
 - Digital signature
- **Others (unkeyed symmetric)**
 - **Cryptographic hash function and more**

Hash function vs. cryptographic hash func



- Hash functions
 - functions that takes inputs of arbitrary length and compress them into short, fixed-length outputs (or **digests**)
 - $H: \{0,1\}^* \rightarrow \{0,1\}^l$
 - evaluation of H is efficient and algorithm is public
- Cryptographic hash functions usually require three properties:
 - Collision resistance (strong collision resistance)
 - Second preimage resistance (weak collision resistance)
 - Preimage resistance (oneway)

Notions of security (informal analysis)

- [P1] *Collision resistance*: it is computationally infeasible to find x, x' such that $x' \neq x$ and $H(x') = H(x)$
- [P2] *Second-preimage resistance*: **given x** , it is computationally infeasible to find $x' \neq x$ such that $H(x') = H(x)$
- [P3] *Preimage resistance*: **given y** , it is computationally infeasible to find x such that $H(x) = y$. (i.e., H is one-way)

Notions of security (informal analysis)

(cont'd)

- [P1] *Collision resistance* implies [P2] *Second-preimage resistance*
 - if adversary can find $x' \neq x$ for given x such that $H(x') = H(x)$, then it can clearly find a colliding pair x and x' .
- [P2] *Second-preimage resistance* implies [P3] *Preimage resistance*:
 - Assume H is second preimage resistance but not preimage resistance, then we show contradiction.
 - For given x (and $y = H(x)$), one can find x' that satisfies $H(x') = y$ because H does not have preimage resistance.
 - When domain is very large, one can find $x' \neq x$ with high probability. Thus, contradiction.

Why are these properties important?

- What if an attacker can easily find a collision?
- Digital signature would not be reliable.
 - Digital signature is often made on the hash value of message

$$m, \sigma = \text{Sign}_{sk}(\mathbf{H}(m))$$

- Finding collision m' , an attacker can send the following!

$$m', \sigma$$

Attack Complexity: *Preimage resistance*

- Assume cryptographic hash function with n -bit output
- Given output y , how many operations does it take to find any x , such that $H(x) = y$?
 - Assumption: best attack is random search
 - For each trial x , probability that output is y is 2^{-n}
 - $P[\text{find } x \text{ after } m \text{ trials}] = 1 - (1 - 2^{-n})^m$
 - Rule of thumb: find x after 2^{n-1} trials on average

Attack Complexity: *Second-preimage resistance*


- Given input x , how many operations does it take to find another $x' \neq x$, s. t. $H(x) = H(x')$?
 - Assumption: best attack is random search
 - For each trial x' , probability that output is equal is 2^{-n}
 - $P[\text{find } x' \text{ after } m \text{ trials}] = 1 - (1 - 2^{-n})^m$
 - Rule of thumb: find x' after 2^{n-1} trials on average

Attack Complexity: Collision Resistance

- How many operations does it take to find x and x' , s. t. $x' \neq x$ and $H(x) = H(x')$?
 - Assumption: best attack is random search
 - Algorithm picks random x' , checks whether $H(x')$ matches any other output value **previously seen**
 - $P[\text{find collision after } m \text{ trials}] =$
 $1 - (1 - 1/2^n)(1 - 2/2^n)(1 - 3/2^n) \dots (1 - (m+1)/2^n)$
 - Rule of thumb: find collision after $2^{n/2}$ trials on average
 - $1.17 * 2^{n/2}$ to be a bit more precise (with 50% probability)
 - $2.146 * 2^{n/2}$ (with 90% probability)

Birthday Paradox

- How many people need to be in a room to have a probability > 50% that at least two people have the same birthday?
- Answer: approximately $1.17 \cdot 365^{1/2} \sim 22.4$



Implication of Birthday Paradox

- Let us consider a hash function that generates 40-bit digest
- How many hash calculations are needed to find collision with probability > 50%?
 - $1.17 * (2^{40})^{(1/2)} = 1.17 * 2^{20} = 1,226,833$
- How many hash calculations are needed to find collision with probability > 90%?
 - $2.146 * (2^{40})^{(1/2)} = 2.146 * 2^{20} = 2,250,244$
- Recent GPU (e.g., Geforce RTX 3080) can calculate 100million hash values per second.
- The size of hash value matters!

SHA-1: Broken! or Broken?

- Retires by the end of 2030



- SHA-1 has a 160-bit message digest and thus has 80-bit strength
- In 2005 an attack shows it requires only 2^{69} operations to find a hash collision
 - Down to 2^{61} in 2010
- According to 2017 report, finding two different PDF files with same SHA-1 hash value required 9,223,372,036,854,775,808 ($=2^{63}$) operations
 - 6,500 years of single-CPU, 110 years of single-GPU

Fingerprinting

- Instead of storing the original data, one can simply store a short hash digest
- Examples:
 - Virus fingerprinting
 - Deduplication
 - P2P file sharing

MySQL Community Server 8.0.27

Select Operating System:
Microsoft Windows

Looking for previous GA versions?

Recommended Download:

MySQL Installer for Windows

All MySQL Products. For All Windows Platforms. In One Package.

Starting with MySQL 5.6 the MySQL Installer package replaces the standalone MSI packages.

Windows (x86, 32 & 64-bit), MySQL Installer MSI [Go to Download Page >](#)

Other Downloads:

Windows (x86, 64-bit), ZIP Archive (mysql-8.0.27-winx64.zip)	8.0.27	209.4M	Download MD5: 9d8e719f18835d8b512c5f8e9182d9a6a Signature
Windows (x86, 64-bit), ZIP Archive Debug Binaries & Test Suite (mysql-8.0.27-winx64-debug-test.zip)	8.0.27	509.6M	Download MD5: fb1a80a91b1aeb95c9be44817978592f Signature

One-Way Hash Chains

- Versatile cryptographic primitive
- Construction
 - Pick random r_N and public one-way function F
 - $r_i = F(r_{i+1})$
 - Secret value: r_N , public value: r_0

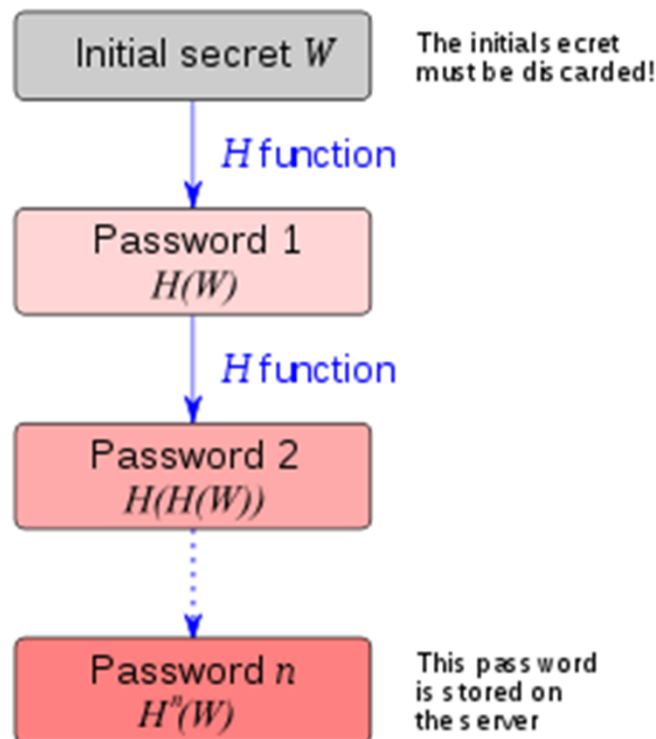


- Properties
 - Use in reverse order of construction: r_0, r_1, \dots, r_N
 - Infeasible to derive r_i from r_j ($j < i$)

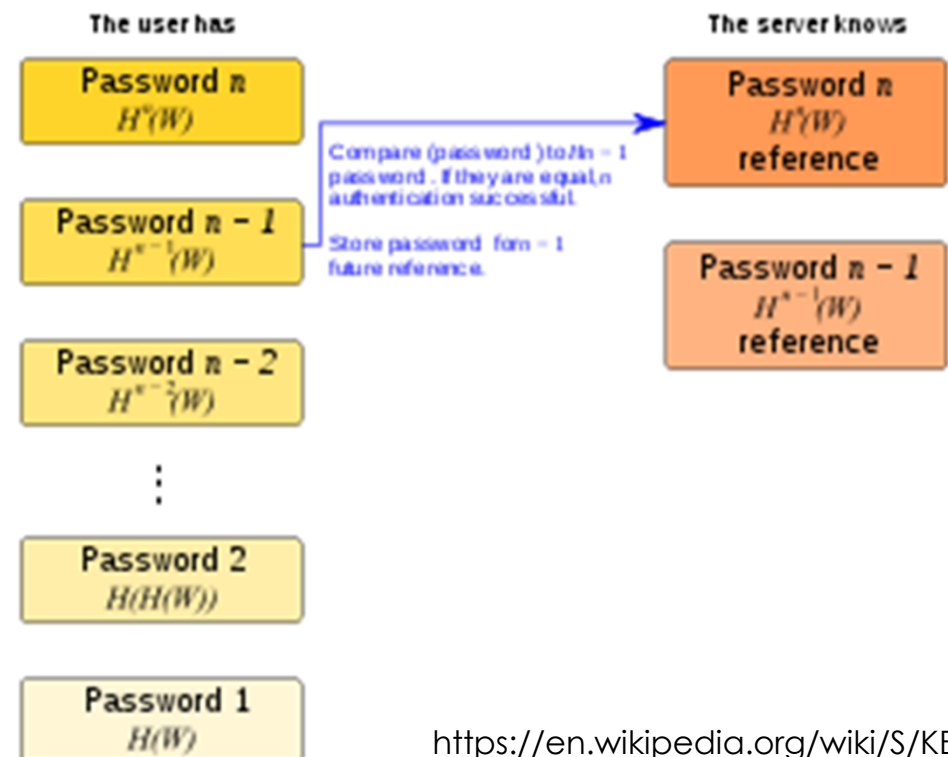
Application: S/KEY

- One-time password authentication
- Consider a setting where user authenticates herself with password W over insecure communication channel
 - Password registration must be done in secure way.

S/KEY password generation



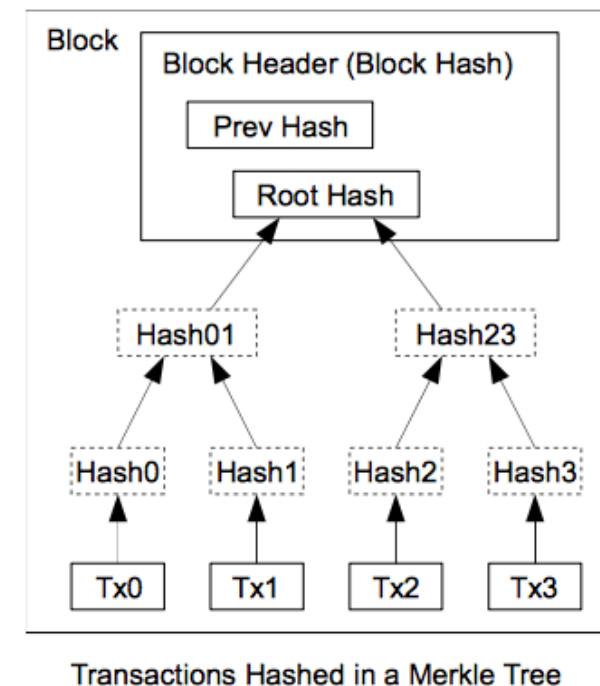
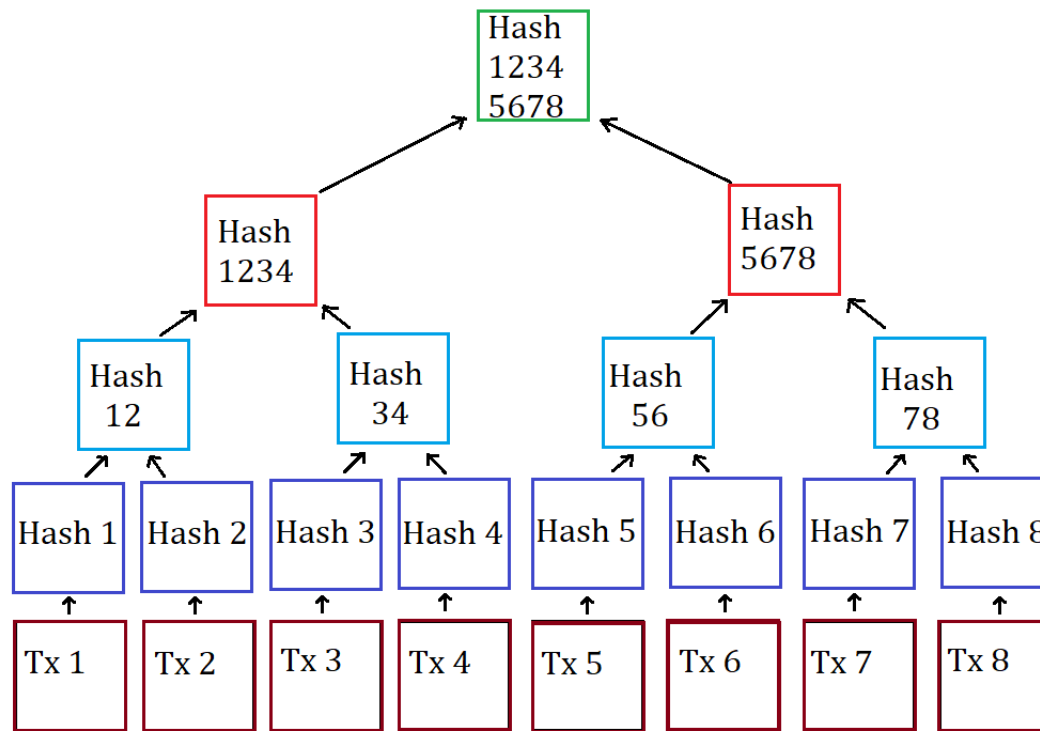
S/KEY authentication



<https://en.wikipedia.org/wiki/S/KEY>

Merkle Hash Trees

- Authenticate a sequence of data values $Tx_1, Tx_2, \dots Tx_N$
- Construct binary tree over data values
- Widely used data structure, e.g., Blockchain



<https://hackernoon.com/achieving-blockchain-scalability-with-sparse-merkle-trees-and-bloom-filters-3b9945f003f>

Password hashing

- Passwords are not stored in plaintext but hashed and stored (e.g., *passwd* file of UNIX system)
- Preimage resistance of H makes inverting it difficult
- Attacks
 - password space is quite limited (e.g., dictionary of English words) \Rightarrow domain of H is small
 - pre-computation makes it even easier (rainbow table)
- Mitigations
 - slow hash functions (e.g., bcrypt)
 - salt: store (h,s) together when $h=H(\text{password}||s)$, s is random
 - pre-computation is hard

Commitment Schemes

- $C(m)$ = commitment to a value m ; e.g., a bid in an auction
 - “hides” the content of m and “binds” the sender to value m
 - value of m revealed later (when auction is over)
 - e.g., $r \leftarrow \{0,1\}^l$ per commitment; $C(m) = h(r||m)$; later reveal r, m .
- Needs secure hash function that has preimage resistance, and collision resistance.

What's next?

- Basics for authentication and secure communication
 - Key Establishment / Distribution
 - Authentication using Symmetric Key
 - Authentication using PKI
 - Authentication using KDC
 - TLS

QUESTIONS?

Reading for Week 4: PKI Security



- **SSL and HTTPS: Revisiting past challenges and evaluating certificate trust model enhancements [IEEE S&P 2013]**
 - A guided tour of history of PKI research and development!
- ***Certificate Transparency* [COMMUNICATIONS OF THE ACM 2014]**
 - Most widely used security mechanisms for certificate misissuance problems