

CS5321 Network Security


Week5: TCP/IP Security

Daisuke MASHIMA

2022/23 Sem 2

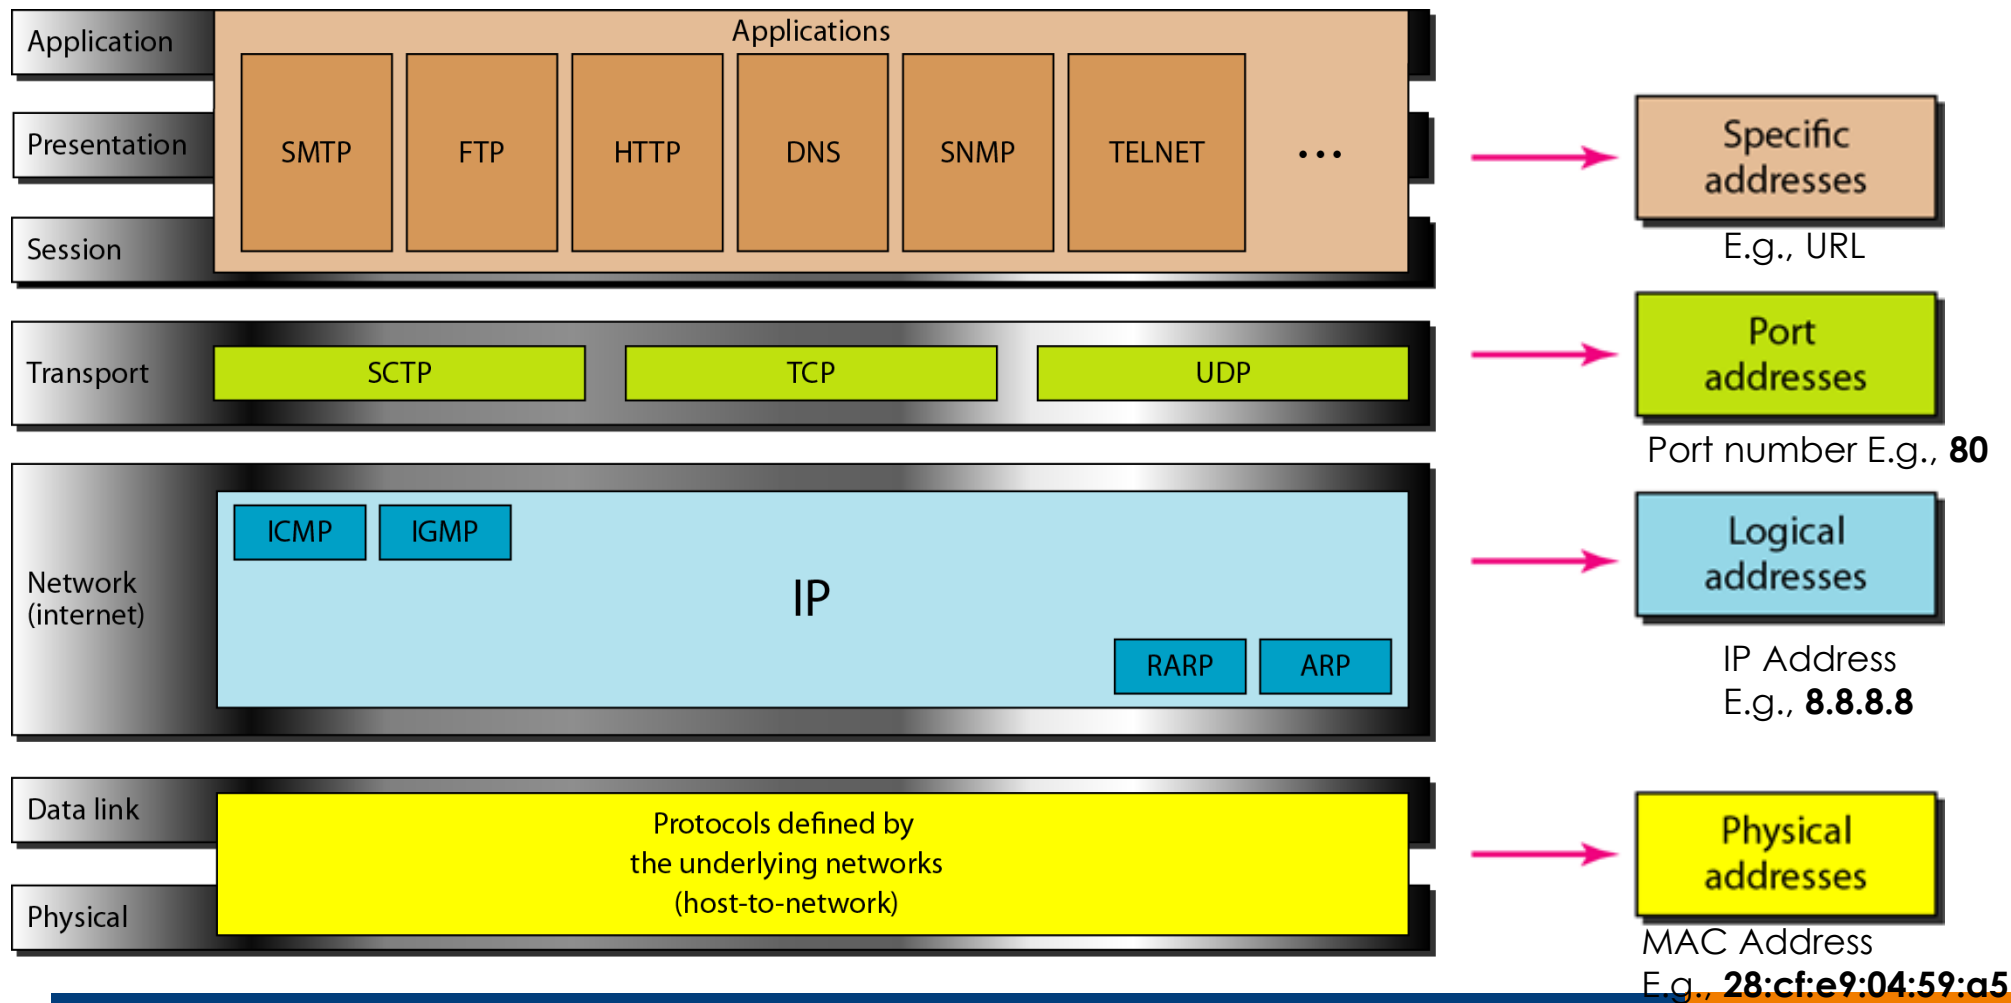
Overview

- In this lecture
 - IP vulnerabilities
 - New IP architecture
 - TCP vulnerabilities
 - TCP Hijacking



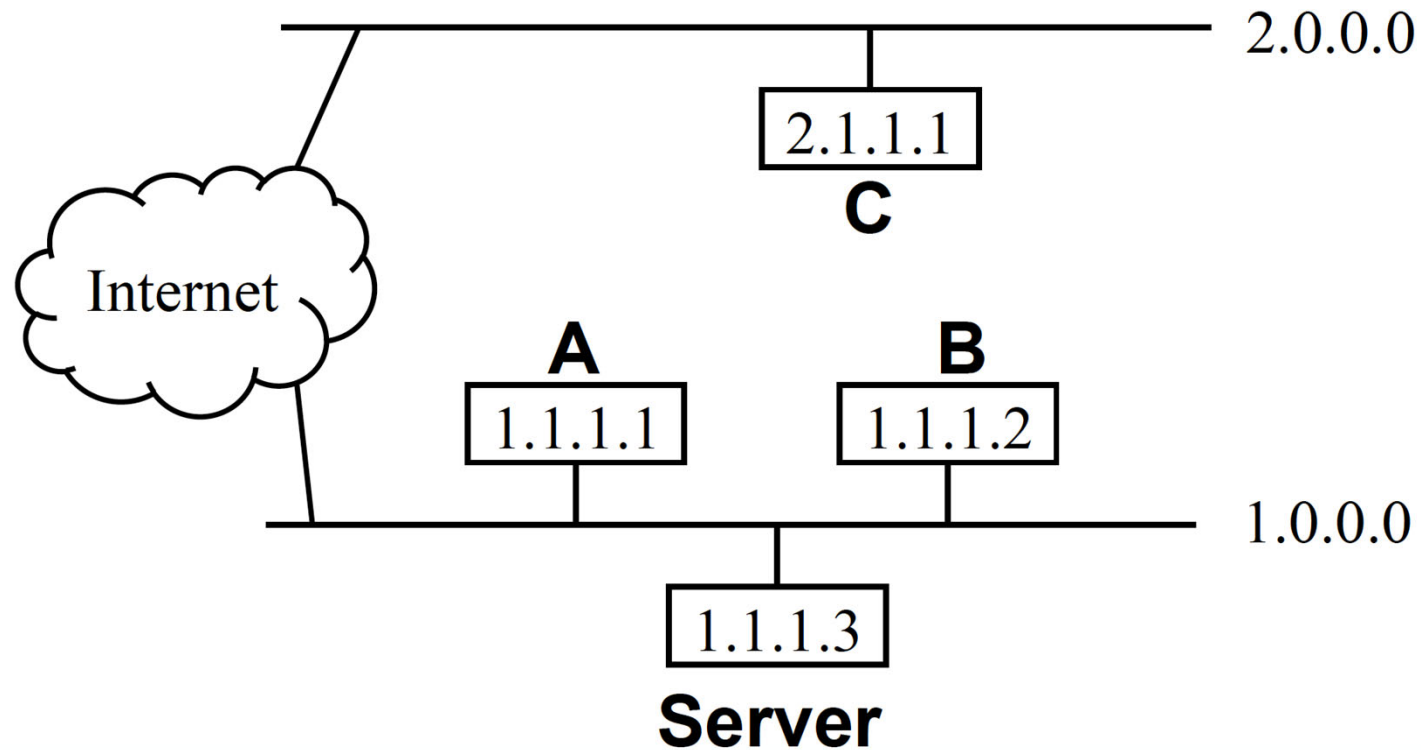
TCP/IP Stack

- Foundation for various Internet-based network services




IP VULNERABILITIES

Security Issues in IP Networks



- Security issues for communication between A, B, C, and Server?

Basic Security Issues



A send S(server) a packet (P)

- $A \rightarrow S: P$ (using the IP protocol)
- How can S know that the packet originated from A?
- Can B overhear P?
- Can B impersonate A to S?
- Can C impersonate A to S?

IP Packet


```
▼ Internet Protocol Version 4, Src: 172.16.8.12, Dst: 172.16.1.41
  0100 .... = Version: 4
  .... 0101 = Header Length: 20 bytes (5)
  > Differentiated Services Field: 0x00 (DSCP: CS0, ECN: Not-ECT)
    Total Length: 113
    Identification: 0x07d8 (2008)
  > Flags: 0x40, Don't fragment
    Fragment Offset: 0
    Time to Live: 64
    Protocol: TCP (6)
    Header Checksum: 0xd159 [validation disabled]
    [Header checksum status: Unverified]
    Source Address: 172.16.8.12
    Destination Address: 172.16.1.41
```

IP Packet captured by Wireshark

- IP packet “**claims**” source and destination IP addresses
- Receiver “**assumes**” that the sender address is the one specified as source address.
- No authentication!


Flaw: Use IP Address for Authentication

- IP source address can be easily spoofed!
- Easy to mount attack for another machine on the same network
- Example: r-utilities (rlogin, rsh, rcp)
 - Consider Server trusts admin's machine A
 - If B spoofs A's address, user on B can log in to Server
- Is it enough for meaningful attack?



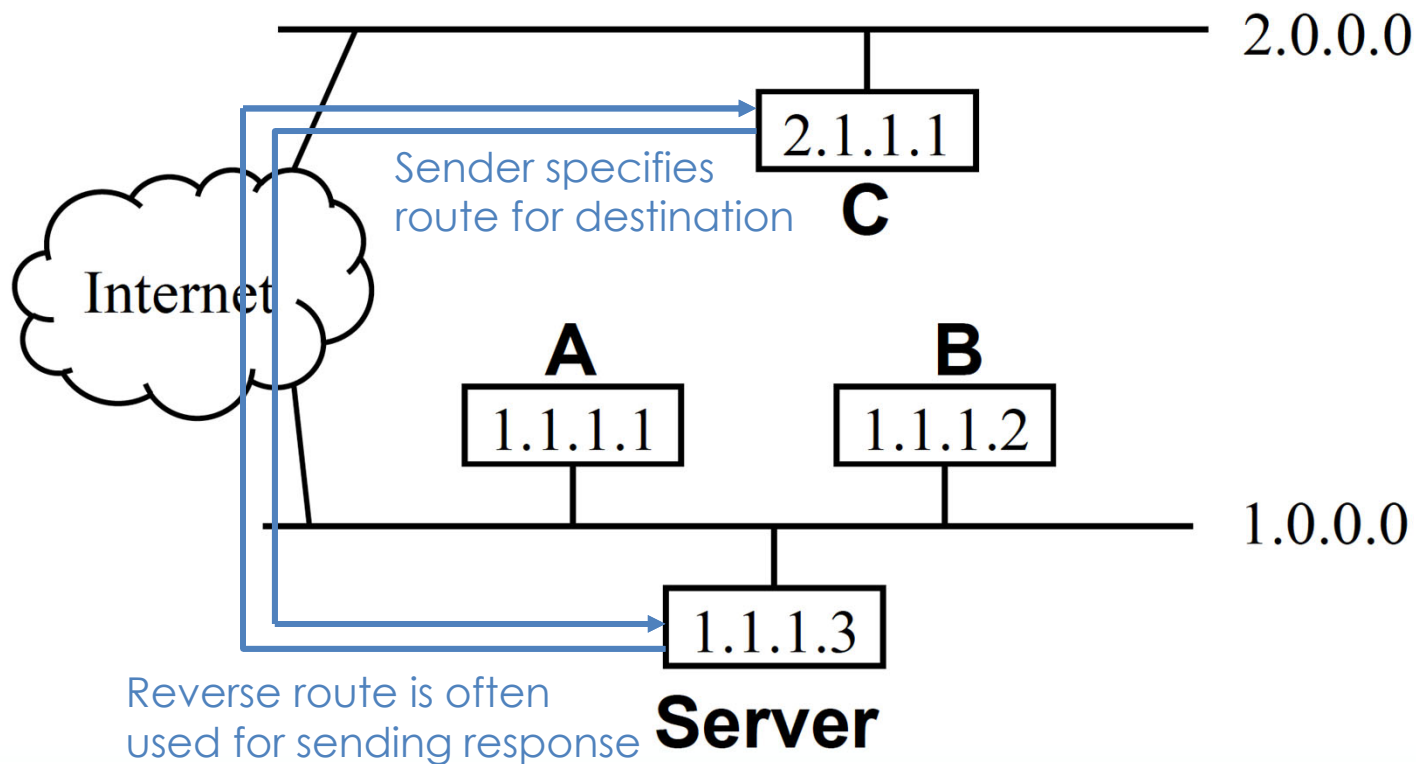
Flaw: Use IP Address for Authentication

- Problem for attacker:
 - A receives S's responses to B's spoofed packets, as the destination address is A!
 - A will respond with a TCP Reset (RST) packet which closes the connection
- Solution for attacker:
 - By overflowing A's queues with connection requests, it is likely that A drops S's replies
 - Note: **DoS attack** is used to enable another attack
 - In the same network (more specifically collision domain), the attacker can still "see" the S's response to A by using "**promiscuous mode**".



Flaw: Use IP Address for Authentication

- How can C impersonate A to S here?



Possible with **source routing!**

Big problem of current IP: IP Spoofing

- **Ingress filtering**
 - Let the upstream network block spoofed IPs
 - Lack of incentive
- **iTrace** (<https://www.cs.columbia.edu/~smb/talks/ietf47/itrace.pdf>)
 - 1 in 20,000 packets triggers a router to send an ICMP packet to a destination with route information for traceback
 - Needs authentication
- **Packet marking**
 - Routers mark 16-bit IP ID field with information that enables traceback
 - Needs changes to routers
- ***Open problem of the Internet!***

New IP Architecture: Accountable Internet Protocol

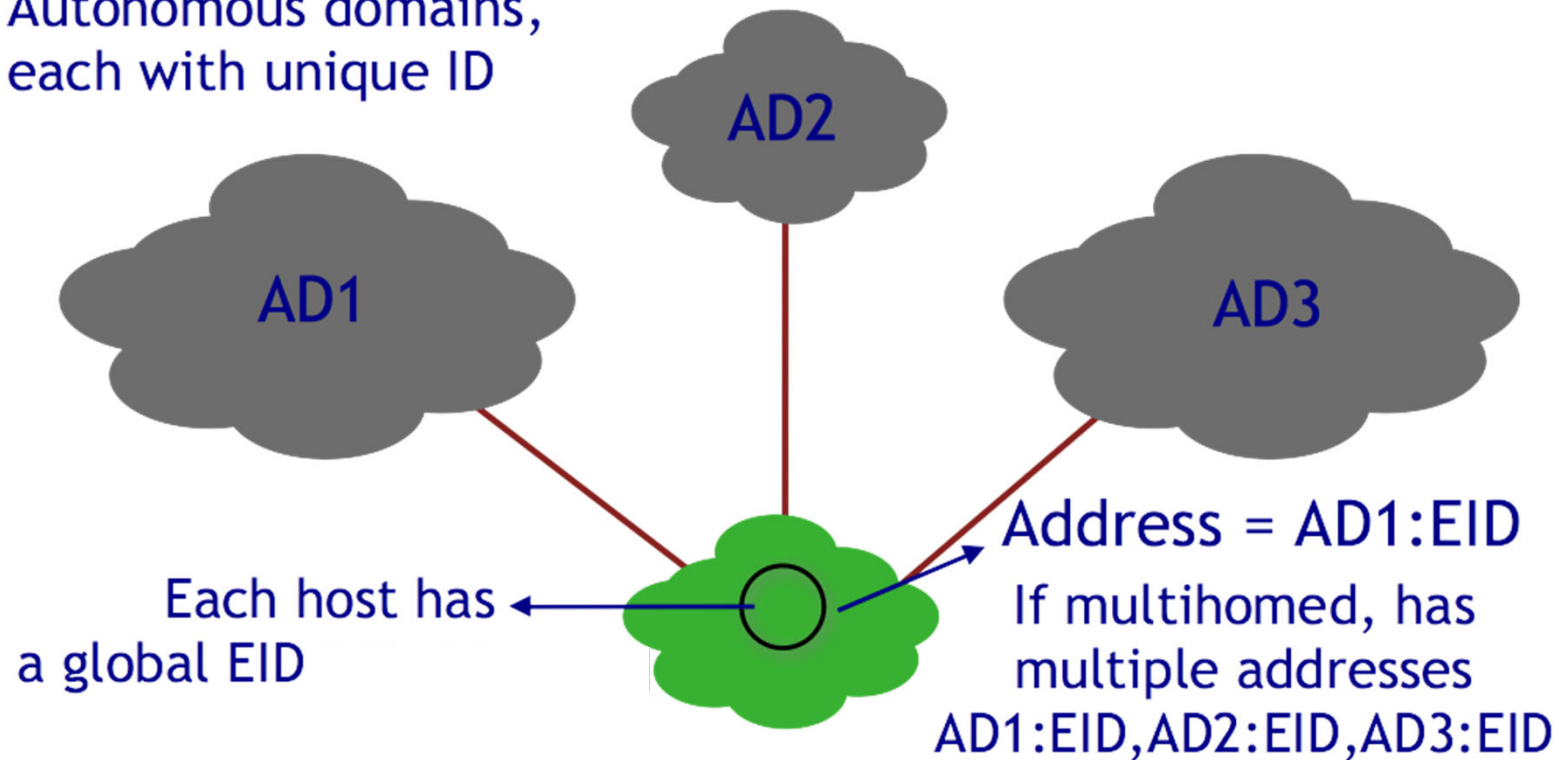
- Paper: **Accountable Internet Protocol** by David Andersen et al.
- Internet Protocol is old. No security in it.
- Lots of security patches but they are not satisfactory:
 - Complicated Mechanisms
 - Many details to circumvent IP weaknesses
 - External Sources of Trust
 - Trusted certificate authorities (e.g., S-BGP)
 - Operator Vigilance
 - Semi-manual configuration (e.g., filters, registries)

IP Layer Names Don't Have Secure Bindings

- Three kinds of IP layer names:
 - IP address, IP prefix, AS (autonomous system) number
- No secure binding of host to its IP addresses
- No secure binding of AS number to its IP prefixes
- Many problems become easier to solve with **network-layer accountability**: *Ability to associate a principal with a message*

AIP Addressing

Autonomous domains,
each with unique ID

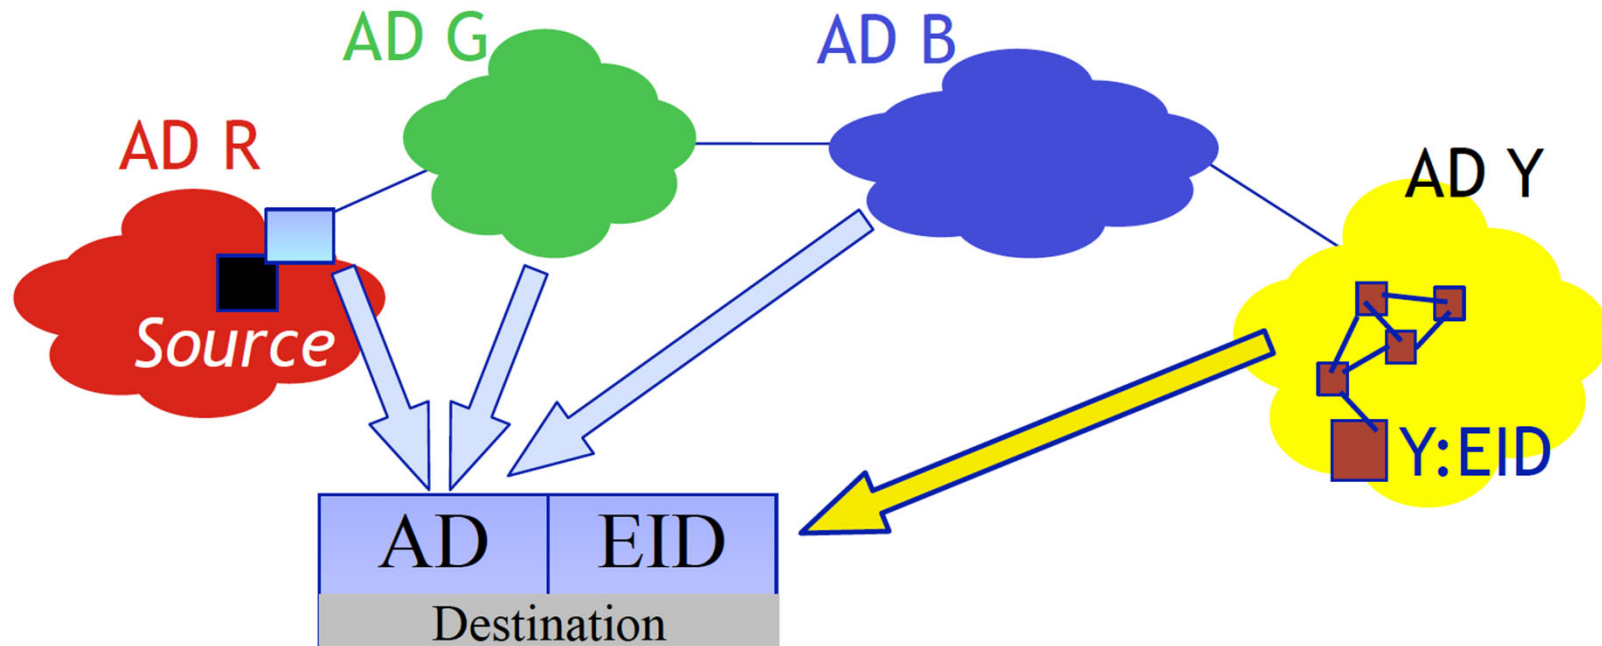


Key Idea:

AD and EID are *self-certifying flat names*

- AD = hash(public_key_of_AD)
- Self-certification binds name to named entity

AIP Forwarding and Routing



- Inter-AD routing & forwarding: AD #s only.
 - When packet is crossing AD boundary, checks are performed
- Intra-AD routing disseminates EIDs.
- Many routing protocols possible - derive security from AIP self-certification

Detecting & Preventing *Spoofing*

- Self-certified entity can prove it sent message:

Let:

rs = Per-router secret, rotated once per minute

$\text{HMAC}_{\text{key}} \langle M \rangle$ = Message authentication code of M

$H \langle P \rangle$ = Hash of P

iface = Interface on which packet arrived

Source $S_{AD} : S_{EID} \rightarrow$ **Dest** $D_{AD} : D_{EID}$
Packet P.

Router R1 \rightarrow **Source**:

Verification packet $V =$

$\text{HMAC}_{rs} \langle S_{AD} : S_{EID} \rightarrow D_{AD} : D_{EID}, H \langle P \rangle, \text{iface} \rangle$

Source \rightarrow **R1**:

$\{\text{accept}, K_{S_{EID}}, V\}_{K_{S_{EID}}^{-1}}$

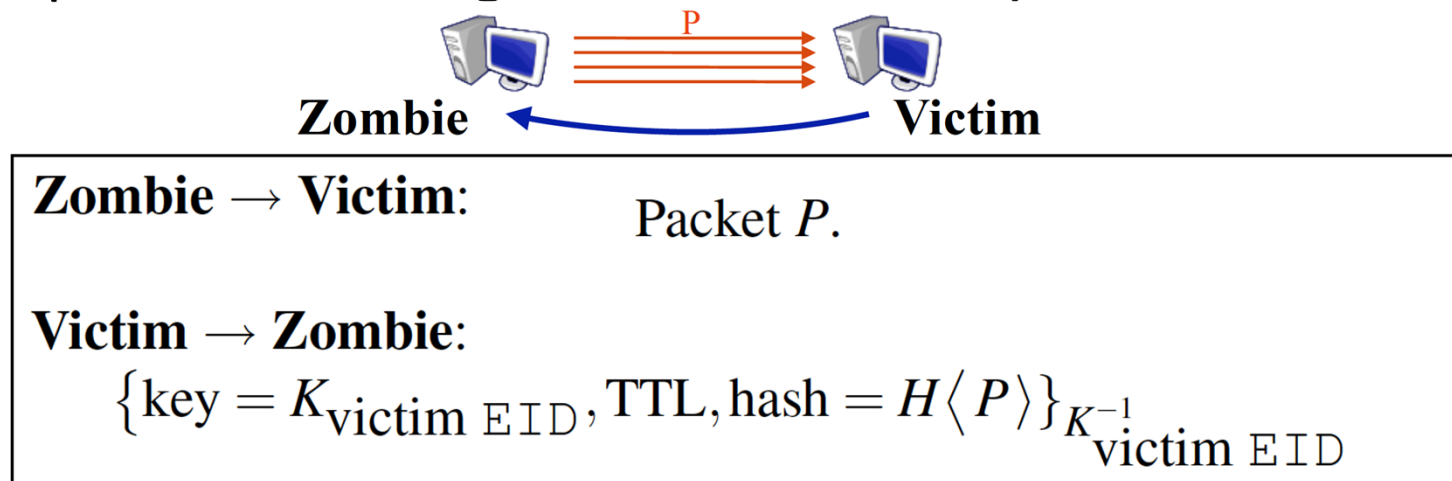
- Routers or hosts seeing packet can check the AD or EID using a challenge-response protocol
- Verification is done at the nearest router as well as intermediate ADs

Address Minting?

- Self-certifying address of AIP enables spoofing detection
- However, it does not prevent malicious hosts from creating large number of EIDs (or minting many identities!)
 - New EIDs can be used for DoS or filter circumvention
- Unfortunately, no clear technical solution
 - Simple solution would be to use public key infrastructure but it violates the core philosophy of AIP
- Engineering solutions (or operational countermeasures):
 - Each first-hop router limits number of new EIDs/second

AIP Enables Secure *Shut-Off*

- Problem: Compromised zombie sending stream of unwanted DoS traffic to victim
 - Zombie is “well-intentioned”, (i.e., owner benign) but compromised owing to its vulnerability.



- Shut-off scheme implemented in “**smart-NIC**” (NIC firmware update requires physical access)
- If smart-NIC saw P that matches $H(P)$, it installs a filter to block it
- Hardware requirements are practical
 - Bloom filter* for replay prevention (8MB SRAM)

* Space-efficient data structure “that is used to test if a certain element is a member of a set” (wikipedia)

Takeaways from AIP paper

- A whole new IP is needed!
 - AIP is one nice proposal: focusing on **accountability**
- Good properties:
 - Spoofing detection
 - Shut-off protocol
 - Secure routing
- Yet, some concerns:
 - Scalability
 - Inability to do CIDR-like aggregation of addresses (e.g., 198.51.100.xxx/24)
 - Perhaps more importantly, lack of backward compatibility

TCP VULNERABILITIES

- Transmission Control Protocol
 - Works at **Transport Layer**, on top of IP
- TCP provides reliable data transfer using the best effort IP service
- Typical TCP packet exchange
 - $A \rightarrow B: \text{SYN}(\text{ISN}_A)$
 - $B \rightarrow A: \text{SYN}(\text{ISN}_B), \text{ACK}(\text{ISN}_A + 1)$
 - $A \rightarrow B: \text{ACK}(\text{ISN}_B + 1)$
 - $A \rightarrow B: \text{data} \dots$

} TCP 3-way handshake

ISN: Initial Sequence Number

TCP network trace

No.	Time	Source	Destination	Protocol	Length	Info
67	3.307390	172.16.2.100	172.16.2.41	TCP	66	1133 → 80 [SYN] Seq=0 Win=8192 Len=0 MSS=1460 WS=256 SACK_PERM=1
69	3.310325	172.16.2.41	172.16.2.100	TCP	66	80 → 1133 [SYN, ACK] Seq=0 Ack=1 Win=14600 Len=0 MSS=1460 SACK_PERM=1 WS=32
71	3.311980	172.16.2.100	172.16.2.41	TCP	54	1133 → 80 [ACK] Seq=1 Ack=1 Win=65536 Len=0
74	3.385125	172.16.2.100	172.16.2.41	HTTP	466	GET / HTTP/1.1
75	3.387219	172.16.2.41	172.16.2.100	TCP	60	80 → 1133 [ACK] Seq=1 Ack=413 Win=15680 Len=0
76	3.389015	172.16.2.41	172.16.2.100	HTTP	189	HTTP/1.1 307 Temporary Redirect

> Frame 69: 66 bytes on wire (528 bits), 66 bytes captured (528 bits) on interface \Device\NPF_{B871E50A-4813-41B4-BA95-95654B00ECB6}, id 0

> Ethernet II, Src: WAGOKont_40:d0:8d (00:30:de:40:d0:8d), Dst: GoodWayI_17:fb:5d (00:50:b6:17:fb:5d)

> Internet Protocol Version 4, Src: 172.16.2.41, Dst: 172.16.2.100

✓ Transmission Control Protocol, Src Port: 80, Dst Port: 1133, Seq: 0, Ack: 1, Len: 0

Source Port: 80

Destination Port: 1133

[Stream index: 1]

[TCP Segment Len: 0]

Sequence Number: 0 (relative sequence number)

Sequence Number (raw): 3189983741

[Next Sequence Number: 1 (relative sequence number)]

Acknowledgment Number: 1 (relative ack number)

Acknowledgment number (raw): 287198956

1000 = Header Length: 32 bytes (0)

> Flags: 0x012 (SYN, ACK)

Window: 14600

[Calculated window size: 14600]

Checksum: 0x6c65 [unverified]

[Checksum Status: Unverified]

Urgent Pointer: 0

> Options: (12 bytes), Maximum segment size, No-Operation (NOP), No-Operation (NOP), SACK permitted, No-Operation (NOP), Window scale

> [SEQ/ACK analysis]

> [Timestamps]

TCP packets captured with Wireshark

TCP ISN Prediction Attack

- Typical TCP packet exchange
 - $A \rightarrow B: \text{SYN}(\text{ISN}_A)$
 - $B \rightarrow A: \text{SYN}(\text{ISN}_B), \text{ACK}(\text{ISN}_A + 1)$
 - $A \rightarrow B: \text{ACK}(\text{ISN}_B + 1)$
 - $A \rightarrow B: \text{data} \dots$
- Attack:
 - $M(A) \rightarrow B: \text{SYN}(\text{ISN}_A)$
 - $B \rightarrow A: \text{SYN}(\text{ISN}_B), \text{ACK}(\text{ISN}_A + 1)$
 - $M(A) \rightarrow B: \text{ACK}(\text{ISN}_B + 1)$
 - $M(A) \rightarrow B: \text{data} \dots$

$M(A)$: Malicious party impersonating A


TCP ISN Prediction

- Are these good choices for next TCP ISN?
 - Always start at same ISN
 - After each connection, $\text{ISN}++$
- **No, attacker can predict next ISN!**
- Better choices for ISN?
 - $\text{ISN} = \text{rand}()$ function of C library?
 - $\text{current ISN} = H(\text{prev ISN}, k)$?
 - $\text{ISN} = \text{DES}_k(\text{counter}++)$?

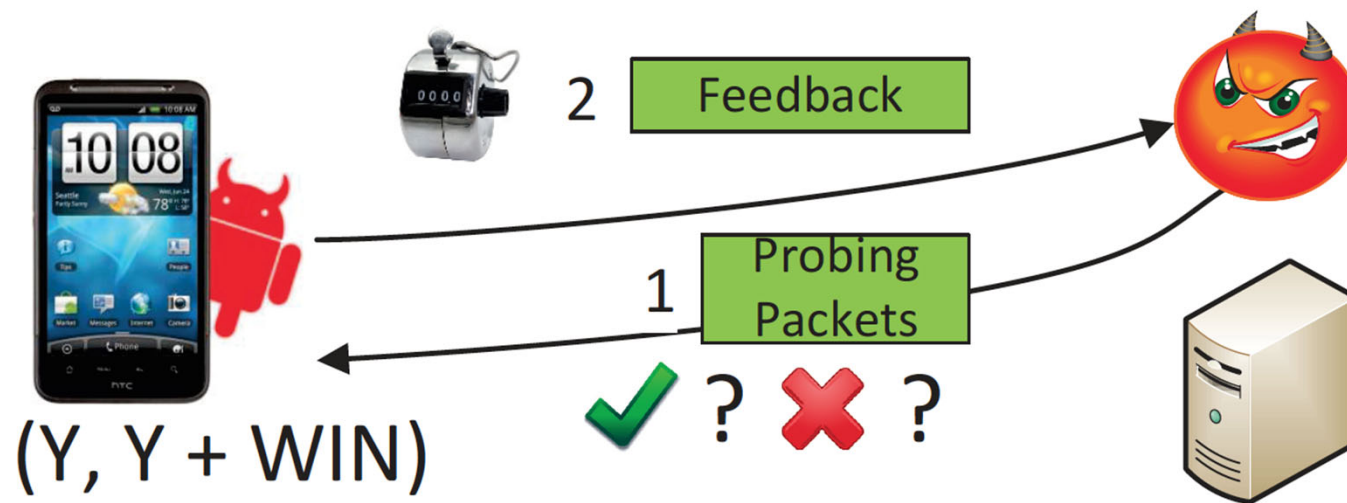
TCP Sequence Inference Attack

- Paper: **“Collaborative TCP Sequence Number Inference Attack — How to Crack Sequence Number Under A Second”** by Zhiyun Qian et al.
- Off-path attacks
 - Can write to existing TCP connection by guessing sequence numbers
 - Works even though Initial sequence number nowadays are randomized (2^{32})

$X = ? \ Y = ?$ 😈



TCP sequence number inference attack

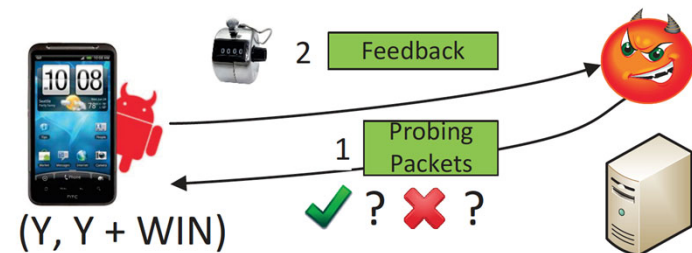


- Required information
 - Target four tuples (source/dest IP, source/dest port)
 - **Feedback** on whether guessed sequence numbers are correct
- “Can an **unprivileged** malware accurately learn if the probing packet is in or out of receive window?”
- “Or, can learn even more useful information?”

Preliminary step: obtaining target four tuples

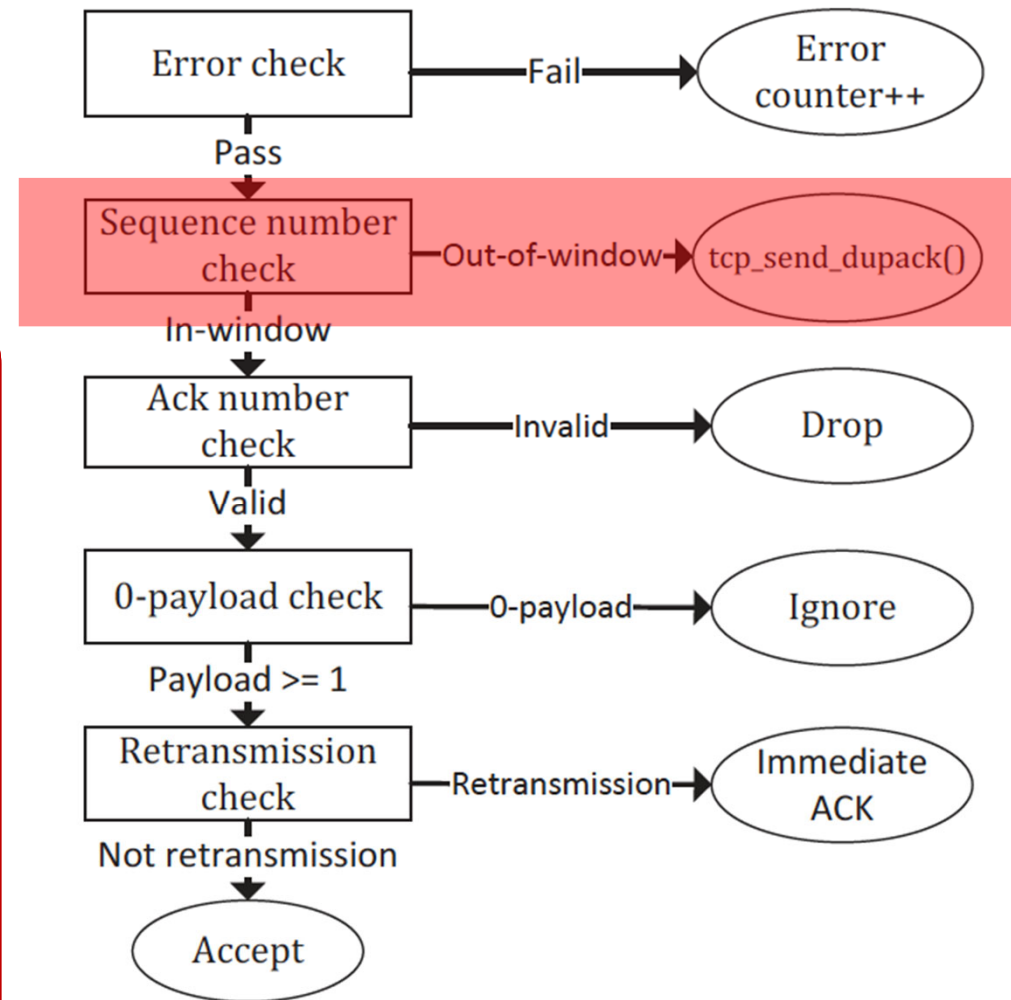
- On-site **unprivileged** malware
 - netstat** (no root required) provides four tuples:
 - (srcIP, dstIP, srcPort, dstPort)

```
netstat -nn
Active Internet connections
Proto Recv-Q Send-Q Local Address Foreign Address
probing
Initiate fake connections
(state)
tcp4 37 0 192.168.1.102.50469 199.47.219.159.443
CLOSE_WAIT
tcp4 37 0 192.168.1.102.50468 174.129.195.86.443
CLOSE_WAIT
tcp4 37 0 192.168.1.102.50467 199.47.219.159.443
CLOSE_WAIT
tcp4 0 0 192.168.1.102.50460 199.47.219.159.443
LAST_ACK
tcp4 0 0 192.168.1.102.50457 199.47.219.159.443
LAST_ACK
tcp4 0 0 192.168.1.102.50445 199.47.219.159.443
```



Linux TCP incoming packet validation logic

- 5 steps to filter out invalid TCP packets (Linux)
- **tcp_send_dupack()**
 - update the global states when out-of-window packet received
 - if received seq# < Y
 - DelayedACKLost+=1
 - otherwise
 - DelayedACKLost+=0
 - DelayedACKLost can be read by any malware!




TCP incoming packet validation logic in Linux 3.2.6


***perfect side-channel
for seq# inference!***

How to find Y efficiently?

- Binary search!
 - total search space: 2^{32} ($\sim 4G$)
 - each iteration, we can eliminate half of the space!



(a). First iteration



(b). Second iteration

- at most 32 iteration (i.e., 32 probing packets) needed
- see the paper for more optimizations

Client-side TCP injection attack

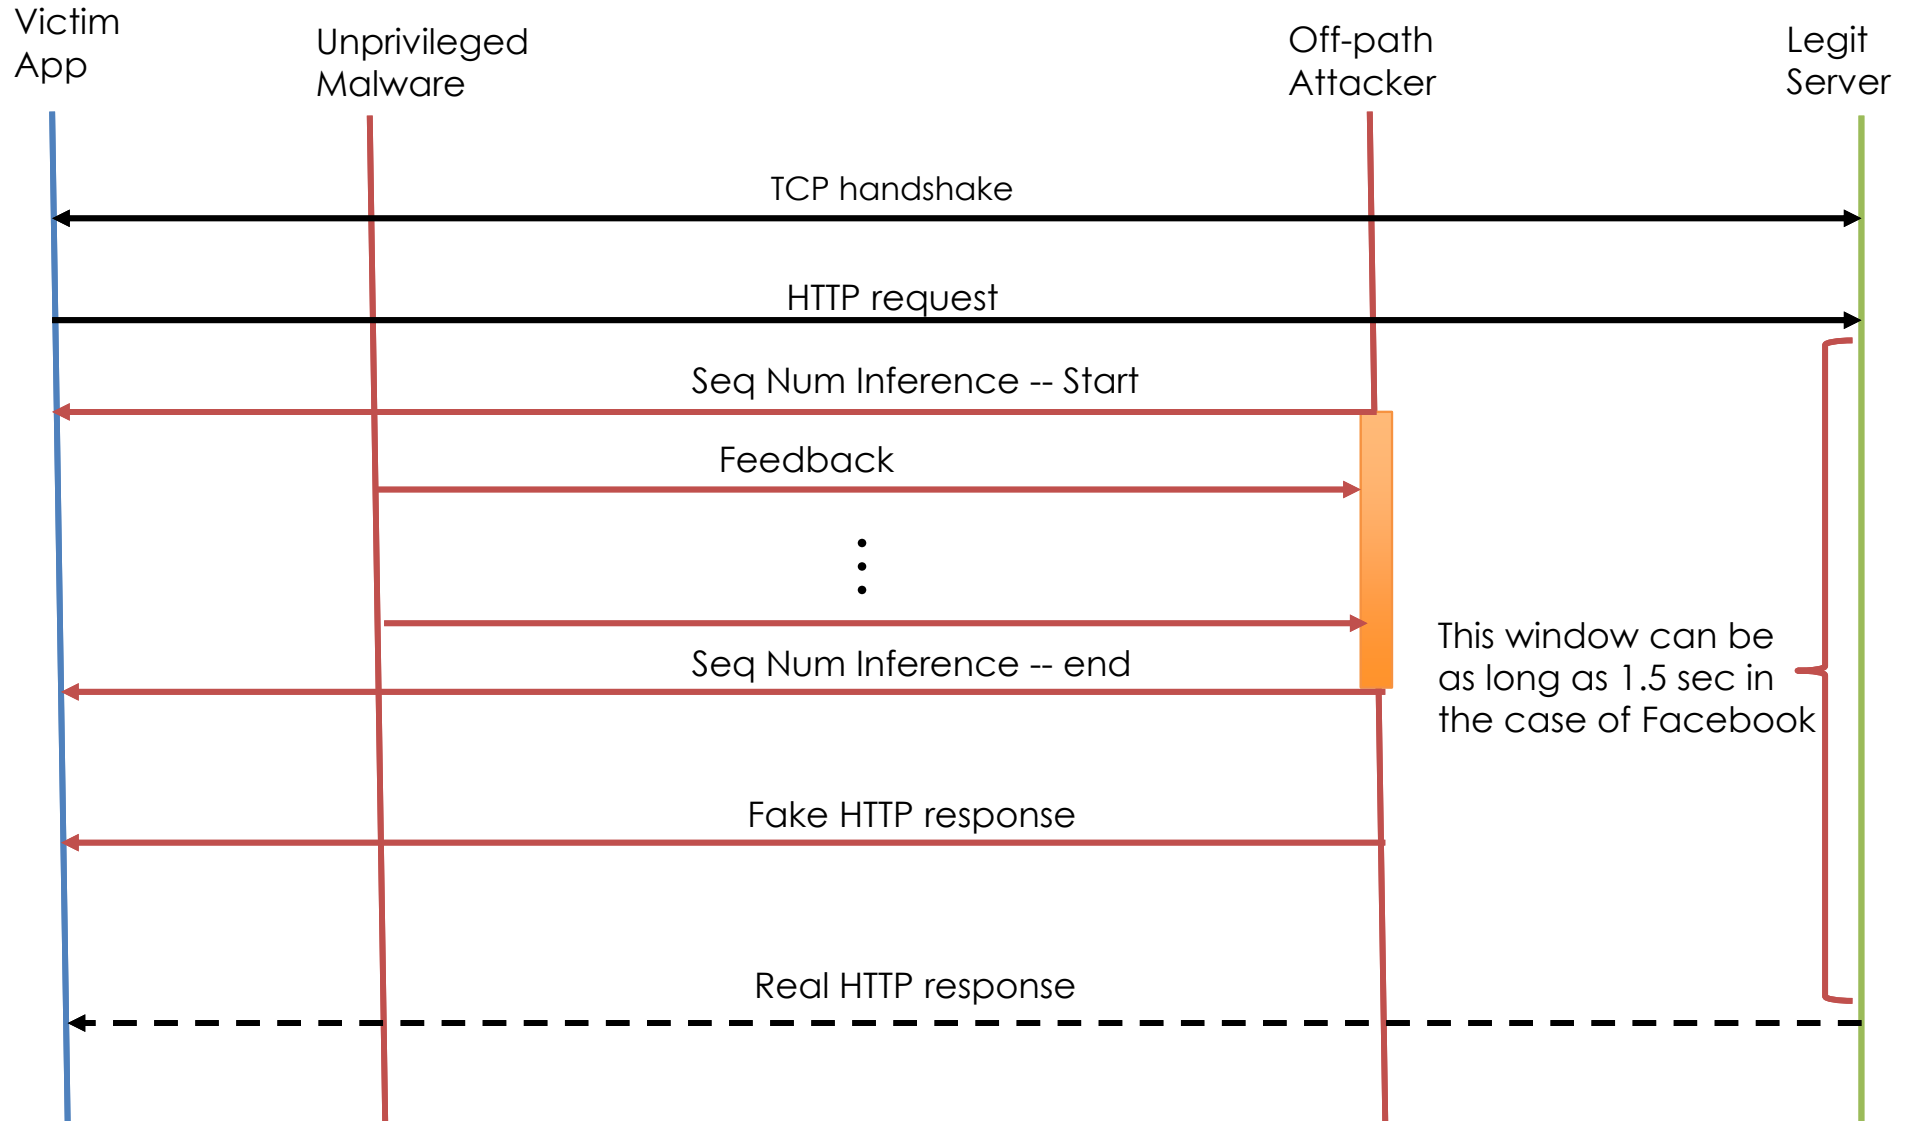
- A mobile phone app initiates a TCP connection to a server
- After TCP handshake, launch TCP seq# inference attack
 - Attacker now knows Y!
- Problem: need to compete with legitimate server
 - valid HTTP response may come first

No.	Time	Source	Destination	Protocol	Length	Info
67	3.307390	172.16.2.100	172.16.2.41	TCP	66	1133 → 80 [SYN] Seq=0 Win=8192 Len=0 MSS=1460 WS=256 SACK_PERM=1
69	3.310325	172.16.2.41	172.16.2.100	TCP	66	80 → 1133 [SYN, ACK] Seq=0 Ack=1 Win=14600 Len=0 MSS=1460 SACK_PERM=1 WS=32
71	3.311980	172.16.2.100	172.16.2.41	TCP	54	1133 → 80 [ACK] Seq=1 Ack=1 Win=65536 Len=0
74	3.385125	172.16.2.100	172.16.2.41	HTTP	466	GET / HTTP/1.1
75	3.387219	172.16.2.41	172.16.2.100	TCP	60	80 → 1133 [ACK] Seq=1 Ack=413 Win=15680 Len=0
76	3.389015	172.16.2.41	172.16.2.100	HTTP	189	HTTP/1.1 307 Temporary Redirect

Window for
ISN inference {

- Demonstration:
 - Facebook takes more than 1 sec to send the first HTTP response back to mobiles
 - With optimizations, attacker can be faster than legit Facebook's response most of times
 - Can inject malicious Facebook Javascripts

Client-side TCP Injection Attack



Passive/Active TCP hijacking

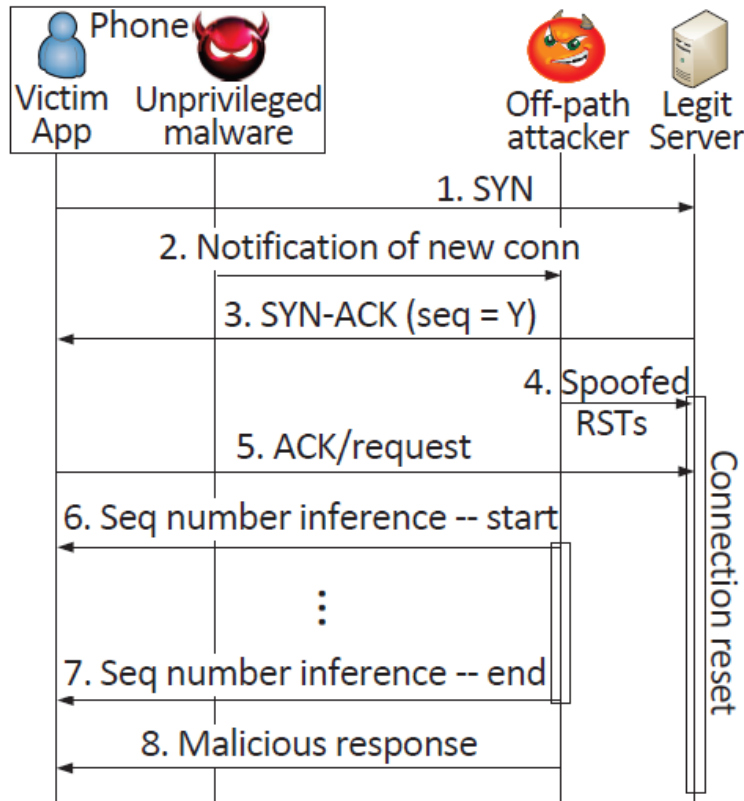


Figure 9: Passive TCP hijacking sequence

(C1). Client-side ISN has only the lower 24-bit randomized. This requirement is necessary so that the malware can roughly predict the range of the ISN of a newly created TCP connection. (*holds for Linux 3.0.2 or earlier*)

(S1). The legitimate server has a host-based stateful TCP firewall. Such a firewall is capable of dropping out-of-state TCP packets.

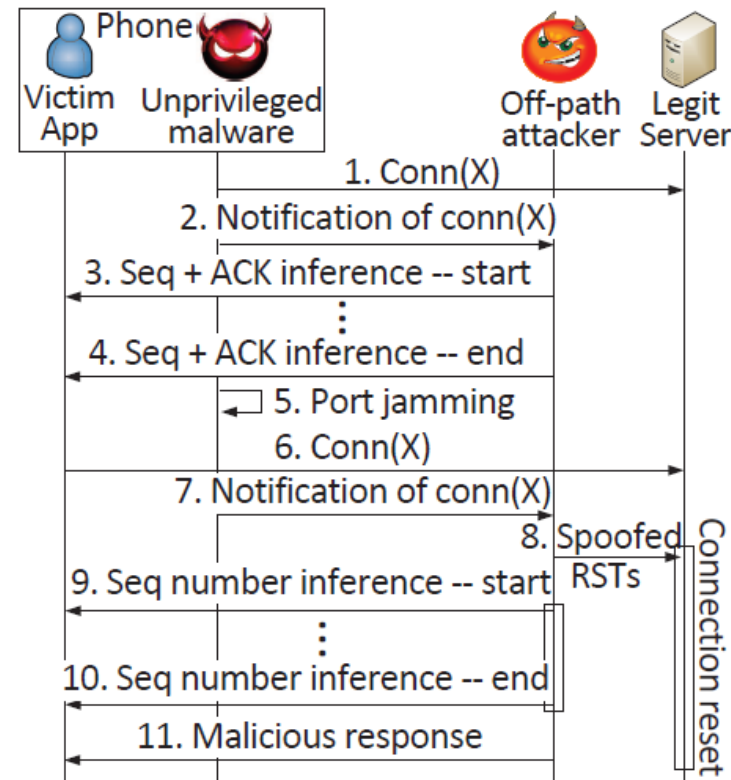


Figure 10: Active TCP hijacking sequence

(C2). Client-side ISN monotonically incrementing for the same four tuples. This client-side requirement is in fact explicitly defined in RFC 793 to prevent packets of old connections, with in-range sequence numbers, from being accepted by the current connection mistakenly.

Takeaways from TCP hijacking paper

- TCP hijacking is still possible!
- Why?
 - Our systems today (Linux, Android, Windows, Mac, etc.) have too much shared state
 - OS aggregated statistics (seemingly harmless) can leak critical internal network states
- Defenses?
 - Use SSL/TLS always
 - Removing unnecessary states
 - Isolating states
- There're more recent attacks (e.g., exploiting Wi-Fi protocol vulnerability to infer TCP sequence numbers)

Questions?