

# Announcement

- Mini-project
  - Due: 16 April 23:59 (Sun)
- Take-home Exam 2
  - Plan to publish 7 April 2023 (Fri)
  - Due: 24 April 2023 23:59 (Mon)
  - Covers Week 7-12
  - Lecture materials as well as papers.
  - Open book.
  - No discussion or collaboration is allowed.

# Blockchain Security

CS 5321 Week 12

Materials are provided by

Dr. Muoi Tran

<https://www.comp.nus.edu.sg/~muoitran/>




# Outline

- Cryptocurrency/blockchain ***overview***
- ***Security*** research of blockchain ***networks***
- ***Partitioning attacks*** against Bitcoin peer-to-peer networks
- ***Summary***

# Cryptocurrency/blockchain *overview*

***Cryptocurrency***: digital currency with strong cryptography and decentralization guarantees

**Cryptocurrency**



- Digital assets (e.g., coins)
- Allow online transactions

- Proof of coin ownership and ownership transfer
- No double-spending

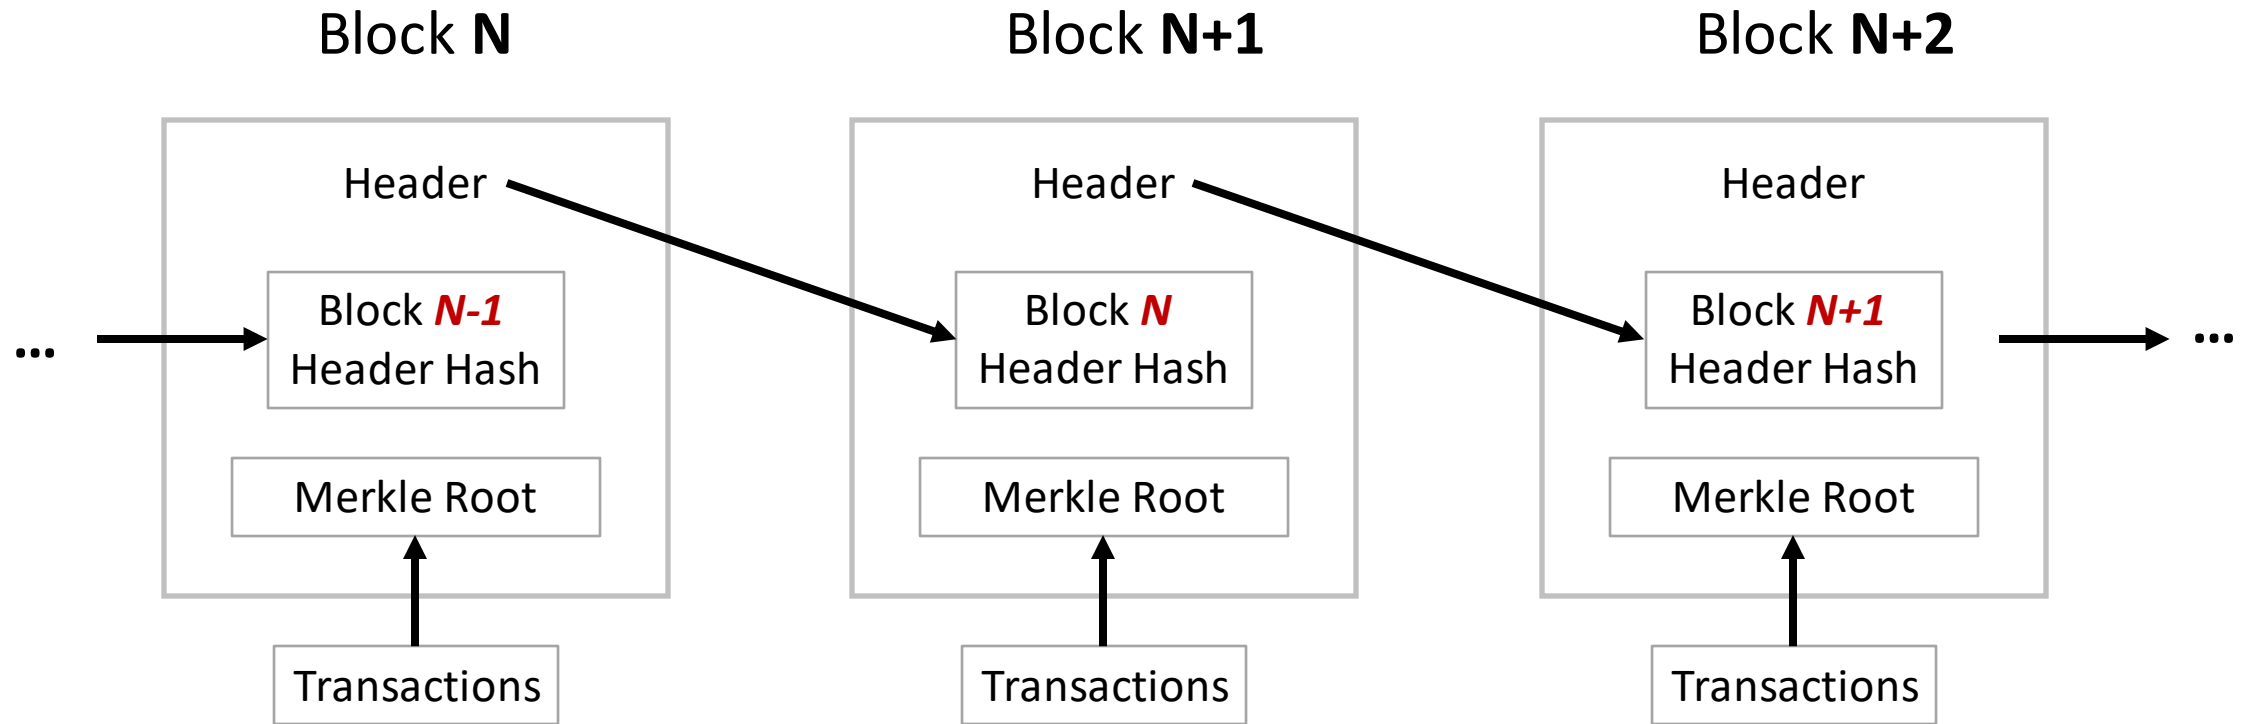
- No centralized authority
- Global transaction database is available to everyone

# Cryptocurrency relies on *blockchain* technologies for decentralization

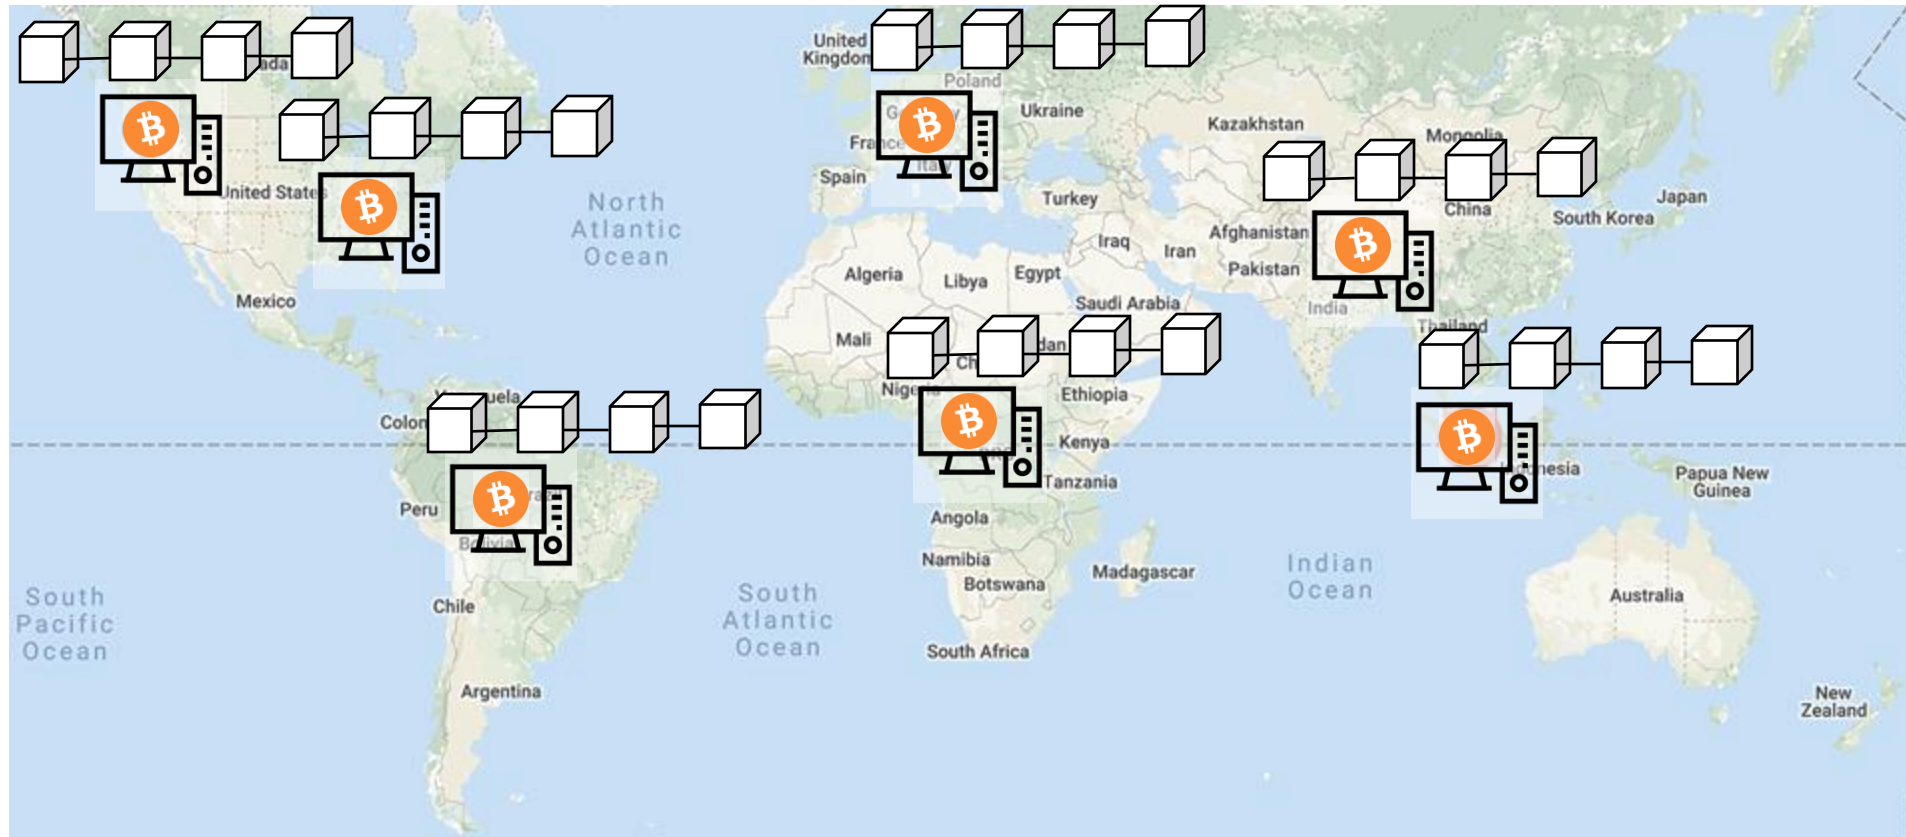
- Digital currencies with cryptography foundations are *not new*
  - ✓ Chaum82: blind signatures for e-cash
  - ✓ Chaum85: retroactive double spender identification
  - ✓ Camenisch05: compact offline e-cash
- **Bitcoin**: the *first* cryptocurrency
  - ✓ Fully decentralized digital currency system
  - ✓ Released by Satoshi Nakamoto in 2008
  - ✓ Innovation: the *blockchain*



# Blockchain is a *linked list* of *blocks* of all transactions



# The blockchain is maintained by many *distributed nodes*






# Nodes follow *pre-defined consensus* to agree on a *single* blockchain

- **All nodes** can generate transactions:
  - ✓ **Unconfirmed** transactions are not in any block
  - ✓ **Confirmed** transactions are included in a block
- **Miners** generate new blocks from unconfirmed transactions
  - ✓ Must own some **powers** (e.g., computational (**PoW**) or stake (**PoS**))
  - ✓ Incentives: new coins + transaction fees
- Nodes follow **consensus** rules:
  - ✓ Only **valid** blocks and transactions are stored
    - ✓ Valid block must contain proof of work
  - ✓ Only **one block** at any height is accepted (e.g., **longest** chain in Bitcoin)


=> a **single** blockchain will eventually remain

# Blockchain relies on underlying *peer-to-peer networks* for data propagation



- Nodes connect via a ***P2P network***
- New transactions/blocks are propagated via ***gossiping*** protocols
- Upon receiving a block /transaction, nodes ***verify*** before forwarding it
- P2P network can be ***permissionless*** or ***permissioned***

# Why blockchain *network security* matters?

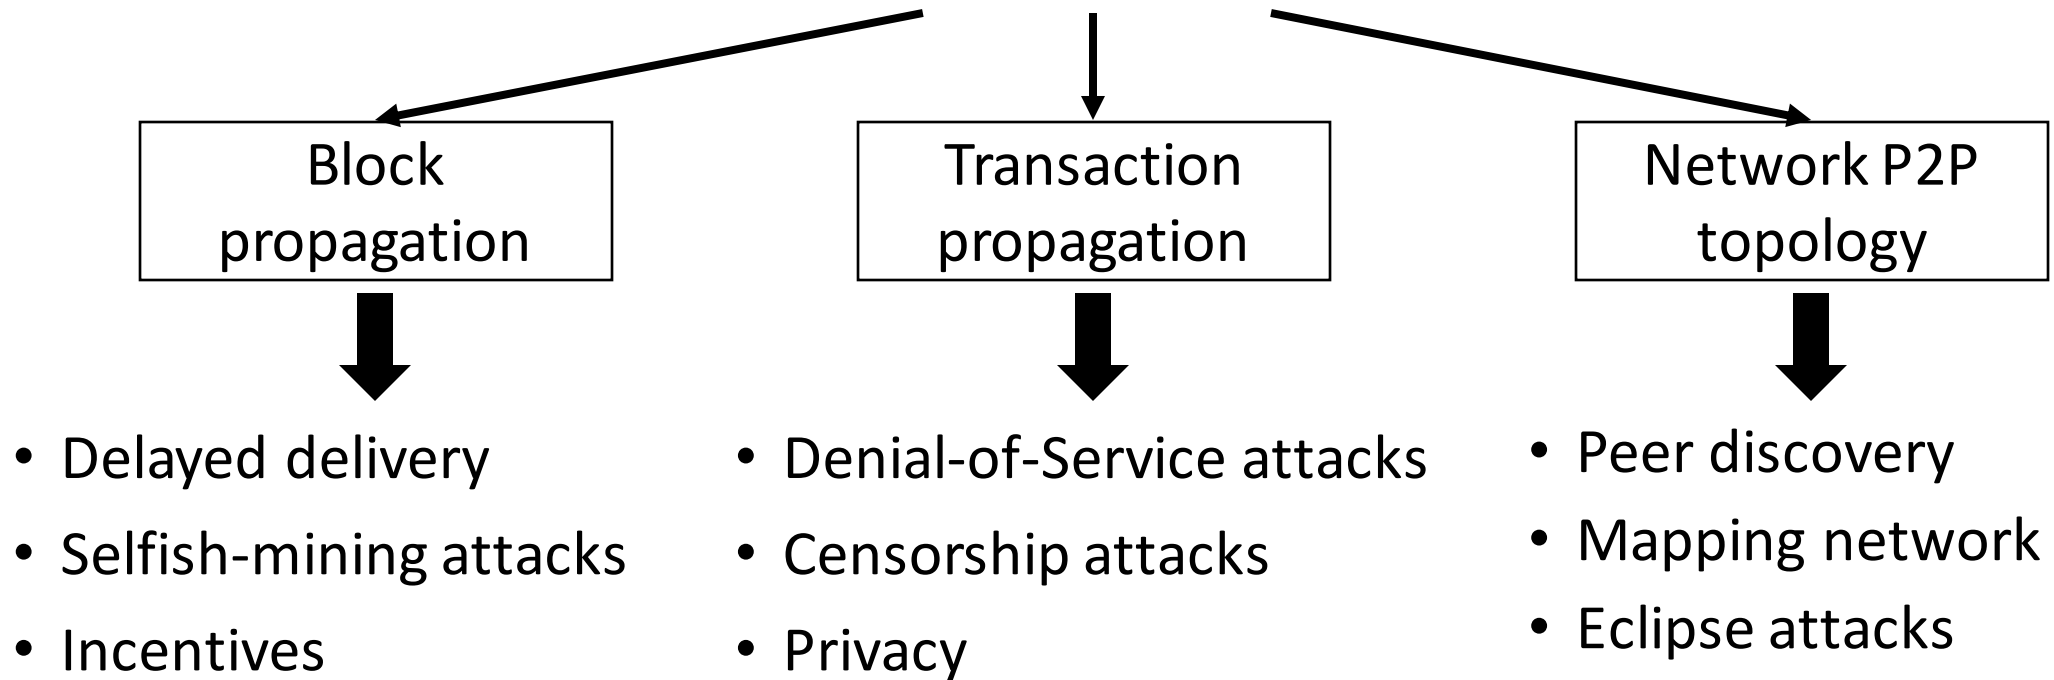


- Blockchain in **three** layers:
  - ✓ Transactions/scripts
  - ✓ Consensus/mining
  - ✓ P2P networks
- Security guarantees of a layer **depend** on a **lower** layer:
  - ✓ Example 1: a valid transaction may be rejected if malicious miners ensor it
  - ✓ Example 2: blockchain will have two versions if network is split into two

# ***Security*** research of blockchain ***networks***

# Topics in blockchain network security

## Blockchain Networks



# Block propagation: *challenges*

- ***Delayed*** blocks lead to different blockchains (i.e., forks)
  - ✓ Non-attacks: block data is large, network latency, churn, ...
  - ✓ MitM attackers purposely delay block delivery
- ***Selfish mining***: Malicious miners do not release new blocks immediately to gain advantages in mining next block
- There is ***no incentives*** for ***non-miner*** nodes to propagate blocks  
=> P2P network becomes more centralized

# Transaction propagation: *challenges*

- Transaction propagation can be abused for ***DoS attacks***
  - ✓ Attackers ***flood transactions*** to the network, which are propagated by ***all nodes***
  - ✓ Only require ***minimal transaction fees***
- Transactions can be ***censored*** or ***not propagated***
  - ✓ Miners do ***not include*** valid transactions that they don't like in new blocks
  - ✓ Miners ***do not propagate*** transactions with high fees to other miners
- The ***origin*** of the transaction can be ***deanonymized***
  - ✓ Transactions are ***grouped*** together and ***linked*** to the origin (e.g., IP address)


# P2P network topology: *challenges*

- How to ***discover*** other peers when a node ***first joins*** the network?
  - ✓ E.g., Bitcoin bootstrap nodes are maintained by developers
  - ✓ E.g., Ethereum uses Distributed Hash Table
  - ⇒ ***not fully decentralized, or no robustness guarantee***
- How to prevent ***mapping*** the full network topology?
  - ✓ Network topology shows which node is connecting to which node  
=> reveal the ***influenced nodes*** to the attackers
- ***Eclipse*** attacks split the P2P network:
  - ✓ Consensus cannot be reached
  - ✓ Users cannot operate on transactions

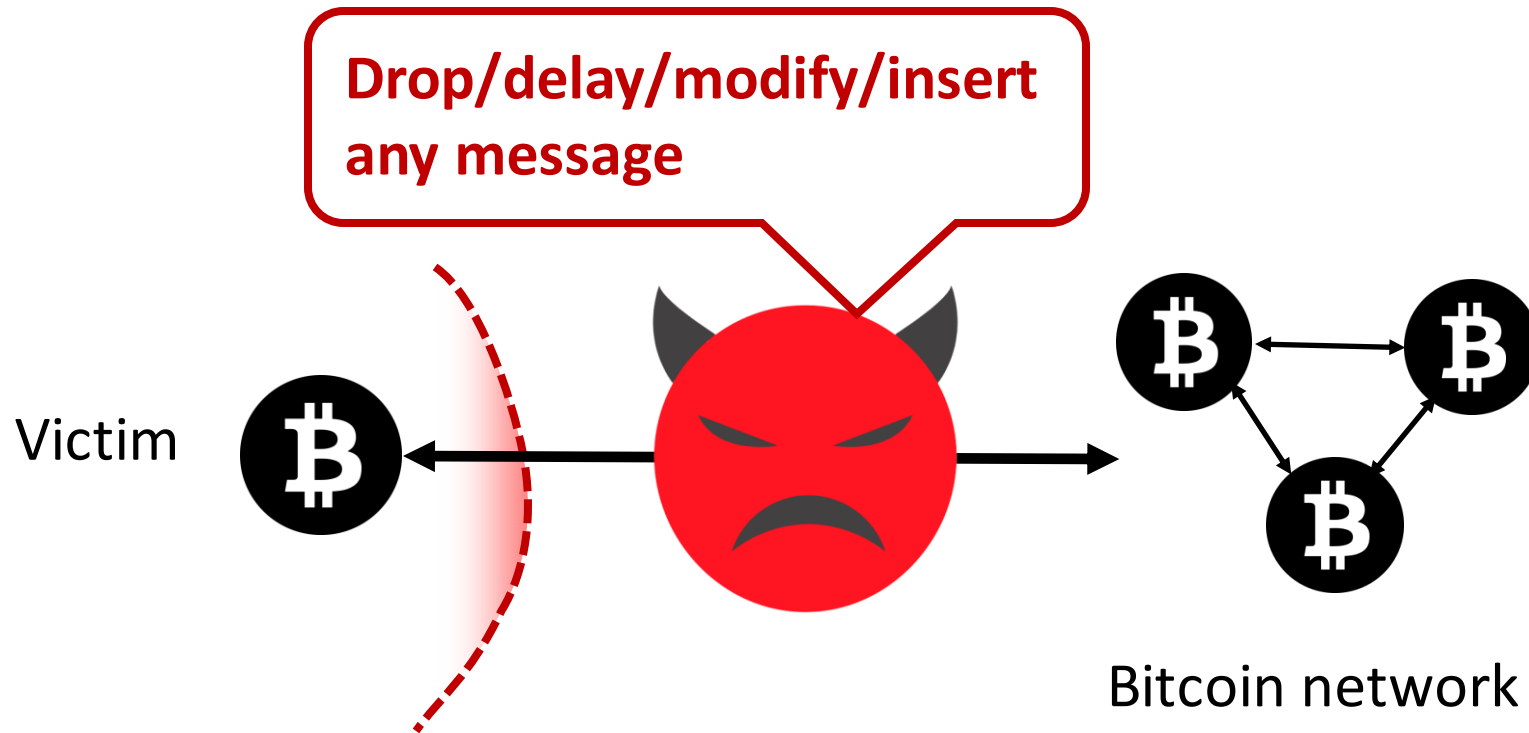


# ***Partitioning attacks*** against Bitcoin peer-to-peer networks

# Bitcoin P2P network under the hood




# Partitioning attacks against Bitcoin network



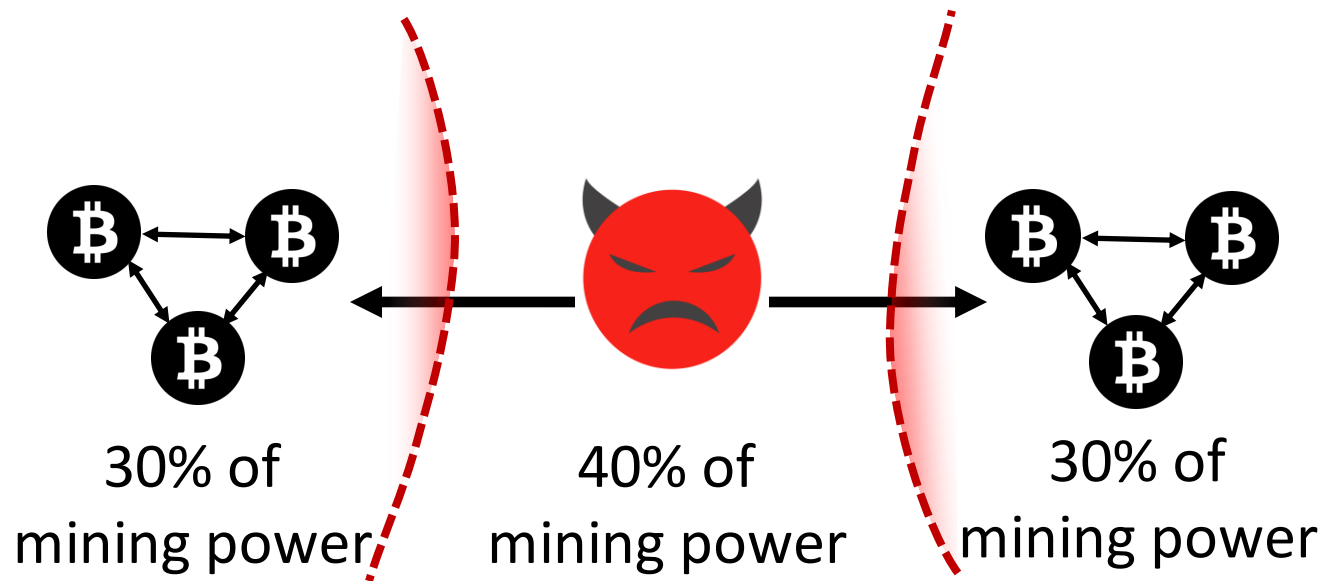
***Partitioning attacks:*** isolate one or more nodes from the rest of network

# Motivation of partitioning attacks




- Double-spending attack
  - ✓ Attacker pays merchant with  $(A \rightarrow B: 10)$  transaction
  - ✓ Attacker pays herself with  $(A \rightarrow C: 10)$  transaction
  - ✓ Mines N blocks to confirm the transaction
  - ✓ Rest of network accepts  $(A \rightarrow C: 10)$  transaction and discards  $(A \rightarrow B: 10)$

# Motivation of partitioning attacks (cont.)




- 51% attack
  - ✓ Attacker outcompetes each partitioned miner, even with only 40% mining power.
  - ✓ Attacker's blockchain will become the longest chain
- Many other attacks:
  - ✓ Selfish-mining
  - ✓ Censoring transactions
  - ✓ Attack layer-2 protocols
  - ✓ ...

# How to partition a Bitcoin node?



- **Strategy 1**: Influence victim to only connect to adversary-controlled peers (e.g., Eclipse, Erebus attack)




- **Strategy 2**: Adversary places herself in all legitimate peering connections (e.g., Bitcoin hijacking attack)

# Eclipse attack

- Paper: “*Eclipse Attacks on Bitcoin’s Peer-to-Peer Network*” by Heilman et al. [USENIX Security 2015]
- Threat model
  - ✓ Attacker’s capability: Control a botnet of ~3,000 IP addresses
  - ✓ Attacker’s goal: Influence all of victim’s connections to be made to adversary-controlled bots
- Assumptions:
  - ✓ Victim run Bitcoin version 0.9.3 or earlier
  - ✓ Victim has a public IP address (i.e., not behind NAT, Tor, VPN, ...)

# How Bitcoin nodes form the P2P network




- Bitcoin node establishes and maintains **8 outgoing connections**
  - ✓ Outgoing peers are selected from internal peer database (i.e., addrman)
- A node also accepts at most **117 incoming connections**
  - ✓ Only full nodes with public IP accept incoming connections
  - ✓ Today network includes ~8K full nodes



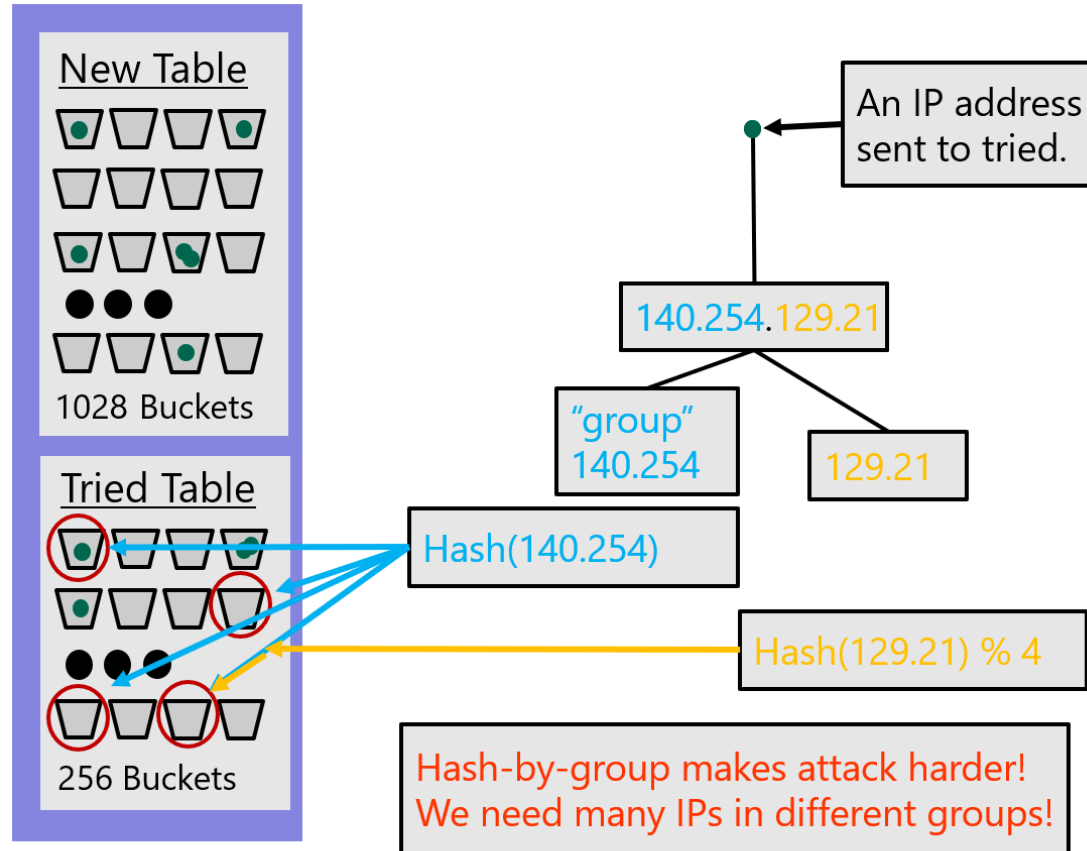
# Address propagation

```
vSeeds.emplace_back("seed.bitcoin.sipa.be"); // Pieter Wuille  
vSeeds.emplace_back("dnsseed.bluematt.me"); // Matt Corallo  
vSeeds.emplace_back("dnsseed.bitcoin.dashjr.org"); // Luke Dashjr  
vSeeds.emplace_back("seed.bitcoinstats.com"); // Christian Decker  
vSeeds.emplace_back("seed.bitcoin.jonasschnelli.ch"); // Jonas Schnelli  
vSeeds.emplace_back("seed.btc.petertodd.org"); // Peter Todd  
vSeeds.emplace_back("seed.bitcoin.sprovoost.nl"); // Sjors Provoost
```



- Boot-strapping: New node gets a list of peers via DNS seeds or hard-coded IP addresses.
- Solicited **ADDR** messages: responses to a **GETADDR** request
- Unsolicited **ADDR** messages: relaying other's advertised IPs
- Each **ADDR** message contains up to 1,000 IP address.

# Storing address in addrman



- `addrman` stores each IP with a timestamp:
  - ✓ New table: IPs from `ADDR` messages
  - ✓ Tried table: IPs of the peered nodes
- Adding an IP to Tried table:
  - ✓ The `/16 group` determines 4 buckets
  - ✓ The `rest` of the IP determines 1 out of 4 buckets
  - ✓ If bucket is full, pick 4 random IPs and delete the oldest one (**bitcoin eviction policy**)

# Attacking steps

- Fill the Tried table
  - ✓ Simply connect to the victim node from botnet
- Fill the New table
  - ✓ Flood unsolicited ADDR messages (contain 1,000 IPs each)
- Restart the victim node
- Make 117 incoming connections to the victim
- All selected outgoing peers are adversary IPs
  - ✓ Outgoing peer is chosen from either New and Tried table
  - ✓ Selected IP is biased toward “fresher” timestamps

# Vulnerabilities exploited

- Vulnerability 1 (**Selection Bias**):
  - ✓ Attacker ensures its IPs are fresher, so they are more likely to be selected  
=> Keep filling the New table to update the timestamp
- Vulnerability 2 (**Try-Try-Again**):
  - ✓ If an attacker IP replaces another attacker IP, she can resend the evicted IP and eventually replace an honest IP
- Vulnerability 3 (**Eviction Bias**):
  - ✓ Attacker ensures its IPs have recent timestamps to avoid being evicted
- IP diversity requirements:
  - ✓ Only 3,000 botnet IPs have sufficient /16 prefix diversity


# Countermeasures to Eclipse attack

- **Random eviction:** No preference to newer IPs
- **Test before evict:** If existing IP is still reachable, do not evict it
- **Feeler connection:** Periodically test an IP in New table and move it to Tried if it is reachable
- **Larger table size:** Both tables are increased by 4 times
- **Remove direct IP inserting to Tried table:** Only IPs, that the node makes outgoing connection to, are stored.
- All above countermeasures are adopted in Bitcoin Core  
=> **Eclipse attack doesn't work against the latest version**

# Bitcoin hijacking


- Paper “*Hijacking Bitcoin: Routing Attacks on Cryptocurrencies*” by Apostolaki et al. [IEEE S&P 2017]
- Threat model:
  - ✓ Attacker’s capability: Adversary is a network attacker (i.e., an ISP)
  - ✓ Attacker’s goals: intercepts all Bitcoin connections of the victim
- Assumptions:
  - ✓ Attacker can launch BGP hijacking attack (not all ASes can do!)
  - ✓ Victim’s IP belongs to a “hijack-able” prefix (i.e., prefix shorter than /24)

# Attack steps



- Nodes of the left and the right side of the network communicate via [Bitcoin connections](#).
- The attacker wishes to split the network into two disjoint components
- The attacker intercepts the traffic destined to the right nodes by performing BGP hijacks

# Attack steps (cont.)



- After the hijack, all traffic sent **from the left to the right side** is forwarded through the attacker.
- The attacker can drop all Bitcoin traffic => partitions are created
- Some connections cannot be intercepted
  - ✓ Nodes within same AS
  - ✓ Miners in the same mining pool



# Attack effectiveness and practicality

- Victim: **93%** Bitcoin nodes' IPs belong to prefixes shorter than /24
- Attacker: Hijack **<100 prefixes** can isolate up to 47% mining power
- ASes (e.g., large ISPs) can launch this Bitcoin hijacking attack
  - Question: Do they really launch this attack in practice?
  - **Yes, but attack is quickly detected!**

[HOME](#)[BLOG](#)[ABOUT US](#)[PRODUCTS AND SERVICES](#)[CLIENT PORTAL](#)

## The Canadian Bitcoin Hijack


Posted by Andree Toonk - August 12, 2014 - [Hijack](#) - [No Comments](#)

A few days ago researchers at [Dell SecureWorks](#) published the details of an attacker repeatedly hijacking BGP prefixes for numerous large providers such as *Amazon, OVH, Digital Ocean, LeaseWeb, Alibaba* and more. The goal of the operation was to intercept data between Bitcoin miners and Bitcoin mining pools. They estimated that \$83,000 was made with this attack in just four months. The [original post](#) has many of details which we won't repeat here, instead will take a closer look at the BGP details of this specific attack. **Attack details** Our friends at Dell SecureWorks decided not to name the network from which the hijacks originated. As a result we won't name the exact Autonomous System either, instead we will suffice by saying that the originator of this hijack is a network operating in Eastern Canada. **Initial experiment** BGPmon

# Erebus attack

- Paper “*A Stealthier Partitioning Attack against Bitcoin Peer-to-Peer Network*” by Tran et al. [IEEE S&P 2020]
- Threat model:
  - ✓ Attacker’s capability: Adversary is a network attacker (i.e., an ISP)
  - ✓ Attacker’s goals: Influence all of victim’s connections to be made to adversary-controlled bots
- Assumptions:
  - ✓ Victim has a public IP address (i.e., not behind NAT, Tor, ...)
  - ✓ Victim runs the latest Bitcoin ver. with Eclipse countermeasures


# Idea: Using “*shadow*” IPs as attack resource



**Observation**: Traffic from victim node to **any** IP addresses at AS **C** and **D** would **traverse attacker AS**

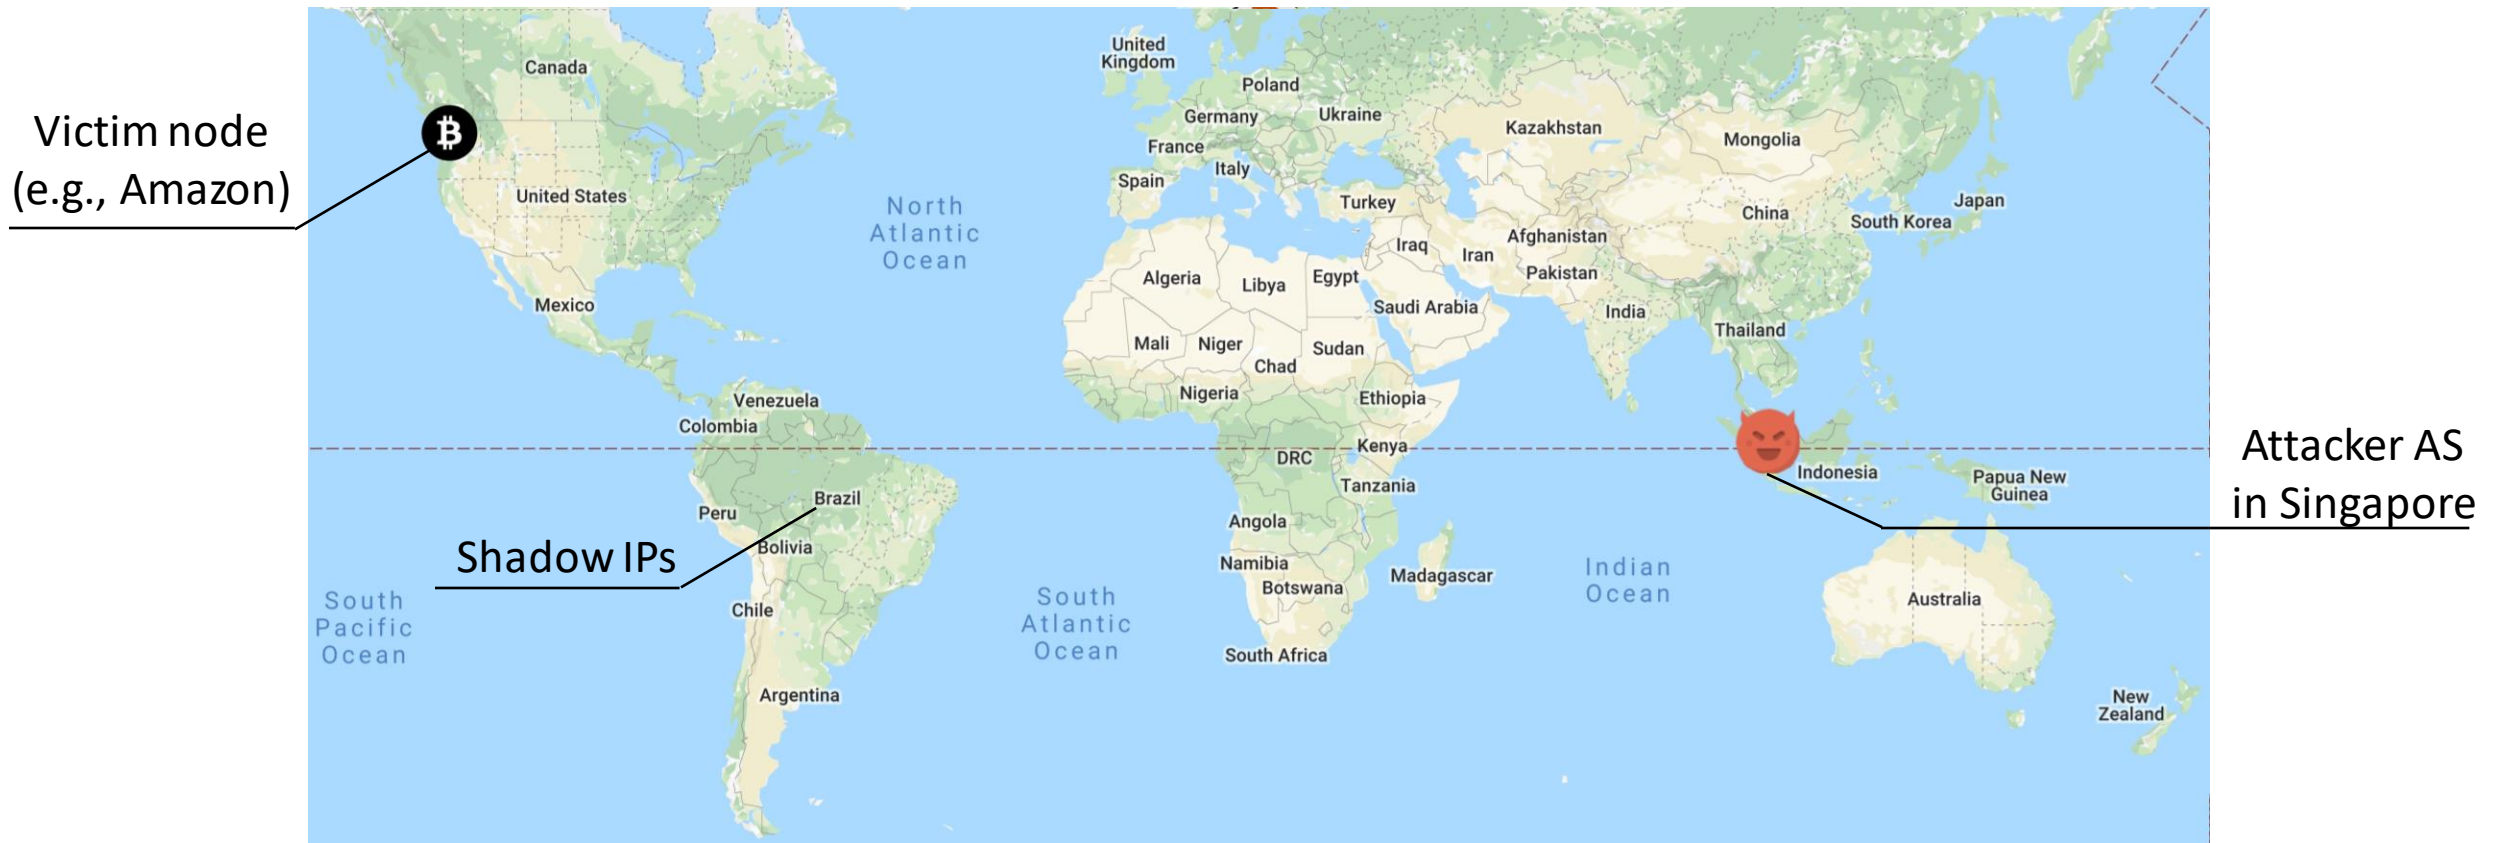
- **Shadow IPs** are valid IPs whose victim-to-IP routes include the attacker AS
- Attacker AS can **spoof** connections with victim node **on behalf** of shadow IPs  
=> Attacker **virtually** controls shadow IPs

# How many shadow IP addresses are available?




When attackers are **top-100 ASes** and victim are hosted at 100 random ASes

# Shadow IPs are *geographically well-distributed*




Shadow IPs are usually *well-distributed* across the world  
=> look normal to cautious Bitcoin nodes


# Dominate *new* table: advertise shadow IPs with *high rate*




- Attacker fills the **new** table by sending shadow IPs with *high rate* to replace legitimate IPs
  - ✓ e.g., an IP is deleted if it is not heard from peers in **30 days**




# Dominate *tried* table: *patiently* wait for IP *trickle-downs*



- Attacker fills the **tried** table by *patiently* waiting for *trickle-downs*
  - ✓ e.g., one IP is moved from **new** to **tried** every *two minutes* via feeler connections



# The Erebus attacker patiently isolates victim Bitcoin nodes in **5 - 6 weeks**

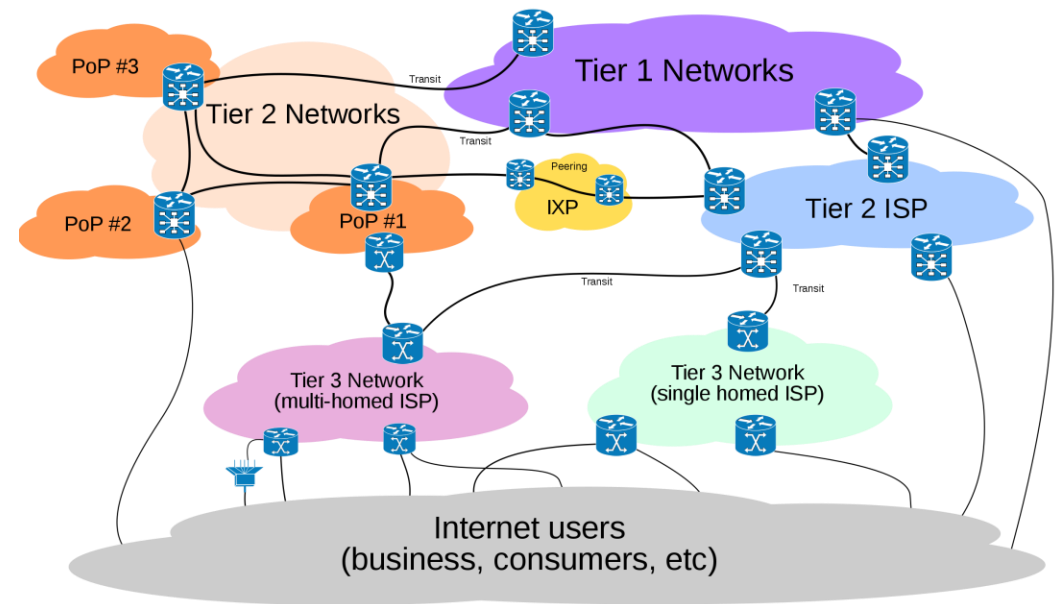


- Attacker keeps sending **low rate** attack traffic (**520 bit/s** or **2 IP/s**) until it controls **all eight** outgoing connections of the victim



# Who can launch the Erebus attack?

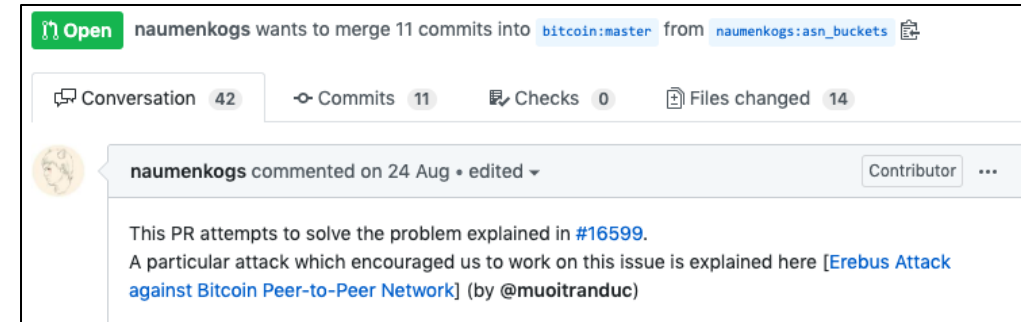
- To attack a targeted node, Erebus attacker needs:
  - ✓ *millions* shadow IP addresses
  - ✓ *several weeks* of attack execution
- **All Tier-1** networks
  - ✓ AT&T, CenturyLink, NTT, ...
  - ✓ Can target *any* Bitcoin node!
- Many **large Tier-2** networks
  - ✓ Singtel, China Telecom, ...
  - ✓ Can target most Bitcoin nodes!
- **Nation-state** adversaries
  - ✓ Some countries are believed to have direct control over their ISPs



By User:Ludovic.ferre - Internet Connectivity Distribution&Core.svg,  
CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=10030716>

# The Erebus attack has a *high impact*

- The Erebus attack works against *latest* Bitcoin core v0.18.1
  - ✓ **Any public Bitcoin node** can be targeted!
  - ✓ Complete mitigation is **difficult**: **No** bugs was exploited but only Internet routing
  - ✓ Deterrence are **being implemented** by Bitcoin
- **34 out of top-100** cryptocurrencies are also potentially vulnerable to the Erebus attack
- Project webpage:  
<https://erebus-attack.comp.nus.edu.sg/>



# Comparing three attacks

	Bitcoin hijacking attack	Erebus attack	Eclipse attack
<b>Attack capabilities</b>			
Network attacker?	Yes	Yes	No
Distributed computing resource required?	No	No	Yes
Route manipulation required?	Yes	No	No
<b>Attack outcomes</b>			
Effective against Bitcoin v0.18.0?	Yes	Yes	No
Attacks are visible and attributable?	Yes	No	No

# Summary

- ***Blockchain*** is the biggest innovation in cryptocurrencies
- Blockchains rely on the ***security*** guarantees of the underlying ***networks***
- ***Partitioning attacks*** are great examples of how a network adversary can attack consensus and/or transactions layers
- ***Network security*** of blockchain is a ***promising*** research direction