

CS5322 Database Security

Query Authentication: Motivation




- A data owner passes a database D to a service provider
- The service provider answers queries from users
- This setting is referred to as *database outsourcing*

Query Authentication: Motivation




- This was the same scenario that we considered in our discussions of *encrypted databases*
- But now, let's assume that we do NOT aim to ensure data confidentiality against the service provider
 - Reason: encrypted databases incurs a lot of overheads
- Instead, we want to ensure that the service provider answers queries *correctly*

Query Authentication: Motivation




- "What do you mean by 'answering queries correctly'?"
- Possible ways to answer queries incorrectly:
 - ❑ Return only a subset of the tuples in the query result (e.g., by stopping the query early)
 - ❑ Return some fake tuples
 - ❑ Return the query result previously obtained from an outdated version of the database

Query Authentication: Motivation



- "Why would the service provider do that? "
- Possible reasons:
 - The service provider is overloaded with a lot of queries on a lot of databases, and does not have enough resource to process all queries in time
 - The service provider wants to compromise the query result for malicious purposes
 - ...


Query Authentication: Motivation



■ Solution:


- When the service provider answers a query, it returns both the query result and a **proof** of the result's **integrity**
- i.e., the proof should let the user verify whether the result is correct and whether it is from the most updated database

Query Authentication: Motivation



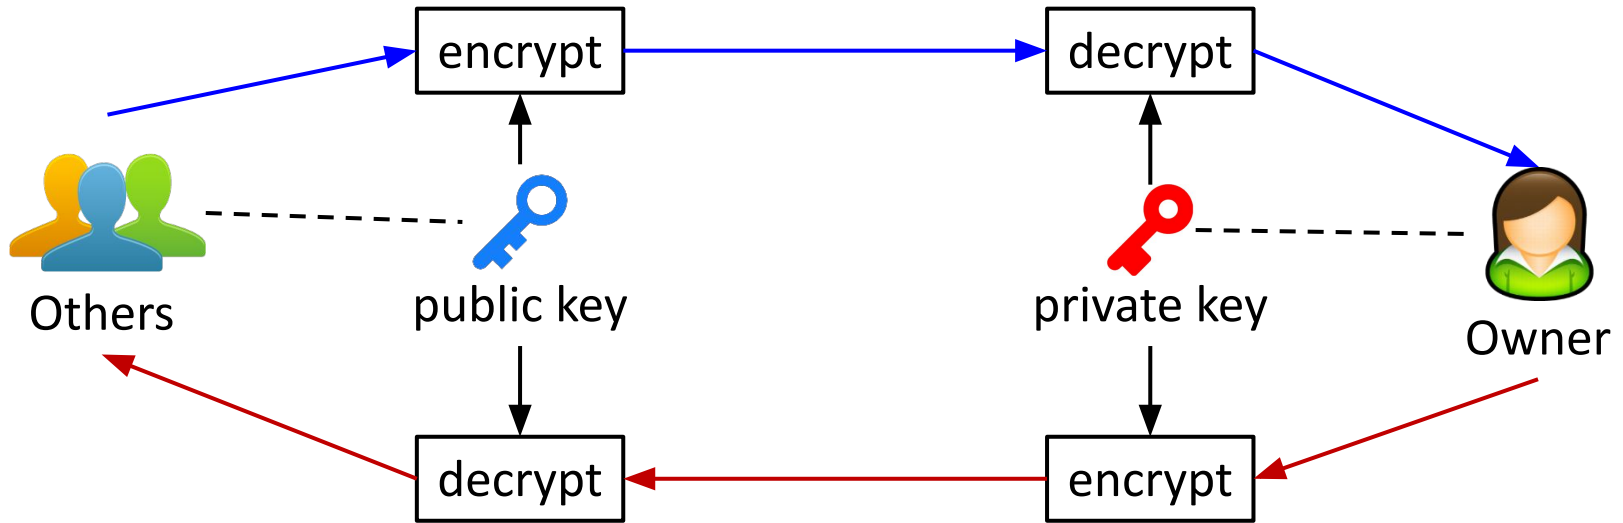
- "I don't think service providers would really do that in practice."
- Amazon has done something based on a similar idea
 - <https://aws.amazon.com/qldb/>
- Microsoft also proposed something along a similar line
 - <https://dl.acm.org/doi/10.1145/3035918.3064030>

Query Authentication: Motivation




- "OK. But how?"
- Use *public-key cryptography*

Public-Key Cryptography




- The data owner has two keys
 - A private key sk , which she keeps to herself
 - A public key pk , which she distributes to the public
- Messages encrypted by pk can only be decrypted by sk
 - So anyone can use pk to send secret messages to the owner
- Messages that can be decrypted by pk must have been encrypted by sk
 - So the owner can send *authenticated* messages to others

Naïve Solution




- Let the owner *sign* each tuple in the database, i.e.,
 - She encrypts each tuple with her private key sk
 - She sends both the encrypted version and the original to the service provider
- Whenever the service provider answers a query by returning some tuples, he also returns the encrypted versions

Naïve Solution



- Whenever the service provider answers a query by returning some tuples, he also returns the encrypted versions
- For each tuple t and its encrypted version t^* , the user can decrypt t^* using the owner's public key, and see if the decrypted tuple is the same as t
 - This prevents the service provider from **modifying** or faking any tuple

Naïve Solution




- Problems:
 - It does not prevent the service provider from dropping any tuples
 - Each tuple needs to be stored twice: encrypted and unencrypted
- We need something better...

Towards An Improved Solution


- Let's focus on a specific type of queries:
equality query on one attribute
 - i.e., `SELECT * FROM T`
`WHERE T.A = X`
 - e.g., `SELECT * FROM Employees`
`WHERE Age = 30`
- How can we outsource this type of queries
and prevent the service provider from faking
or dropping results?

Towards An Improved Solution




- Idea: We let the data owner
 - ❑ Divide the tuples into groups according to their Age values
 - ❑ Encrypt each group separately using her private key sk
 - ❑ Send the encrypted groups to the service provider

Towards An Improved Solution




- Suppose that the user issues the following query:
 - `SELECT * FROM Employee WHERE Age = 21`
- The service provider just returns the encrypted group of tuples for Age = 21
- The user then decrypts the tuple group using the data owner's public key pk
 - If the tuples have Age = 21, then it is guaranteed that there is no fake or missing tuples

Towards An Improved Solution




- Question:
 - Suppose that there is no employee with Age = 21
 - What does the data owner do?
- Option 1:
 - Just omit the Age = 21 group
 - Only create encrypted groups for those age values that have at least one tuple
- Does this work?

Towards An Improved Solution



- Option 1:
 - Only create encrypted groups for those age values that have at least one tuple
- Does this work?
- No
 - If we do this, then the service provider can drop any tuple group from the query result, by pretending that there is no tuple with that age value


Towards An Improved Solution



■ Option 2:


- If Age = 21 does not have any tuple, then the data owner encrypts a message “Age = 21 is empty” using her private key
- Do this for every empty Age group
- If a user asks for an Age group that is empty, just return the corresponding encrypted message

Towards An Improved Solution



- Option 2: Create an encrypted message for each empty group
- Problem: There could be too many empty groups
 - Suppose that the attribute to be queried is Salary
 - Creating an encrypted group or message for each possible Salary value is rather expensive in terms of time and space


Towards An Improved Solution



■ Option 3:


- For each non-empty group, we not only encrypt the tuples, but also create an encrypted message to indicate which is the next non-empty group
- E.g., in the above example, we not only encrypt the tuple group with Age = 20, but also create an encrypted message: “After Age = 20, the next non-empty group is Age = 22”
- The total number of encrypt messages equals the total number of non-empty groups

Towards An Improved Solution



- If the user queries a non-empty group, then returns the corresponding encrypted group
- If the user queries an empty group, return the encrypted message that can prove the empty results
 - E.g., if the user enquires “Age = 21”, then return the encrypted message “After Age = 20, the next non-empty group is Age = 22”


Towards An Improved Solution



■ Any problem?

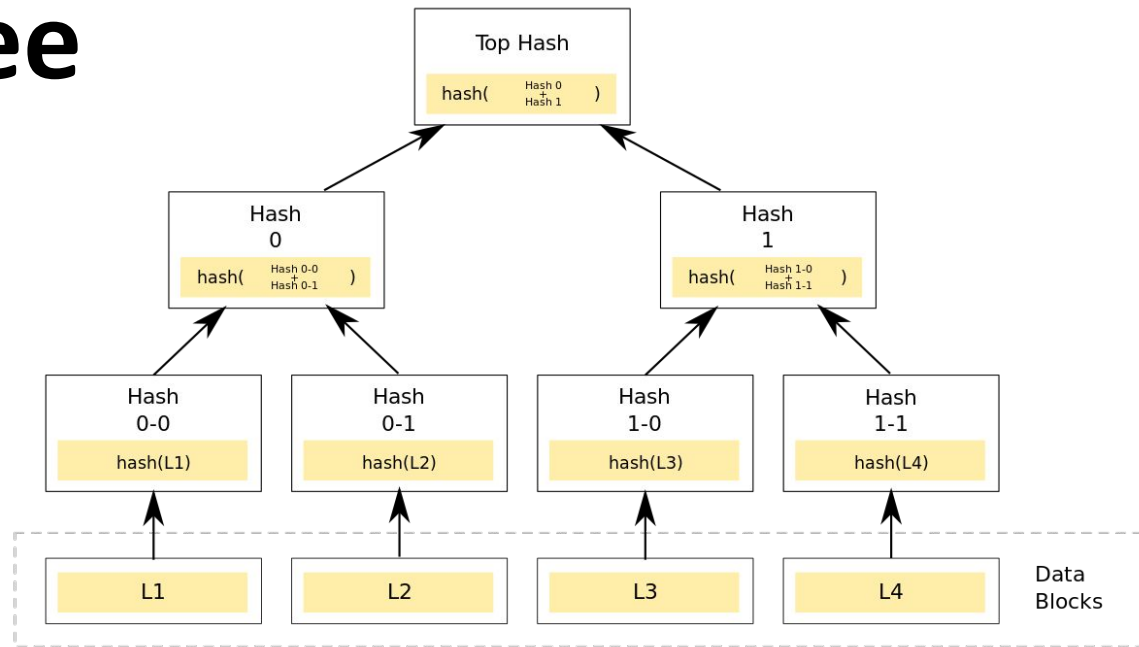
- ❑ There could be too many non-empty groups to encrypt
 - Think about Salary
- ❑ Only works well for equality query
 - Does not work for range queries, e.g., “Salary > 100k and Salary < 200k”

Towards An Improved Solution



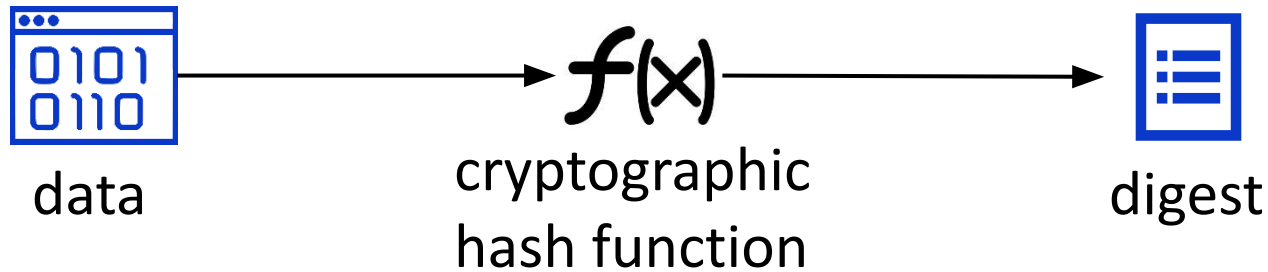
- We will introduce a solution that only require the data owner to create one encrypted “message”
 - No need to encrypt any tuple at all
 - And it can support range queries
- It is called the *Merkle tree*

Merkle Tree




- Invented by Ralph Merkle in 1979
- A tree structure that allows efficient verification of its content
- Key ingredient: a cryptographic hash function

Cryptographic Hash Function



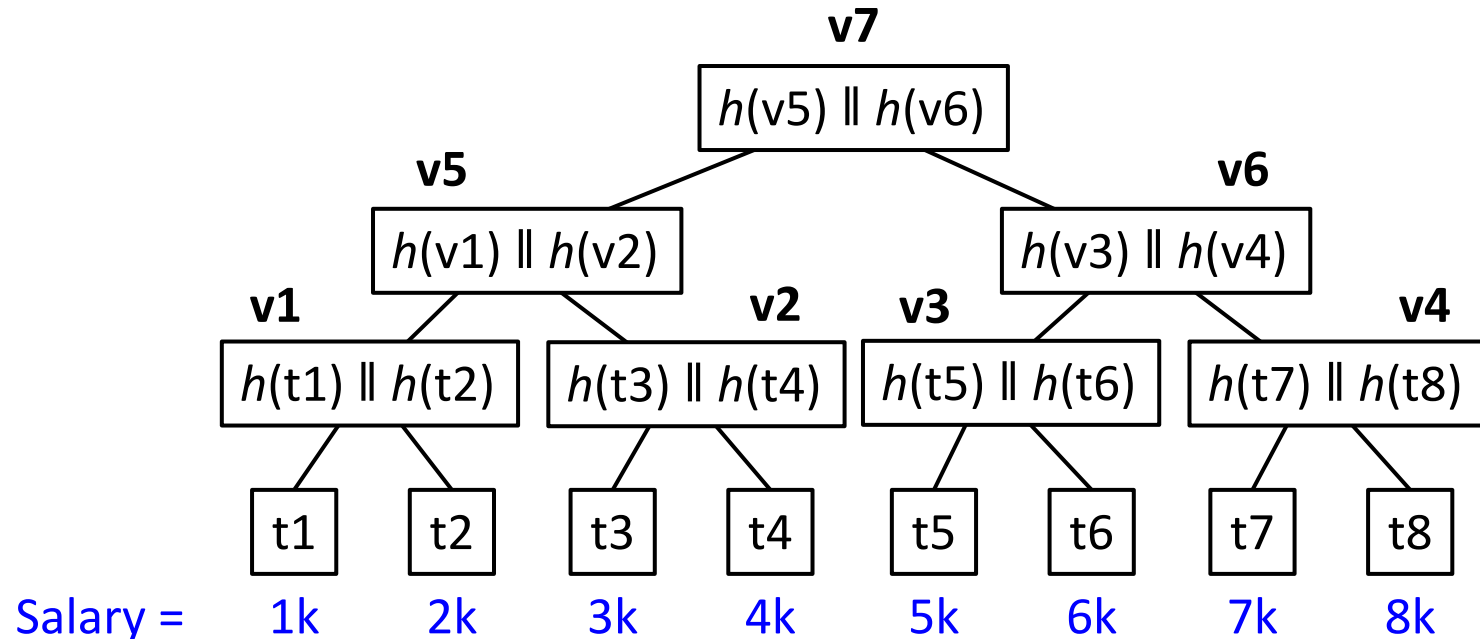
- A function that maps data of arbitrary size to a bit string of a fixed size, with the following properties:
 - ❑ It is deterministic (so the same message always results in the same hash)
 - ❑ It is efficient to compute the hash value for any given message
 - ❑ It is infeasible to generate a message from its hash value except by trying all possible messages
 - ❑ A small change to a message should change the hash value so extensively that the new hash value appears uncorrelated with the old hash value
 - ❑ It is infeasible to find two different messages with the same hash value

Merkle Tree: Example




- Suppose that we are to build a Merkle tree to support range queries on Salary
- First, sort all tuples by Salary
- Second, build a binary tree on the sorted sequence

Merkle Tree: Example




- Third, materialize the non-leaf nodes in a bottom up manner, using a cryptographic hash function h
 - For each non-leaf v , its content equals $h(v_{\text{left}}) \parallel h(v_{\text{right}})$, where v_{left} and v_{right} are v 's left and right children, respectively, and \parallel denotes concatenation

Merkle Tree: Example



- Finally, let the data owner encrypts $h(\text{root})$ using her private key sk
 - In the above example, the root is $v7$
- Then, the data owner sends the encrypted digest to the service provider


Merkle Tree: Example



■ Intuition:

- The encrypted $h(v7)$ ensures that the service provider cannot make any change to the sorted sequence $t1, t2, \dots, t8$


Merkle Tree: Example



■ Why?


- Since $h(v7)$ is signed by the data owner, the service provider will get caught if he changes $v7$
- Since $v7$ cannot be changed, the service provider will get caught if he changes $v5$ or $v6$
- Since $v5$ and $v6$ cannot be changed, the service provider will get caught if he changes $v1$, $v2$, $v3$, or $v4$, and so on...

Merkle Tree: Example



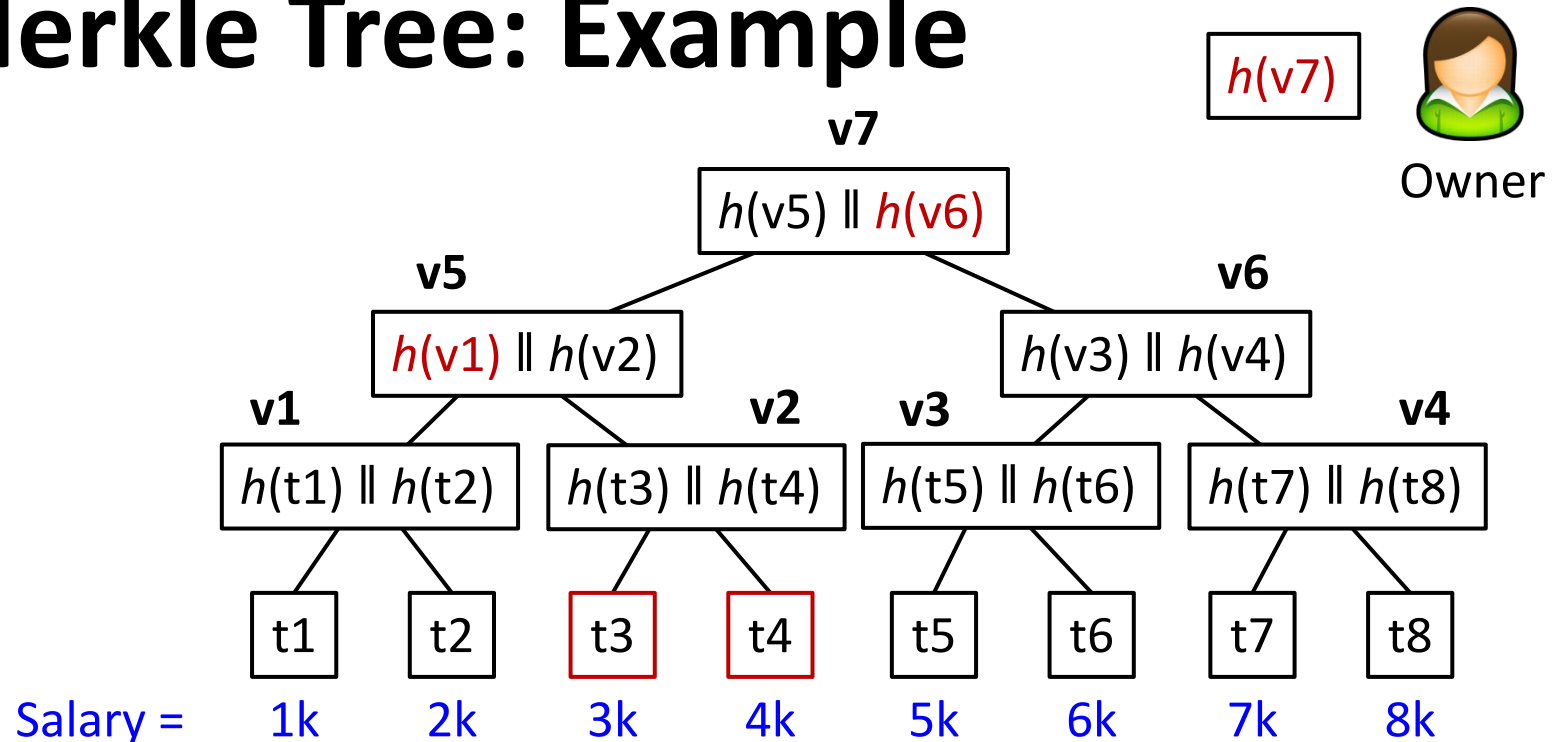
- In other words, the service provider can answer any query as follows
 - Return the sorted sequence $t1, t2, t3, \dots, t8$, along with the signed $h(v7)$
 - Ask the user to verify the correctness of the sorted sequence, and then answer the query herself using the sorted sequence
- Problem: this approach returns too many irrelevant tuples
- Solution: return some hash values instead of tuples

Merkle Tree: Example





- Suppose that a user searches for “Salary = 3.5k”
- The service provider would return the following
 - $t3$, $t4$, $h(v1)$, $h(v6)$, and the encrypted $h(v7)$

Merkle Tree: Example





- Suppose that a user searches for “Salary = 3.5k”
- The service provider would return the following
 - $t3, t4, h(v1), h(v6)$, and the encrypted $h(v7)$
- Why? We will explain using a simpler tree

Merkle Tree: Example




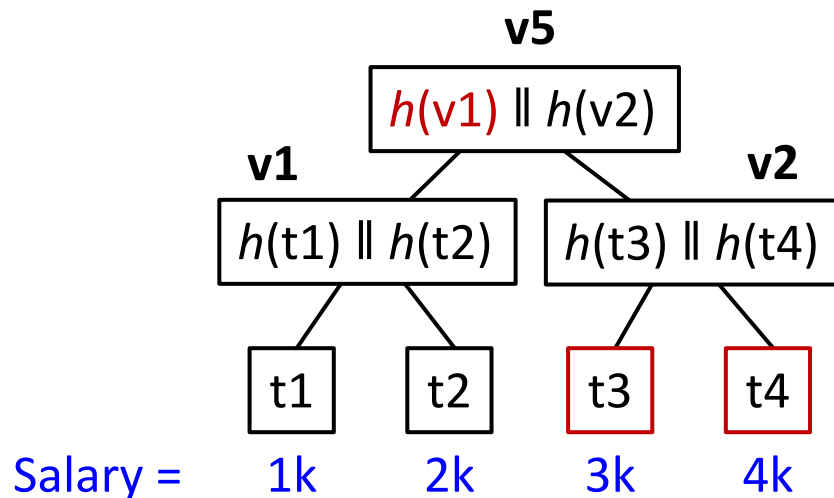
- Consider the above Merkle tree
- Re-consider the query on “Salary = 3.5k”
- Option 1:
 - The service provider could return $t1$, $t2$, $t3$, $t4$, as well as the encrypted $h(v5)$
 - The user could then compute $h(t1)$, $h(t2)$, $h(t3)$, $h(t4)$
 - Based on that, she computes $h(v1)$ and $h(v2)$
 - Then she can compute the hash of $h(v1) \parallel h(v2)$ and verify it against the encrypted $h(v5)$
 - Then she can be sure that the data has only $t1$, $t2$, $t3$, and $t4$; so no “Salary = 3.5k”

Merkle Tree: Example




- Option 1: The user computes
 - ▣ $h(t1), h(t2), h(t3), h(t4)$,
 - ▣ and then $h(v1) \parallel h(v2)$
 - ▣ and then verify it again $h(v5)$
- Question: does the user really need $t1$ and $t2$?
 - ▣ No; She only needs $h(v1)$
- That is, given $t3, t4$, and $h(v1)$, the user can already verify the query result against $h(v5)$

Merkle Tree: Example




- So the service provider could simply return $t3$, $t4$, and $h(v1)$, as well as the encrypted $h(v5)$
- This is sufficient for the user to verify the answer for “Salary = 3.5k”

Merkle Tree: Example



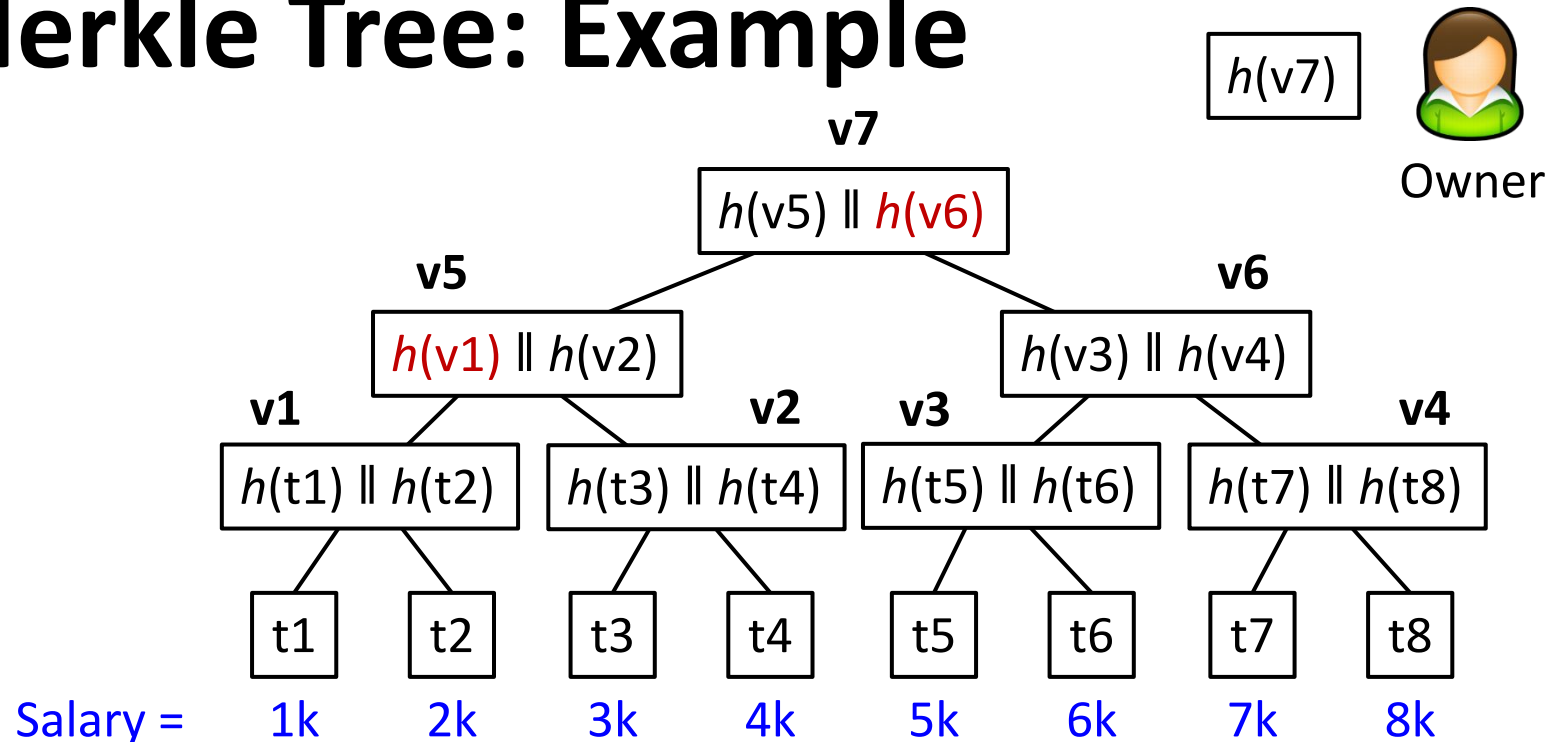
- Now reconsider the query on “Salary = 3.5k”
- The service provider could return the following:
 - $t1, t2, t3, t4, t5, t6, t7, t8$, and the encrypted $h(v7)$
- But $t5, t6, t7, t8$ could be replaced by $h(v6)$

Merkle Tree: Example



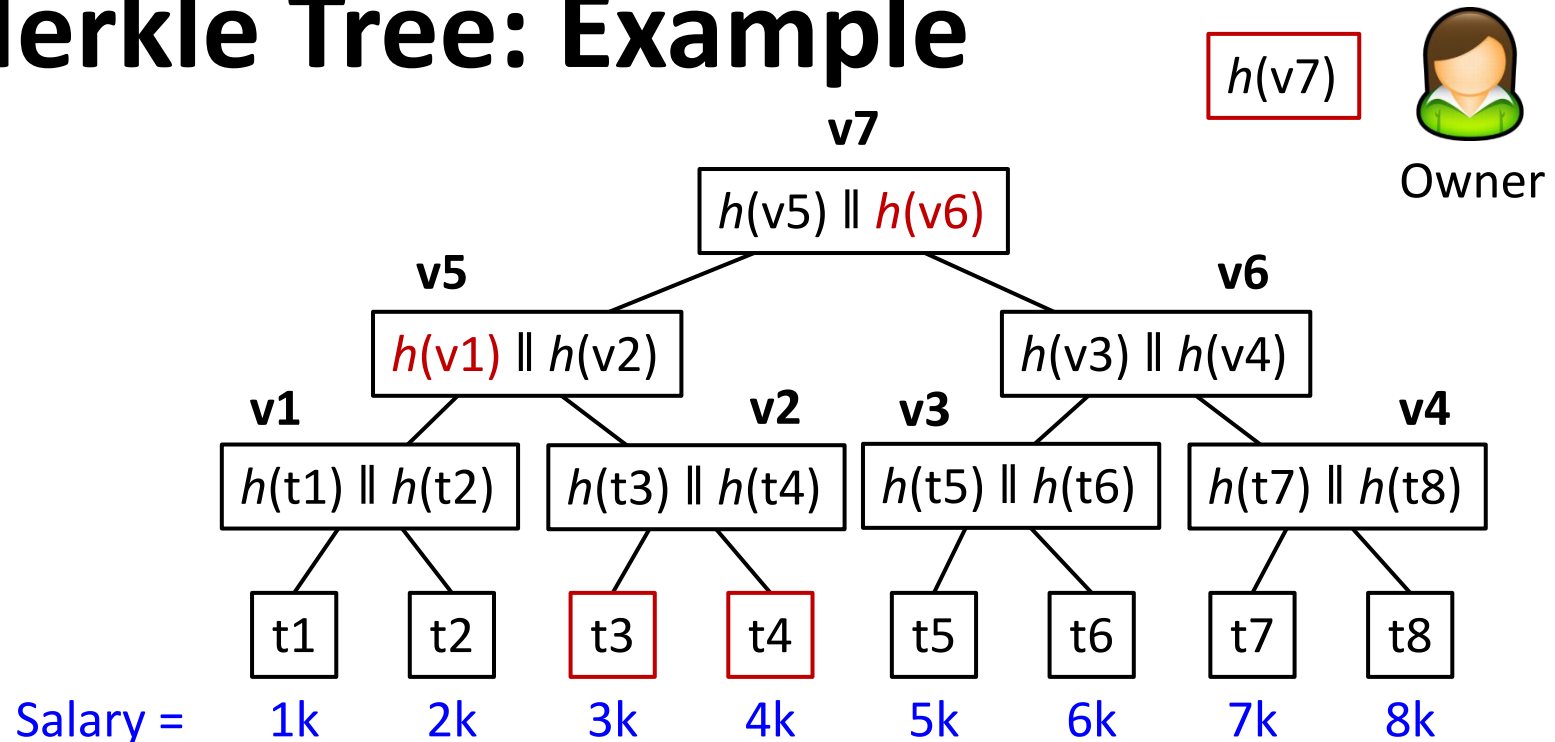
- Now reconsider the query on “Salary = 3.5k”
- The service provider could return the following:
 - $t1, t2, t3, t4, h(v6)$, and the encrypted $h(v7)$
- But $t1, t2$ could be replaced by $h(v1)$

Merkle Tree: Example




- Now reconsider the query on “Salary = 3.5k”
- The service provider could return the following:
 - $h(v1)$, $t3$, $t4$, $h(v6)$, and the encrypted $h(v7)$
- Any more replacement?
 - No; we definitely need $t3$, $t4$, and the encrypted $h(v7)$

Merkle Tree: Example




- Now reconsider the query on “Salary = 3.5k”
- Final answer: The service provider returns the following:
 - $h(v1)$, $t3$, $t4$, $h(v6)$, and the encrypted $h(v7)$

Merkle Tree: Example



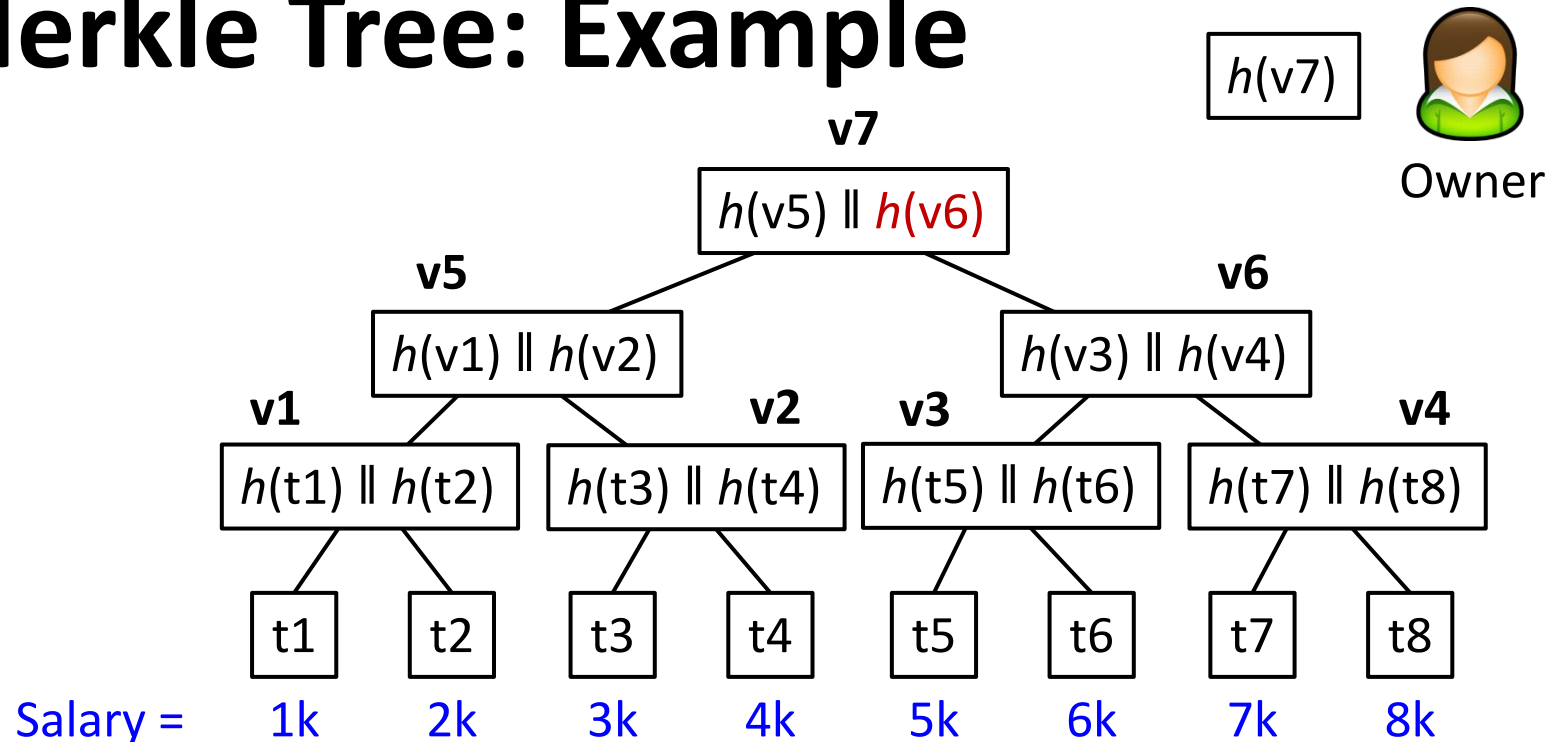
- Now consider a query on “Salary > 2.5k and Salary < 3.5k”
- The service provider could return the following:
 - $t1, t2, t3, t4, t5, t6, t7, t8$, and the encrypted $h(v7)$
- Any replacement possible?
 - $t5, t6, t7, t8$ could be replaced by $h(v6)$

Merkle Tree: Example



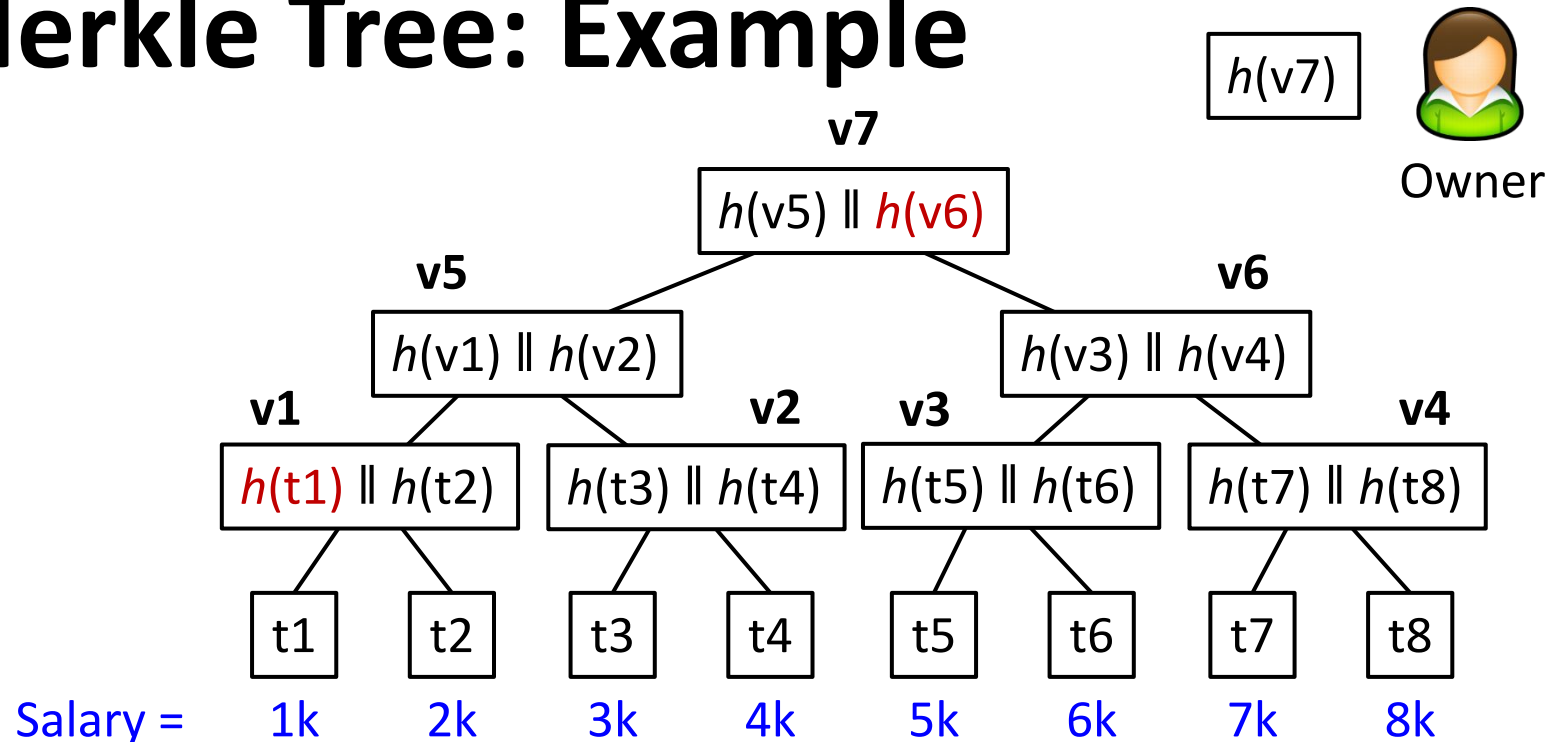
- Now consider a query on “Salary > 2.5k and Salary < 3.5k”
- The service provider could return the following:
 - t1, t2, t3, t4, $h(v6)$, and the encrypted $h(v7)$
- Could we replace t1, t2 with $h(v1)$?
 - No; otherwise, the user cannot verify whether the service provider has hidden a tuple with Salary = 2.6k

Merkle Tree: Example



- Now consider a query on “Salary > 2.5k and Salary < 3.5k”
- The service provider could return the following:
 - $t1, t2, t3, t4, h(v6)$, and the encrypted $h(v7)$
- Could we replace $t1$ with $h(t1)$?
 - This is OK

Merkle Tree: Example

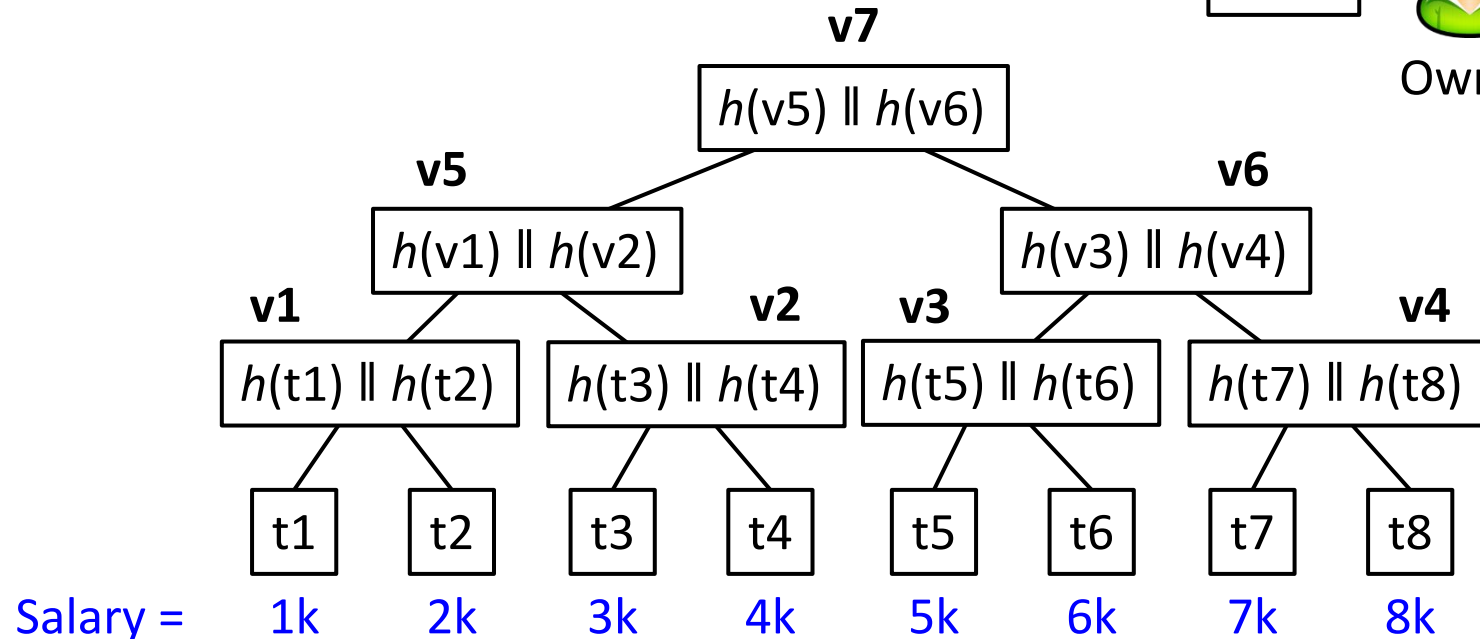



- Now consider a query on “Salary > 2.5k and Salary < 3.5k”
- The service provider could return the following:
 - $h(t1)$, $t2$, $t3$, $t4$, $h(v6)$, and the encrypted $h(v7)$
- Any more replacement?
 - No; we definitely need $t2$, $t3$, $t4$ to prove correctness

Merkle Tree: General Algorithm

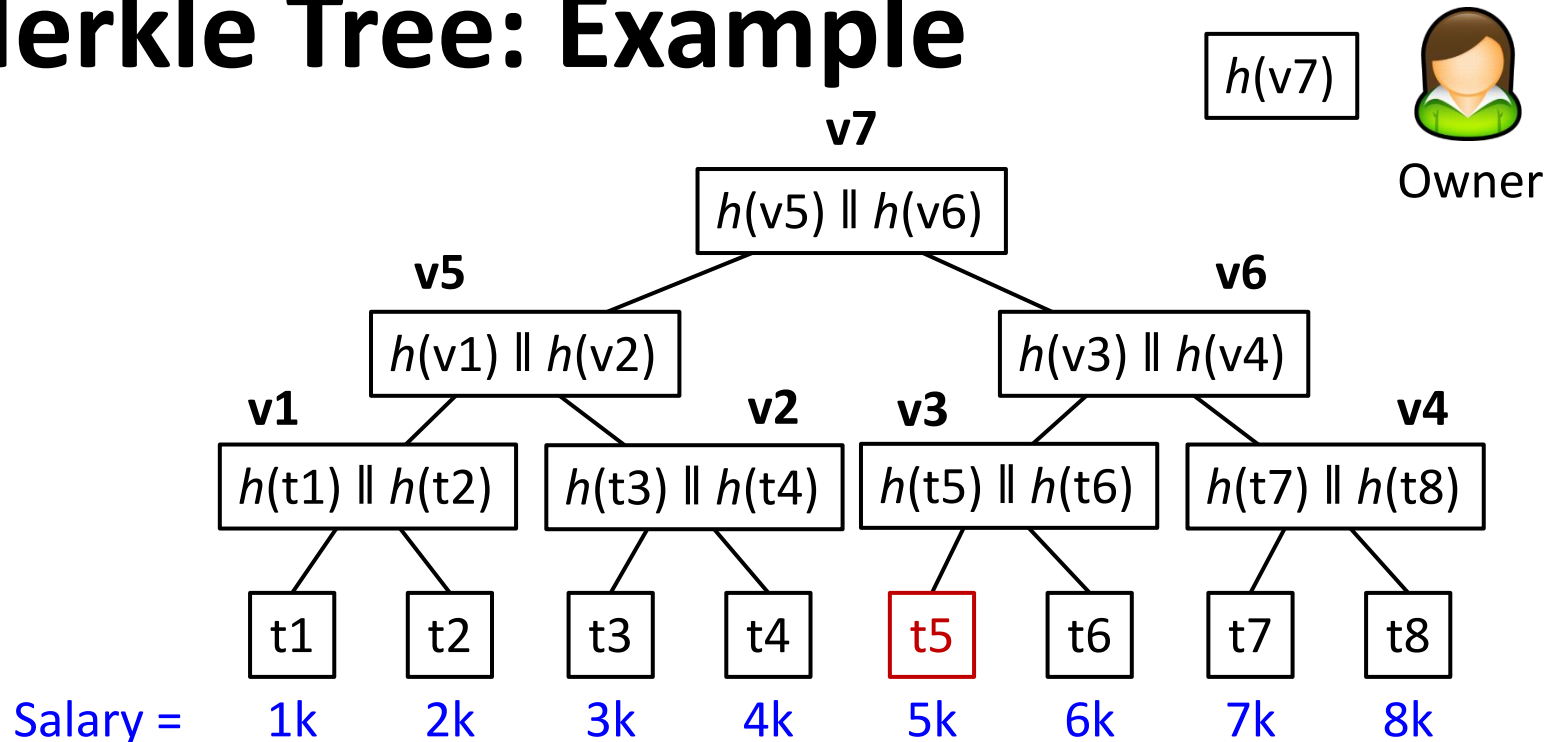
- Consider a query on T.A in $[x, y]$
- Among the tuples t with $t.A < x$, find the tuple t_x whose A value is the largest
 - Identify the path from t_x to the root of the Merkle tree
 - For every “left branch” on the path, collect the hash value of the branch
- Among the tuples t with $t.A > y$, find the tuple t_y whose A value is the smallest
 - Identify the path from t_y to the root of the Merkle tree
 - For every “right branch” on the path, collect the hash value of the branch
- Return t_x , t_y , and all tuples between them, and all hash values collected, as well as the encrypted Merkle root

Merkle Tree: Example



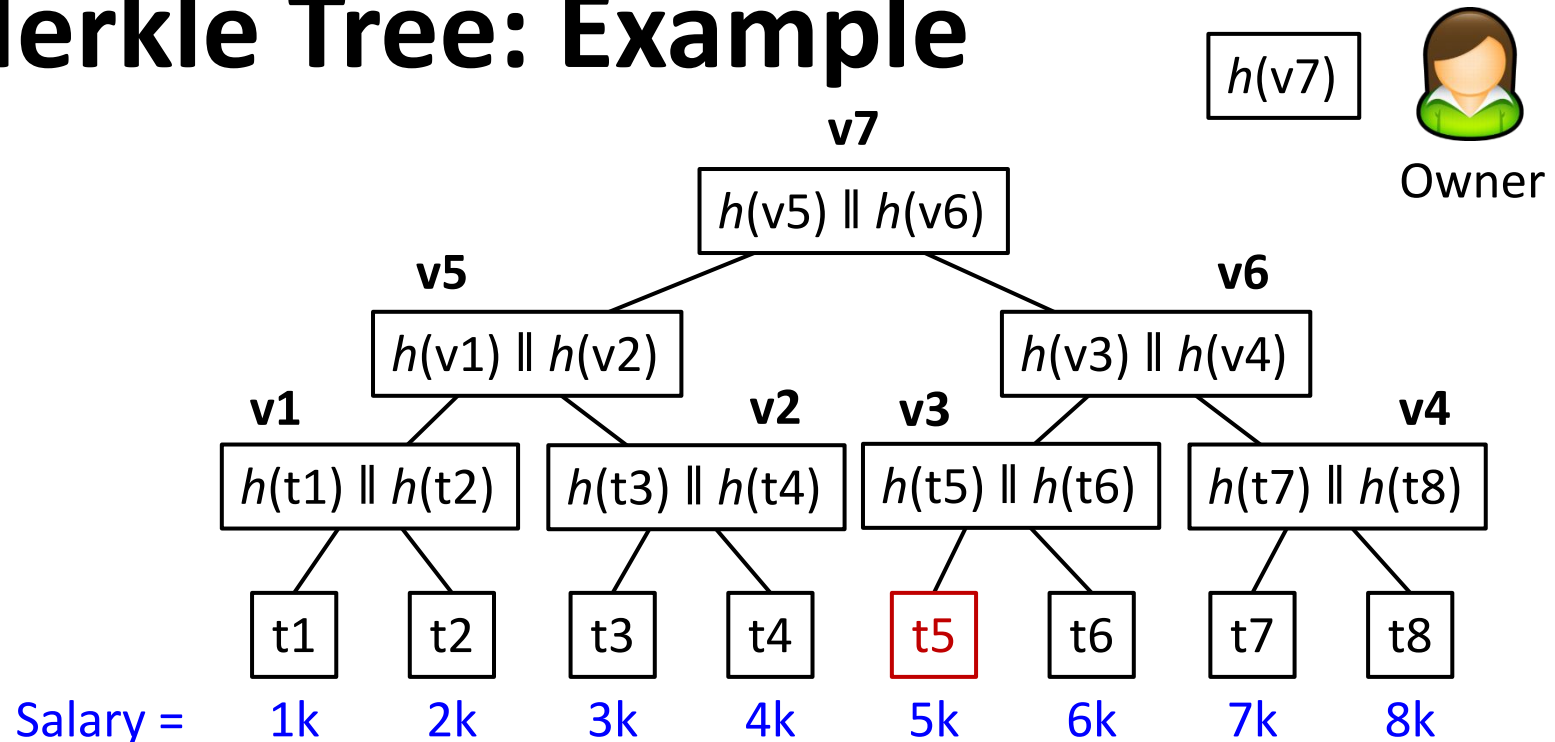
- Consider a query on “Salary in [5.5k, 6.5k]”
 - i.e., A is Salary, and $[x, y] = [5.5k, 6.5k]$
- “Among the tuples t with $t.A < x$, find the tuple **tx** whose A value is the largest”
 - **tx** is t5

Merkle Tree: Example



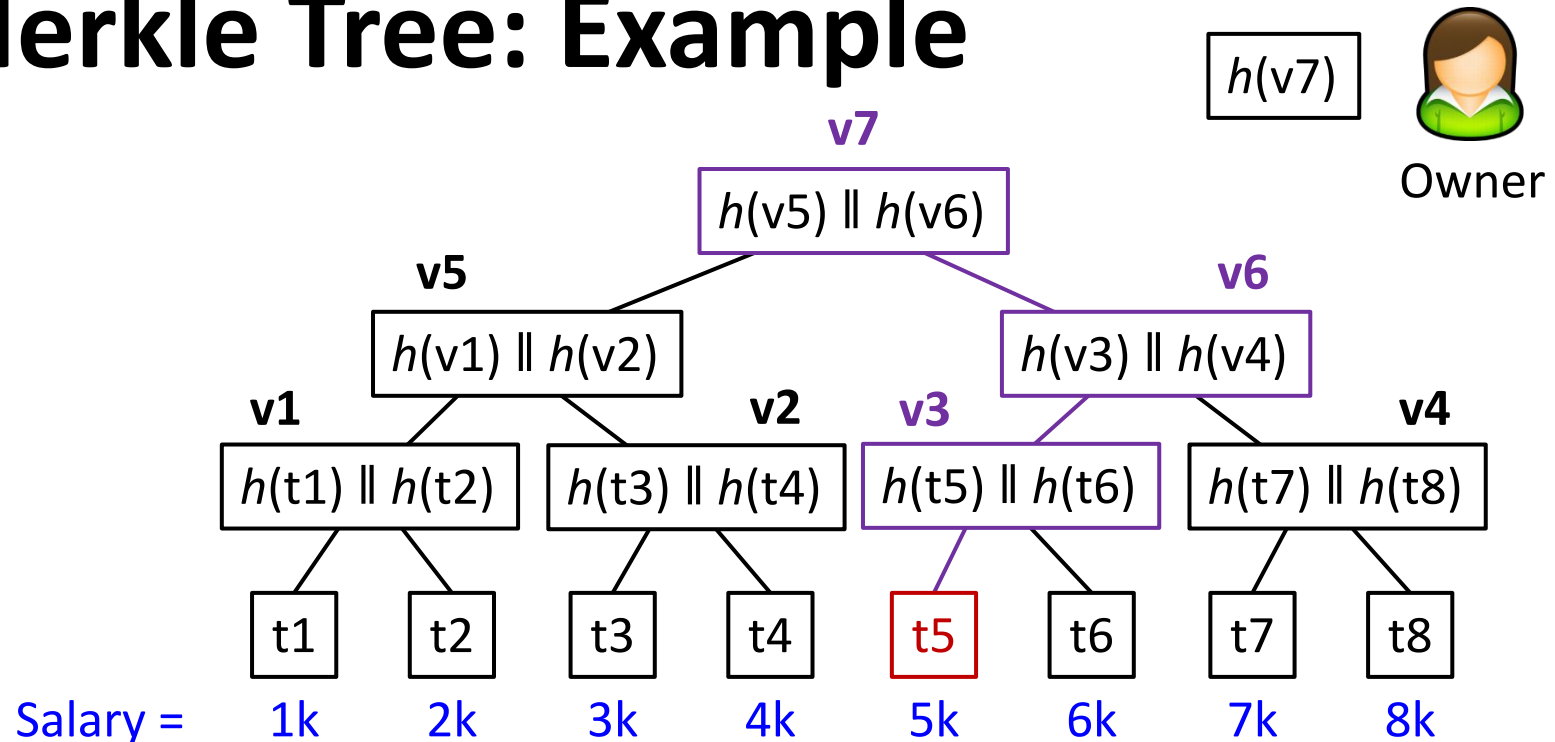
- Consider a query on “Salary in [5.5k, 6.5k]”
 - i.e., A is Salary, and $[x, y] = [5.5k, 6.5k]$
- “Among the tuples t with $t.A < x$, find the tuple tx whose A value is the largest”
 - tx is $t5$

Merkle Tree: Example




- Consider a query on “Salary in [5.5k, 6.5k]”
 - i.e., A is Salary, and $[x, y] = [5.5k, 6.5k]$
- “Identify the path from **tx** to the root of the Merkle tree”
 - The **path** from t5 to v7

Merkle Tree: Example



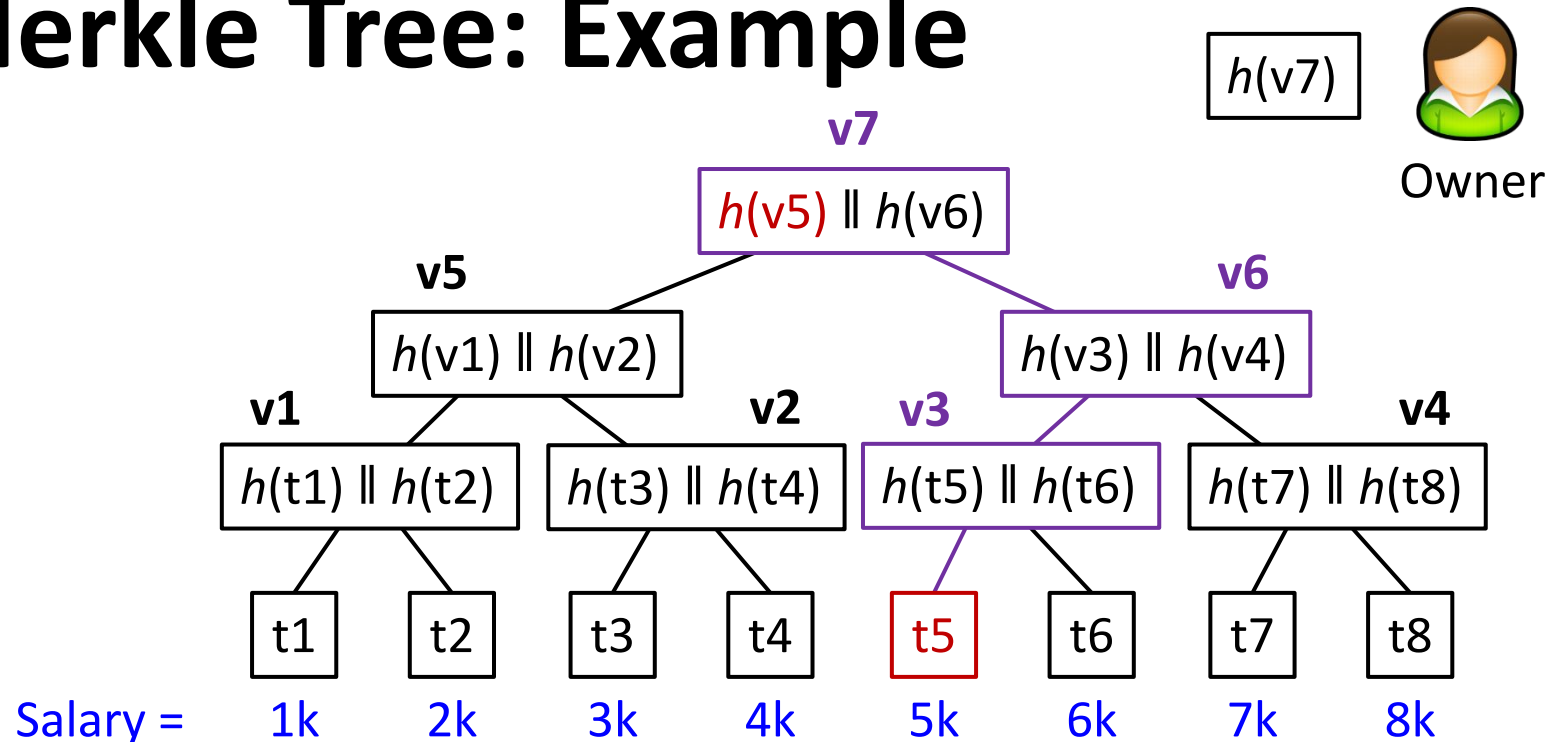
- Consider a query on “Salary in [5.5k, 6.5k]”
 - i.e., A is Salary, and $[x, y] = [5.5k, 6.5k]$
- “Identify the path from tx to the root of the Merkle tree”
 - The path from $t5$ to $v7$

Merkle Tree: Example




- Consider a query on “Salary in [5.5k, 6.5k]”
 - i.e., A is Salary, and $[x, y] = [5.5k, 6.5k]$
- “For every left branch on the **path**, collect the hash value of the branch”
 - Collected hash: $h(v5)$

Merkle Tree: Example




- Consider a query on “Salary in [5.5k, 6.5k]”
 - i.e., A is Salary, and $[x, y] = [5.5k, 6.5k]$
- “For every left branch on the **path**, collect the hash value of the branch”
 - Collected hash: $h(v5)$

Merkle Tree: Example



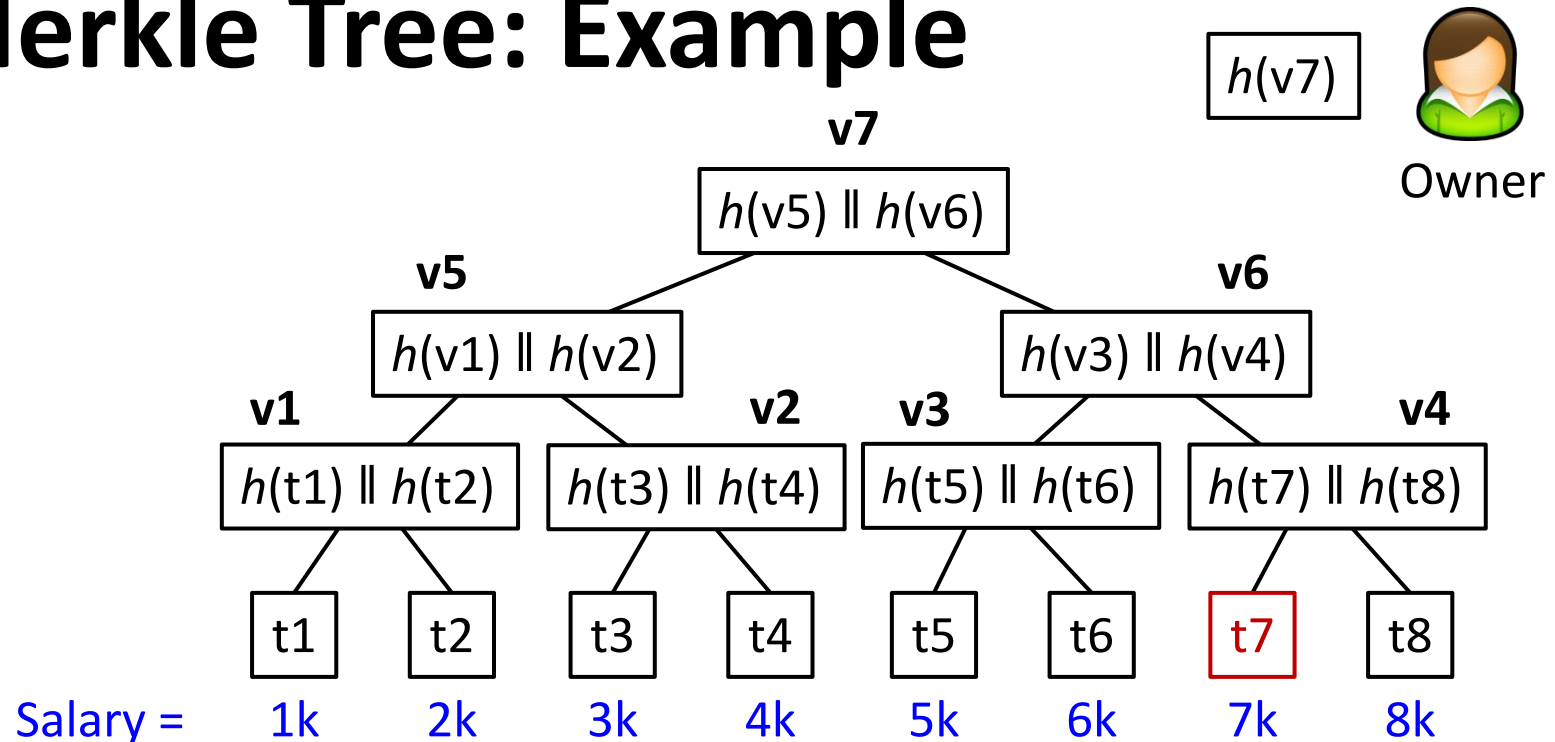
- Consider a query on “Salary in $[5.5k, 6.5k]$ ”
 - i.e., A is Salary, and $[x, y] = [5.5k, 6.5k]$
- “Among the tuples t with $t.A > y$, find the tuple **ty** whose A value is the smallest”
 - **ty** is $t7$

Merkle Tree: Example




- Consider a query on “Salary in [5.5k, 6.5k]”
 - i.e., A is Salary, and $[x, y] = [5.5k, 6.5k]$
- “Among the tuples t with $t.A > y$, find the tuple **ty** whose A value is the smallest”
 - **ty** is t7

Merkle Tree: Example



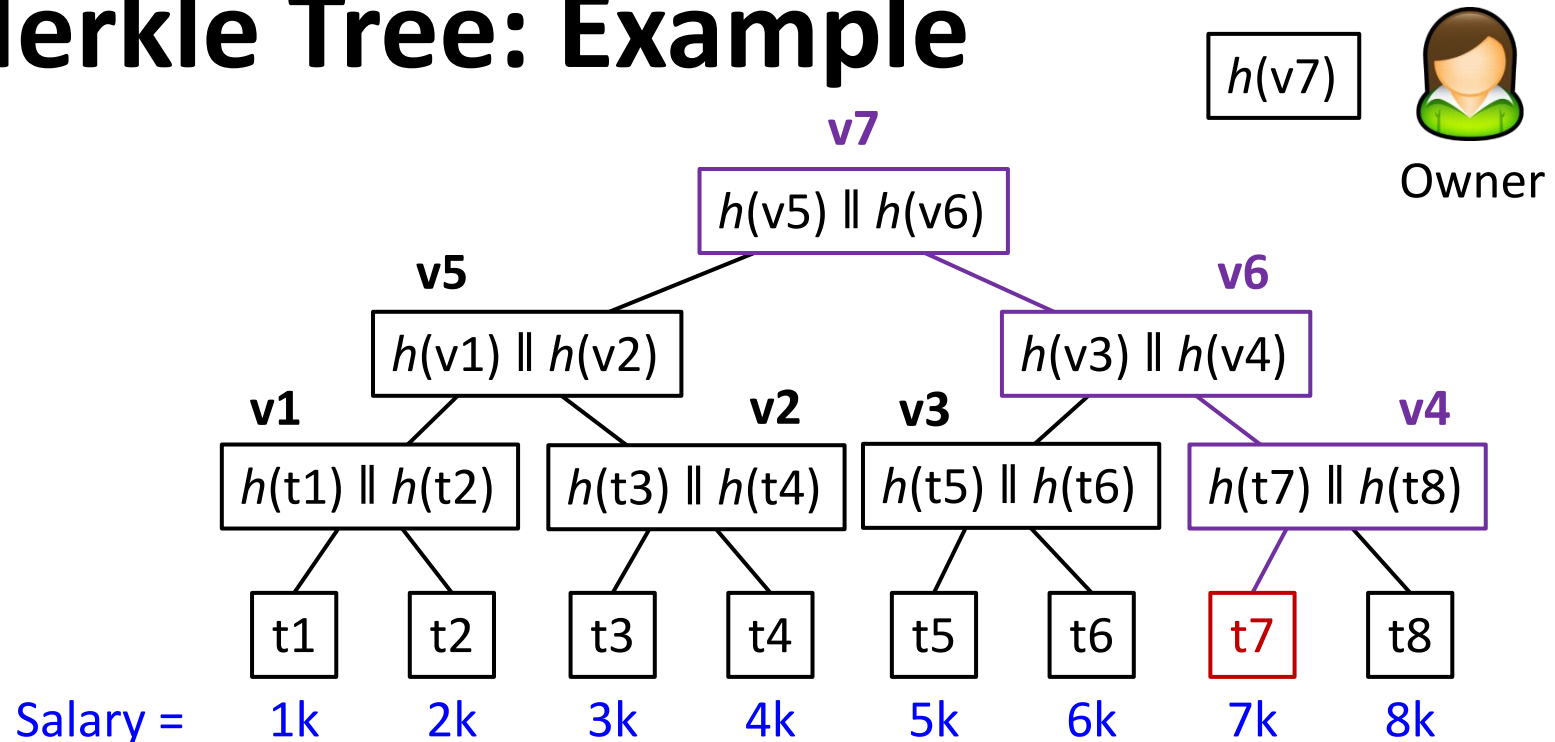
- Consider a query on “Salary in [5.5k, 6.5k]”
 - i.e., A is Salary, and $[x, y] = [5.5k, 6.5k]$
- “Identify the **path** from t_7 to the root of the Merkle tree”
 - the **path** from t_7 to v_7

Merkle Tree: Example



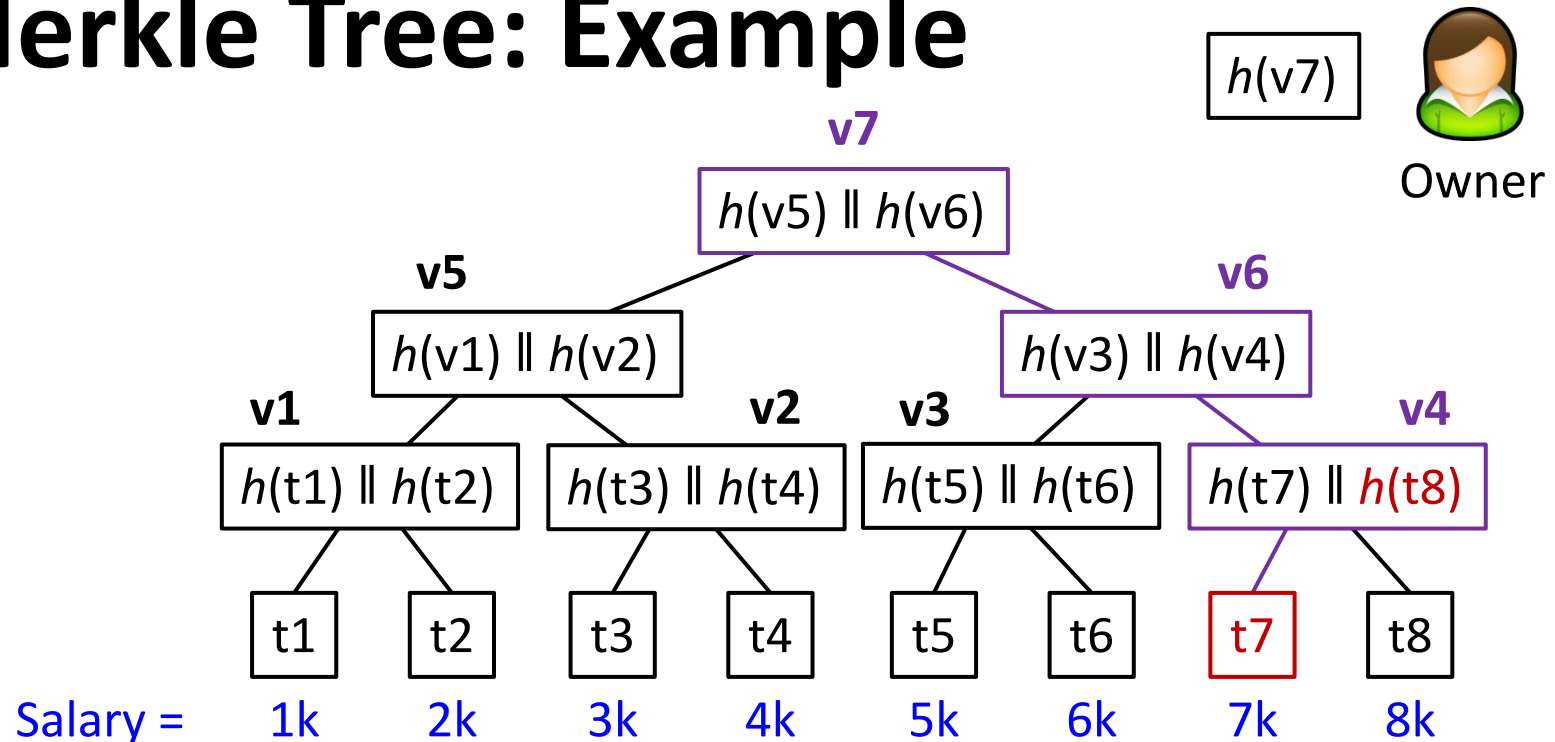
- Consider a query on “Salary in [5.5k, 6.5k]”
 - i.e., A is Salary, and $[x, y] = [5.5k, 6.5k]$
- “Identify the **path** from t_7 to the root of the Merkle tree”
 - the **path** from t_7 to v_7

Merkle Tree: Example



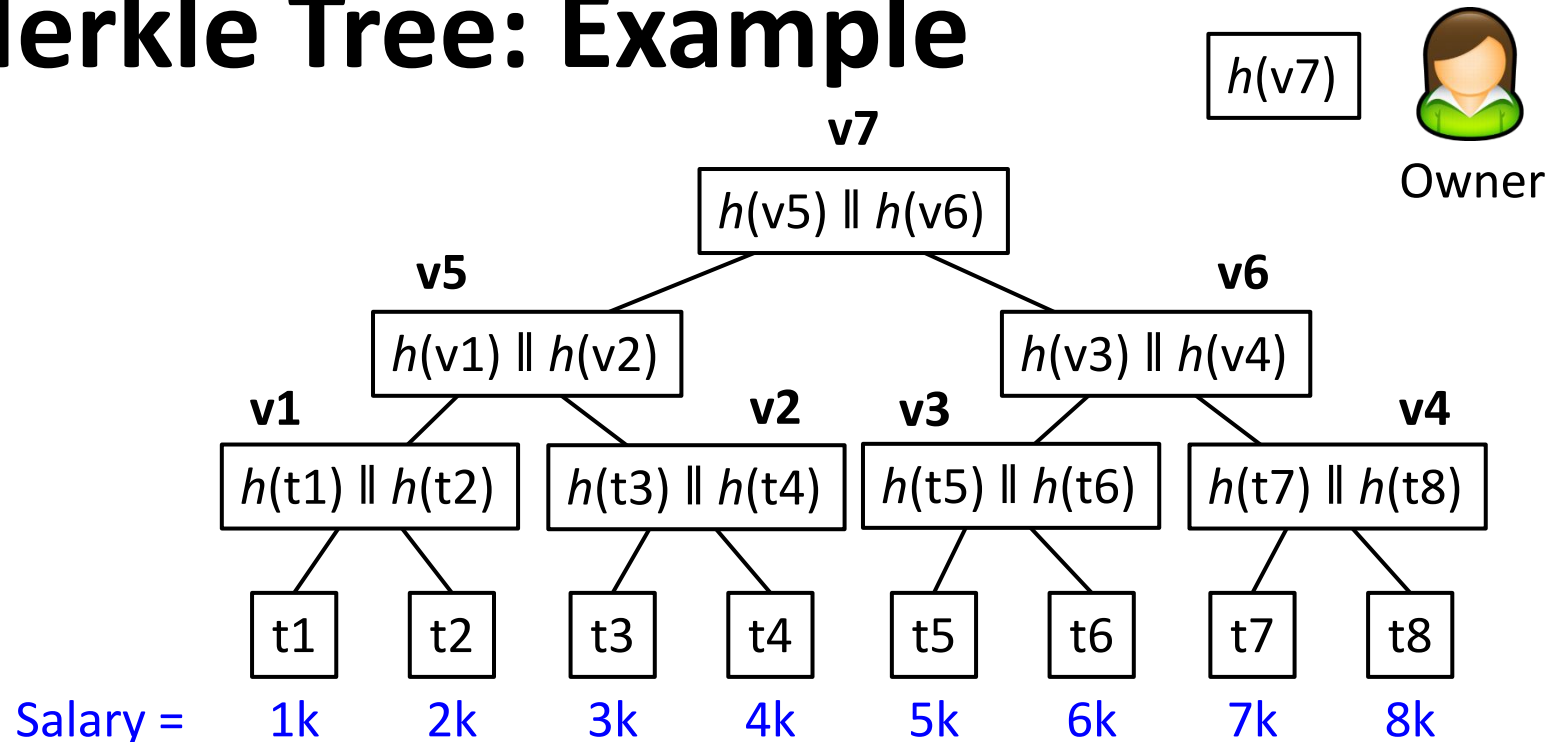
- Consider a query on “Salary in [5.5k, 6.5k]”
 - i.e., A is Salary, and $[x, y] = [5.5k, 6.5k]$
- “For every right branch on the path, collect the hash value of the branch”
 - Collected hashes: $h(t8)$

Merkle Tree: Example



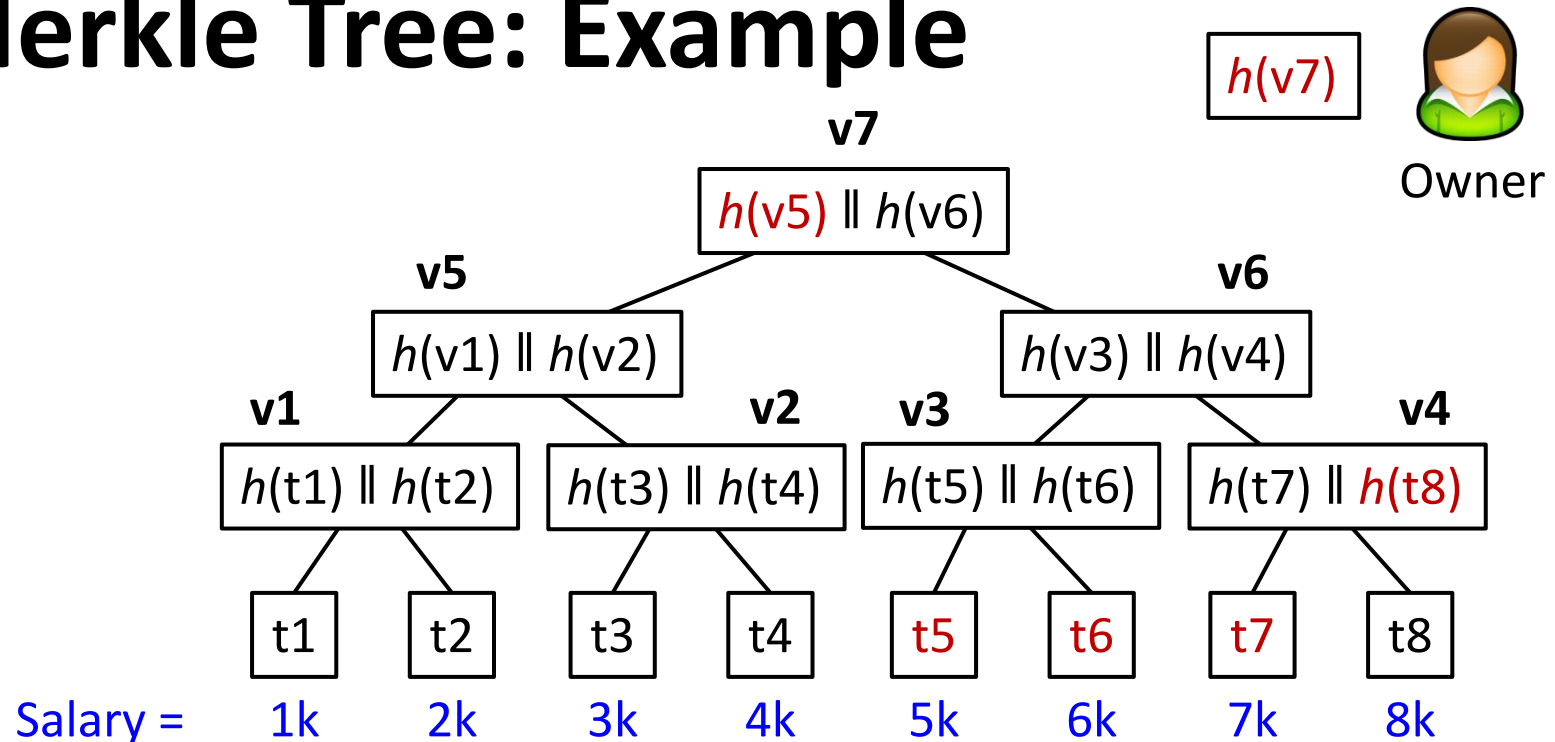
- Consider a query on “Salary in $[5.5k, 6.5k]$ ”
 - i.e., A is Salary, and $[x, y] = [5.5k, 6.5k]$
- “For every right branch on the path, collect the hash value of the branch”
 - Collected hashes: $h(t8)$

Merkle Tree: Example



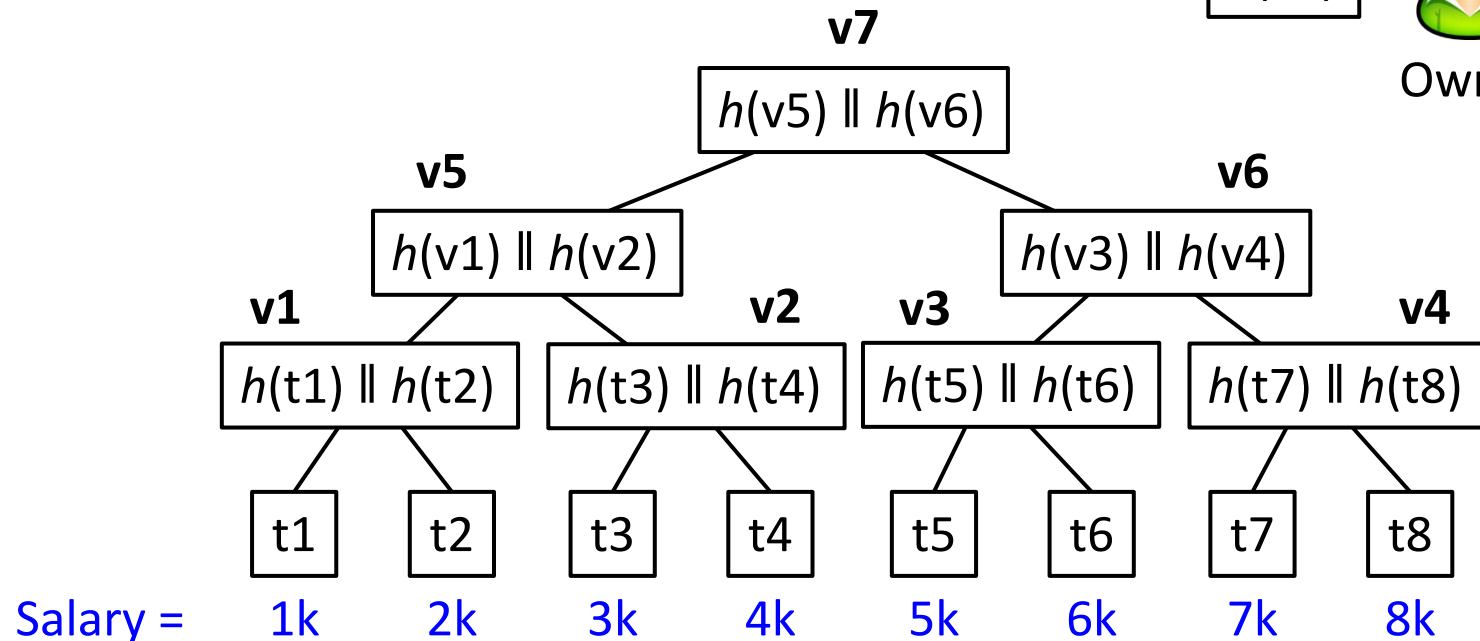

- Consider a query on “Salary in [5.5k, 6.5k]”
 - i.e., A is Salary, and $[x, y] = [5.5k, 6.5k]$
- “Return t_x , t_y , and **all tuples between them**, and all hash values collected, as well as the encrypted Merkle root”
 - i.e., t_5 , t_7 , and **t_6** , and $h(v_5)$ and $h(t_8)$, and the encrypted $h(v_7)$

Merkle Tree: Example





- Consider a query on “Salary in [5.5k, 6.5k]”
 - i.e., A is Salary, and $[x, y] = [5.5k, 6.5k]$
- “Return tx, ty, and all tuples between them, and all hash values collected, as well as the encrypted Merkle root”
 - i.e., t5, t7, and t6, and $h(v5)$ and $h(t8)$, and the encrypted $h(v7)$

Merkle Tree: Exercise



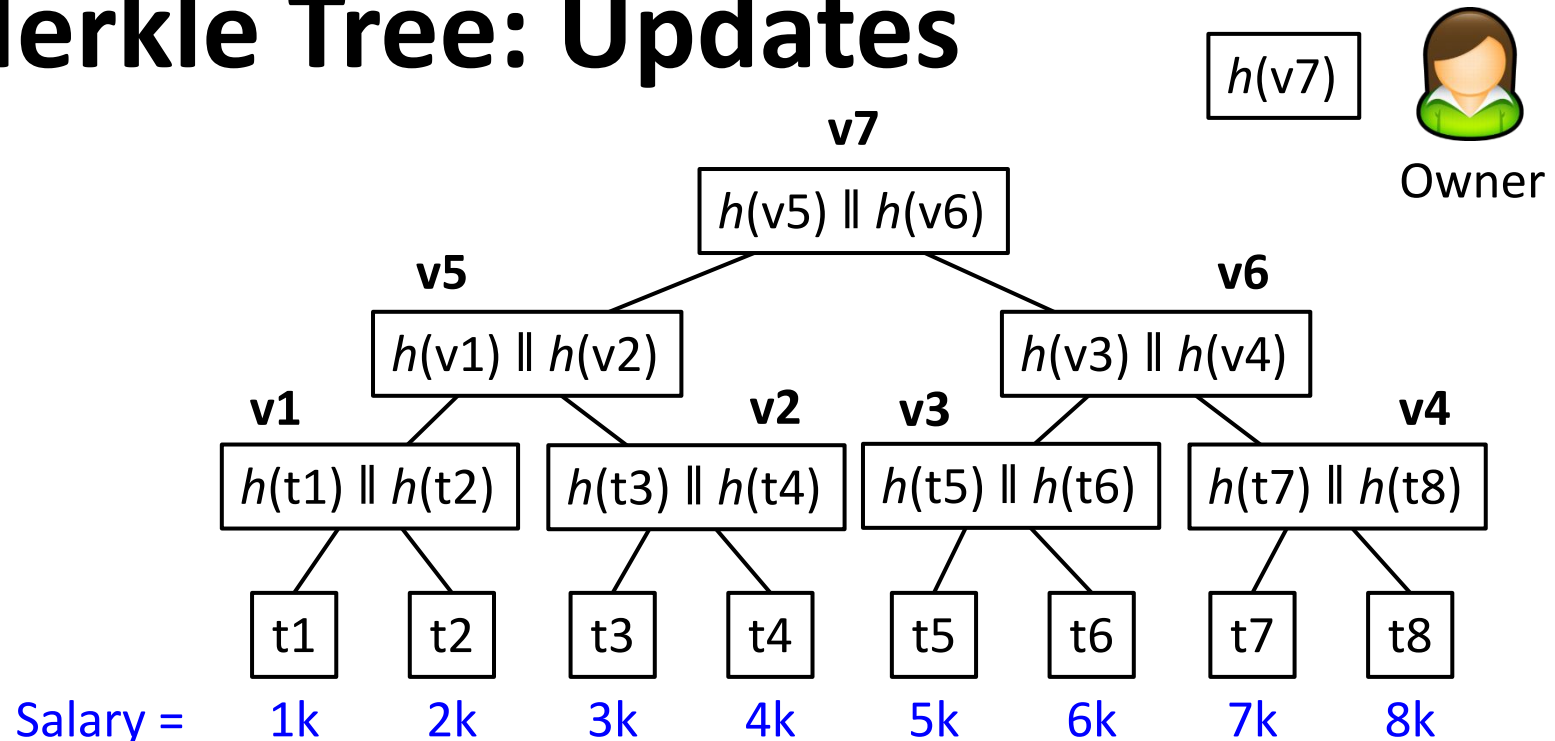
- Consider a query on “Salary = 4k”
- What should the service provider return?

Merkle Tree: Exercise



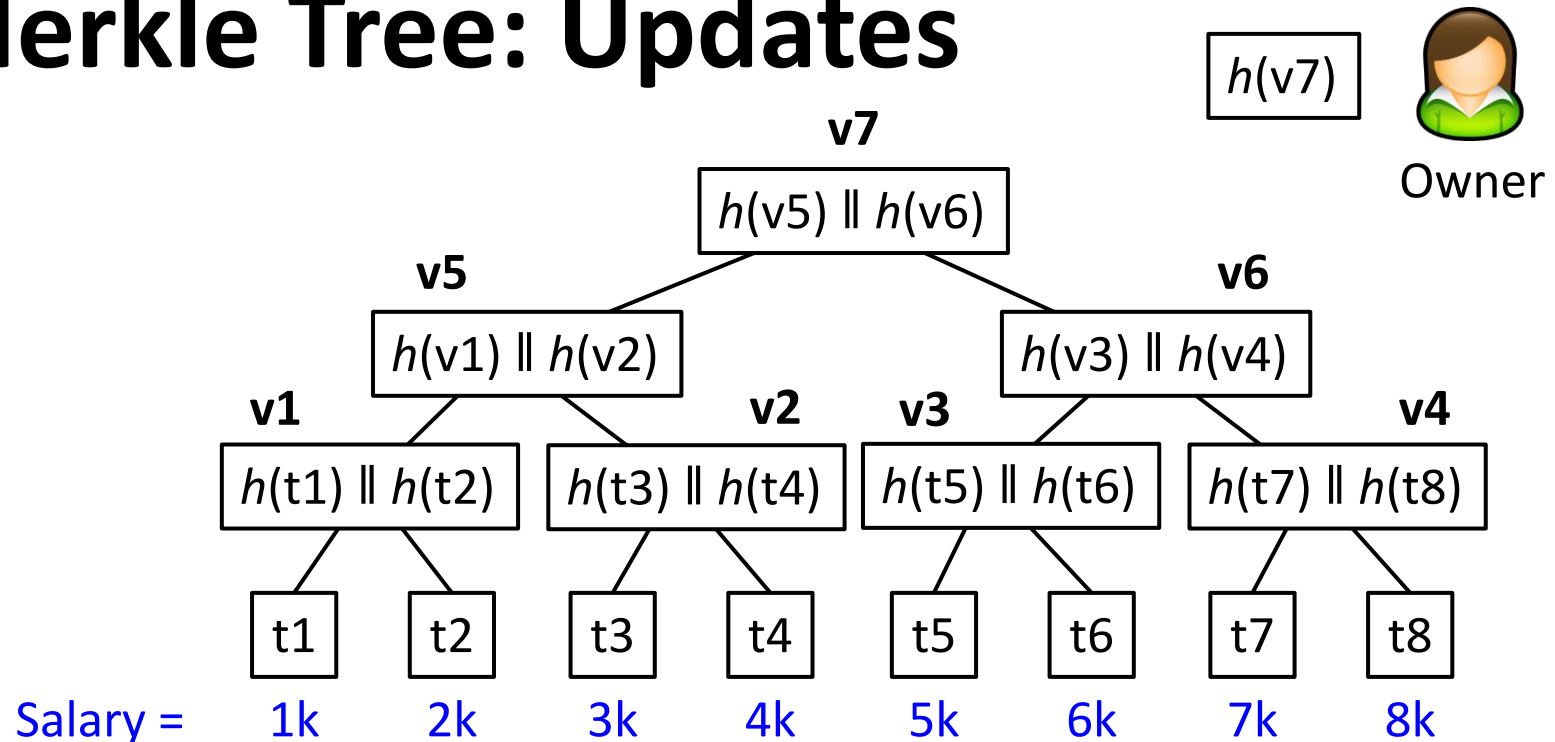
- Consider a query on “Salary in (3k, 4k]”
- What should the service provider return?

Merkle Tree: Updates



- What if the tuples are updated?
 - E.g., the tuple with Salary = 6k is deleted
- Option 1:
 - The data owner reconstructs the binary tree, re-computes $h(\text{root})$, signs it with a timestamp, and sends it to the service provider
 - The data owner announces the timestamp to users

Merkle Tree: Updates



- Problem with Option 1:
 - Reconstructing the whole binary tree once per update is time-consuming
- Improved solution:
 - Use an update-friendly tree structure, e.g., a red-black tree
 - Any tree structure could be signed by the data owner like a Merkle tree

Extension to Multi-Dimensional Data

- The previous discussions focus on one-dimensional queries
- How about multi-dimensional queries? e.g.,
 - `SELECT * FROM Employee
WHERE Age > 30 AND Salary > 10000`
- Option 1:
 - Build two Merkle trees on Age and Salary, respectively
- Problem:
 - No benefit from using two trees simultaneously
 - We need to choose either the tree on Age or the one on Salary to answer a query

Extension to Multi-Dimensional Data

- The previous discussions focus on one-dimensional queries
- How about multi-dimensional queries? e.g.,
 - `SELECT * FROM Employee`
`WHERE Age > 30 AND Salary > 10000`
- Option 2:
 - Use multi-dimensional indices
 - E.g., an R-tree

R-tree

name	semester	credits
A	8	100
B	4	10
C	6	35
D	1	10
E	6	40
F	5	45
G	7	85
H	3	20
I	10	70
J	2	30
K	8	50
L	4	50

